

S.O.L.A.R.I.S.

Solar-Optimized Land Autonomous Recharge & Intelligence System

Authors:
(Group #12)

Rafael Puig
Electrical Engineering

Ryan Markowitz
Computer Engineering

Nathan Hammond
Electrical Engineering

Garrett Fortier
Computer Engineering

Reviewer Committee:

Dr. Maddox
Dr. Rakshan
Dr. Sahawneh

ECE
ECE
ECE

Adjunct Professor
Assistant Professor
Lecturer

Mentor:
Dr. Suboh



Table of Contents

List of Tables	x
List of Figures	xi
1 Introduction.....	1
2 Project Description.....	2
2.1 Project Background and Motivation	2
2.2 Current Commercial Technologies and Existing Projects	3
2.2.1 Jackery Solar Mars Bot.....	3
2.2.2 SolarX: DIY STEM Solar Powered Robot	4
2.2.3 Summary	4
2.3 Goals and Objectives	5
2.3.1 Basic Goals Summarized	5
2.3.2 Advanced Goals Summarized.....	5
2.3.3 Stretch Goals Summarized.....	5
2.3.4 Basic Goals	6
2.3.5 Advanced Goals	6
2.3.6 Stretch Goals	7
2.3.7 Basic Objectives.....	8
2.3.8 Advanced Objectives	8
2.3.9 Stretch Objectives	8
2.4 Required Specifications	9
2.4.1 Overall System Specifications	9
2.4.2 Component System Specifications	10
2.4.3 Hardware and Software Diagrams + Prototype Illustration.....	11
2.5 Design Requirements	15
2.5.1 Marketing & Engineering Requirements	15
2.5.2 House of Quality	17
3 Hardware & Software Research.....	18
3.1 Hardware Technologies	18
3.1.1 Brief	18
3.1.2 Control Unit	18

3.1.2.1	What is a MCU?	18
3.1.2.1.1	ESP32-S3.....	19
3.1.2.1.2	ESP32-P4.....	20
3.1.2.1.3	MSP430FR6989.....	21
3.1.2.1.4	STM32U0	21
3.1.2.2	MCU Vs. MPU	22
3.1.2.3	MCU Vs. FPGA.....	23
3.1.2.4	Control Unit Part Selection.....	24
3.1.3	Drive Motors.....	24
3.1.3.1	Brushed DC Motor	24
3.1.3.2	Brushless DC Motor.....	24
3.1.3.3	Stepper Motor	25
3.1.3.4	AC Motor.....	25
3.1.3.5	AC Motors Vs. DC Motors	26
3.1.3.6	Drive Motor Type Selection.....	26
3.1.3.7	Brushed Vs. Brushless DC Motors.....	27
3.1.3.8	DC Motor Type Selection.....	27
3.1.3.9	Drive DC Motor Selection	28
3.1.3.10	StepperOnline Brushed Gear Motor.....	29
3.1.3.11	E-S Motor Planetary Gear Motor	30
3.1.3.12	goBILDA Yellow Jacket Planetary Gear Motor	30
3.1.3.13	Drive DC Motor Selection.....	30
3.1.4	Solar Panel Pan/Tilt Motor Selection	31
3.1.4.1	Solar Panel Pan/Tilt Motor Selection	31
3.1.4.2	goBILDA Yellow Jacket Planetary Gear Motor	31
3.1.4.3	E-S Motor Worm Gear Motor w/ Encoder	32
3.1.4.4	FeeTech Digital Servo Motor	32
3.1.4.5	Solar Panel Pan/Tilt Motor Part Selection.....	32
3.1.5	Motor Drivers.....	33

3.1.5.1	Unidirectional Drivers	33
3.1.5.2	Half-Bridge Drivers	33
3.1.5.3	Relay-Based Driver:	33
3.1.5.4	H-Bridge Driver.....	33
3.1.5.5	Motor Driver Selection	34
3.1.5.6	Motor Driver Components	35
3.1.5.7	Motor Driver Component Selection	36
3.1.6	Solar Panel Technologies.....	36
3.1.6.1	What is a Solar Panel?.....	36
3.1.6.2	Monocrystalline Silicon.....	37
3.1.6.3	Polycrystalline Silicon	38
3.1.6.4	Thin-Film Silicon.....	38
3.1.6.5	Solar Panel Technology Selection	39
3.1.7	Solar Panel Parts	39
3.1.7.1	ACOPOWER HY020-12M	40
3.1.7.2	Newpowa NPA20S-12J	41
3.1.7.3	Ameresco C1D2-420J	41
3.1.7.4	SLD TECH ST-20P-12	41
3.1.7.5	Solar Panel Part Selection	42
3.1.8	Solar Panel Charge Controller Technologies.....	42
3.1.8.1	What is a Solar Panel Charge Controller?	42
3.1.8.2	Maximum Power Point Tracker	43
3.1.8.3	Pulse Width Modulation	43
3.1.8.4	Solar Panel Charge Controller Technology Selection	43
3.1.9	Solar Panel Charge Controller Parts	44
3.1.9.1	GV-5-LI-14.2V	45
3.1.9.2	GV-5-MOD-LFP.....	45
3.1.9.3	Victron SmartSolar MPPT 75/10	46
3.1.9.4	RNG-CTRL-RVR20-US.....	46
3.1.9.5	Solar Panel Charge Controller Part Selection	47

3.1.10	Light Sensors	48
3.1.10.1	Photoresistor	48
3.1.10.2	Phototransistor	50
3.1.10.3	Photodiode	52
3.1.10.4	Digital Light Sensor (LUX, I2C)	53
3.1.10.5	Light Sensing Technology Selection	53
3.1.10.6	Phototransistor Structural Type Comparison	54
3.1.10.6.1	BJT-Based Phototransistors	54
3.1.10.6.2	FET-Based Phototransistors	54
3.1.10.6.3	Structural Type Selection	55
3.1.10.7	Phototransistor Current Gain Configuration	55
3.1.10.7.1	Standard Phototransistor	55
3.1.10.7.2	Photodarlington	55
3.1.10.7.3	Gain Configuration Selection.....	55
3.1.10.8	Phototransistor Hardware Comparison.....	56
3.1.10.8.1	VTT9812FH	56
3.1.10.8.2	VTT9814FH	56
3.1.10.8.3	TEPT5600.....	57
3.1.10.8.4	TEPT4400.....	57
3.1.10.9	Phototransistor Selection	57
3.1.11	Energy Monitoring Technologies	57
3.1.11.1	Coulomb Counting.....	58
3.1.11.2	Hybrid (Coulomb Counting + Open Circuit Voltage)	58
3.1.11.3	Energy Monitoring Technology Selection	58
3.1.12	Energy Monitoring Parts.....	59
3.1.12.1	LTC2944	59
3.1.12.2	INA226	59
3.1.12.3	INA228	59
3.1.12.4	Energy Monitoring Part Selection	60
3.1.13	Inertial Measurement Unit (IMU).....	60

3.1.13.1	Accelerometer	60
3.1.13.2	Gyroscope	61
3.1.13.3	Magnetometer	62
3.1.13.4	Degree of Axis	63
3.1.13.4.1	3-DoF	64
3.1.13.4.2	6-DoF	64
3.1.13.4.3	9-DoF	64
3.1.13.5	IMU Comparison	65
3.1.13.5.1	BNO-085	65
3.1.13.5.2	MPU-6050	65
3.1.13.5.3	ICM-20948	66
3.1.13.6	IMU Selection	66
3.1.14	Ultrasonic Sensors	66
3.1.14.1	How Ultrasonic Sensors Work	66
3.1.14.2	Calculation for Beam Cone Coverage	67
3.1.14.3	Ultrasonic Sensor Comparisons	72
3.1.14.3.1	Note on Limited Selection	72
3.1.14.3.2	MB1020-000	73
3.1.14.3.3	MB1232-000	73
3.1.14.3.4	MB1013-000	74
3.1.14.4	Ultrasonic Sensor Selection	74
3.1.15	Chassis Technology Comparison	75
3.1.15.1	3D Printed Filament	75
3.1.15.2	Aluminum	76
3.1.15.3	Carbon Fiber	76
3.1.15.4	Chassis Technology Selection	76
3.1.16	Chassis Part Comparison	77
3.1.16.1	ASA Filament	77
3.1.16.2	PETG Filament	78

3.1.16.3	ABS Filament	78
3.1.16.4	PC Filament	78
3.1.16.5	5052-H32 Aluminum	78
3.1.16.6	6061-T6 Aluminum	78
3.1.16.7	7075-T6 Aluminum	78
3.1.16.8	Chassis Part Selection.....	79
3.2	Programming Languages	79
3.2.1	High-Level Languages.....	79
3.2.1.1	C/C++.....	79
3.2.1.2	Python	80
3.2.1.3	Java	80
3.2.1.4	JavaScript/Typescript.....	81
3.2.1.5	Swift.....	81
3.2.1.6	Dart.....	82
3.2.1.7	Code Selection.....	82
3.2.2	Algorithms	82
3.2.2.1	Reactive Obstacle Avoidance	82
3.2.2.2	Kalman Filters	83
3.2.2.3	Solar Position	83
3.3	Frameworks.....	84
3.3.1	What are Frameworks?	84
3.3.2	Frontend Phone App Frameworks	85
3.3.2.1	React Native	85
3.3.2.2	SwiftUI.....	85
3.3.2.3	Flutter.....	86
3.3.2.4	Frontend Phone App Framework Selection	86
3.3.3	Backend Phone App Frameworks.....	87
3.3.3.1	FastAPI	87
3.3.3.2	Django	88
3.3.3.3	SpringBoot.....	88

3.3.3.4	Express.js	89
3.3.3.5	Backend Phone App Framework Selection.....	89
3.3.4	Embedded Development Frameworks	89
3.3.4.1	Platform I/O	90
3.3.4.2	ESP IDF.....	90
3.4	Wireless Communication Protocols.....	91
3.4.1	Wi-Fi.....	91
3.4.2	Bluetooth Classic	91
3.4.3	Bluetooth Low Energy (BLE).....	92
3.4.4	LoRaWAN	92
3.4.5	Zigbee	93
3.4.6	Wireless Communication Protocol Selection	93
3.5	Power Supply	94
3.5.1	Brief	94
3.5.2	Battery Considerations.....	95
3.5.3	Power Distribution	95
3.5.4	Battery Technology Comparison	96
3.5.4.1	Lithium Ion (NMC)	96
3.5.4.2	Lithium Iron Phosphate	97
3.5.4.3	Lead-Acid	97
3.5.4.4	Battery Technology Selection	97
3.5.5	Battery Part Comparison.....	98
3.5.5.1	Dr.Prepare DBT1220LFP	98
3.5.5.2	Bioenno Power BLF-1220AS.....	98
3.5.5.3	Renogy RBT1220LFP-TM-G1	99
3.5.5.4	Ultralife URB12200	99
3.5.5.5	Battery Part Selection	99
3.5.6	Voltage Regulator Technology Comparison	99
3.5.6.1	Voltage Regulator Considerations	99
3.5.6.2	Switching Regulator	100

3.5.6.3	Linear Regulator	101
3.5.6.4	Voltage Regulator Technology Selection	101
3.5.7	Voltage Regulator Part Comparison	102
3.5.7.1	LM2679.....	103
3.5.7.2	LM2576.....	103
3.5.7.3	LM2596.....	103
3.5.7.4	LT8650S.....	103
3.5.7.5	LTC3633.....	103
3.5.7.6	Voltage Regulator Part Selection	104
3.5.8	Circuit Safety	104
4	Standards and Design Constraints	104
4.1	Standards.....	104
4.1.1	Brief	104
4.1.2	Communication Protocol Standards	105
4.1.2.1	I2C	105
4.1.2.2	UART	106
4.1.2.3	SPI.....	107
4.1.3	Wireless Communication Standards.....	108
4.1.3.1	Bluetooth Low Energy (BLE).....	108
4.1.3.2	IEEE 802.11	109
4.1.4	Health & Safety Standards.....	110
4.1.4.1	IPC-2221.....	110
4.1.5	Documentation Standards	112
4.1.5.1	ISO-5807	112
4.2	Design Constraints	113
4.2.1	Brief	113
4.2.2	Power Constraints	113
4.3	Practical Constraints	114
4.3.1	Economic	114
5	Comparison of ChatGPT with other Similar Platforms.....	115

5.1	Brief	115
5.2	Case Study 1	116
5.3	Case Study 2	117
5.4	Case Study 3	117
5.5	Case Study 4	118
6	Hardware Design	120
6.1	Brief	120
6.2	Schematic Diagrams	121
6.2.1	MCU	121
6.2.2	Sensors	124
6.2.3	Power	126
6.3	Mechanical Design.....	129
7	Software Design.....	131
7.1	Brief	131
7.2	Case and State Diagrams	132
7.2.1	User State Diagram	132
7.2.2	Autonomous Navigation State Diagram	134
7.2.3	Software Case Diagram	135
7.3	Data Transfer	136
7.3.1	BLE Data Transfer	136
7.4	User Interface.....	137
8	System Fabrication/Prototype Construction	142
8.1	Brief	142
8.2	Sensor Boards	143
8.3	Power Boards	146
8.4	MCU Board.....	149
9	System Testing and Evaluation.....	151
9.1	Brief	151
9.2	Hardware and Software Integration	152
9.2.1	ESP32-S3-WROOM-1/1U MCU Dev Kits	152
9.2.2	TEPT5600 Phototransistors	152

9.2.3	MB1020-000 Ultrasonic Sensors.....	153
9.2.4	BNO085 IMU	154
9.2.5	INA228EVM Energy Monitor Unit Dev Kit.....	155
9.2.6	Multiplexer.....	156
9.2.7	Motors	156
9.2.8	NPA20S-I2J Solar Panel.....	157
9.2.9	Charge Controller.....	158
9.3	Phone Application Stack.....	159
9.3.1	NimBLE Setup on ESP32-S3-WROOM-1/1U	159
9.3.2	Phone App.....	159
9.3.2.1	Frontend Integration	159
9.3.2.2	Database Setup and Connection	160
9.3.2.3	Backend Weather Services.....	161
10	Administrative Content.....	161
10.1	Budget.....	161
10.2	Bill of Materials	162
10.3	Distribution of Work.....	163
10.4	Milestones	164
11	Conclusion	164
	Appendix A – Reference.....	166
	Appendix B – AI References	178

List of Tables

Table 2.1	Energy Specifications	9
Table 2.2	Mobility Specifications.....	9
Table 2.3	Component Specifications	10
Table 2.4	Justification Chart for Requirements	16
Table 3.1	MCU Comparison.....	19
Table 3.2	MCU Vs. MPU	22
Table 3.3	MCU Vs. FPGA.....	23
Table 3.4	AC Motors Vs. DC Motors.....	26

Table 3.5 DC Motor Comparison	29
Table 3.6 Brushed DC Motor Comparison	31
Table 3.7 DC Motor Driver Models	35
Table 3.8 Solar Panel Technology Comparison	36
Table 3.9 Solar Panel Part Comparison	39
Table 3.10 Solar Panel Charge Controller Technology Comparison	42
Table 3.11 Solar Panel Charge Controller Part Comparison	44
Table 3.12 Light Sensing Technology Comparison	48
Table 3.13 Phototransistor Comparison.....	56
Table 3.14 Energy Monitoring Technology Comparison	57
Table 3.15 Energy Monitoring Part Comparison.....	59
Table 3.16 Types of Measurement Devices.....	60
Table 3.17 Degrees of Freedom (DoF) Comparison	64
Table 3.18 IMU Unit Comparison	65
Table 3.19 Beam Angle and Offset Calculation Comparison.....	71
Table 3.20 Ultrasonic Comparison Table	72
Table 3.21 Chassis Technology Comparison.....	75
Table 3.22 Chassis Part Comparison	77
Table 3.23 Kalman Filter Reference Terms.....	83
Table 3.24 Frontend Framework Comparison	85
Table 3.25 Backend Framework Comparison	87
Table 3.26 Wireless Communication Protocol Comparison.....	91
Table 3.27 Power Consumption Table.....	95
Table 3.28 Battery Technology Comparison	96
Table 3.29 Battery Part Comparison.....	98
Table 3.30 Voltage Regulator Technology Comparison	100
Table 3.31 Voltage Regulator Part Comparison	102
Table 4.1 Communication Modes for I2C	105
Table 4.2 Documentation Hierarchy for IPC-2221.....	111
Table 4.3 General Symbols Guide for ISO-5807.....	113
Table 7.1 SOLARIS Software Case Overview	135
Table 10.1 Bill of Materials	163
Table 10.2 Distribution of Work.....	163
Table 10.3 Milestone Progress.....	164

List of Figures

Figure 2.1 Jackery Solar Mars Robot	3
Figure 2.2 DIY STEM Solar Powered Robot	4
Figure 2.3 Hardware Block Diagram.....	11

Figure 2.4 SOLARIS Prototype Illustration	12
Figure 2.5 SOLARIS Main Function Software Diagram	13
Figure 2.6 SOLARIS Robot Driving Software Diagram.....	13
Figure 2.7 SOLARIS Phone App Software Diagram	14
Figure 2.8 House of Quality.....	17
Figure 3.1 Microcontroller Hardware Layout.....	18
Figure 3.2 Two-pole DC Motor Model.....	28
Figure 3.3 H-Bridge Topology	34
Figure 3.4 Photovoltaic Cell Diagram	37
Figure 3.5 Genasun GV-5-LI 5A MPPT Solar Buck Controller	45
Figure 3.6 Genasun GV-5-MOD-LFP 5A MPPT Solar Charge Controller	46
Figure 3.7 Victron SmartSolar MPPT 75/10 Charge Controller	46
Figure 3.8 Rover Li 20 Amp MPPT Solar Charge Controller.....	47
Figure 3.9 LDR Light Sensing Circuit.....	48
Figure 3.10 LDR Resistance Vs. Temperature	49
Figure 3.11 LDR Non-Linear relationship between LUX and Resistance	50
Figure 3.12 Phototransistor Light Sensing Circuit	51
Figure 3.13 Photodiode Light Sensing Circuit	52
Figure 3.14 Coulomb Counting Formula.....	58
Figure 3.15 Degrees of Freedom Illustrated	63
Figure 3.16 Two Ultrasonic Sensor Beam Cone Coverage	67
Figure 3.17 One Ultrasonic Sensor Flush with X-Axis	68
Figure 3.18 Ultrasonic Sensor Beam Vs. Ground.....	69
Figure 3.19 Ultrasonic Sensor Angle Pointing Towards the Ground	70
Figure 3.20 MB1020-000 Sensor Chaining Diagram	73
Figure 6.1 MCU Schematic	122
Figure 6.3 MCU Programming Bridge Schematic	123
Figure 6.5 MCU Connectors Schematic	124
Figure 6.6 Energy Monitor Schematic.....	125
Figure 6.7 Light Sensor Schematic.....	126
Figure 6.8 3.3V Regulator Schematic.....	127
Figure 6.9 5V Regulator Schematic.....	128
Figure 6.10 Motor Driver Operating Truth Table.....	128
Figure 6.11 Motor Driver Schematic	129
Figure 6.12 SOLARIS Chassis Without Covers.....	130
Figure 6.13 SOLARIS Chassis With Covers.....	131
Figure 7.1 Software State Diagram with User Interaction.....	132
Figure 7.2 Software State Diagram with Phone User Interface.....	133
Figure 7.3 Autonomous Navigation Sub-User State Diagram.....	134
Figure 7.4 Bluetooth Low Energy (BLE) Data Transfer (Client to ESP).....	136

Figure 7.5 SOLARIS Phone App Home Screen	137
Figure 7.6 SOLARIS Phone Sign Up Screen	138
Figure 7.7 SOLARIS Phone Verification Functionality	139
Figure 7.8 SOLARIS Phone Sign In Page	140
Figure 7.9 SOLARIS Telemetry Page	141
Figure 7.10 SOLARIS Mode Switch Page	142
Figure 8.1 2D Energy Monitor Board Layout	144
Figure 8.2 3D Energy Monitor Board Layout	144
Figure 8.3 2D Light Sensor Board Layout.....	145
Figure 8.4 3D Light Sensor Board Layout.....	146
Figure 8.5 2D Regulator Board Layout	147
Figure 8.6 3D Voltage Regulator Board Layout.....	148
Figure 8.7 2D Motor Driver Board Layout.....	149
Figure 8.8 3D Motor Driver Board Layout.....	149
Figure 8.9 2D MCU Board Layout	150
Figure 8.10 3D MCU Board Layout	151

1 Introduction

The Solar-Optimized Land Autonomous Recharge & Intelligence System (*SOLARIS*) is an autonomously moving robot designed specifically for solar farming. As solar energy adoption accelerates alongside improvements in battery storage and panel efficiency, *SOLARIS* aims to merge renewable energy generation with autonomous robotic technologies to create a more practical and adaptable solution. The primary motivation behind *SOLARIS* is to explore the power-conscious design of a completely self-sufficient system. Traditional stationary solar panels suffer from inefficiencies due to changing environmental conditions and shade. *SOLARIS* addresses this by evaluating the tradeoffs between the energy expended for movement and the additional energy gained by actively repositioning itself toward a light source. To continue, the project serves as a comprehensive educational lesson for Electrical and Computer Engineering students, providing hands-on experience in balancing system costs, low-power constraints, and net-positive charging capabilities.

Currently, there is no commercially realized product that directly compares to the full scope of *SOLARIS*. Some existing concepts that come close are educational level implementations which only demonstrate sensor-driven panel alignment using servos. *SOLARIS* pushes this further by making solar tracking fully autonomous behavior focused on power efficiency. Instead of merely tilting a panel the robot evaluates its environment and determines if simply changing the panel's orientation is insufficient, then it physically navigates to a brighter location, all while ensuring the energy spent moving is less than the energy gained. To accomplish this goal the project implements several distinct engineering disciplines to create a cohesive robot platform. These consist of a control unit to govern the robot. This "brain" must balance low power and decision making. For energy storage we also must consider energy harvesting which directly relies on the power and charging system for the robot. Lastly the robot must receive input from the outside world and move through it at the same time via sensors and other electronic devices.

The remainder of this report provides a comprehensive technical breakdown of the *SOLARIS* platform. It begins with a detailed analysis of the hardware selection process, justifying the choice of the control system and the power supply chemistry based on power efficiency and safety requirements. Following the hardware specifications, the document outlines the structural design and fabrication of the chassis, the development of the custom PCB and power distribution system, and the logic behind the autonomous navigation algorithms. Finally, the report details the integration of the mobile application and summarizes the testing procedures used to validate the system's ability to achieve a net-positive energy yield. Through this multidisciplinary approach, *SOLARIS* demonstrates a viable path toward self-sustaining, mobile renewable energy systems.

2 Project Description

2.1 Project Background and Motivation

In senior design, constraints define a project's limitations in terms of functionality and capabilities. Our project aims its sights on industry-relevant processes and standards, which serve as a learning opportunity. Our project, Solar-Optimized Land Autonomous Recharge & Intelligence System (*SOLARIS*), is defined by two primary technical pillars: solar optimization and integrated autonomous intelligence. With these ideas in mind, we plan to tackle power-conscious design of a system intended to be self-sufficient. Budgetary limitations represent a primary constraint, as the project is entirely self-funded. By splitting the cost between all team members, it remains important to find the balance between expenditure, low-power, and net positive charging capability. The key educational value remains an important portion as it relates heavily to our academic studies and our own personal interests as Electrical Engineer (EE) and Computer Engineering (CpE) majors. As we dive further into motivation for the project, we see software and hardware equally important to achieving this final goal.

Software design, firmware included, is one such interest that drives a large portion of the project. Embedded architecture has long since dominated consumer electronics markets and continues to prove itself valuable time and time again. An embedded system would be nothing without firmware or a software interface to connect it to the internet. That is why it is our motivation to not only design a robust control system but also interact with it through a wireless connection. This not only provides an interesting learning experience but also adds valuable and relevant experience to the large toolkit expected from engineers. The combination provides an interaction between hardware constraints and software design in real-world applications. Integration of wireless communications reflects industry-relevant practices where modularity, scalability, and maintainability are essential. To emphasize, mobile and desktop applications for control and statistics serve as valuable experience in developing a bridged embedded system. These embedded applications showcase engineering skills that are indicative of common workflows, frequent pairing with companion applications, and automated control systems.

Solar energy adoption continues to accelerate as improvements in battery storage and solar panel efficiency make renewable systems increasingly practical. This project offers hands-on experience designing and implementing a solar-powered system in a real-world context. At the same time, the growing fields of robotics and automation demand engineers who can effectively integrate electromechanical hardware with reliable control strategies. Through this project, the team will gain experience in motor and actuator selection, control system design and tuning, and system-level power management. By addressing efficiency, variable loading, and practical implementation constraints, the project will develop skills applicable to both energy systems and robotics and automation engineering roles.

2.2 Current Commercial Technologies and Existing Projects

2.2.1 Jackery Solar Mars Bot



Figure 2.1 Jackery Solar Mars Robot

Currently, there is no commercially realized product directly comparable to *SOLARIS*. This is primarily due to the high cost of solar panel technology, the reduced effectiveness of mobile robots burdened with large-area solar panels, and the greater efficiency of traditional stationary solar panel installations for battery charging. However, one company that specializes in clean, portable power solutions, Jackery, has proposed a concept for a future autonomous solar-farming mobile robot known as the “Solar Mars Bot” [1].

While the Solar Mars Bot has not yet been released as a consumer product, its existence as a concept is significant. Jackery describes the system as an autonomous robot capable of navigating toward regions of stronger sunlight and repositioning itself when light conditions deteriorate. The design integrates foldable photovoltaic panels with onboard energy storage, allowing the robot to both generate and manage its own power. When deployed, the panels are rated to deliver up to 600 W of charging power to the onboard 5 kWh battery.

Although technical details regarding its sensing, control algorithms, and navigation methods are not publicly available, the concept itself demonstrates a growing commercial interest in autonomous, mobile solar collection. This proposal reinforces the idea that *SOLARIS* occupies an emerging design space. The Solar Mars Bot suggests that industry is beginning to explore robotic platforms that can actively seek out optimal light conditions rather than relying on fixed installations. Even in its conceptual form, it validates the direction of this project and highlights the relevance of developing efficient, autonomous systems that can adapt their position in response to changing environmental conditions.

2.2.2 SolarX: DIY STEM Solar Powered Robot




Figure 2.2 DIY STEM Solar Powered Robot

The SolarX: DIY STEM Solar Powered Robot is one of many consumer products designed as learning tools for directing a solar panel toward a light source, typically the sun, as effectively as possible. It uses photo resistive sensors to measure incoming light and determine how the panel should be oriented. Two servo motors, mounted on perpendicular axes, control the horizontal and vertical position of the solar panel, allowing it to continuously adjust until it is facing the brightest direction. [2]

Although this product represents only a small subset of what *SOLARIS* aims to accomplish, it felt important to include it as a design reference for that specific subpart of the project. The core goal of *SOLARIS* is to improve power generation by responding to changing light conditions, and this kit provides a concrete example of how sensor-driven solar alignment can be implemented in practice.

One major difference is in how motion is achieved. The SolarX relies on servos, which require continuous power to hold their position. For *SOLARIS*, the intent is to explore a brushless motor actuator or similar mechanism that can reposition the panel without needing constant power to maintain its orientation. Since power efficiency is a primary constraint of the system, reducing energy draw is critical. While the actuation method will differ, the underlying design principle remains the same: use light feedback to determine how the panel should be oriented and adjust it accordingly to maximize available power.

2.2.3 Summary

Currently, there is no consumer product that fully matches the scope of *SOLARIS*. However, existing concepts and educational devices show that the core ideas behind the project are both feasible and relevant. The Jackery Solar Mars Bot highlights a growing commercial interest in autonomous, mobile solar platforms that can reposition themselves in response to changing light conditions. Even as a concept, it supports the idea that mobile robots can be used to actively seek out better solar exposure rather than relying on fixed installations.

On a smaller scale, products such as the SolarX DIY STEM Solar Powered Robot demonstrate how light sensors and actuators can be used to orient a solar panel toward the brightest direction. While these systems are limited and primarily educational, they provide

a clear and proven example of sensor-driven solar alignment, which directly aligns with one of *SOLARIS*' core goals.

SOLARIS takes these existing ideas and pushes them further by turning solar tracking into a fully autonomous, mobile behavior. Instead of only adjusting the angle of a panel, the system is designed to decide when orientation alone is no longer sufficient and physically move to a better location. Power efficiency is treated as a core constraint, so every action is driven by the goal of gaining more energy than it consumes. By combining light sensing, autonomous decision making, and low-power actuation, *SOLARIS* explores how a robot can actively respond to its environment to improve its own energy generation.

2.3 Goals and Objectives

2.3.1 Basic Goals Summarized

- Create a robot capable of autonomously moving out of shade on flat ground.
- Rotate the solar panel about two axes to orient itself in peak sunlight conditions.
- Energy gained from moving the robot must be greater than an equivalent, fixed solar panel over a full daylight cycle.
- Measure and report information to the end user such as distance travel, net energy, state of charge, etc.
- Create a phone app to serve as the primary means of communication between the user and the robot.

2.3.2 Advanced Goals Summarized

- Allow users to manually select between three different modes through the phone app, stationary, autonomous, and manual.
 - In stationary, the robot will not move, however it will still move the solar panel.
 - In autonomous mode, full control is given to the robot.
 - In manual, it receives throttle and steer information and uses that to move.
- The phone app should display all pertinent information we collect from the robot to the user.
- The robot should be capable of avoiding collisions.
- The robot will be able to reliably deliver regulated electrical power from its onboard energy storage system to a user's personal electronic device.

2.3.3 Stretch Goals Summarized

- Create a website that mimics the phone app's interface. It will not be capable of connecting to the robot, but if users want to see their data collected from the robot, they can sign in and see it.
- Use IR sensors to detect going onto a curb / road.
- Add GPS functionality to track position throughout the day, and use that information to have a more robust model of where good spots are at given times.

2.3.4 Basic Goals

At the project's core, our main goal is to create an autonomously moving robot capable of solar farming. The solar farming robot should be able to rotate and angle a solar panel for peak sunlight positioning. *SOLARIS* needs to be capable of operating in outdoor environments like flat yards or open fields, as opposed to traditional solar harvesting methods. As it moves, there will be a low-level obstacle avoidance protocol causing the robot to traverse non-ideal environments. With both basic goals in mind, we want to expand capabilities that support our broader goal, maximizing solar energy collection without requiring external intervention. The overarching goal is what supports the basic idea to design a new kind of solution that is more practical.

Another central goal that is explored in the senior design project is providing a net positive energy benefit when compared to stationary solar harvesting approaches. This robot will evaluate the tradeoffs seen between energy expended for movement and the additional energy gained by repositioning itself. It is without question that aiming the solar panel at the source of sunlight will increase electrical intake; however, it is yet unknown whether movement can also play a role in this collection process. Therefore, success is defined by achieving greater net energy generation over a full daylight cycle compared to an equivalently sized fixed solar panel.

Though autonomous in nature, our goals do not completely neglect human interaction with the project. For these reasons, we need observable data that gives us insight into efficiency and mobility. The robot should also serve as a platform for observing some of its most basic operating behaviors. Key metrics like battery life, charging rate, and distance traveled should be available to the user through a wireless connection. This system should also support safe operation, as it is meant to be autonomous in nature but observable by a user. The robot should be able to identify an event, such as a collision, and perform a maneuver to avoid the continuation of the event. These goals ensure autonomous operation, which can remain practical and robust for a real-world setting while also keeping the project innovative and fresh.

2.3.5 Advanced Goals

An advanced goal of the project is to ultimately enhance user interaction through the usage of a mobile application. The mobile application will serve as the primary interface between the user and the robot. With this platform, we can provide performance metrics, system behavior, and influence robot operation without touching the robot physically. In the mobile application, users should be able to select between one of two modes. These two modes exist transparently but functionally work the same. One mode allows the bot to become stationary, and the bot will function as a traditional solar harvesting setup with a rotating solar panel. The other mode would be default, as the robot's full autonomous nature would allow it to navigate on its own and relocate if necessary.

This fully autonomous behavior highlights an advanced goal for this project as the robot is given complete control over its own decisions. The robot will be responsible for its energy management, navigation, and overall decision-making. In this mode, the robot will

systematically evaluate conditions to determine if movement is necessary and then avoid obstacles on its way to a perceived optimal solution. All while a user can watch over system metrics via a wireless connection to the robot. This presents the last advanced goal of our project, which is to present relevant data to the user's device. On the device, metrics like distance traveled, temperature, and charging rate can all be displayed. Displaying this data allows a user to understand the robot's behavior while answering an important question about the robot's functionality.

2.3.6 Stretch Goals

To expand upon our advanced goals, we would like to introduce additional functionality that can enhance user control, system intelligence, and long-term data analysis. With these features, we could apply a more sophisticated system to explore decision-making capabilities that expand upon the original core system requirements for functionality. One such stretch goal is to expand upon a manual control system from the user interface of the application. This could include finer control over movements, adjustments for speed, and the ability to directly interact with solar panel orientation. Ultimately, these adjustments would provide users with greater flexibility in their own use of the platform. Similarly, this would require a much more robust design, allowing for a much more fine-tuned system entirely.

In addition to mobile application enhancements, there could be an integration of higher-level control logic or the mounting of an artificial intelligence to improve autonomous behavior. This next level of autonomy could enhance decision-making behavior capable of strategizing navigation. The robot would also be capable of a higher level battery management with adaptive battery level thresholds and movement statistics. These upgrades would allow the system to better respond to changing environmental conditions, which could produce increased performance over time. To expand on connectivity, the robot could alternatively connect to a web-based interface where the structure of the app would be mirrored, but would apply a more practical interaction. The website would not provide direct communication with the robot, but would alternatively allow for more detailed logging of events and statistics for later consumption. All stretch goals mentioned support a long-term analysis and provide a more accessible platform for interaction with the robot.

In addition to mobile application enhancements, there could be an integration of higher-level control logic or the mounting of an artificial intelligence to improve autonomous behavior. This next level of autonomy could enhance decision-making behavior capable of strategizing navigation. The robot would also be capable of a higher level battery management with adaptive battery level thresholds and movement statistics. These upgrades would allow the system to better respond to changing environmental conditions, which could produce increased performance over time. To expand on connectivity, the robot could alternatively connect to a web-based interface where the structure of the app would be mirrored, but would apply a more practical interaction. The website would not provide direct communication with the robot, but would alternatively allow for more detailed logging of events and statistics for later consumption. All stretch goals mentioned

support a long-term analysis and provide a more accessible platform for interaction with the robot.

2.3.7 Basic Objectives

- Design and fabricate a mobile robotic platform driven by DC motors and a motor driver capable of traversing flat outdoor terrain such as yards or open fields.
- Implement a mechanically actuated solar panel mounting system using a motorized joint to enable controlled rotation and tilt for sunlight alignment.
- Integrate obstacle detection sensors (ultrasonic sensors) and low-level motor control logic to support autonomous avoidance of environmental obstacles.
- Incorporate onboard energy storage, charge regulation, and efficient power distribution circuitry to support simultaneous solar charging and system operation.
- Develop embedded control firmware on a microcontroller to manage motion control, sensor input, and autonomous decision-making.
- Measure and record battery voltage, current, and net power using onboard sensing circuitry.
- Implement wireless communication via bluetooth to transmit system metrics and operational status to an external user device.
- Ensure the system can operate autonomously without external power or physical user input after deployment for at least 12 hours.

2.3.8 Advanced Objectives

- Support multiple operational modes selectable via wireless communication, including:
 - A stationary mode where only the solar panel actuation hardware is active.
 - A fully autonomous mode where the drive motors and sensors are continuously engaged.
 - A manual mode where the microcontroller receives throttle and steering commands over bluetooth from the user and uses it to drive the robot.
- Integrate motion sensing to quantify distance traveled during movement operation.
- Design the electronics and firmware architecture with modularity.
- Detect collision or near-collision events using physical or proximity sensors and trigger corrective motor commands to prevent sustained contact.
- Integrate a regulated power delivery system capable of safely converting stored onboard energy into standardized output to charge consumer electronic devices with reliable performance and protection mechanisms.

2.3.9 Stretch Objectives

- Develop a web-based interface to visualize logged system data independent of direct robot communication.
- Create a weather API to get sunlight patterns to optimize robots algorithms.
- Use hall effect sensors to detect collisions as a bumper instead of checking if an axle has difficulty moving through the encoder.

2.4 Required Specifications

2.4.1 Overall System Specifications

Specification	Description	Target Value and Unit(s)
Energy for User	After a sunny, optimal day, the system shall provide the end user with at least 15 watt-hours (Wh) of usable energy from the battery.	≥ 15 Wh/day
Energy Generated	After a sunny, optimal day, the system shall generate at least 75 watt-hours (Wh) of electrical energy through solar harvesting.	75 Wh/day
Power Output Connection	A USB-A connection shall be provided for user energy extraction from the system.	USB-A Connection

Table 2.1 Energy Specifications

Specification	Description	Target Value and Unit(s)
Max Forward Robot Speed	Maximum designed robot speed (forward).	0.3 m/s
Average Robot Speed	Normal forward robot speed.	0.1 m/s
Max Reverse Robot Speed	Maximum designed robot speed (reverse).	0.3 m/s
Weight	Total weight of robot	<50 lbs.
Solar Panel Rotation	Pan axis range of motion	360 Degrees
Solar Panel Tilt	Tilt axis range of motion	30 Degrees


Table 2.2 Mobility Specifications

2.4.2 Component System Specifications

Component	Parameter	Subsystem and Description	Target Value and Unit(s)
Geared BDC Motor w/ Planetary Gearbox	Torque	4 Motors that provide the means of rotating the wheels for movement.	24 V, 48 W, 23 RPM
Solar Panel	Power	Primary source of energy for battery charging.	20 Watts
Solar Panel	Voltage	Voltage level that solar panel operates at.	12 V
Battery	Capacity	Powers all devices on the robot system.	256 Wh
Pneumatic Knobby Tires	Size	Wheels that each motor will be linked to via a shaft.	4"
High Torque Worm Gear BDC Motor w/ Self Locking	Torque	Motor for pan-axis rotation of solar panel.	24 V or 12 V 5 RPM
Bluetooth Low Energy (BLE)	Wireless Communication Range	Communication subsystem. Used to send data from the robot to a mobile app.	1-6 meters
Linear Actuator	Torque	Tilt-axis motion for solar panel	200 N-M, 4" Stroke
PWM H-Bridge Motor Drive	Amperage	3 boards for driving the motors and linear actuator	10 A
5 V Regulator	Voltage	Regulator for sensors (ultra-sonic sensor)	5 V
3.3 V Regulator	Voltage	Regulator for photoresistors	3.3 V

Table 2.3 Component Specifications

2.4.3 Hardware and Software Diagrams + Prototype Illustration



SOLARIS Block Diagram

Figure 2.3 Hardware Block Diagram




Figure 2.4 SOLARIS Prototype Illustration




Figure 2.5 SOLARIS Main Function Software Diagram




Figure 2.6 SOLARIS Robot Driving Software Diagram




Figure 2.7 SOLARIS Phone App Software Diagram

2.5 Design Requirements

2.5.1 Marketing & Engineering Requirements

Marketing Requirements	Engineering Requirements	Justification
1, 2	1. Rover should reposition itself autonomously toward higher light intensity using sensors and the control system.	The ability to autonomously reposition corrects inefficiencies in fixed solar harvesting systems while reducing the need for manual intervention.
1, 2, 4	2. The system should operate in low-power sleep mode while charging, consuming <300mW in the idle state.	Low-power operation lets the system maximize energy gain while extending operational lifetime; the bot can outlast suboptimal weather conditions entirely.
1, 2, 3	3. The rover should detect obstacles within a minimum range of 20-50 cm accurately using ultrasonic sensors to avoid collision.	Automated obstacle avoidance ensures reliability in dynamic real-world environments, whether outdoors or indoors.
2, 3	4. The system should provide real-time battery voltage, charging current, and charge state data.	Being able to monitor energy metrics allows users to evaluate system performance like distance traveled, battery level, and charging statistics.
2, 3, 4	5. The rover should communicate wirelessly with a smartphone via Bluetooth Low Energy (BLE).	Bluetooth Low Energy (BLE) enables low-power communication with user devices, making it suitable for field deployment and demonstrations.
1, 3	6. Rover should have four wheels with encoders to estimate distance traveled and improve distance traveled accuracy.	Feedback from encoders provides improved navigational efficiency while allowing precise movements to be made for travel.
3	7. The system should require no external infrastructure for normal operation.	An independent infrastructure makes deployability a priority, allowing for remote or agricultural environments.
4	8. Total system hardware cost shall not exceed \$1000	Cost constraints enable feasibility for deployment while also aligning with the competitive market for low-cost energy monitoring.

3, 4	9.	The system should be implemented on a single custom PCB with modular sensor connections.	Custom PCB demonstrates design project integrity, while the modularity of sensors allows for an easier debugging process.
4, 5	10.	The charging system for solar-battery should exceed an efficiency of >50% under nominal operating conditions.	An efficient charging system is achievable using commercial MPPT-based Li-ion charging ICs and demonstrates the balance between complexity, cost, and performance.
5	11.	The System should operate reliably in outdoor lighting and temperature without user intervention.	The reliability of an unattended rover in a working environment increases the number of practical applications.
1, 2, 4	12.	The rover should operate for a minimum of 12 hours without solar input under normal conditions.	Ensuring continued operation even during sub-optimal conditions demonstrates robust power design and applicability in differing environments.
Marketing Requirements 1. The system should maximize solar energy harvesting through autonomous behavior 2. The system should operate reliably with minimal human intervention 3. The system should be easy to deploy and monitor 4. The system should be low-cost and energy-efficient 5. The system should be reliable in a real-world operating condition			

Table 2.4 Justification Chart for Requirements

2.5.2 House of Quality




Figure 2.8 House of Quality

3 Hardware & Software Research

3.1 Hardware Technologies

3.1.1 Brief

Up until this point we have shared our goals and requirements for this project. Mentions of hardware usage have been stated but the reasoning for their choice has not been stated in detail. In this section we look at the valuable research used in making those decisions. It is not only valuable to investigate the technologies we want to use but to also be open-minded as to why we did not choose others given sufficient investigation. This presents itself as a strong opportunity to deepen our understanding of the hardware technologies available to us for best fit within our project *SOLARIS*.

3.1.2 Control Unit

3.1.2.1 What is a MCU?




Figure 3.1 Microcontroller Hardware Layout

Our control unit consideration begins with understanding microcontrollers (MCU). Designed as a small computer on a single chip, MCUs are made for embedded systems without requiring an operating system. This is important because this combines all effective traditional hardware into one compact integrated circuit (IC). As compared to a microprocessor (MPU), MCUs integrate processing, memory, and input/output (I/O) peripherals like timers, counters and analog-to-digital converters (ADCs) into one standalone unit [3]. The general makeup of a microcontroller goes as follows: it receives one or more inputs, processes the data, and generates system outputs. These inputs and outputs are generally voltages which the microcontroller can recognize and output. While the I/O interacts with the external environment, code generally written in C/C++ is stored in local memory alongside stored data from a running program [4]. This allows for a bare

metal connection of user control for device programming. In a nutshell, MCUs are high speed devices which can operate at low power or on battery powered systems for self-contained systems like telephones, appliances, vehicles and other computer systems.

Specification	ESP32-S3	ESP32-P4	MSP430FR6989	STM32U0
Architecture	Xtensa LX7 Dual-core	RISC-V Dual HP + LP	16-bit RISC	Cortex-M0+
Max Clock Speed	240 MHz	400 MHz	16 MHz	56 MHz
Flash Memory	up to 16 MB	up to 32 MB	128 KB	64KB ≤ 256KB
Bluetooth/Wi-Fi	BLE/Wi-Fi 4 (802.11 b/g/n)	None	None	None
ADC	2 x 12-bit, 20 ch	2x 12-bit, 8 ch	12-bit, 16 ch	12-bit, 10 ch
UART / SPI / I2C	3/4/2	3/8/3	2/2/1	4/3/2
Low Power Mode	Modern Sleep, Light/Deep Sleep	Multiple LP modes	LPM0-LPM4.5	Stop / Standby / Shutdown
Active Current	~340 mA peak	~123 mA peak	~275 µA	~25uA/MHz
Low Power Current	~7 µA	~25 µA	~0.5 µA	Varies < ~0.4 µA
Operating Temp	-40 to 85°C	-40 to 85°C	-40 to 85°C	-40 to 85°C
Operating Voltage	3.0-3.6V	3.0-3.6V	1.8-3.6V	1.71-3.6V
Development Ecosystem	ESP-IDF, Arduino	ESP-IDF	TI Code Composer	STM32CubeIDE
GPIO Count	~45	~54	~83	~29
Unit Cost	~\$6.10	~\$10.50	~\$11.20	~\$3.50

Table 3.1 MCU Comparison

Source: [5]-[10]

3.1.2.1.1 ESP32-S3

The ESP32-S3 can be considered the low-power family of processors from Espressif Systems. This processor utilizes a dual-core Xtensa LX7 processor which runs at 240 MHz.

The integration of a 2.4 GHz Wi-Fi and Bluetooth 5 LE radio makes this chip highly capable for Internet of Things (IoT) applications. Designed for workloads in artificial intelligence (AI) and machine learning (ML), the processor uses a five-stage pipeline which supports varieties of code density. The ultra-low power (ULP) coprocessors are complemented by a ULP-FSM used with deep sleep operations making this chip incredibly power efficient and responsive. The ESP32-S3 comes with an internal USB controller meaning it can act as a native USB device without using any external silicon. Opening a range of possibilities for flashing and debugging the device. The microcontroller is rich in I/O with almost 45 GPIO ready for configuration in a multitude of designs.

Major considerations that make the ESP32-S3 stand out compared to the other options in research are the integrated RF handling and low power settings. There are four operating modes defined by a deep-sleep drawing around 7 μ A which makes it competitive for a battery powered device. Espressifs support for open-source design goes as far as their framework ESP-IDF and integration for use with Arduino. The ESP32-S3 receives a large portion of community support with online access to projects and libraries to make programming a simpler process. Configurability is not an issue since it has 3 hardwired UART lines and 4 hardwired SPI lines with 2 multiples configurable I2C lines to be used. It is also important to have on board storage, and this microcontroller comes in around the middle of the pack with different variations of ROM and SRAM which reduce complexity and allow for storage of data components for export wirelessly or wired.

3.1.2.1.2 ESP32-P4

Another one of Espressif Systems offerings, this option comes in similarly to the S3 model. Some key differences here are the lack of Bluetooth and Wi-Fi support but a plethora of GPIO and horsepower. The absence of integrated wireless means that the P4 is wired only, relying on Ethernet MAC for network connectivity or pairing with an external wireless module. Sitting above traditional MCUs the P4 uses integrated image processing pipelines to reduce component count and bill of materials (BOM) cost for camera and display applications. Making its strongest cases in multimedia applications and edge vision, the combinations of image processing hardware and generous integrated memory simplify this process to one chip. One big issue with design is the lack of community support due to this microprocessor's unique use case.

Where the ESP32-P4 makes sense for SOLARIS is its strong dual-core RISC-V architecture which constitutes real time motor and sensor control. The processing power also would support high level autonomous logic, communication, and decision making without blocking each other. The P4 also comes with two full Motor Control PWM (MCPWM) blocks. Each of these blocks has three timers, fault detection, and capture submodules. For a vehicle with motors and panel positioning this would be purpose built rather than making a more customized motor control system. In the case of computer-vision based shadow detection or for more complicated processing the P4 is a perfect option. Similar hardware like the ADC with 14 analog input channels which can be routed across two 12-bit SAR gives headroom for current/voltage sensing, light detection, and battery monitoring at the same time.

3.1.2.1.3 MSP430FR6989

The MSP430FR6989 is purpose built for minimizing energy consumption. When running in active mode it only draws around 275 μ A. Beyond this it offers seven distinct low-power modes (LPM0 - LPM4.5) which allows for an application to sleep between sensing intervals and then wake instantly on interrupts. In the application of a solar vehicle like SOLARIS, these key power saving features are critical. Like many other MCUs it supports a 12-bit ADC for applications like an array of phototransistors or reading battery state-of-charge. Onboard timer peripherals provide hardware level pulse width modulation (PWM) generation to drive motor controls. This avoids CPU overhead with hardware specific modules for performance scenarios. Bluetooth communication is not native but highly supported with use of a Bluetooth-to-serial module. Via a UART interface, the Bluetooth module appears as a simple serial port as the app sends commands in ASCII or as binary packets while the MSP430 parses and executes them.

Something to appreciate about Texas Instrument as the developers of this MCU is the toolset for use with the MSP430FR6989. TI Code Composer Studio is industry developed software with support for a multitude of libraries making it easy to set up and flash a microcontroller with ease. Features like per line of code energy usage statistics make it a massive win for designing low power systems. Another major win for this MCU is its non-volatile FRAM storage. This local storage allows for unlimited write cycles for logs of solar stats, distance traveled, and battery charge level even after power loss. The MSP430FR6989 could be a well-matched choice for this project due to its low power drawn with flexible sleep modes. The main limitation is its lack of native Bluetooth or Wi-Fi and 16MHz processor speed capacity.

3.1.2.1.4 STM32U0

The STM32U0 is a family of entry level low power microcontrollers from the company STMicroelectronics. Built on the ARM Cortex-M0+, this family boasts static power consumption of 160 nA in standby with as low as 16nA in full shutdown. Between the different models they go from 20 to 81 GPIO pins and offer between 64 KB and 256 KB of flash memory. In active mode the chip draws 25 μ A/MHz at 3.0 V. With a tiered sleep architecture, the MCU can be put into sleep mode for low power draw and then woken quickly for a superior wake-up performance. From the analog perspective this MCU uses one 12-bit ADC with up to 16 channels, a digital to analog converter (DAC), an analog comparator, and an operational amplifier all on one chip. For example the comparator can be used to watch an array of phototransistors for shadow detection while the CPU sleeps.

Like the MSP430FR6989, the STM32U0 family includes a 16-bit PWM timer which is dedicated specifically for motor control. For wireless communication like Bluetooth the STM32U0 uses UART connections to remain active in receiving Bluetooth commands. This means a separate chip for serial from Bluetooth is not necessary but a receiver is still required for use. ST has also developed an ecosystem for programming their devices which stands tall around many. Key functionality like a graphical pin configuration and code generation tools make programming the device visually simple. Overall this MCU contains a large timer functionality, dedicated motor control hardware, a mature software ecosystem for development, and hierarchical low power design.

3.1.2.2 MCU Vs. MPU

	Microprocessor(MPU)	Microcontroller(MCU)
Description	CPU integrated into a single chip, designed for general purpose computing, requiring external memory and peripherals.	Small computer on a single chip with CPU, RAM, and program memory.
Memory	External memory RAM and storage (scalable).	On-chip RAM and program memory (limited).
Speed	High (for complex tasks)	Moderate (for dedicated tasks)
Power	High (may need cooling)	Low (energy efficient)
Cost	Higher (external components)	Low (inexpensive)
Programming	Runs bare-metal firmware.	Supports full operating system (OS)
Applications	Personal Computers (PCs), laptops, and high-performance computing (HPC).	Embedded systems, IoT devices, and robotics.
Strengths	High performance, flexible, and scalable.	Integrated features, low power, and low cost.
Weaknesses	High power consumption and more expensive.	Limited processing power and memory.

Table 3.2 MCU Vs. MPU

Source: [11]

When comparing these technologies it is important to note that MCUs are more similar to MPUs than FPGAs are to either. While similar they all differ in architecture as well as performance capability. Between MCUs and MPUs, the key difference lies in MCUs containing all additional peripherals with integrated memory, input/output (I/O) components and other varied peripherals [11]. Typically more cost effective, small in size, and low in power, microcontrollers are optimized for an all in one functionality. On the other side MPUs are designed to support higher performance in demanding applications like personal computing and graphics processing [11]. Therefore, a MCU can succeed in an area of real-time and high-speed computing tasks while an MPU sees itself as reliable and scalable. This leaves microcontrollers to be typically used for handling fast signal processing while microprocessors are favored for more demanding high-performance computing (HPC) related tasks [11].

3.1.2.3 MCU Vs. FPGA

	<i>Microcontroller(MCU)</i>	<i>Field-Programmable Gate Array (FPGA)</i>
Description	Small computer on a single chip with CPU, RAM, and program memory.	A configurable integrated circuit composed of programmable logic blocks designed for reprogrammable hardware
Key Components	Single chip computer with CPU, memory, and built in peripherals	Configurable Logic Blocks (CLBs), programmable interconnects, and external memory interface
Speed	Sequential Processing (executes instructions one at a time)	Parallel processing (can perform multiple operations simultaneously)
Power	Low (energy efficient)	Higher (configuration dependant)
Cost	Low (inexpensive)	Higher (hardware and development)
Programming	Software-based (C/C++, Python, Assembly)	Hardware Description Languages (Verilog VHDL)
Applications	Embedded systems, IoT devices, and robotics.	Digital signal processing (DSP), rapid prototyping, and ASIC development.
Strengths	Integrated features, low power, and low cost.	Highly flexible, parallel-processing, hardware optimization.
Weaknesses	Limited processing power and memory.	Requires special knowledge with a longer development cycle.

Table 3.3 MCU Vs. FPGA

Source: [12]

Though all can be thought of as small computers it is still important to compare FPGAs as well. Unlike a microcontroller, FPGAs offer a level of reprogrammability at a hardware level. MCUs operate in a sequential pattern by reading lines of code whereas FPGAs are unique with capabilities of handling parallel inputs simultaneously. An FPGA can be programmed to perform the function of a microcontroller; however, a microcontroller cannot be reprogrammed to perform as an FPGA [12]. While FPGAs appear to be a clear winner it would certainly be overlooking many constraints like cost, ease of programming, energy efficiency, and best fit for application. MCUs are cheaper and require basic conceptual knowledge of software languages. MCUs also consume much less power than FPGAs allowing them to be suitable for battery powered applications and embedded

systems [12]. This means that small, affordable, and non-volatile microcontrollers are ubiquitous in modern electronics [12]. For our application it may be safe to assume that an FPGA would be overkill as the major processing will certainly see itself reduced to a simple control system with calculations.

3.1.2.4 Control Unit Part Selection

ESP32-S3: After consideration between microcontroller, microprocessor, and field-programmable gate arrays it is a smart decision to implement a microcontroller as the brains of our system. The ESP32-S3 is one of the best choices for three main reasons. Built in Bluetooth and Wi-Fi capabilities make this a strong choice for pairing with a potential companion app and controlling the system wirelessly. This eliminates the need for an external module which could potentially complicate design constraints. When it comes to self-autonomy the balance between computational power and amount of GPIO also make this microcontroller shine as the best option. The ESP32-S3s dual core processor is ample for semi-autonomous logic, sensor fusion, and PWM motor control. Lastly the ability to program this microcontroller easily is a huge advantage over other options seeing as the Espressifs line is strictly open source. There are many libraries online with large communities of people making the support for programming this device overwhelmingly helpful. The ESP32-S3 also utilizes low power modes like other microcontrollers. For our project *SOLARIS* this option excels in all areas relevant to our project and does not lack basic feature sets common with other MCUs.

3.1.3 Drive Motors

3.1.3.1 Brushed DC Motor

Brushed DC motors are commonly used in robotic systems due to their simple control requirements, low cost, and high starting torque, making them practical for drive mechanisms and auxiliary actuators. From a system design perspective, these motors are advantageous because their speed can be controlled directly using voltage or pulse-width modulation without complex control electronics, reducing controller overhead and hardware complexity. The torque produced is proportional to the armature current, allowing straightforward current based torque regulation for motion control if desired. However, mechanical commutation via brushes introduces friction losses, electrical noise, and gradual component wear due to more moving parts, which reduce efficiency and limit operational lifetime. These are factors of concern for *SOLARIS* since energy losses translate directly into reduced runtime which reduced desired efficiency. Consequently, brushed motors are most suitable for applications where simplicity, cost, and ease of integration outweigh long-term efficiency or durability considerations.

3.1.3.2 Brushless DC Motor

Brushless DC motors provide significant advantages in robotic systems where efficiency, reliability, and dynamic performance are critical design priorities. Unlike brushed motors, commutation is performed electronically using a motor driver that energizes stator windings in sequence based on rotor position feedback, creating a rotating magnetic field that drives the rotor. This eliminates brush friction and electrical arcing, resulting in higher efficiency, reduced thermal losses, and longer service life. In robotic platforms with limited power resources, these efficiency improvements can substantially extend operating time which is a primary goal with *SOLARIS*. Brushless motors also offer higher torque density

and smoother operation, which improves motion precision and reduces mechanical vibration transmitted to the chassis. The primary tradeoff is increased system complexity, as they require dedicated electronic speed controllers and more sophisticated firmware for commutation and feedback processing. Despite this added complexity, brushless motors are a good choice due to their increased overall efficiency and longer lifespan.

3.1.3.3 Stepper Motor

Stepper motors are small motors that produce movement to a desired angular position that is precisely controlled by the given control system. The main difference is that these changes are given in small step increments. These motors are typically used in small applications such as robot arms due to their small size, weight and highly precise nature. They were not selected due to their small nature and inability to handle *SOLARIS*' load demands.

3.1.3.4 AC Motor

AC motors are widely used in industrial automation and large scale robotic systems because of their robustness, high efficiency, and ability to operate continuously under heavy loads. They generate torque through a rotating magnetic field produced by alternating current, which induces rotor motion either through electromagnetic induction or synchronous magnetic coupling. From a design standpoint, AC motors are quite advantageous in applications that require sustained high power output, such as industrial robotic arms or automated manufacturing equipment. However, they typically require variable frequency drives for speed control, which significantly increases complexity and cost for just the control system. In small or mobile robotic platforms minimizing power consumption requirements, an AC motor is less desirable due to its larger size, weight and overall power requirements when compared to DC motors.

3.1.3.5 AC Motors Vs. DC Motors

	Microcontroller(MCU)	Field-Programmable Gate Array (FPGA)
Description	Small computer on a single chip with CPU, RAM, and program memory.	A configurable integrated circuit composed of programmable logic blocks designed for reprogrammable hardware
Key Components	Single chip computer with CPU, memory, and built in peripherals	Configurable Logic Blocks (CLBs), programmable interconnects, and external memory interface
Speed	Sequential Processing (executes instructions one at a time)	Parallel processing (can perform multiple operations simultaneously)
Power	Low (energy efficient)	Higher (configuration dependant)
Cost	Low (inexpensive)	Higher (hardware and development)
Programming	Software-based (C/C++, Python, Assembly)	Hardware Description Languages (Verilog VHDL)
Applications	Embedded systems, IoT devices, and robotics.	Digital signal processing (DSP), rapid prototyping, and ASIC development.
Strengths	Integrated features, low power, and low cost.	Highly flexible, parallel-processing, hardware optimization.
Weaknesses	Limited processing power and memory.	Requires special knowledge with a longer development cycle.

Table 3.4 AC Motors Vs. DC Motors

Sources: [13-15]

3.1.3.6 Drive Motor Type Selection

DC Motor: While AC motors have a higher lifespan and efficient continuous operation, the advantages that these properties afford come with significant drawbacks for our application. AC motors have inherently larger sizes and would require an inverter to convert the DC from the battery to AC and then back to DC in the variable frequency drive to produce the required signal to rotate the motor. This chain of equipment raises cost and,

in particular the variable frequency drive, will require lots of space inside of the robot and add to the weight. DC motors are directly compatible with our already low voltage battery without requiring additional power conversion. The high torque at low speeds is desirable for efficiency and for outdoor terrain. Additionally, their compact size and straightforward integration with embedded microcontrollers simplify electrical and mechanical implementation. For a battery powered robotic system where efficiency, controllability, and minimal hardware overhead are the focus, DC motors offer a more practical and scalable solution than AC alternatives.

3.1.3.7 Brushed Vs. Brushless DC Motors

Brushed and brushless DC motors differ in how they rotate and control rotation. Brushed motors work by switching current to the windings using a physical brush on a commutator meaning they will spin naturally when voltage is applied which is significantly easier to implement and control using a device such as an H-bridge controller. The main drawback of brushed DC motors stems from the inherent mechanical contact of the brush riding on the commutator which introduces electrical noise, brush wear and efficiency. Brushless motors use solid-state switching through the windings which requires the implementation of an external controller that creates the desired signal for rotation and control of the motor. While more complex in its implementation and control, this allows brushless motors to have smoother torque, more precise speed regulation and minimal wear. While brushed motors are simpler and able to respond very predictably for a given input, they demonstrate lower efficiency over time and performance under heavier loads and prolonged use. Brushless motors can better adapt to varying load conditions because the external controller actively regulates the current and timing. Overall, brushed motors offer simplicity and direct operation while brushless motors offer controller performance, adaptability and consistency and higher efficiency.

3.1.3.8 DC Motor Type Selection

Brushed Motor: Brushed DC motors were selected for driving the robot because the efficiency and precision gains offered by brushless motors were not in the project scope for *SOLARIS*. While brushless motors allow for greater electrical efficiency, their implementation and operation requires dedicated controllers, more complex control, and more integration concerns. All of these factors raise the risk of increasing development time and system operation risk. Brushed DC motors are typically driven with simpler pulse width modulation (PWM) control methods which allows for rapid prototyping, reliable operation with minimal setup required. Since *SOLARIS* does not operate continuously and only in short bursts, the efficiency difference is not likely to significantly impact the battery life or efficiency of the robot. The simplified wiring and reduced control requirements with ease of predictability is desired to allow for ease of troubleshooting and allow focus on higher level system design and implementation. The benefits of brushed DC motors will allow us to achieve our desired robustness, and system reliability while meeting all of our performance requirements.




Figure 3.2 Two-pole DC Motor Model

Source: [16]

3.1.3.9 Drive DC Motor Selection

When selecting specific motors for *SOLARIS* the first major factor is ensuring that the motors can handle the load of the entire robot. To meet this end, there are several factors that must be taken into consideration. The first factor was finding a gearbox output speed that could ensure a substantial torque margin. The wheel size that has been selected were 4-inch diameter wheels and a desired speed of 0.3 m/s. The required wheel speed is approximately 53 RPM after gearbox reduction. A motor RPM of 60-80 RPM was determined to be the target range that would meet the desired performance and margins.

The second major design consideration is the torque under load that the motors must be able to handle during routine, expected operation. *SOLARIS* must be able to roll on various grass types and heights, startup inertia, and friction losses from a differential steering scheme; all of these obstacles will affect the desired torque requirements. The next consideration is stall torque. Stall torque is the maximum capability torque that a motor can sustain. Running torque should be at a margin approximately 20-30% below this maximum torque to ensure efficient operation and avoid high current spikes. The motors were considered based on their rated running torque with an error margin to what is the minimum required for *SOLARIS*.

The third major design consideration was the gear ratio and gearbox construction for the robot for the required performance. Higher gearbox ratios will increase the output torque while reducing speed, making them better for heavier, lower speed robots. Too high of a reduction scheme can yield undesirable mechanical losses and backlash which would affect responsiveness and efficiency. Motors with planetary gearboxes were considered most

heavily because they distribute the load evenly on the gears which offers improved torque density and robustness when compared to similar gearbox configurations such as spur gears.

The final major consideration were the electrical attributes of the motor. A 12 V motor was selected because it is the voltage level that the battery for *SOLARIS output* and thus will not require any additional power conditioning to allow the motors to operate. The current rating must also be able to handle current surges from motor stalls under unexpected conditions to work with the motor drivers, wiring and battery. Since 4 motors are desired with a differential steering configuration, startup inrush current and turning currents can place high stress on the power delivery systems so the motor driver must be able to accommodate these worst-case events. Additionally, most of the considered motors come with encoders already built in which is very desirable for ease of implementation with simplified control design.

<i>Brushed DC Motor Model</i>	<i>Voltage</i>	<i>Rated Torque</i>	<i>Rated Speed</i>	<i>Gear Ratio</i>	<i>Gearbox</i>	<i>Encoder</i>	<i>Cost</i>
StepperOnline F5D60-12GN-30S-5GN50K	12V	6.27 Nm	60 RPM	50:1	Spur	No	\$52
Duma Dynamics 13PPR Encoder 60RPM DC Gearmotor	12V	5 Nm	60 RPM	104:1	Planetary	Yes	\$59
goBILDA 5204 Series Yellow Jacket Planetary Gear Motor	12V	18.14 Nm	60 RPM	139:1	Planetary	Yes	\$56.99

Table 3.5 DC Motor Comparison

Sources: [17], [18-19]

3.1.3.10 StepperOnline Brushed Gear Motor

This motor is a relatively inexpensive, simple and conservative option that is an industrial adjacent drive train motor for *SOLARIS*. The motors 12V rating with a torque of 6.27 NM at 60 RPM post-reduction is a very adequate fit for the robot and will more than carry the robot's estimated weight of 15-20 lb across grass at the desired speed of ~0.3 m/s. This motor is the strong and robust option making it a good choice where a lot of resistance is expected for the environment or terrain. With a priority being simplicity, the spur gearbox

makes the drivetrain integration easier due to the conventional gearing inside of the box at the expense of efficiency. The main drawback of this motor is that it does not come with an encoder which is a massive consideration when selecting motors. Including an encoder is strongly desired to minimize extra wiring and reduce complexity.

3.1.3.11 E-S Motor Planetary Gear Motor

This motor offers a strong balance between torque capability and control functionality. With an output torque around 5 Nm and a gearbox ratio of 104:1 and a speed near the desired operating range, it provides more than enough mechanical capability for grass traversal while remaining within budget constraints. Its integrated encoder simplifies implementation of feedback control, allowing for more precise speed regulation and improved navigation accuracy without additional hardware. Functionally, this makes it attractive for autonomous systems that rely on wheel feedback for distance estimation or correction. Compared to StepperOnline's motor, it provides slightly less torque but still maintains ample safety margin for real-world operation.

3.1.3.12 goBILDA Yellow Jacket Planetary Gear Motor

This motor stands on its own as a highly capable choice for a small outdoor robotic drivetrain. Its planetary gearbox architecture provides a strong torque density, meaning it can deliver a large output torque relative to its size and weight along with providing more efficiency, and the 60 RPM gearbox output aligns well with a robot's desired cruising speed when paired with a wheel size of both 4 and 6 inches. This motor also includes an integrated quadrature encoder that runs on 3.3–5 V logic, making it easy and straightforward to interface with the ESP32's PWM pins for closed-loop speed control and odometry without needing external sensors which would drive up the integration complexity. The documented stall torque and stall current give ample headroom for traction on grass and small disturbances such as rocks or uneven terrain, and goBILDA's product pages include downloadable specifications that support formal engineering justification. Additionally, the Yellow Jacket series is widely used in competitive robotics (FTC, FIRST, VEX communities), so its mechanical behavior and integration patterns are well understood, making it a reliable and well-supported choice.

3.1.3.13 Drive DC Motor Selection

goBILDA 5204 Series Yellow Jacket Planetary Gear Motor: The goBILDA strikes a strong balance between performance, documentation, and system integration. The StepperOnline F5D60-12GN-30S-5GN50K motor has a respectable rated torque point but is a spur gearbox design without an integrated encoder, which adds complexity for encoder feedback and makes closed-loop control harder to implement directly. On the other hand, the Duma Dynamics 13PPR Encoder 60 RPM DC Gearmotor does offer encoder feedback and decent torque, but it tends to be more expensive and less commonly used in the robotics community, which can make finding mounting options and support more difficult. The Yellow Jacket's planetary gearbox gives it an inherent mechanical advantage in torque density and smoothness over simpler spur gear alternatives, and its integrated encoder reduces wiring and sensor overhead relative to motors without feedback. For *SOLARIS* needs, budget, and ESP32 architecture, the Yellow Jacket provides the best blend of torque margin, encoder support, documentation quality, and community familiarity.

3.1.4 Solar Panel Pan/Tilt Motor Selection

Brushed DC Motor Model	Voltage	Rated Torque	Rated Speed	Gearbox	Encoder	Cost
5203 Series Yellow Jacket Planetary Gear Motor (188:1 Ratio, 24mm Length 8mm REX® Shaft, 30 RPM, 3.3 - 5 V Encoder)	12 V	24.5 Nm	30 RPM	Planetary	Yes	\$54.99
E-S Motor 12 V 105 RPM Worm Gearmotor with Encoder	12 V	3.92 Nm	54 RPM	Worm (Self Locking)	Yes	\$27.50
FeeTech FT5330M Digital Servo (180°)	7.4V	3.48 Nm	0.3 sec/60 degrees	Servo	Yes	\$20.74

Table 3.6 Brushed DC Motor Comparison

Sources: [20] [21-22]

3.1.4.1 Solar Panel Pan/Tilt Motor Selection

When selecting a motor for the solar panel pan-and-tilt system, the main goals were to ensure the motor could provide enough torque to move and hold the panel, operate from the project's electrical system, and remain practical for a senior design build. Since the panel is mounted near its center, the torque demand is lower than a worst-case edge-mounted estimate, but the motor still needs enough margin to account for imbalance, friction, and wind loading. The three options considered here were chosen because they each provide some combination of geared torque multiplication, position feedback, and manageable size or cost. Other important considerations included whether the motor could be controlled in a servo-like manner, whether it could be integrated easily into the mechanical design, and whether reliable documentation was available for design justification. Because the same motor is desired for both pan and tilt, torque margin, controllability, and documentation quality became more important than simply choosing the cheapest possible option.

3.1.4.2 goBILDA Yellow Jacket Planetary Gear Motor

This is a strong candidate because it offers a high torque rating of 24.5 N·m at 12 V while still maintaining a moderate speed of 30 RPM. Its planetary gearbox gives it much better torque density than a small un-geared motor and helps reduce the effort required from the motor itself during holding and positioning. The built-in encoder also makes it well suited

for closed-loop control, allowing it to be used in a servo-style system with a motor driver and controller. This motor greatly exceeds the estimated torque requirement, which gives a design margin for imbalance, wind disturbance, and startup torque. At the same time, it is not so large that it becomes unrealistic for the project's scale, making it a very balanced choice at an affordable price.

3.1.4.3 E-S Motor Worm Gear Motor w/ Encoder

This motor is appealing because its worm gearbox offers a self-locking characteristic, which is desirable for holding position with reduced back driving. This is desired because of efficiency. With a self-locking mechanism, there will be no need to send a hold signal to the motors and thus waste power. Its listed torque of 3.92 N·m is close to the original design target, so on paper it appears capable of handling the required load with a modest margin. The built-in encoder is also useful for feedback and position control, which makes it possible to implement controlled motion on both the pan and tilt axes. However, its higher 54 RPM speed is somewhat faster than desired, and more importantly, motors in this class are often limited by weaker documentation and fewer supporting ecosystem parts.

3.1.4.4 FeeTech Digital Servo Motor

This is the simplest option from a control perspective because it is already a packaged servo, meaning that it includes internal feedback and electronics control. Its torque rating of 3.48 N·m is within the general range of the project's needs, and its rated speed of 0.3 sec/60 degrees indicates that it is fast enough for slow solar tracking adjustments. Because it is a servo, it is easy to command directly without requiring a separate feedback loop implementation, which can reduce software and integration effort. However, it runs at 7.4 V instead of 12 V which increases integration complications, which require an additional regulator, and servos typically hold position by continuously applying torque rather than by mechanically resisting motion.

3.1.4.5 Solar Panel Pan/Tilt Motor Part Selection

goBILDA 5203 Series Yellow Jacket Planetary Gear Motor: Overall, this is the best choice because it offers the best combination of torque margin, control flexibility, documentation quality, and mechanical practicality. It is not a worm gearmotor, but truly well-documented 12 V worm gear motors with encoder support within a reasonable price range are very difficult to find. The 188:1 gear ratio is valuable because it means the motor only needs to generate a small fraction of the output torque at the motor shaft, which reduces the current needed to hold the panel in place, so the remaining power draw is mostly electrical loss, which is minimized by the large gear ratio. A mechanical brake can be easily and cheaply added to assist the holding of the motor if desired. On top of that, goBILDA provides strong documentation, and a large mounting and shaft ecosystem making integration much easier and troubleshooting than other options.

3.1.5 Motor Drivers

3.1.5.1 *Unidirectional Drivers*

Unidirectional motor drivers control speed but do not allow reversal of rotation. This method usually utilizes a ground side N-channel MOSFET switch that uses the PWM voltage signal to regulate the motor speed. High-side switch drivers operate in a similar manner but control the supply side of the motor instead which adds safety to the design. Linear drivers control voltage by operating the MOSFET in the linear region which provides smoother control but at the cost of power losses as heat. While these approaches are electrically simple and cheap they do not offer the bidirectional functionality required by the robots' driving system.

3.1.5.2 *Half-Bridge Drivers*

A half-bridge driver consists of two switching devices that control a single side of the motor and is a part of the H-bridge circuit topology. Half-bridge drivers do not offer bidirectional motion on their own but can be used in more specialized drive configurations due to their modularity. Half bridges provide good speed control with PWM voltage but are insufficient alone for the desired bidirectional movement.

3.1.5.3 *Relay-Based Driver:*

Relay motors drivers are drivers that use physical switching relays to reverse polarity and control the motor direction. While relays are excellent for discrete control in many industrial applications due to their robustness and simplicity, Relays switch slowly compared to solid state options and take up much more space within the robot to implement the same task. Slower switching is not generally well suited for speed control and mechanical wear is also a factor with continuous rapid switching. They can be combined with FET transistors for speed control improvements, but the added cost and complexity make the simplicity of the option hard to consider.

3.1.5.4 *H-Bridge Driver*

H-Bridge Drivers are common for brushed DC motor control due to their efficient and fast switching speeds, bidirectional motion and good speed control. The driver works by having 4 N-channel MOSFET transistors arranged in an H-shape that has the motor commutator in the center of the bridge. The MOSFETS switch so that the polarity of the voltage applied to the motor can be reversed which allows for bidirectional motion of the motor. H-Bridge drivers also use pulse-width modulation to regulate the voltage supply to the motor which enables precise motor speed control. H-Bridge drivers also provide electrical braking which allows faster stopping of the motor when its usage is no longer desired.




Figure 3.3 H-Bridge Topology

Source: [16]

3.1.5.5 Motor Driver Selection

H-Bridge Driver: The H-Bridge driver was selected for *SOLARIS* due to the balance of cost, efficiency and functional requirements. While unidirectional drivers are efficient and allow for precise speed control, they do not allow for the reversal of the voltage polarity thus lack the ability to reverse the direction of the motor. The H-Bridge's switch configuration allows for bidirectional motion which enables the robot to reposition itself as required for optimal sunlight exposure. Relay based drivers did not provide the switching speed required to precisely control the speed of the motor as desired and also introduced the unreliability of mechanical wear into the driver system. The PWM based speed control allows for precision and electrical braking which will enable more precise movement which will meet the requirements for robot navigation over the variable terrain. The H-Bridge driver is the most balanced solution that allows for all of the control, simplicity and ease of integration requirements that the project demands.

3.1.5.6 Motor Driver Components

Driver Model	Channels	Voltage Range	Current	Logic Voltage	PWM/C control	Cost
Pololu TB6612FNG Dual Motor Driver Carrier	2	4.5-13.5 V	1 A/Channel continuous, 3 A peak	2.7-5.5 V	Up to 100 kHz PWM	\$4.95
STMicroelectronics VN5019ATR-E	1	5.5 - 24 V	30 A	3.3/5 V	In/Out control 20 kHz	\$9.22
Pololu VN5019 Motor Driver	1	5.5 - 24 V	12 A Continuous, 30 A peak	2.5 - 5 V	Up to 20 kHz PWM, current-sense output	\$29.95

Table 3.7 DC Motor Driver Models

Source: [23]

When selecting an H-bridge motor driver for our DC motor, the most critical consideration is current handling capacity. The driver must safely support both the motor's continuous operating current and its higher startup surge or stall current without overheating or triggering protection faults. In multi-motor robots like *SOLARIS*, cumulative current draw during acceleration or turning maneuvers can momentarily spike, so driver selection must account for worst-case conditions rather than nominal cruising current. Under sizing the driver risks voltage drop, thermal shutdown, or permanent damage, while oversizing may increase cost and board footprint unnecessarily. Therefore, a balanced selection ensures sufficient headroom while remaining efficient and cost effective.

Another important factor is logic compatibility and control interface to ensure compatibility with the ESP32. The driver must accept 3.3 V logic signals or explicitly support mixed-voltage systems to avoid the need for level shifting hardware. Control schemes such as PWM + direction or locked antiphase control affect how easily the system can implement speed and directional control in firmware. Additionally, maximum PWM frequency support influences smoothness of operation. Selecting a driver that integrates cleanly with the ESP32 simplifies firmware programming and improves overall system robustness.

Protection features and electrical robustness also play a significant role in driver selection. Features such as over-current protection, thermal shutdown, under-voltage lockout, and reverse polarity protection can prevent catastrophic failure during testing or unexpected operating conditions. In outdoor mobile robotics, transient loads and battery voltage fluctuations are common, making integrated safeguards particularly valuable. Some drivers

also include current-sense outputs, which can be leveraged for monitoring motor load or detecting stall conditions. These protections improve safety, increase system durability, and reduce the risk of damage during iterative development.

Finally, practical integration considerations must be evaluated alongside electrical performance. The number of motor channels per board affects packaging, wiring complexity, and overall system layout, especially in a four-motor drivetrain. Board size, heat dissipation requirements, and mounting options influence mechanical integration within the chassis. Cost per channel must also be considered to maintain project budget constraints. Ultimately, driver selection requires balancing electrical capacity, control compatibility, protection features, packaging practicality, and cost to ensure reliable and efficient operation of the robot's drivetrain.

3.1.5.7 Motor Driver Component Selection

VNH5019ATR-E: The STMicroelectronics VNH5019ATR-E was selected over the other drivers due to its better power performance and system flexibility. The TB6612FNG is limited to low-current applications, making it an unideal for *SOLARIS*, whereas the VNH5019 supports much higher current and a much wider voltage range. Compared to the Pololu VNH5019 board, the standalone IC provides the same electrical performance but allows for a custom PCB design with improved thermal management and optimized layout for higher current functionality. Overall, the VNH5019ATR-E offers the robustness, scalability, and integration flexibility required for reliably driving motors in a 20–30 lb mobile robot without concerns for current and overheating or efficiency.

3.1.6 Solar Panel Technologies

<i>Solar Panel Technology</i>	<i>Efficiency</i>	<i>Weight</i>	<i>Cost</i>	<i>Temperature Performance</i>	<i>Degradation Rate</i>
Monocrystalline Silicon	Excellent	Good	High	Good	Low
Polycrystalline Silicon	Good	Good	Moderate	Moderate	Moderate
Thin-Film Silicon	Low	Excellent	Low	Best	High

Table 3.8 Solar Panel Technology Comparison

Sources: [24], [25]

3.1.6.1 What is a Solar Panel?

Solar photovoltaic (PV) modules operate by orienting an array of semiconductor cells toward incident solar irradiance. Upon absorption of photons within the p-n junction of the semiconductor material (typically silicon-based), electron-hole pairs are generated, producing a photovoltage via the photovoltaic effect. This voltage is scaled up by interconnecting multiple cells in series to achieve the desired output potential [24]. Commercially available PV technologies primarily differ in cell composition, efficiency,

and encapsulation methods, including the protective front glass and back sheet materials. Key selection criteria for PV modules include module efficiency, specific power (W/kg), mechanical and environmental durability (e.g., hail impact resistance and degradation rate), form factors (e.g., dimensions, flexibility, and mounting compatibility), and cost [25].

Inside a photovoltaic cell




Figure 3.4 Photovoltaic Cell Diagram

Source: [25]

3.1.6.2 Monocrystalline Silicon

Monocrystalline silicon solar panels are widely regarded as the highest-performance mainstream photovoltaic technology. Because each cell is cut from a single, continuous crystal structure, electron flow is more uniform, resulting in higher efficiencies typically in the 20-23% range for commercial modules. This means more power output per square meter compared to polycrystalline or thin-film alternatives, making them ideal when roof space is limited. They also tend to have slightly lower weight per watt because fewer panels

are needed to achieve the same system size, though individual modules weigh similarly to other silicon panels. The tradeoff is cost: monocrystalline panels are generally more expensive due to the energy-intensive crystal growth process and material purity requirements.

In terms of performance in real-world conditions, monocrystalline panels offer strong temperature characteristics, with lower temperature coefficients (often around -0.3 to $-0.4\%/^{\circ}\text{C}$). This means they lose less efficiency as temperatures rise, which is critical in hot climates where rooftop surfaces can become extremely warm. Additionally, they exhibit low annual degradation rates allowing them to retain over 85-90% of their original output after 25 years.

3.1.6.3 Polycrystalline Silicon

Polycrystalline silicon solar panels are a widely used and cost-effective photovoltaic technology, characterized by cells made from multiple silicon crystal fragments melted together. This manufacturing process is simpler and less energy-intensive than that of monocrystalline silicon, which translates to lower overall cost per panel. In terms of efficiency, polycrystalline panels typically operate in the 15-18% range, meaning they produce less power per unit area and require more space to achieve the same output as monocrystalline panels. Their weight per panel is generally comparable to other silicon-based modules, but because more panels are needed for a given system size, the total system weight can be slightly higher for equivalent power output.

Regarding performance and longevity, polycrystalline panels have slightly worse temperature performance, with temperature coefficients often around -0.38 to $-0.45\%/^{\circ}\text{C}$. As a result, their output drops more noticeably in high-temperature environments, making them somewhat less ideal for very hot climates. However, their degradation rates are still respectable, typically around 0.5-0.8% per year, allowing systems to maintain reliable output over decades.

3.1.6.4 Thin-Film Silicon

Thin-film silicon solar panels use very thin layers of photovoltaic material, deposited onto substrates like glass, metal, or flexible polymers. Because the semiconductor layer is much thinner than in crystalline panels, manufacturing can be less material-intensive and potentially lower at scale costs. However, efficiency is significantly lower, typically in the 10-13% range for silicon-based thin-film modules. This means more surface area is required to generate the same power as crystalline panels. On the positive side, thin film panels are generally lighter and can even be flexible.

In real-world operating conditions, thin-film silicon panels often perform well at high temperatures, with relatively favorable temperature coefficients compared to polycrystalline modules. Their output tends to degrade less sharply as temperatures rise, which can partially offset their lower nominal efficiency in hot climates. However, degradation rates can be somewhat higher in the early years due to light-induced degradation effects, though performance typically stabilizes afterward. Over the long term, annual degradation may range around 0.5-1% depending on manufacturing quality.

3.1.6.5 Solar Panel Technology Selection

Monocrystalline silicon: With industry-leading conversion efficiencies typically exceeding 20%, monocrystalline silicon panels generate more electrical power per unit area than polycrystalline or thin-film alternatives. For a mobile robotic platform like *SOLARIS*, surface area is limited and every square inch of panel space must be maximized. Higher efficiency directly translates to greater energy harvest under the same sunlight conditions, enabling longer operation time, improved battery charging capability, and more reliable performance in partially shaded or lower-irradiance environments. Simply put, monocrystalline technology allows *SOLARIS* to extract the maximum possible energy from the space available.

Beyond efficiency, monocrystalline panels also offer strong temperature performance and low long-term degradation, which reinforces their suitability for an outdoor robotic system. Their lower temperature coefficient ensures that power output remains more stable in hot conditions, preventing significant efficiency losses during peak sunlight hours when energy generation matters most. While thin-film technology often exhibits more favorable temperature coefficients than monocrystalline, the increased area of solar panel that would be required for equivalent power would not be worthwhile. Additionally, their low annual degradation rate ensures sustained performance over years of operation, protecting the long-term reliability of *SOLARIS*. While the upfront cost is higher, the improved energy density, durability, and system efficiency justify the investment by delivering consistent, high-quality power generation critical to the robot's functionality and mission success.

3.1.7 Solar Panel Parts

<i>20W 12V Solar Panel Model</i>	<i>Temperature Coefficient of Pmax</i>	<i>Size</i>	<i>Weight</i>	<i>Power Density (mW/in²)</i>	<i>Cost</i>
ACOPOWER HY020-12M	-0.45%/°C	13.6 × 18.5 × 0.8 in	6.2lbs	79.5	\$39.50
Newpowa NPA20S-12J	-0.36%/°C	15.6 × 12 × 0.9 in	3.16lbs	106.8	\$39.99
SLD TECH ST-20P-12	-0.5%/°C	15.4 × 14.2 × 1.2 in	4lbs	91.5	\$63.80
Ameresco C1D2-420J	-0.45%/°C	16.7 × 19.8 × 2 in	6lbs	60.5	\$167.70

Table 3.9 Solar Panel Part Comparison

Sources: [26]

The key voltage and current parameters when selecting a solar panel are the values at the maximum power point (V_{mp} and I_{mp}). A solar panel exhibits a nonlinear current–voltage (I – V) characteristic, meaning power output ($P = V \times I$) varies across the curve [26]. There exists a specific operating voltage at which the product of voltage and current is maximized, this is known as the maximum power point (MPP). Operating at higher voltage

is advantageous because, for a given power level, higher voltage corresponds to lower current. Since resistive cable losses are proportional to I^2R , reducing current significantly decreases conduction losses in wiring. Therefore, panels that allow higher operating voltage can improve system efficiency.

The temperature coefficient of P_{max} quantifies the rate at which a solar module's maximum power output decreases as cell temperature rises above Standard Test Conditions (25°C cell temperature). It is typically expressed in $\%/^{\circ}\text{C}$ and is negative for silicon modules, indicating power loss with increasing temperature. A coefficient closer to zero (i.e., less negative) means the panel experiences a smaller reduction in power as temperature increases, allowing it to maintain higher output under real-world operating conditions.

The Nominal Operating Cell Temperature (NOCT) specifies the expected cell temperature under defined conditions: 800 W/m^2 irradiance, 20 °C ambient temperature, and moderate wind speed. It provides a realistic operating temperature reference. The difference between NOCT and the 20 °C ambient used in its definition represents the typical temperature rise of the cells above ambient under those conditions. This value can then be used to estimate operating cell temperature at other ambient conditions.

For example, assume a solar panel module has a temperature coefficient of $-0.45\%/^{\circ}\text{C}$ and a Nominal Operating Cell Temperature (NOCT) of 45°C at 20°C ambient. This corresponds to a 25°C temperature rise above ambient under NOCT conditions. Assuming a 35°C ambient temperature on a sunny July day and similar irradiance, the expected cell temperature would be approximately 60°C. Since STC is defined at 25°C cell temperature, this represents a 35°C increase. Applying the temperature coefficient yields a total power reduction of 15.75%, resulting in an expected maximum power output of approximately 16.85 W.

The physical size, weight, and resulting power density of the solar panel directly impact on the compactness, mobility, and mechanical integration of the robot. A larger or heavier panel increases structural load, raises the center of mass, and requires greater torque from any panning or tilting mechanisms used for solar tracking. This can increase actuator power consumption, reduce stability, and complicate mechanical design. Therefore, minimizing panel size and mass is critical for maintaining efficient mobility and responsive directional control.

3.1.7.1 ACOPOWER HY020-12M

The ACOPOWER HY020-12M offers solid mechanical resilience for portable and small-scale applications, with an aluminum alloy frame, low-iron tempered glass, and EVA/TPT encapsulation. Its compact dimensions (13.6×18.5×0.8 in) and weight (6.2lbs) provide decent portability, while electrical performance includes V_{mp} of 17.5 V and I_{mp} of 1.14A at STC, with a module efficiency of 10.89% and $\pm 3\%$ power tolerance. Temperature coefficients are balanced (P_{max} $-0.45\%/^{\circ}\text{C}$, V_{oc} $-0.35\%/^{\circ}\text{C}$, I_{sc} $+0.05\%/^{\circ}\text{C}$), and NOCT of $45\pm 5^{\circ}\text{C}$ supports operation from -40°C to $+85^{\circ}\text{C}$. The 0-100% relative humidity rating listed for this panel is particularly advantageous for deployment in central Florida, where summer conditions routinely drive relative humidity levels to 90-100%. On the downside,

its lower efficiency translates to modest power density (about 79.5 mW/in²), necessitating larger arrays for higher energy needs, and the specific power (3.23 W/lbs.) is unremarkable for weight-sensitive robotics.

3.1.7.2 Newpowa NPA20S-12J

The Newpowa NPA20S-12J delivers exceptional efficiency and compactness, boasting 22% module efficiency with 33 Class A high-efficiency monocrystalline cells, yielding the highest power density (106.8 mW/in²) and specific power (6.33 W/lbs.) among the options. It features solid encapsulation for potential-induced degradation protection and offers a 0-+5W positive tolerance with V_{mp} of 19.08V and I_{mp} of 1.05A at STC. The panel's high reliability in harsh environments make it suitable for outdoor, mobile use in Florida's variable weather.

Weaknesses are minimal but include a slightly higher NOCT (45±2°C) that could affect peak-hour performance, though mitigated by the best temperature coefficient for power (-0.36%/°C), reducing heat-related losses. Its dimensions (15.6×12×0.9 in) and light weight (3.16lbs) are advantages. Its high performance-to-cost ratio makes this solar panel a strong candidate for *SOLARIS*.

3.1.7.3 Ameresco CID2-420J

The Ameresco 420J excels in durability and reliability for off-grid applications, featuring a robust design proven in extreme environments like Antarctica and the Australian outback. Its mechanical construction includes a thick, scratch-resistant white polyester back sheet for enhanced insulation and protection against rough handling, along with a silver anodized aluminum frame and high-transmission 3.2mm glass front cover. Electrically, it offers a nominal 12V output with V_{mp} of 16.8V and I_{mp} of 1.19A at STC, supported by certifications including IEC 61215/61730, UL1703, and suitability for Class 1 Division 2 hazardous locations.

However, the Ameresco 420J has limitations in efficiency and compactness, with a module efficiency of only 9.4% due to its 36 monocrystalline silicon cut cells, resulting in lower power density (60.5 mW/in²) and requiring more surface area for equivalent output. Its dimensions (16.7×19.8×2 in) and weight (6lbs.) make it bulkier and heavier than competitors, reducing specific power (3.33 W/lbs.) and suitability for space-constrained mobile applications. Temperature performance is moderate, with a P_{max} coefficient of -0.45%/°C and NOCT of 47±2°C, leading to noticeable derating in hot climates like Florida, while the ±10% power tolerance introduces variability in performance.

3.1.7.4 SLD TECH ST-20P-12

The SLD TECH ST-20P-12 achieves a module efficiency of 20% with 32 monocrystalline cells, enabling strong power density (about 91.5 mW/in²) and a positive tolerance of 0-+5W for reliable output. It features heavy-duty anodized frames, high-transparent low-iron tempered glass, and high potential-induced degradation resistance. This panel is resistant to salt mist, per IEC 61701, making it ideal for coastal or harsh weather. Electrical specs include V_{mp} of 18.53V and I_{mp} of 1.08A at STC.

Weaknesses include a relatively poor temperature coefficient for power (-0.5±0.05%/°C), the most negative among comparators, causing greater output loss in high temperatures. Having a NOCT of 47±2°C also indicates this solar panel operates hotter than the other

ones, which further reduces power efficiency. Its dimensions (15.4×14.2×1.2 in) and weight (4lbs) yield moderate specific power (5 W/lbs.).

3.1.7.5 Solar Panel Part Selection

Newpowa NPA20S-12J: This panel was selected primarily for its superior power density and optimal temperature coefficient of Pmax, making it the most suitable for the *SOLARIS* mobile robotic platform operating in Central Florida’s hot and humid climate. With a power density of approximately 106.8 mW/in², it maximizes energy harvest from limited mounting space, ensuring extended autonomous operation without expanding the robot's footprint. Its Pmax temperature coefficient of -0.36%/°C, the least negative among the panels, minimizes derating during peak solar hours when ambient temperatures often exceed 35°C (cell temperatures reaching 50-65°C), preserving 5-10% more output compared to panels with -0.45% to -0.5%/°C coefficients. This combination, alongside high specific power (6.33 W/lbs) for reduced system weight, justifies the choice. Delivering reliable, high-density power in variable thermal conditions is critical for mission success, and the Newpowa NPA20S-12J panel accomplishes that goal effectively.

3.1.8 Solar Panel Charge Controller Technologies

<i>Solar Panel Charge Controller Technology</i>	<i>Power Efficiency</i>	<i>Daily Energy Capture</i>	<i>Safety</i>	<i>Complexity</i>	<i>Cost</i>
Multi Power Point Tracker	95-99%	High	High	High	High
Pulse Width Modulation	70-80%	Moderate	Moderate	Moderate	Moderate

Table 3.10 Solar Panel Charge Controller Technology Comparison

Sources: [27]

3.1.8.1 What is a Solar Panel Charge Controller?

A solar panel charge controller regulates the power flowing from a photovoltaic panel to a battery, ensuring the battery charges safely, efficiently, and without overvoltage or overcurrent damage. It prevents overcharging, limits reverse current at night, and manages charging stages to protect battery lifespan. The two primary types are PWM (Pulse Width Modulation) and MPPT (Maximum Power Point Tracking). PWM controllers are simpler and less expensive, but they force the panel to operate near battery voltage, which reduces efficiency and lowers daily energy capture. MPPT controllers are more complex and costly because they use DC-DC conversion to continuously track the panel’s maximum power point, allowing higher efficiency and greater daily energy harvest. In short, PWM offers lower cost and simplicity, while MPPT offers higher efficiency and better overall energy performance [27].

Constant current (CC) and constant voltage (CV) charging are critical for LiFePO₄ batteries because they ensure the battery is charged efficiently without exceeding safe electrical

limits. During the CC stage, the battery receives a controlled, steady current that allows it to charge quickly while its voltage rises in a predictable manner. Once the battery reaches its specified upper voltage limit, the charger must transition to the CV stage, holding the voltage constant while the current tapers off as the battery approaches full capacity. This prevents overvoltage stress, overheating, and long-term degradation of the cells. A solar panel charge controller should ideally manage this CC–CV profile automatically, regulating both current and voltage to match the battery’s specifications while also coordinating with the battery management system (BMS). Without proper CC and CV control, the battery’s lifespan, safety, and performance can be significantly compromised.

3.1.8.2 Maximum Power Point Tracker

An MPPT controller uses a DC-DC power converter to actively regulate the solar panel operating voltage at its maximum power point. Because a solar panel has a nonlinear I–V curve, there is a specific voltage (V_{mp}) at which the product of voltage and current is highest. The converter then steps the panel voltage down to the required battery charging voltage while conserving power (minus conversion losses). By allowing the panel to operate at its optimal voltage and current combination, the MPPT controller increases the total energy delivered to the battery compared to other solar panel charge controllers, particularly when environmental conditions vary throughout the day.

3.1.8.3 Pulse Width Modulation

A PWM solar charge controller regulates battery charging by rapidly switching the connection between the solar panel and the battery on and off at a high frequency, effectively controlling the average current flowing into the battery. Instead of converting voltage, a PWM controller essentially forces the solar panel to operate near the battery’s voltage by directly coupling them during each pulse. When the battery is low, the controller allows longer “on” pulses to deliver more current. As the battery voltage approaches its target charge voltage, the controller shortens the pulse width to taper the charging current and prevent overcharging. This simple switching method approximates CC and CV charging behavior through duty cycle modulation in practice, but because the panel voltage is pulled down to the battery voltage rather than held at its maximum power point, overall energy harvest and efficiency are lower compared to MPPT systems.

3.1.8.4 Solar Panel Charge Controller Technology Selection

Multi Power Point Tracker: For the *SOLARIS* platform, an MPPT (Maximum Power Point Tracking) charge controller was selected as the interface between the photovoltaic module and the 12 V LiFePO₄ battery system. The primary design driver for this project is maximizing usable energy generation within strict size and power constraints. Operating the solar panel at battery voltage (as in PWM control) does not generally correspond to its maximum power point. An MPPT controller continuously adjusts the panel operating voltage to maintain operation at V_{mp} while performing DC–DC conversion to deliver the appropriate CC/CV charging profile to the battery. This enables higher instantaneous conversion efficiency and increased total daily energy harvest, particularly under elevated cell temperatures and varying irradiance conditions typical of outdoor operation.

Although PWM controllers offer lower cost, reduced circuit complexity, and simpler implementation, their inability to track the maximum power point results in reduced energy extraction. Even a 10–15% improvement in daily energy capture can have a significant

impact on the mission duration, recharge time, and system energy balance. Given that *SOLARIS* operates under a constrained energy budget, prioritizing energy optimization over minimal cost or simplicity is justified. Therefore, the performance gains and improved system efficiency provided by MPPT technology outweigh the incremental increase in cost and design complexity.

3.1.9 Solar Panel Charge Controller Parts

<i>Solar Panel Charge Controller Part</i>	<i>Electrical Efficiency</i>	<i>Operating Power Consumption</i>	<i>Key Features</i>	<i>Weight</i>	<i>Cost</i>
GV-5-LI-14.2V	94-99.85%	0.19 mA	Fixed charge algorithm, 10-year warranty, IP40, 6-position terminal block connection	0.176lbs	\$99
GV-5-MOD-LFP	94-99.85%	0.15 mA	Fixed charge algorithm, 10-year warranty, 24-pin connection	0.062lbs	\$80
Victron SmartSolar MPPT 75/10	≤98%	20 mA	Built-in Bluetooth, programmable charge algorithm, 5-year warranty, IP22, screw terminal connection	1.102lbs	\$73
RNG-CTRL-RVR20-US	≥95%	≤100 mA	Optional Bluetooth add-on, programmable charge algorithm, 3-year warranty, IP32, screw terminal connection	3.1lbs	\$116

Table 3.11 Solar Panel Charge Controller Part Comparison

Sources: [28]-[31]

The solar panel charge controller in *SOLARIS* is essential for regulating the power from the solar panel to safely charge the 12V LiFePO₄ battery, preventing overcharging, undercharging, or damage. Selecting an optimal solar panel charge controller requires a balanced evaluation of electrical efficiency, power consumption, and integrated features to ensure system longevity. As shown in the comparison of various models, electrical efficiency typically ranges from 95% to as high as 99.85%. While efficiency is critical for maximizing energy harvest, the operating power consumption is equally vital for maintaining battery state of charge during periods of low solar input.

The final selection often hinges on specific application needs, such as connectivity and physical durability. Programmable controllers like the Victron SmartSolar and the RNG-CTRL-RVR 20-US offer built-in or optional Bluetooth capabilities, allowing for real-time monitoring and custom charge algorithms. Conversely, the Genasun models prioritize simplicity and reliability with fixed charge algorithms and an impressive 10-year warranty, compared to the 3-to-5-year warranties found on other units. Physical constraints must also be considered; weights vary significantly from the ultra-lightweight GV-5-MOD-LFP at 0.062 pounds to the more robust RNG-CTRL-RVR 20-US at 3.1 pounds.

3.1.9.1 *GV-5-LI-14.2V*

The Genasun GV-5-Li-14.2V is a compact MPPT solar charge controller designed for 12V lithium batteries, boasting an industry-leading 99.85% peak efficiency that extracts up to 30% more power than PWM controllers in varying conditions, such as low light or partial shading [28]. Key features include electrolytic-free ceramic capacitors for exceptional durability, ultra-fast true MPPT tracking at 15 Hz, excellent low-light performance, and a fixed CC/CV charging profile at 14.2V. What sets it apart is its military-grade reliability suited for harsh environments like marine, RV, or off-grid applications, backed by a 10-year warranty and a lightweight 0.176 pounds form factor that ensures long-term performance without wear-out components, making it ideal for small, high-efficiency setups where longevity and maximal energy harvest are priorities.




Figure 3.5 Genasun GV-5-LI 5A MPPT Solar Buck Controller

Source: [28]

3.1.9.2 *GV-5-MOD-LFP*

The Genasun GV-5-MOD-LFP is the more modular version of the GV-5 series, optimized for 12V (4S) Lithium Iron Phosphate batteries with a fixed CC/CV charging profile at 14.2V, offering the same 99.85% peak efficiency and 5A rated output for up to 65W panels [29]. Notable features include a 24-pin header for seamless PCB integration, low self-consumption of 0.15mA, and protections against overvoltage, reverse polarity, and short circuits, while supporting custom voltage variations for diverse battery types. On the contrary, this component is limited by its lack of ingress protection making it require a proper enclosure to operate safely, long-term. It stands out for its OEM-focused design that reduces assembly labor and system costs in solar-powered products, providing plug-and-play reliability in embedded applications like industrial or street lighting, with a 10-year warranty and compact dimensions (3.28" x 1.88" x 0.48") that differentiate it from standard controllers by prioritizing integration ease over user-facing interfaces.



Figure 3.6 Genasun GV-5-MOD-LFP 5A MPPT Solar Charge Controller

Source: [29]

3.1.9.3 Victron SmartSolar MPPT 75/10

The Victron SmartSolar MPPT 75/10 is a high-quality MPPT charge controller for 12 V/24 V systems, delivering up to 10 A output with 98% peak efficiency, supporting up to 145 W at 12 V and 75 V max PV open circuit voltage, and featuring advanced adaptive 3-stage charging with flexible algorithms, automatic equalization, and temperature compensation via internal or external sensors [30]. Notable features include built-in Bluetooth for wireless configuration via the VictronConnect app, VE.Direct port for data communication, VE.Smart Networking for synchronized operation with other devices, a 15 A load output with BatteryLife algorithm, and extensive protections like over-temperature and PV short circuit. It distinguishes itself through superior smart connectivity and ecosystem integration within the Victron lineup, enabling remote monitoring, firmware updates, and scalability for professional installations in boats, vans, or off-grid homes, backed by a 5-year warranty and robust build quality that emphasizes long-term reliability and user customization over basic plug-and-play simplicity.



Figure 3.7 Victron SmartSolar MPPT 75/10 Charge Controller

Source: [30]

3.1.9.4 RNG-CTRL-RVR20-US

The Renogy Rover Li 20A is a versatile MPPT solar charge controller tailored for lithium batteries in 12V/24V systems, featuring a 20A charge current, 100V maximum PV input, and up to 260W power handling at 12V, with 98% efficiency and programmable parameters via an integrated LCD screen for easy monitoring and customization [31]. Key features include auto voltage recognition, full system protection (overcharge, overload, short circuit, reverse polarity), RS485 communication for optional Bluetooth monitoring,

and compatibility with a variety of battery chemistries, along with a 4-stage charging algorithm (bulk, boost, float, equalization). What sets it apart is its user-friendly interface with built-in diagnostics and error codes, making it accessible for DIY enthusiasts in RV or home setups, combined with a higher power capacity for medium-scale systems at an affordable price point, though with a shorter 3-year warranty compared to premium competitors.



Figure 3.8 Rover Li 20 Amp MPPT Solar Charge Controller

Source: [31]

3.1.9.5 Solar Panel Charge Controller Part Selection

GV-5-LI-14.2V: The GV-5-LI-14.2V was selected as the solar panel charge controller for *SOLARIS* due to its exceptional suitability for this compact robotic application, where reliable solar harvesting is key. Weighing just 0.176 pounds, its low weight minimizes the overall mass of the robot, enhancing mobility and energy efficiency without compromising performance. The controller's compact dimensions (4.3" x 2.2" x 1.14") and simple screw-terminal design make it incredibly easy to mount directly onto the robot's chassis, reducing installation complexity and integration time. Furthermore, its ultra-low self-consumption of only 0.19 mA ensures minimal parasitic drain on the battery during standby, preserving precious energy in a mobile system like *SOLARIS*. Finally, boasting the highest peak efficiency in its class at 99.85%, it maximizes power extraction from the solar panel even in partial shade or low-light conditions common in variable outdoor environments, providing up to 50% more energy harvest compared to PWM alternatives and backed by a robust 10-year warranty for long-term reliability.

3.1.10 Light Sensors

<i>Sensor Type</i>	<i>Ease of Implementation</i>	<i>Can Detect Sun vs Shade?</i>	<i>Multi-Sensor Scalability (4-8 sensors)</i>	<i>Detecting Small Directional Differences</i>	<i>Cost</i>
Photoresistors (LDR)	Very Easy	Yes	No issues	Moderate (works well with shading geometry)	Very Low
Phototransistor	Medium	Yes	No issues	Good (more sensitive than LDRs, but still benefit from shading geometry)	Low
Photodiode	Medium / Hard	Yes	No issues	Very good (with an amplifier. Raw signal is small)	Low - Medium
Digital Light Sensor (LUX, I2C)	Medium	Yes	Potential address issues (using multiple of the same sensor)	Excellent. Gives exact LUX values. In terms of accuracy it is the best.	Medium

Table 3.12 Light Sensing Technology Comparison

Sources: [32]-[35]

3.1.10.1 Photoresistor




Figure 3.9 LDR Light Sensing Circuit

Source: [36]

Photoresistors, also called light dependent resistors (LDRs), are one of the simplest and most common ways to measure light in embedded systems. An LDR changes its resistance based on the amount of light hitting its surface. The most straightforward way to interface an LDR with a microcontroller is by using a voltage divider. The LDR is placed in series with a fixed resistor, and the midpoint voltage is connected to an ADC pin on the MCU. This makes the overall hardware and software implementation very simple, since the system only needs to sample an analog voltage and compare it to the other sensor readings.




Figure 3.10 LDR Resistance Vs. Temperature

Source: [37]

One major drawback of photoresistors is that they are not designed for accurate or repeatable light measurement. Their response to changing light is relatively slow compared to semiconductor-based sensors, and their readings can fluctuate due to temperature changes. This makes them less ideal for applications where exact or consistent intensity measurements are required.




Figure 3.11 LDR Non-Linear relationship between LUX and Resistance

Source: [37]

Another limitation of photoresistors is that their response to light is not linear. This means that a change in resistance does not correspond to a proportional change in illumination, especially across different lighting conditions. As a result, comparing measurements between multiple LDRs can be less consistent than with semiconductor-based sensors. While this does not prevent LDRs from being used for directional tracking, it does make them less ideal for applications where consistent scaling and predictable sensor matching are desired.

However, for *SOLARIS*, the goal is not to measure absolute illuminance in lux, but instead to detect which direction contains the strongest sunlight so the robot can orient itself accordingly. In this case, LDRs are still considered “good enough,” especially when used in a geometry where the chassis or a small divider structure creates shaded and unshaded regions. The sensor array can reliably detect sunlight versus shade and provide enough directional information for tracking, and the slower response time is not a major concern since the system does not need to react instantly to light changes.

Overall, photoresistors offer an extremely low-cost and low-complexity solution that can provide directional tracking information with minimal hardware overhead. Even though they are not the most accurate sensing option, their simplicity makes them a strong baseline sensor choice for the *SOLARIS* light tracking system, and they are widely used in low-cost solar tracking designs for this reason.

3.1.10.2 Phototransistor

Phototransistors can be thought of as a photodiode and a bipolar junction transistor (BJT) combined into a single device. Like a normal BJT, they have an NPN or PNP structure, but instead of driving the base terminal, the base region is exposed to incoming light. When

light hits the device it generates a small photo-induced current, like a photodiode. The key difference is that this small current is then amplified by the transistor's internal gain, which produces a much larger output current than a standalone photodiode would typically provide. Because of this built-in gain, phototransistors often do not require an external amplifier stage to generate a useful signal, making them a practical middle ground between photodiodes and photoresistors.




Figure 3.12 Phototransistor Light Sensing Circuit

Source: [38]

A simple phototransistor sensing circuit can be implemented using a single resistor as a load, as shown in the figure. In this configuration, the phototransistor is connected to the supply voltage and the resistor is connected to ground, with the output voltage measured at the junction between the two. As the light level increases, the phototransistor conducts more current. This increases the voltage drop across the resistor, causing the measured output voltage to rise. The output voltage can be read directly by the microcontroller's ADC, making the software side straightforward and similar to reading an LDR voltage divider.

This circuit is also safe to interface with a 3.3 V microcontroller as long as the circuit is powered from the same 3.3 V supply. Since the output voltage cannot exceed the supply voltage, the ADC input will remain within the valid 0–3.3 V range. The main design choice is selecting the resistor value so the signal does not saturate at 3.3 V too easily. A larger resistor increases sensitivity but can cause the output to peg near the supply voltage in bright sunlight, while a smaller resistor provides more headroom at the cost of reduced sensitivity. In the context of *SOLARIS*, phototransistors provide a strong, easy-to-read

signal with minimal additional circuitry, while still enabling reliable detection of light differences between different directions around the robot.

3.1.10.3 Photodiode




Figure 3.13 Photodiode Light Sensing Circuit

Source: [39]

Photodiodes are semiconductor light sensors that convert light into a small current. In most practical sensing applications, photodiodes are operated in reverse bias, meaning a reverse voltage is applied across the diode. Reverse bias increases the size of the depletion region, which reduces the junction capacitance. This improves response time significantly and also helps produce a more predictable, linear relationship between light intensity and output current.

One drawback of photodiodes is that the raw output signal is typically very small. While the device itself produces a photocurrent proportional to the light level, this current is often too low to directly interface with a microcontroller ADC in a stable way. As a result, photodiodes are commonly paired with an amplifier stage (such as a transimpedance amplifier) to convert the photocurrent into a readable voltage signal. This increases overall circuit complexity compared to photoresistors or phototransistors, but it also allows for improved signal-to-noise ratio and more consistent measurements.

Photodiodes also provide the fastest response time of the sensor types considered, often operating in the nanosecond range depending on the device and biasing method. This makes them ideal for high-speed optical communication and precision measurement applications. However, for *SOLARIS*, the system does not require extremely fast light sensing since solar tracking is a slow process and does not need rapid reaction to change in illumination.

An additional benefit of photodiodes is that they can be selected based on semiconductor material and spectral response, meaning specific devices can be chosen for ultraviolet (UV), visible, or infrared (IR) detection. While this capability can be useful in specialized sensing applications, *SOLARIS* primarily needs to detect general sunlight intensity for directional tracking, so broad visible-light response is sufficient for this project.

Overall, photodiodes offer excellent speed and linearity and can provide very high-quality measurements when paired with the proper amplifier circuitry. However, the additional analog design effort required to obtain a stable, readable output makes them less practical for the *SOLARIS* tracking system compared to simpler sensor options.

3.1.10.4 Digital Light Sensor (LUX, I2C)

Digital light sensors are integrated devices designed to measure ambient light intensity and report it as a digital value, often in units of lux. Unlike photoresistors, photodiodes, or phototransistors, these sensors typically include internal circuitry such as amplification, filtering, and an analog-to-digital converter (ADC), allowing the microcontroller to receive a stable numeric measurement rather than interpreting a raw analog voltage. This makes digital light sensors one of the most robust and repeatable options for measuring light intensity in embedded systems.

A major advantage of digital light sensors is that they provide an actual lux-based output rather than a relative signal. This can be useful for system logging, long-term analysis, and more consistent comparisons between sensors. Because the measurement and conversion are handled internally, digital sensors are also less affected by ADC noise, resistor tolerance, and unit-to-unit variation compared to analog light sensing approaches.

However, digital light sensors typically communicate over I2C, meaning they share the same two-wire bus (SDA and SCL) with other I2C peripherals. While this simplifies wiring for a single sensor, it can create scalability issues when multiple identical sensors are required. Many digital light sensors use a fixed I2C address or only provide a limited number of selectable addresses. As a result, using multiple identical devices on the same I2C bus can become problematic without additional hardware such as an I2C multiplexer. This makes digital light sensors less convenient for multi-sensor directional tracking configurations where four or more sensors are required around the robot.

Overall, digital lux sensors provide the most complete and accurate light measurement capability of the sensor types considered, and they are especially useful for systems that require precise measurement and data logging. For *SOLARIS*, digital light sensors remain a strong option for single-point illuminance measurement, but their reliance on I2C addressing can complicate the design if many sensors are needed for directional tracking.

3.1.10.5 Light Sensing Technology Selection

Phototransistors: Phototransistors were selected as the light sensing technology for *SOLARIS* because they provide high sensitivity and a strong analog output signal while still requiring minimal circuitry. Unlike photoresistors, phototransistors have built-in current gain, meaning the light-induced signal is amplified internally. This allows them to produce a larger and more stable output signal without requiring an external amplifier stage.

Phototransistors were chosen over photodiodes primarily due to implementation complexity. While photodiodes can provide fast and highly linear sensing, their current output is typically very small and often requires a dedicated amplifier circuit (such as a

transimpedance amplifier) to generate a usable voltage signal. For *SOLARIS*, solar tracking is a slow process and does not require nanosecond-scale response times, so the added analog complexity of photodiodes is not justified for this application.

Digital I2C light sensors were also considered because they provide direct lux readings and can offer very repeatable measurements. However, they introduce additional software and communication complexity due to the use of the I2C bus, and multi-sensor configurations can become problematic due to device addressing limitations. Since *SOLARIS* requires multiple sensors around the robot for directional tracking, relying on multiple identical I2C light sensors would likely require extra hardware such as an I2C multiplexer.

In addition to these design tradeoffs, a published study [34] comparing LDR-based solar tracking to phototransistor-based tracking found that phototransistors produced approximately a 40% increase in voltage gain under the same sensing setup and software approach. This supports the decision to use phototransistors for *SOLARIS*, since stronger sensor output improves the ability to detect directional light differences and reduces sensitivity to noise. Overall, phototransistors provide the best balance of sensitivity, simplicity, scalability, and tracking performance for *SOLARIS* directional sunlight detection.

3.1.10.6 Phototransistor Structural Type Comparison

3.1.10.6.1 BJT-Based Phototransistors

[40] Most discrete phototransistors are implemented using a bipolar junction transistor (BJT) structure. In this configuration, incident light generates a small photo-induced base current within the device. This base current is then amplified by the transistor's internal current gain (β), producing a larger collector current. Because of this built-in gain, BJT phototransistors can produce a usable analog signal with only a simple load resistor, making it easy to interface with a microcontroller ADC. The operational description provided earlier in this section assumes this standard BJT-based structure, as it represents the dominant architecture used in general-purpose light sensing devices. BJT phototransistors are widely available in silicon form, are inexpensive, and are well suited for low-frequency or DC light sensing applications.

3.1.10.6.2 FET-Based Phototransistors

Photo-FET devices work differently than BJT phototransistors. Instead of light creating a small base current that gets amplified, the incoming light affects the electric field at the gate region, which changes how much current can flow through the channel. In other words, the device behaves more like a voltage-controlled element rather than a current-amplified one. Photo-FETs are typically used in applications requiring high input impedance, low noise, or higher frequency optical detection, such as optical communication systems and precision instrumentation. They are less common in simple discrete, through-hole packages intended for basic analog light sensing. For low-frequency sunlight detection applications, the additional complexity and performance characteristics of photo-FET devices provide no meaningful advantage.

3.1.10.6.3 Structural Type Selection

BJT: For *SOLARIS*, the light signal we are measuring is essentially DC. Sunlight does not change rapidly, and even when it does (for example, passing clouds), the change happens very slowly relative to electronic time scales. Because of this, the circuit does not need high bandwidth or special AC performance. What matters more is having a stable, predictable analog output that can be read reliably by the 3.3 V ADC. Due to their internal gain, ease of biasing with a single resistor, and suitability for low-frequency light detection, silicon BJT-based phototransistors are selected for this design. All further analysis and component comparison assume a standard BJT phototransistor structure.

3.1.10.7 Phototransistor Current Gain Configuration

3.1.10.7.1 Standard Phototransistor

[41] A standard phototransistor consists of a single bipolar junction transistor with a light-sensitive base region. When light strikes the base–collector junction, a photocurrent is generated and then amplified by the transistor’s internal current gain, producing a collector current that increases with illumination. This configuration results in moderate gain and relatively faster response compared to multi-stage devices. Because only one transistor stage is used, standard phototransistors are a common choice for general light sensing applications where a balance of sensitivity and dynamic response is desired.

3.1.10.7.2 Photodarlington

A photodarlington is a phototransistor configuration in which the emitter of a light-sensitive transistor stage drives the base of a second transistor. This provides two stages of current gain in series, resulting in overall gain that is the product of the gains of each transistor. The two-stage design yields significantly higher sensitivity, making photodarlingtons useful in low-light applications. However, the higher gain means that under high illumination the device will generate a much larger collector current, making it more prone to producing output voltages near the supply limit in bright light, essentially clipping the output too early.

3.1.10.7.3 Gain Configuration Selection

Standard: For the *SOLARIS* project, the light sensors will be exposed to bright outdoor sunlight and used to compare relative illumination across directions. In this context, extreme internal gain is not required and can be counterproductive. With a fixed 3.3 V supply and load resistor, a phototransistor with excessively high gain will produce large collector current under bright sunlight, which can push the output voltage toward the supply and reduce the usable range, meaning small changes in brightness become harder to distinguish. A standard phototransistor, with its lower gain and faster response, is less likely to saturate the output in strong light and provides more usable headroom for ADC readings. For these reasons, a standard (non-Darlington) phototransistor is selected for *SOLARIS*.

3.1.10.8 Phototransistor Hardware Comparison

	<i>Collector Light Current</i>	<i>Viewing Angle</i>	<i>Wavelength Range</i>	<i>Dark Current</i>	<i>Max Operating Voltage</i>	<i>Cost per unit</i>
VTT9812FH	0.1 mA	56°	450-700 nm	100 nA	30 V	\$0.57
VTT9814FH	0.1 mA	50°	450-700 nm	50 nA	30 V	\$0.63
TEPT5600	20 mA	20°	440-800 nm	50 nA	6 V	\$0.58
TEPT4400	20 mA	30°	440-800 nm	50 nA	6 V	\$0.40

Table 3.13 Phototransistor Comparison

Sources: [42]-[44]

To select an appropriate light sensing device for *SOLARIS*, several through-hole phototransistors were evaluated. Through-hole packaging was intentionally prioritized over surface-mount options because the sensors will be mounted externally on the robot chassis rather than directly soldered to the main PCB. Using through-hole components allows flexible placement and simplifies connection through jumper wiring while maintaining mechanical robustness in outdoor conditions.

Additionally, devices with spectral responses centered within the visible light range were preferred. Since *SOLARIS* performs relative brightness comparisons rather than absolute irradiance measurements, selecting a phototransistor responsive to visible wavelengths improves consistency with perceived light intensity and reduces potential distortions caused by excessive infrared sensitivity. These constraints limited the amount of phototransistors we could choose from. The options in the comparison table were the best we found.

3.1.10.8.1 VTT9812FH

The VTT9812FH offers a maximum operating voltage of 30 V and a moderate angular response of approximately 56 degrees. Its dark current is low, and its through-hole package makes it mechanically compatible with the *SOLARIS* mounting requirements. However, its wider viewing angle increases susceptibility to reflected light from surrounding surfaces, which can reduce contrast between opposing sensors. Additionally, its light current sensitivity is lower compared to the TEPT series under similar illumination conditions.

3.1.10.8.2 VTT9814FH

The VTT9814FH shares similar characteristics with the VTT9812FH, including a 30 V operating margin and through-hole packaging. It maintains a relatively wide angular response near 50 degrees. Although the higher breakdown voltage provides additional margin, this benefit is not meaningful in a 3.3 V system. Like the VTT9812FH, its broader field of view increases exposure to indirect light, which may reduce tracking contrast. As a result, it does not provide a strong advantage for the intended application.

3.1.10.8.3 TEPT5600

The TEPT5600 features a significantly narrower angular response of approximately 20°, which improves directional sensitivity and enhances contrast between illuminated and shaded regions of the robot chassis. Its spectral response range of 440 to 800 nm is centered within the visible spectrum, aligning well with natural sunlight conditions. Although its maximum collector-emitter voltage is limited to 6V, this remains fully compatible with the 3.3 V operating environment of *SOLARIS*.

3.1.10.8.4 TEPT4400

The TEPT4400 was also evaluated because it closely matches the desired characteristics for *SOLARIS*. It provides a visible-centered spectral response similar to the TEPT5600 and maintains low dark current, making it well suited for relative sunlight comparison. Its viewing angle is wider than the TEPT5600 (around 30 degrees rather than 20 degrees), which can be beneficial for tolerance to mounting misalignment and small changes in orientation while still preserving strong directional contrast. In other words, it offers a good middle ground between being directional enough to improve “brighter side vs darker side” discrimination while not being so narrow that the sensor becomes finicky. However, despite being a strong candidate, the TEPT4400 was not selected due to current availability limitations.

3.1.10.9 Phototransistor Selection

TEPT5600: The TEPT5600 was ultimately selected for *SOLARIS* based on its mechanical compatibility, spectral response, and directional characteristics. Its narrower viewing angle improves contrast between illuminated and shaded sides of the robot, which is important since the control system relies on comparing relative light intensity rather than measuring absolute values. Although the device is sensitive and could approach saturation under direct sunlight, this can be managed through proper resistor selection in the signal conditioning stage. Its visible-spectrum response aligns well with natural sunlight, and the through-hole package makes it practical for external mounting and jumper-wire integration on the chassis.

3.1.11 Energy Monitoring Technologies

<i>Energy Monitoring Technology</i>	<i>Accuracy</i>	<i>Power Loss</i>	<i>Complexity</i>	<i>Key Features</i>
Coulomb Counting	High (short-term)	Low	Moderate	Drift accumulates
Hybrid (Coulomb Counting + Open Circuit Voltage)	Very High	Moderate	High	Most accurate long-term

Table 3.14 Energy Monitoring Technology Comparison

Sources: [45]-[47]

3.1.11.1 Coulomb Counting

Coulomb counting provides high short-term accuracy by integrating bidirectional current over time to track the net charge entering or leaving the battery [45]:

$$SOC = SOC_{initial} \pm \frac{\int I dt}{Capacity}$$

Figure 3.14 Coulomb Counting Formula

This method requires precise current sensing and periodic recalibration to correct drift caused by measurement error, battery self-discharge, and temperature variation. Implementation complexity is moderate, as it requires accurate current sensing hardware and processing to accumulate charge over time. Power loss is minimal but present due to the shunt resistor's I^2R dissipation, which can be made negligible by selecting a very low resistance value. The primary limitation of coulomb counting is accumulated drift during long-term operation if no correction mechanism is used [47].

3.1.11.2 Hybrid (Coulomb Counting + Open Circuit Voltage)

The hybrid approach combines coulomb counting for real-time charge tracking with periodic open-circuit voltage (OCV) calibration to correct accumulated drift. During periods of low load or rest, the battery voltage can be compared against a pre-characterized LiFePO_4 OCV–SOC curve to recalibrate the estimated state of charge [46]. Although this approach increases implementation complexity due to additional sensing and processing requirements, it provides significantly improved long-term SOC accuracy. Hybrid methods are particularly effective in systems where the battery may experience simultaneous charging and discharging conditions, such as solar applications.

3.1.11.3 Energy Monitoring Technology Selection

Hybrid: The hybrid energy monitoring approach is selected for *SOLARIS* due to its ability to provide reliable long-term SOC estimation in environments with fluctuating solar input and variable load demand. In this system, coulomb counting will continuously track charge flow to provide real-time energy monitoring, while the microcontroller periodically performs OCV-based recalibration when stable voltage conditions are detected, such as after charging cycles or during low-load periods. By referencing a pre-characterized LiFePO_4 OCV curve, the system can correct any accumulated drift in the coulomb counter and maintain accurate SOC estimation over extended operation. Although this approach introduces additional system complexity, the improved accuracy and reliability of the energy monitoring data provided to the user outweighs the added implementation overhead.

3.1.12 Energy Monitoring Parts

<i>Energy Monitoring Part</i>	<i>Current Accuracy</i>	<i>Bus Voltage Range</i>	<i>Accessible Measurements</i>	<i>Cost of One IC</i>
LTC2944	±1%	3.6-60 V	Battery Voltage, Current, Die Temperature, Accumulated Charge	\$12.82
INA226	±0.02-0.1%	0-36 V	Shunt Voltage, Battery Voltage, Current, Power	\$2.99
INA228	±0.01-0.05%	-0.3-85 V	Shunt Voltage, Battery Voltage, Die Temperature, Current, Power, Accumulated Energy, Accumulated Charge	\$5.27

Table 3.15 Energy Monitoring Part Comparison

Sources: [48]-[50]

3.1.12.1 LTC2944

The LTC2944 monitors the battery's energy by measuring the differential voltage across a high-side sense resistor located between the battery's positive terminal and the common junction of the solar charger and loads [48]. Its bidirectional analog integrator is specifically designed to accommodate current polarity; it "counts up" as the solar panels provide a charging current and "counts down" as the system draws power for the loads. This "true" coulomb counting occurs in real-time with 1% accuracy, allowing the chip to maintain a continuous record of the accumulated battery charge regardless of whether the system is net-charging or net-discharging.

3.1.12.2 INA226

The INA226 features an industry-standard 16-bit resolution and its ability to handle bus voltages up to 36V [49]. In a solar charging scenario, the INA226 monitors the current flow bidirectionally. When the solar panels are generating excess power, the sensor detects a positive current flowing into the battery, and when the loads exceed solar production, it detects a negative current. Because it features a dedicated alert pin, it can be programmed to interrupt the microcontroller if the battery voltage hits a critical low or if the solar charging current exceeds the battery's 20A continuous charging limit, providing an essential hardware-level safety layer.

3.1.12.3 INA228

The INA228 represents a significant technological leap, utilizing a 20-bit delta-sigma ADC to provide extreme precision for long-term energy tracking [50]. In solar applications where light intensity, and therefore charging current can fluctuate minutely, the higher resolution of the INA228 ensures that even the smallest energy gains from a solar panel

during low-light conditions are captured without falling into the "noise" of the sensor. Furthermore, the INA228 includes an on-chip energy and charge accumulator. Unlike the INA226, which requires the microcontroller to constantly poll and integrate current data to calculate Amp-hours, the INA228 can perform these calculations internally. This reduces the processing overhead required of the *SOLARIS* MCU.

3.1.12.4 Energy Monitoring Part Selection

INA228: The INA228 was selected as the primary energy monitor over the INA226 and LTC2944 due to its combination of high-resolution data acquisition and self-contained computational features. While the INA226 is a capable industry standard, its 16-bit ADC is surpassed by the INA228's 20-bit delta-sigma ADC, which provides the superior precision required to detect minute fluctuations in solar charging current that might otherwise be lost to sensor noise. Furthermore, the INA228 offers a significant architectural advantage by integrating dedicated hardware for energy and charge accumulation. Unlike the INA226, which requires the microcontroller to constantly poll current values to manually calculate Amp-hours, the INA228 performs these complex integrations internally on-chip. This not only reduces the processing overhead and power consumption of the microcontroller but also eliminates the risk of "integration drift" caused by timing variances in software-based polling. By offloading these calculations to the sensor, the INA228 provides a more stable and accurate State of Charge (SoC) tracking system, making it the most robust choice for managing a 12.8V LiFePO₄ battery system in a variable solar-charging environment.

3.1.13 Inertial Measurement Unit (IMU)

<i>Specifications</i>	<i>Accelerometer</i>	<i>Gyroscope</i>	<i>Magnetometer</i>
What it Measures	Linear acceleration & gravity	Angular velocity (rotation)	Magnetic field strength
Primary Unit	$m/s^2 = gravity$	deg/s (degrees per second)	μT (microteslas)
What it Detects	Which way is down	How fast it's spinning	Which way is north
Strength	Great at sensing tilt & impact	Very precise for fast motion	Provides an absolute heading
Weakness	Noisy due to vibrations	Drifts over time	Easily distorted by metal/magnets

Table 3.16 Types of Measurement Devices

Sources: [51], [52]

3.1.13.1 Accelerometer

An accelerometer is a key piece of hardware used to measure the inertial movement of devices like cell phones and even vehicles like cars. While a measurement like speed tells you how fast something is going, acceleration tells you how quickly that speed is changing. An accelerometer detects three things. It measures the inclination of mostly stationary

objects by detecting the acceleration of gravity [53]. It also detects the acceleration of an object when it is vibrating [53]. Lastly, it detects the acceleration of an object when it experiences an impact [53]. To accomplish this, the accelerometer does not calculate speed over time but instead measures force using an elementary equation.

$$F = m \times a$$

$$a = \frac{F}{m}$$

Though this equation covers the most basic kinds of accelerometers, there exist different "flavors" of sensors used for measuring acceleration. One such type is piezoelectric, which relies on piezoelectric material like a quartz crystal to sense changes in acceleration [54]. When acceleration is applied, the material responds with a compression or strain relative to its mass. Another technology, known as piezoresistive, operates in a similar fashion but measures changes in resistance rather than electrical charge. Manufactured as a semiconductor, each axis includes multiple piezo resistors that decrease in resistance when force is applied [54]. The last accelerometer technology worth mentioning for *SOLARIS* is capacitive. Like piezoresistive, it is generally a semiconductor sensing device, but the key difference is a tether system that collectively acts like a spring, allowing the displacement of the mass to represent the measured acceleration [54]. Other accelerometer types are more specialized in their use cases and therefore are not relevant to *SOLARIS* as a potential technology selection.

3.1.13.2 Gyroscope

A gyroscope (gyro) is a microelectromechanical system (MEMS) sensor designed to measure angular velocity [55]. In a traditional sense, gyroscopes are generally heavy spinning wheels used to maintain orientation. The MEMS application used in electronics instead relies on tiny vibrating structures that shift when the device is rotated. While an accelerometer measures linear force, a gyroscope detects angular velocity, which can be defined as the rate at which an object rotates around an axis [55]. Within MEMS gyroscopes there are two main types: vibrating structure and the tuning fork. Vibrating structure gyroscopes utilize the Coriolis force, which causes a secondary oscillation that can be measured to determine angular velocity [56]. The tuning fork configuration operates similarly but uses a pair of vibrating masses. When the Coriolis force acts on the tines, it can be measured in the same way to determine angular velocity [56].

$$\omega = \frac{V_{out} - V_{bias}}{S}$$

A gyroscope can detect three major movements: pitch, roll, and yaw. Using the equation shown above, the gyroscope receives a bias voltage (V_{bias}) from its power source and, given its own sensitivity (S), calculates angular velocity against its output voltage (V_{out}). This equation is generalized since MEMS gyroscopes can come in different configurations. Some of the main advantages of MEMS gyroscopes are their small size, low cost, low power consumption, and reliability [56]. These qualities make them useful across a wide

range of application areas including consumer electronics, automotive, defense, and healthcare [56].

3.1.13.3 Magnetometer

Just like an accelerometer and a gyroscope, a magnetometer is another way for a digital device to interface with an analog world. A magnetometer measures the magnetic fields present in a surrounding environment [52]. Measuring magnetic fields is not only useful for *SOLARIS* but is also crucial for research, navigation, and exploration. A magnetometer uses a sensor capable of measuring either magnetic flux density (B) or magnetic field strength (H) [52]. Like the sensors discussed previously, these devices also come in different configurations but operate on the same mathematical principles shown below.

$$V_H = \frac{IB}{nte}$$

The equation above references the most common type of magnetometer, the Hall Effect. When a current-carrying conductor is placed in a magnetic field, a voltage is generated perpendicular to both the direction of the field and the current flow [52]. In this equation, (I) represents current, (B) represents magnetic flux density, (n) represents the charge carrier density of the material, (t) represents the thickness of the plate, and (e) represents the elementary charge of an electron.

$$|M| = \sqrt{M_x^2 + M_y^2 + M_z^2}$$

$$\psi = \arctan^2(M_y, M_x)$$

The two equations above define the mathematical principles behind another popular magnetometer technology known as vector based. The first equation combines three directional readings to calculate total magnetic intensity using a three-dimensional Pythagorean theorem. The second equation converts these raw magnetic readings into a compass heading using trigonometry. Together, these equations mean that vector magnetometers measure magnetic flux density (B) in a specific direction within three-dimensional space [52]. This makes the magnetometer particularly valuable for its ability to provide a true heading for spatial awareness.

3.1.13.4 Degree of Axis




Figure 3.15 Degrees of Freedom Illustrated

In the grand scheme, the three major components that make up an inertial measurement unit (IMU) are the accelerometer, gyroscope, and magnetometer. On their own, each of these sensors measures three degrees of movement. An IMU is categorized by its Degrees of Freedom (DoF), which represents how many ways it can sense motion [57]. The DoF of a device expresses how many ways something can move in three-dimensional space [57]. The accelerometer is responsible for three degrees of translational movement, the gyroscope measures three degrees of rotational movement, and the magnetometer measures three degrees of direction. When referring to the DoF of a particular IMU, we are referring to how many of these sensors it contains. To obtain usable orientation information from an IMU, a process called sensor fusion is used to combine the raw readings from each sensor, which are otherwise output in different units and scales, into an accurate representation of an object's orientation in three-dimensional space [57].

<i>DoF</i>	<i>Sensors Included</i>	<i>What it Measures</i>	<i>Best For</i>
3-DoF	Accelerometer	Linear acceleration, static tilt	Screen rotation, step counting, impact sensing
6-DoF	Accelerometer + Gyroscope	Translation and rotation	VR controllers, gamepads, balancing robots
9-DoF	Accelerometer + Gyroscope + Magnetometer	Translation, rotation, and magnetic heading	Drones, autonomous vehicles, 3D mapping.

Table 3.17 Degrees of Freedom (DoF) Comparison

Sources: [57]

3.1.13.4.1 3-DoF

A 3-DoF IMU can refer to any singular sensor on its own but generally refers to an accelerometer. This means it can only sense translational movement along the X, Y, and Z axes [57]. However, it cannot detect yaw since rotating the sensor around a vertical axis does not change the force of gravity acting on it.

3.1.13.4.2 6-DoF

A 6-DoF IMU generally refers to a 3-axis accelerometer paired with a 3-axis gyroscope. With this combination, the device can sense both linear movement and angular velocity [57]. It is called 6-DoF since it combines the three axes of the accelerometer with the three axes of the gyroscope. The key advantage of this setup is the ability to calculate orientation while in motion, where the gyroscope tracks real-time rotation and the accelerometer provides a reference point.

3.1.13.4.3 9-DoF

A 9-DoF IMU adds a 3-axis magnetometer to the previous 6-DoF setup, allowing the device to sense translational movement, rotation, and direction simultaneously. What makes this configuration particularly powerful is that sensor fusion across all three sensors enables a form of self-correction within the system [57]. By sensing the Earth's magnetic field, the device has a fixed reference point for the horizontal plane, allowing for accurate heading calculation [57].

3.1.13.5 IMU Comparison

<i>Specifications</i>	<i>BNO-085</i>	<i>MPU-6050</i>	<i>ICM-20948</i>
Degrees of Freedom	9-DoF	6-DoF	9-DoF
Onboard Processing	Full Sensor Fusion (ARM Cortex-M0)	Basic Digital Motion Processor (DMP)	Advanced DMP (to offload calculations)
Power Profile	Low Power Mode	Sleep Register	Low Power Mode
Stability	High (Absolute Orientation)	Medium (No Magnetometer)	Medium-High (Relative Compass)
Use Case	Autonomous navigation and high-accuracy robotics	Budget-friendly tilt sensing and simple motion detection	Low-power mobile robots and high-speed telemetry

Table 3.18 IMU Unit Comparison

Sources: [58-60]

3.1.13.5.1 BNO-085

The BNO085 is a specialized 9-axis IMU that combines the features of a classical IMU while also functioning as a System in Package (SiP). Unlike typical IMUs that output raw data for the microcontroller to process, the BNO085 features an onboard 32-bit ARM Cortex-M0+ microcontroller that handles all sensor fusion internally [58]. Another key highlight is its ability to self-calibrate, continuously estimating and compensating for magnetometer distortions and gyroscope zero-rate offset. It also supports three communication interfaces: I2C, SPI, and UART. I2C and SPI can be used to directly configure the device, enabling pre-implemented sleep modes for low power applications and allowing individual sensors like the accelerometer, gyroscope, or magnetometer to be toggled on or off as needed [58]. This means any combination of sensors can be used interchangeably or together for maximum sensor fusion accuracy. For an application where power efficiency and data accuracy are critical, the BNO085 offloads sensor processing from the ESP32-S3, freeing it to focus on controlling the rest of the system.

3.1.13.5.2 MPU-6050

The MPU-6050 is a 6-axis sensor hub that outputs raw motion tracking data and is typically more affordable than other selections, making it a common choice for projects where the microcontroller handles all sensor fusion processing. It is marketed as one of the world's first devices to combine a 3-axis gyroscope and a 3-axis accelerometer with an onboard Digital Motion Processor (DMP) [59]. It uses six dedicated 16-bit ADCs to sample all axes simultaneously, ensuring no time skews between the two peripherals. One key feature is its auxiliary I2C master port, which allows it to act as a controller for external sensors [59]. This means it can be paired with a third-party magnetometer to automatically read compass data and fuse it with its own output for full heading calculation. The main downsides compared to other options are its lack of native magnetometer support and onboard

processing, though its low price point and strong community support due to its legacy status remain notable advantages.

3.1.13.5.3 ICM-20948

The ICM-20948 is an advanced 9-axis motion tracking device designed as a successor to legacy hardware [60]. It is specifically built for low power consumption and high-performance sensor fusion, making it a strong candidate for *SOLARIS*. The device can consume as little as 2.5mW in full 9-axis mode, making it one of the lowest power 9-axis IMUs available [60]. It also supports a native sleep mode where it does not draw power upon startup until it is needed. By combining a 3-axis gyroscope, 3-axis accelerometer, and 3-axis compass, it can provide a full absolute heading to magnetic north. Like the other options discussed, its Digital Motion Processor (DMP) supports high-frequency data transfers of up to 200Hz, offloading motion processing from the ESP32-S3 [60]. The DMP reliably fuses data from all nine axes into stable output, which can reduce both code complexity and the likelihood of unreliable readings. Like the MPU-6050, the ICM-20948 features an auxiliary I2C master port and can act as a controller for other sensors [60]. Onboard firmware also supports runtime calibration, helping maintain accuracy as environmental conditions and temperature change throughout the day.

3.1.13.6 IMU Selection

For *SOLARIS*, the BNO085 was selected as the IMU of choice due to its reliability, power efficiency, and software simplicity. While the MPU-6050 and ICM-20948 are excellent options for raw data output, the BNO085 functions as a sensor hub that simplifies autonomous navigation. By offloading sensor fusion math onto its onboard ARM Cortex-M0+, the ESP32-S3 is free to handle other critical system tasks while the BNO085 runs mathematically intensive fusion algorithms in the background [58]. Its dynamic calibration algorithms continuously monitor the magnetic environment and filter out distortions in real time, which is particularly important for keeping the vehicle on course as it traverses uneven terrain. The chip is designed with an always-on feature but excels in low-power states where proper configuration is key. With its rich feature set and zero processing load on the main controller, the BNO085 is a natural fit for a solar tracking robot. This selection reflects *SOLARIS*'s priority for data integrity and system stability, ensuring the vehicle can accurately navigate toward the sun or unshaded regions.

3.1.14 Ultrasonic Sensors

3.1.14.1 How Ultrasonic Sensors Work

[61] Ultrasonic sensors measure distance by transmitting sound waves and detecting their reflections. A typical ultrasonic sensor module includes four primary connections: VCC and GND for power, a Trigger input, and an Echo output. When a short pulse is applied to the Trigger pin, the sensor emits a burst of sound waves from its transmitter. These waves travel outward in a cone-shaped pattern, commonly referred to as the beam cone, which defines the sensor's field of view. Any object located within this cone can reflect the sound waves back toward the sensor.

Once the reflected wave returns, the receiver detects it and the sensor drives the Echo pin high for a duration equal to the time it took for the sound to travel to the object and back. A microcontroller measures this time interval and calculates the distance by using the speed

of sound and dividing the total travel time by two, since the signal makes a round trip. The width of the beam cone is important because it determines the coverage area and influences how accurately objects can be localized in front of the sensor.

For *SOLARIS*, ultrasonic will be used for object avoidance. By continuously sending trigger pulses and monitoring the returned echoes, *SOLARIS* can determine how close it is to obstacles. If an object is detected within a predefined threshold distance, the control system can respond by slowing down, stopping, or changing direction to prevent collisions.

3.1.14.2 Calculation for Beam Cone Coverage




Figure 3.16 Two Ultrasonic Sensor Beam Cone Coverage

For *SOLARIS*, we plan to use four ultrasonic sensors mounted at the corners of the robot, two facing forward and two facing backward. When selecting an ultrasonic sensor, the beam cone coverage is therefore an important consideration. As shown in Figure 3.1.14.1, there are three primary regions that *SOLARIS* will be unable to detect: the far left side, the far right side, and a triangular region directly in front of the robot. While the beam cone plays a significant role in defining these blind areas, sensor placement also influences coverage. The sensors can be mounted with a slight angular offset relative to the chassis. Although not explicitly labeled in the figure, Figure 23 was created using a 10-degree offset, with the two sensors angled 10 degrees away from each other. The benefit of introducing this offset is that it reduces the undetectable angle on the far left and right sides by 10 degrees. The smaller this angle becomes, the greater the turning flexibility *SOLARIS* can be given. However, the tradeoff is that the maximum blind distance in the central region increases.




Figure 3.17 One Ultrasonic Sensor Flush with X-Axis

As already established, the two values we are looking for are the maximum blind distance in the middle region of *SOLARIS*'s blind spot and the angle at which *SOLARIS* will be unable to see on its left and right sides, which we will call α . Figure 3.16 can be used to determine this angle. Using angle-sum rules, we know the three angles shown in the figure must add up to 180. Since the beam cone is centered, the angles on both sides are equal. In the case where the offset is 0:

$$\alpha = \frac{(180-\beta)}{2}$$

If we add an offset angle to the left sensor in this case, the angle on the left side increases by that amount, while the angle on the right-side decreases by the same amount. Therefore, when we account for an offset, the angle at which *SOLARIS* will be unable to see becomes:

$$\alpha = \frac{(180 - \beta)}{2} - offset$$

Now we can create an equation to determine the maximum blind distance in the middle of the robot. Looking back at Figure 3.14, the central blind-spot region can be split into two right triangles. The angle formed on the inside of the beam cones (not at the intersection) will be called δ . This angle can be solved in the same way as α , but in this case it increases with the offset. Therefore, the equation becomes:

$$\delta = \frac{(180 - \beta)}{2} + offset$$

Before continuing, we should also define the distance between the sensors, which we will call d . We expect this distance to be 2 feet, but for the purpose of deriving the equation, we will keep it as d in case the spacing changes later. We now have everything needed to form an equation for the maximum blind-spot distance. Since the interior region forms a right triangle and we know the relevant angle while the distance between the sensors is known, we can use the following equation:

$$\tan(\delta) = \frac{MaxBlindDistance}{(d/2)}$$

Thus max blind distance is:

$$MaxBlindDistance = \frac{2}{d} * \tan(\delta)$$




Figure 3.18 Ultrasonic Sensor Beam Vs. Ground

The last consideration is that the beam cone can reach the ground and create a false reading, making *SOLARIS* think there is an obstacle ahead when it is actually detecting the ground. This can be mitigated by adding a vertical (y-axis) tilt so the beam cone does not intersect the ground until a chosen distance. A reasonable target is around 5 to 6 meters (16 to 20 feet), since many ultrasonic sensors reach their effective limit around that range anyway. An example of this is shown in Figure 3.17, where the sensor is tilted upward so the lower edge of the cone does not hit the ground until approximately the desired distance.

One more consideration is how high off the ground the sensor should be mounted. For the figure, we chose a height of about 4 inches (or 1/3 of a foot). This height was selected because the assumption is that objects below that level are likely small enough for *SOLARIS* to drive over without issue, such as a rock, a slight incline, or a branch. This choice will likely be adjusted during testing as we determine what works best in real conditions. For now, however, assuming a sensor height of 4 inches above the ground is a reasonable starting point.




Figure 3.19 Ultrasonic Sensor Angle Pointing Towards the Ground

We now have everything needed to determine what offset we should use, given that our sensor height is 1/3 of a foot and a selected beam angle. Although it is difficult to see in the figures, the portion of the beam closest to the ground forms a right triangle with the ground surface. Like how α was derived using Figure 3.18, we can determine the angle the beam makes within this right triangle. This angle is shown as already derived in Figure 3.17, and we will call it θ . θ is:

$$\theta = \frac{(180 - \beta)}{2} + offset$$

We can again use the properties of right triangles to develop an equation that solves the desired offset. Let the target distance at which the beam first touches the ground be 5 meters, or 16.4 feet.

$$\tan(\theta) = \frac{16.4}{(1/3)}$$

$$\tan\left(\frac{(180 - \beta)}{2} + offset\right) = \frac{16.4}{(1/3)}$$

$$\frac{(180 - \beta)}{2} + offset = \arctan\left(\frac{16.4}{(1/3)}\right)$$

$$offset = \arctan\left(\frac{16.4}{(1/3)}\right) - \frac{(180 - \beta)}{2}$$

Now that we have finished deriving the calculations we need, we can compare our constraints and determine how the sensors should be positioned, including whether offsets

are necessary, for ultrasonic sensors with different beam angles. This comparison will help us narrow down which sensor is the best fit for *SOLARIS*.

	$x_offset = 0^\circ$	$x_offset = 5^\circ$	$x_offset = 10^\circ$	$x_offset = 15^\circ$	$x_offset = 20^\circ$
$\beta = 15^\circ$	$\alpha = 82.5^\circ$ MaxBlindDist = 7.6' $y_offset = 6.3^\circ$	$\alpha = 77.5^\circ$ MaxBlindDist = 22.9' $y_offset = 6.3^\circ$	$\alpha = 72.5^\circ$ MaxBlindDist = und. $y_offset = 6.3^\circ$	$\alpha = 67.5^\circ$ MaxBlindDist = und. $y_offset = 6.3^\circ$	$\alpha = 62.5^\circ$ MaxBlindDist = und. $y_offset = 6.3^\circ$
$\beta = 20^\circ$	$\alpha = 80^\circ$ MaxBlindDist = 5.7' $y_offset = 8.8^\circ$	$\alpha = 75^\circ$ MaxBlindDist = 11.4' $y_offset = 8.8^\circ$	$\alpha = 70^\circ$ MaxBlindDist = und. $y_offset = 8.8^\circ$	$\alpha = 65^\circ$ MaxBlindDist = und. $y_offset = 8.8^\circ$	$\alpha = 60^\circ$ MaxBlindDist = und. $y_offset = 8.8^\circ$
$\beta = 25^\circ$	$\alpha = 77.5^\circ$ MaxBlindDist = 4.5' $y_offset = 11.3^\circ$	$\alpha = 72.5^\circ$ MaxBlindDist = 7.6' $y_offset = 11.3^\circ$	$\alpha = 67.5^\circ$ MaxBlindDist = 22.9' $y_offset = 11.3^\circ$	$\alpha = 62.5^\circ$ MaxBlindDist = und. $y_offset = 11.3^\circ$	$\alpha = 57.5^\circ$ MaxBlindDist = und. $y_offset = 11.3^\circ$
$\beta = 30^\circ$	$\alpha = 75^\circ$ MaxBlindDist = 3.7' $y_offset = 13.8^\circ$	$\alpha = 70^\circ$ MaxBlindDist = 5.7' $y_offset = 13.8^\circ$	$\alpha = 65^\circ$ MaxBlindDist = 11.4' $y_offset = 13.8^\circ$	$\alpha = 60^\circ$ MaxBlindDist = und. $y_offset = 13.8^\circ$	$\alpha = 55^\circ$ MaxBlindDist = und. $y_offset = 13.8^\circ$
$\beta = 35^\circ$	$\alpha = 72.5^\circ$ MaxBlindDist = 3.2' $y_offset = 16.3^\circ$	$\alpha = 67.5^\circ$ MaxBlindDist = 4.5' $y_offset = 16.3^\circ$	$\alpha = 62.5^\circ$ MaxBlindDist = 7.6' $y_offset = 16.3^\circ$	$\alpha = 57.5^\circ$ MaxBlindDist = 22.9' $y_offset = 16.3^\circ$	$\alpha = 52.5^\circ$ MaxBlindDist = und. $y_offset = 16.3^\circ$
$\beta = 40^\circ$	$\alpha = 70^\circ$ MaxBlindDist = 2.7 $y_offset = 18.8^\circ$	$\alpha = 65^\circ$ MaxBlindDist = 3.7' $y_offset = 18.8^\circ$	$\alpha = 60^\circ$ MaxBlindDist = 5.7' $y_offset = 18.8^\circ$	$\alpha = 55^\circ$ MaxBlindDist = 11.4' $y_offset = 18.8^\circ$	$\alpha = 50^\circ$ MaxBlindDist = und. $y_offset = 18.8^\circ$
$\beta = 45^\circ$	$\alpha = 67.5^\circ$ MaxBlindDist = 2.4' $y_offset = 21.3^\circ$	$\alpha = 62.5^\circ$ MaxBlindDist = 3.1' $y_offset = 21.3^\circ$	$\alpha = 57.5^\circ$ MaxBlindDist = 4.5' $y_offset = 21.3^\circ$	$\alpha = 52.5^\circ$ MaxBlindDist = 7.6' $y_offset = 21.3^\circ$	$\alpha = 47.5^\circ$ MaxBlindDist = 22.9' $y_offset = 21.3^\circ$
$\beta = 60^\circ$	$\alpha = 60^\circ$ MaxBlindDist = 1.7' $y_offset = 28.8^\circ$	$\alpha = 55^\circ$ MaxBlindDist = 2.1' $y_offset = 28.8^\circ$	$\alpha = 50^\circ$ MaxBlindDist = 2.7' $y_offset = 28.8^\circ$	$\alpha = 35^\circ$ MaxBlindDist = 3.7' $y_offset = 28.8^\circ$	$\alpha = 40^\circ$ MaxBlindDist = 5.67' $y_offset = 28.8^\circ$
$\beta = 75^\circ$	$\alpha = 52.5^\circ$ MaxBlindDist = 1.3' $y_offset = 36.3^\circ$	$\alpha = 47.5^\circ$ MaxBlindDist = 1.5' $y_offset = 36.3^\circ$	$\alpha = 42.5^\circ$ MaxBlindDist = 1.9' $y_offset = 36.3^\circ$	$\alpha = 37.5^\circ$ MaxBlindDist = 2.4' $y_offset = 36.3^\circ$	$\alpha = 32.5^\circ$ MaxBlindDist = 3.17' $y_offset = 36.3^\circ$

Table 3.19 Beam Angle and Offset Calculation Comparison

A beam angle of 75° may initially appear to be the best option since it produces the smallest α and minimizes the maximum blind distance. However, a wider beam is not always the most practical choice. As beam angle increases, range and precision typically decrease, and the likelihood of false readings rises due to reflections from unintended objects within the sensing cone.

On the other hand, a very narrow beam provides improved range, higher precision, and fewer false detections. The tradeoff is reduced environmental awareness, as the sensor covers a smaller portion of the robot's surroundings. If the beam angle is too small, the robot may lack sufficient spatial information for reliable obstacle detection.

For *SOLARIS*, a balanced approach is preferred. The design goal is to keep the maximum blind distance near five feet while ensuring that α remains less than or equal to 60°. Based on this analysis, the most effective compromise lies in the mid-range beam angles. The selected configuration is an ultrasonic sensor with a beam angle between 40° and 45°, paired with an x_{offset} of 10°. This combination provides a strong balance between coverage, precision, and detection reliability, making it well-suited for the system's design requirements.

3.1.14.3 Ultrasonic Sensor Comparisons

	<i>Operating Voltage</i>	<i>Typical Current Draw</i>	<i>Beam Angle</i>	<i>Range</i>	<i>Minimum Time Between Readings</i>	<i>Communication Interface</i>	<i>Cost per unit</i>
MB1020-000	2.5-5.5 V	2mA	42°	.15-6.5m	50ms	UART + ADC	\$19.95
MB1232-000	3-5.5 V	Not Specified	44°	.2-7.7m	25ms	I2C	\$37.44
MB1013-000	2.5-5.5 V	2.5-3.1mA	44°	.3-5m	100ms	UART + ADC	\$31.19

Table 3.20 Ultrasonic Comparison Table

Sources: [62]-[64]

3.1.14.3.1 Note on Limited Selection

For *SOLARIS*, the ultrasonic sensor requirements of a 3–5 V operating voltage and a 40–45° beam angle significantly narrowed the available options. Many sensors that met these specifications were priced at over \$100 per unit, which is not practical for a four-sensor configuration. After evaluating the available market options, the only reasonably priced sensors that satisfied both the electrical and beam angle constraints were manufactured by MaxBotix®.

MaxBotix offers multiple sensor models, many of which belong to the same product families with similar electrical characteristics and performance specifications. While the

company sells additional variants, several were either out of stock or did not align with our exact design requirements. Therefore, the sensors included in the comparison table were selected to represent their respective families while also being currently available for purchase. This approach ensures that the comparison remains both technically relevant and practically feasible for implementation in *SOLARIS*.

3.1.14.3.2 MB1020-000

One ultrasonic sensor under consideration for *SOLARIS* is the MB1020-000. This model operates within the required 3–5 V range and provides a beam angle that aligns with our 40–45° design constraint, making it electrically and mechanically compatible with the system architecture.

The sensor offers sufficient range for *SOLARIS*' expected operating speed and obstacle detection requirements. Its current drawing is lower than many other commercially available sensors that meet similar beam and voltage specifications, which supports the platform's low-power, solar-driven design goals. Cost is also competitive, which is important given that multiple units will be required.




Figure 3.20 MB1020-000 Sensor Chaining Diagram

Source: [62]

An additional advantage is the documented sensor chaining capability. The ability to configure sensors to trigger sequentially helps mitigate crosstalk without requiring complex timing logic in firmware. This simplifies integration in a multi-sensor configuration. Potential drawbacks include a lower maximum range and a slower update rate compared to some other ultrasonic sensors.

3.1.14.3.3 MB1232-000

Another ultrasonic sensor under consideration for *SOLARIS* is the MB1232-000. This model operates within the required 3–5 V range and is electrically compatible with the system architecture. Among the sensors being compared, the MB1232-000 provides the highest maximum range and the fastest time between readings. This could allow for more frequent environmental updates and improved responsiveness in autonomous navigation. The sensor uses an I2C interface. One consideration to make is that using multiple identical I2C sensors introduces potential addressing limitations. If the devices share a fixed address or limited configurable address range, additional hardware solutions such as I2C

multiplexers may be required. This adds design complexity compared to sensors that can be independently triggered and read without bus conflicts.

Cost is another significant factor. The MB1232-000 is the most expensive option among the sensors being evaluated. Because *SOLARIS* requires multiple units, this increases total system cost substantially. The primary tradeoff with this sensor is improved range and update speed versus higher cost and potential integration complexity related to I2C multi-device communication. Final suitability will depend on whether the performance benefits justify the added expense and system complexity.

3.1.14.3.4 MB1013-000

The MB1013-000 is also being evaluated as a potential ultrasonic sensor for *SOLARIS*. It meets the required 3–5 V operating range, making it compatible with the existing power design. In comparison to the other sensors under review, this model has the shortest maximum range and the longest time between readings. This could reduce system responsiveness in situations where faster environmental updates are needed. Its current draw is also higher than the MB1020-000, which is a relevant consideration for a solar-powered platform where minimizing power consumption is important.

The pin configuration differs from the other candidates, which may require adjustments in PCB layout or firmware handling. However, like the MB1020-000, it supports sensor chaining, allowing multiple units to be triggered sequentially to reduce crosstalk without complex timing logic.

From a cost standpoint, the MB1013-000 is less expensive than the MB1232-000 but still significantly more expensive than the MB1020-000. Since multiple sensors are required for *SOLARIS*, this cost difference scales at the system level. Overall, this model presents a tradeoff between cost, power consumption, range, and updating speed when compared to the other options being considered.

3.1.14.4 Ultrasonic Sensor Selection

MB1020-000: The MB1020-000 was selected as the ultrasonic sensor for *SOLARIS*. Its support for sensor chaining, lowest current draw among the evaluated options, and significantly lower cost made it the most practical and well-balanced choice for a multi-sensor implementation.

While its maximum range is lower than that of the MB1232-000, it comfortably exceeds our design requirement of a minimum 5 m maximum detection distance. The reduced range does not limit system performance within the intended operating environment.

The MB1020-000 also has a slower minimum time between readings compared to the MB1232-000. For high-speed robotic systems, this could be a critical limitation due to increased latency. However, *SOLARIS* is designed as a slow-moving platform, meaning the additional latency does not meaningfully impact navigation performance.

When balancing performance, power consumption, cost, and system complexity, the MB1020-000 provides the strongest overall fit for *SOLARIS*.

3.1.15 Chassis Technology Comparison

<i>Chassis Technology</i>	<i>Strength-to-Weight</i>	<i>Pros</i>	<i>Cons</i>	<i>Durability</i>	<i>Cost</i>
3D Printed Filament	Moderate	Fast iteration, Extremely low cost, Complex geometries easily achievable, Lightweight	Weak in structural applications, Heat sensitivity	Poor(UV+ heat degrade)	Low
Aluminum	High	Strong and predictable behavior, Excellent for outdoor robotics, Ductile	Heavier than composites, Requires tools	Excellent	Moderate
Carbon Fiber	Very high	Best strength-to-weight ratio, Extremely stiff, Professional-grade performance	Expensive, Brittle failure, Hard to machine safely, Difficult to repair	Excellent	High

Table 3.21 Chassis Technology Comparison

Sources: [65]-[67]

The mechanical design of *SOLARIS* requires a structured approach to selecting an appropriate chassis fabrication method to house all system components. A key parameter in this trade study is strength-to-weight ratio, as higher values enable a structurally robust design while minimizing overall mass. Reducing system weight is critical for improving energy efficiency and extending operational runtime.

Another important consideration is outdoor durability. Since *SOLARIS* is intended for outdoor operation, materials must withstand environmental exposure, particularly UV radiation, which can degrade certain polymers over time and lead to premature failure. Therefore, material selection must account for long-term reliability in sun-exposed conditions.

Finally, cost is a significant factor, especially during the prototyping phase. Multiple design iterations are expected, and material and manufacturing costs can accumulate quickly. Selecting cost-effective materials and fabrication methods allows for rapid iteration without excessive resource expenditure, enabling a more efficient development cycle.

3.1.15.1 3D Printed Filament

3D printed filament is well-suited for rapid prototyping of a robot chassis due to its low cost, fast turnaround, and ability to produce complex geometries without specialized tooling [65]. It enables tight integration of features such as mounting bosses, wire routing channels, and sensor housings directly into the structure. However, its structural

performance is limited by relatively low strength and stiffness, as well as anisotropic properties caused by layer-by-layer fabrication. Additionally, common filaments are susceptible to UV degradation and thermal softening, which can compromise reliability in outdoor environments. As a result, filament-based construction is best used for non-load-bearing components.

3.1.15.2 Aluminum

Aluminum is a highly practical and widely used material for small robot chassis due to its strong balance of strength, weight, manufacturability, and durability [66]. For a robot chassis, aluminum sheet or extrusion can provide a rigid and robust frame capable of supporting motors, batteries, and electronics without significant deflection. It is well-suited for outdoor use, offering good corrosion resistance and stable mechanical properties across a wide temperature range. Fabrication requires basic machining processes such as cutting, drilling, and fastening, which are accessible but more involved than 3D printing. Overall, aluminum offers a reliable and repeatable solution for a functional, field-ready chassis.

3.1.15.3 Carbon Fiber

Carbon fiber provides an exceptional strength-to-weight ratio and very high stiffness, making it attractive for applications where minimizing mass is critical [67]. In a small robot chassis, carbon fiber plates or tubes can produce a lightweight yet rigid structure that enhances dynamic performance and efficiency. However, the material introduces challenges in fabrication, including difficulty in machining, the need for specialized tools, and safety concerns related to carbon fiber dust. Additionally, carbon fiber exhibits a brittle failure mode and is more difficult to repair compared to metals. Its higher cost and manufacturing complexity typically limit its use to performance-critical designs rather than general-purpose prototyping.

3.1.15.4 Chassis Technology Selection

3D Printed Filament and Aluminum: Aluminum was selected as the primary material for load-bearing and structural components due to its strong balance of mechanical performance and reliability. For a robot of this scale, the chassis must support the weight of batteries, electronics, and drive systems while maintaining rigidity under dynamic loads such as movement over uneven terrain. Aluminum provides high strength and stiffness with predictable, ductile behavior, meaning it will deform rather than fail catastrophically under stress. It also offers excellent durability for outdoor environments, with good resistance to corrosion and temperature variation. These characteristics make aluminum a dependable choice for the core structural frame where integrity and long-term performance are critical.

In addition to aluminum, 3D printed filament is incorporated to enable rapid iteration and customization of non-load-bearing components. Additive manufacturing allows for quick design changes and fabrication of complex geometries such as mounts, enclosures, and cable management features without the need for specialized tooling. This is especially valuable during development, where multiple design revisions are expected. Filament-based parts are also lightweight, which helps minimize the overall system mass without compromising the primary structural integrity provided by the aluminum frame. By combining aluminum for strength and 3D printing for flexibility, the design achieves both robustness and adaptability.

3.1.16 Chassis Part Comparison

<i>Chassis Part</i>	<i>Density (g/cm³)</i>	<i>Yield/Tensile Strength (MPa)</i>	<i>UV Resistance</i>	<i>Cost</i>
ASA Filament	1.07	~45	Excellent	Low
PETG Filament	1.27	~50	Moderate	Low
ABS Filament	1.04	~40	Poor	Moderate
PC Filament	1.20	~65	Moderate	High
5052-H32 Aluminum	2.68	~193	Excellent	Low-Moderate
6061-T6 Aluminum	2.70	~255	Excellent	Moderate
7075-T6 Aluminum	2.81	~450	Excellent	High

Table 3.22 Chassis Part Comparison

Sources: [65],[68]-[70]

The next step is selecting the specific 3D printing filament and aluminum alloy for use in *SOLARIS*. The 3D printed material will primarily serve in non-load-bearing roles such as component mounting, enclosures, and wire routing, where geometric flexibility and rapid fabrication are advantageous. In contrast, aluminum will be used for the primary chassis framework and base plates, where structural integrity and stiffness are critical to supporting system loads and maintaining alignment.

Several key material properties are considered in this selection. Density provides a measure of mass per unit volume and directly impacts the overall weight of the robot, which influences energy efficiency and mobility. Tensile strength refers to the maximum tensile (pulling) stress a material can withstand before failure, typically measured in units of pressure such as MPa. It characterizes how well a material resists being stretched or pulled apart under load, which is especially important for structural components subjected to forces during operation. Lastly, UV resistance is a critical factor for outdoor deployment, as prolonged exposure to sunlight can degrade certain materials (especially polymers) leading to embrittlement, discoloration, and reduced mechanical performance over time.

3.1.16.1 ASA Filament

ASA (Acrylonitrile Styrene Acrylate) is specifically engineered for high-performance outdoor applications due to its exceptional UV stability and weather resistance. Unlike its relative ABS, ASA does not yellow or become brittle when exposed to prolonged sunlight, making it the ideal 3D printing filament for solar-integrated projects like *SOLARIS*. It offers high mechanical strength and a high glass transition temperature of approximately 95–100°C, ensuring that external sensor mounts and solar panel clips maintain their structural integrity even in intense thermal environments.

3.1.16.2 PETG Filament

PETG (Polyethylene Terephthalate Glycol) serves as a versatile "all-rounder" filament that combines the ease of printing found in PLA with the functional durability of ABS [65]. It is characterized by its excellent layer adhesion and low shrinkage, which allows to produce dimensionally accurate parts without the warping issues common in other technical polymers. While it is moderately resistant to chemicals and impacts, its relatively lower heat deflection temperature (approx. 70–80°C) means it is best suited for internal structural components or electronics housings that are shielded from direct solar heating.

3.1.16.3 ABS Filament

ABS (Acrylonitrile Butadiene Styrene) is a legacy engineering plastic valued for its toughness and high heat resistance. It is highly impact-resistant and has a "ductile" failure mode, meaning it tends to bend or deform before snapping, which provides a safety margin for robotic chassis components. However, ABS is notoriously difficult to print without a heated enclosure due to its high thermal expansion coefficient, and its poor UV resistance makes it unsuitable for long-term outdoor use unless it is painted or treated with a protective coating.

3.1.16.4 PC Filament

Polycarbonate (PC) is the premier choice for 3D printed parts requiring extreme impact resistance and high-temperature tolerance. Often used in industrial applications where components must withstand significant mechanical stress, PC filament can handle operating temperatures up to 110–140°C. In a hybrid robot chassis, it is best utilized for high-torque motor housings or protective battery enclosures, as its superior tensile strength (60–72 MPa) allows it to survive collisions that would shatter standard filaments.

3.1.16.5 5052-H32 Aluminum

5052-H32 Aluminum is an alloy primarily defined by its excellent formability and high resistance to corrosion, particularly in marine or outdoor environments [68]. As a non-heat-treatable alloy, the "H32" temper indicates that it has been strain-hardened and stabilized for a balance of strength and ductility. Its most significant feature for robotics is its ability to be bent at 90-degree angles with a small radius without cracking, making it the standard material for fabricating "U-channel" frames, base plates, and folded sheet-metal enclosures.

3.1.16.6 6061-T6 Aluminum

6061-T6 Aluminum is the most widely used structural aluminum alloy due to its excellent strength-to-weight ratio and superb machinability [69]. The "T6" temper is achieved through solution heat-treating and artificial aging, resulting in a yield strength that is significantly higher than the 5000-series alloys. Because it can be easily drilled, tapped, and CNC-machined with high precision, it is the industry standard for 2020 T-slot extrusions and the primary load-bearing "skeleton" of most mobile robotic platforms.

3.1.16.7 7075-T6 Aluminum

7075-T6 Aluminum is a high-strength "aerospace-grade" alloy that utilizes zinc as its primary alloying element [70]. It offers a yield strength that rivals many steels while maintaining the low density of aluminum, providing the highest specific strength of the materials considered. While it is more expensive and has lower corrosion resistance than

the 6000-series, its extreme hardness makes it indispensable for critical, high-wear components such as drive axles, custom wheel hubs, and high-reduction gearbox plates where mechanical failure is not an option.

3.1.16.8 Chassis Part Selection

PETG Filament and 6061-T6 Aluminum: PETG filament was selected as the preferred 3D printing material due to its balanced mechanical properties and improved environmental resistance compared to other common filaments. It offers greater toughness and layer adhesion than ASA, reducing the likelihood of brittle failure in functional parts such as mounts and enclosures. Additionally, PETG has better thermal stability, making it less prone to softening under elevated temperatures that may occur during outdoor operation. Its enhanced UV and moisture resistance relative to many other consumer filaments also supports longer service life in sun-exposed environments. These characteristics make PETG a reliable choice for non-load-bearing components that still require durability and repeated handling during prototyping and operation.

6061-T6 aluminum was selected for the primary structural components due to its strong combination of strength, weight efficiency, and manufacturability. This alloy provides high tensile strength and good stiffness, allowing the chassis to maintain structural integrity under mechanical loads while minimizing unnecessary mass. It is also widely available and easy to machine, making it practical for fabrication processes such as cutting, drilling, and fastening. In addition, 6061-T6 exhibits excellent corrosion resistance, which is important for outdoor use where exposure to moisture and environmental conditions is expected. Overall, it offers a well-proven, cost-effective solution for building a rigid and durable chassis framework.

3.2 Programming Languages

3.2.1 High-Level Languages

3.2.1.1 C/C++

C and C++ were considered for *SOLARIS* because they are both widely used in embedded systems and provide the level of hardware access needed for a robotics project. In particular, these languages allow the software to interact closely with the microcontroller, including configuring peripherals, handling real-time I/O, and managing memory in a predictable way. This matters for *SOLARIS* because the system needs reliable low-level control for sensors, motor drivers, and communication while still being maintainable over a long development cycle.

C is the primary language considered for this project due to its simplicity, efficiency, and strong compatibility with embedded toolchains. Most microcontroller vendors provide their best support, examples, and documentation in C, and many low-level driver libraries are designed with C as the default. C is also highly predictable in terms of performance and memory usage, which is important for a system that may need to run continuously and operate under power constraints.

C++ was also considered because some embedded libraries and frameworks provide better support or cleaner interfaces when using C++. In addition, C++ can improve code organization through features like stronger type checking and better modular structure. However, since *SOLARIS* does not require complex object-oriented design, and because simplicity is a major advantage in embedded development, C remains the preferred option for most of the system. Overall, C/C++ provides an ideal balance of control, performance, and compatibility, with C serving as the primary language and C++ used only when it offers clear practical benefits.

3.2.1.2 Python

Python is a strong language to consider for *SOLARIS* for a few key reasons. First, it is simply an easy language to work with. Python is high-level, dynamically typed, and supports object-oriented programming, which makes development fast and helps keep code organized as the project grows. This makes it especially useful for higher-level software tasks where readability and development speed matter more than raw execution speed.

Second, Python has extremely strong library support across a wide range of tools. Even though we are not currently planning on using computer vision or advanced AI for this project, it is still worth considering as a future option if our initial approach is not robust enough. Python has one of the best ecosystems for intelligent systems, including libraries for computer vision, machine learning, and sensor data processing. It is also very useful for analyzing and visualizing collected data, which could be helpful if we want to evaluate performance metrics like energy gained, energy used, distance traveled, and overall system efficiency.

Lastly, Python is attractive because of its usefulness outside of the embedded side of the project. It is widely used for backend development and automation, and it can support tasks like building APIs, handling stored data, and integrating with a phone application or website. For *SOLARIS*, the performance tradeoffs of Python are negligible for these higher-level tasks, since real-time motor control and low-level peripheral access would still be handled directly on the microcontroller in a lower-level language such as C. Familiarity also plays a major role, since the entire team is comfortable with Python, and some members already have experience using it in backend web development frameworks.

3.2.1.3 Java

Java is another language that could be considered for *SOLARIS*, mostly due to how structured and portable it is. Java is a strongly object-oriented and class-based language, which can be useful for keeping a large project organized and scalable. For team-based development, Java's design naturally encourages cleaner code separation, reuse, and easier collaboration since different components can be developed as separate classes and modules.

However, Java is not a strong fit for the low-level embedded side of this project. Since *SOLARIS* requires direct interaction with microcontroller peripherals, real-time motor control, and hardware-level I/O, languages like C tend to be far more common and better supported. Java's runtime environment and memory overhead also make it less practical for most microcontroller targets.

Where Java becomes more relevant is on the software side outside of the microcontroller. Java is commonly used for backend development and is also the primary language historically used for Android app development. If *SOLARIS* ever required a native Android application or a backend service that manages stored data, system metrics, or communication between devices, Java could support those parts of the project well. Overall, Java is a strong general-purpose option, but for *SOLARIS* it is more applicable to application-level software than embedded firmware.

3.2.1.4 JavaScript/Typescript

JavaScript and TypeScript are another strong option for *SOLARIS*, and they are grouped together since they are used in the same development ecosystem and typically share the same frameworks and tooling. The main difference is that TypeScript is essentially JavaScript with static typing added, which can help reduce bugs and improve maintainability in larger projects. In practice, both languages are commonly used for the same types of applications, especially when working with modern frontend frameworks.

For *SOLARIS*, the main reason to consider JavaScript and TypeScript is for the phone application. These languages are widely used with React Native, which provides a cross-platform development approach that supports both iOS and Android. This is a major advantage since it allows one shared codebase instead of needing to develop separate native applications for each platform. For a senior design project with limited time and resources, that cross-platform support can significantly reduce development overhead. React Native also has strong library support for Bluetooth Low Energy communication, which is important since BLE is a realistic option for connecting the phone application to the robot.

Another benefit is familiarity. While the team may not have direct experience with React Native specifically, React Native is based on the same general React development model used in web development. Since React and React-based frameworks are already familiar from web projects, it makes the learning curve more manageable compared to starting from a completely new mobile development environment. Overall, JavaScript and TypeScript are practical options for building a mobile interface for *SOLARIS* due to strong ecosystem support, cross-platform compatibility, and existing team familiarity.

3.2.1.5 Swift

Swift is another language considered for the *SOLARIS* phone application. Swift is Apple's primary language for iOS development and is commonly used for building native iPhone applications. One major advantage of Swift is that it provides strong performance, direct access to iOS features, and reliable support for Bluetooth Low Energy communication.

However, the main drawback is that Swift is limited to Apple devices. Since *SOLARIS* is intended to support both iOS and Android users, using Swift would require a separate Android application to be developed in a different language, which would increase development time and complexity. For this reason, Swift is less practical as the primary language for mobile applications, even though it is a strong option for iOS-only development.

3.2.1.6 Dart

Dart is another language considered for the *SOLARIS* phone application, mainly because it is used with the Flutter framework. Flutter is a cross-platform mobile development tool that supports both iOS and Android from a single codebase, similar to React Native. This makes Dart a practical option for a senior design project since it reduces the need to maintain separate mobile applications.

Dart and Flutter also have strong support for building modern user interfaces and can support Bluetooth Low Energy communication through existing libraries. However, one downside is that Dart is less commonly used than JavaScript and TypeScript, and it is less familiar to the team. Because of this, Dart is still a valid option, but it would likely have a steeper learning curve compared to React Native development.

3.2.1.7 Code Selection

C/C++: C/C++ will be used for the embedded firmware running directly on the microcontroller. This portion of the system requires direct hardware access, predictable timing, and efficient power usage. C is expected to be the primary language for the embedded side due to its strong vendor support, large ecosystem of driver libraries, and predictable performance. C++ may be used in specific cases where libraries provide better support or cleaner integration, but the firmware will remain mostly C to keep the embedded code lightweight and reliable.

Python: Python will be used for backend development through the FastAPI framework (as further discussed in Section 3.3.3). This allows the system to expose a clean API for storing and retrieving telemetry, robot status, and logged data. Python is a strong fit for this role because it supports rapid development, has excellent library support, and integrates easily with database tools such as PostgreSQL libraries. This choice also aligns well with the team's existing familiarity with Python.

JavaScript/TypeScript: JavaScript and TypeScript will be used for the mobile application frontend through React Native (as further discussed in Section 3.3.2). This provides a practical cross-platform solution for supporting both iOS and Android users from a single codebase. TypeScript also improves maintainability through static typing, which is useful as the application grows. React Native has strong ecosystem support for BLE communication, making it a good match for *SOLARIS*'s wireless protocol selection.

3.2.2 Algorithms

3.2.2.1 Reactive Obstacle Avoidance

A reactive obstacle avoidance algorithm is a navigational strategy for mobile robots that relies on local sensor data to detect and avoid collisions in real time. Rather than depending on a global system with a complete map of the environment, a reactive system makes quick decisions based on immediate sensor input, allowing it to function effectively in dynamic environments. For sensory input, reactive systems most commonly use range sensors such as ultrasonic sensors and depth sensors [71]. Ultrasonic sensors measure distance by emitting sound signals, while depth sensors use infrared lasers to generate a 3D representation of the environment. To implement this type of algorithm, two low-level AI

systems are traditionally used: a Fuzzy Logic Inference System (FIS) and an Artificial Neural Network (ANN). A FIS uses defined "IF-THEN" rules to handle uncertainty and ambiguity, making soft and gradual adjustments based on sensor features [71]. An ANN is more straightforward in practice, as it can be trained via supervised learning to associate sensor inputs with specific motor actions [71]. Reactive algorithms like these allow robots to learn from experience and adapt to new situations without requiring large, rigid programs for every possible scenario. Since they do not require prior knowledge of the environment, these models are highly flexible and capable of semi-autonomous operation.

3.2.2.2 Kalman Filters

<i>Variable</i>	<i>Description</i>	<i>Use Case</i>	<i>Input/Output</i>
x	State variable	$n \times 1$ column vector	Output
P	State covariance matrix	$n \times n$ matrix	Output
z	measurement	$m \times 1$ column vector	Input
A	State transition matrix	$n \times n$ matrix	System Model
H	State-to-measurement matrix	$m \times n$ matrix	System Model
R	Measurement covariance matrix	$m \times m$ matrix	Input
Q	Process noise covariance matrix	$n \times n$ matrix	System Model
K	Kalman Gain	$n \times m$	Internal

Table 3.23 Kalman Filter Reference Terms

Source: [72]

In applications that make use of inertial measurement units, a Kalman Filter is a common approach for combining multiple degrees of freedom through sensor fusion. The Kalman Filter works by estimating the true state of a system from two noisy sources of information. More specifically, it operates on the principle of Bayesian inference, where a probability distribution is used to represent uncertainty in both predictions and measurements [73]. The filter maintains two key pieces of information: the state estimate and the uncertainty of that estimate. Through a recursive process, the filter begins with a prediction step where the current state is estimated based on the previous state estimate [73]. An update step then follows, incorporating new sensor measurements to refine the prediction and produce an optimal estimate that balances the prediction against the reliability of the incoming data.

3.2.2.3 Solar Position

The Solar Positioning Algorithm (SPA) is a step-by-step procedure for tracking solar radiation defined by the National Renewable Energy Laboratory (NREL) [74]. The primary goal of the SPA is to calculate the solar zenith and azimuth angles based on a given location and time [74]. The algorithm operates in five steps: time scale conversion, earth position, solar coordinates, observer-specific data, and the final output of angles. The first step

calculates variables based on Universal Time (UT) and Terrestrial Time (TT) to convert raw time data into a usable format [74]. Earth position is then calculated by taking the heliocentric longitude, latitude, and radius vector into account [74]. Solar coordinates are calculated in the same manner, followed by observer-specific data where sun ascension, declination, and local hour angle are determined based on the observer's elevation and location [74]. Finally, a topocentric zenith and azimuth angle are output, which can be used in a control system to define motor behavior [74]. This algorithm has been shown to produce results within an uncertainty range of only $\pm 0.0003^\circ$, making it one of the most precise solar positioning methods available [74].

3.3 Frameworks

3.3.1 What are Frameworks?

Before selecting frameworks for *SOLARIS*, it is important to understand the difference between the main software layers in the system and why multiple frameworks may be needed. *SOLARIS* includes both an embedded system running directly on the robot and an application layer that allows the user to monitor and control the robot. Because these parts of the system have very different requirements, they often rely on different development tools and frameworks.

The application layer is split into frontend and backend development. The frontend refers to the part of the system the user directly interacts with, such as the mobile app screens, buttons, and overall user interface. For *SOLARIS*, the frontend is also responsible for direct communication with the robot, such as connecting over Bluetooth Low Energy and sending or receiving live data. The backend refers to software hosted separately from the phone, typically on a server. The backend is responsible for storing and organizing collected information, processing requests from the app, and providing a consistent way to access stored data later. In the context of *SOLARIS*, the backend could store robot logs, historical energy metrics, usage data, and other telemetry that the phone application collects. This is important because the phone application is not intended to permanently store large amounts of data, and backend storage also makes it easier to access data from other devices such as a website, which is a stretch goal of the project.

In addition to the phone application, *SOLARIS* also requires embedded development on the robot itself. This includes the firmware running on the microcontroller, which is responsible for tasks such as reading sensors, controlling motors, managing power, handling wireless communication, and logging data. Embedded development frameworks are important because they provide the build system, drivers, libraries, and debugging tools needed to develop reliable firmware efficiently.

A framework is a set of tools and built-in structure that helps developers build software faster and more consistently. Frameworks reduce the need to write everything from scratch by providing common features such as project organization, communication libraries, hardware drivers, and platform-specific support. For *SOLARIS*, choosing the right frameworks is important because it affects development time, cross-platform

compatibility, available libraries, and how easily the system can be developed and maintained.

3.3.2 Frontend Phone App Frameworks

	<i>React Native</i>	<i>SwiftUI</i>	<i>Flutter</i>
Android & iOS Support	Yes	No: iOS Only	Yes
Language	Javascript/Typescript	Swift	Dart
BLE Support	Yes (libraries available)	Yes (via CoreBluetooth)	Yes (via plugins)
Team Familiarity	Medium (react experience)	Low	Low
Library Ecosystem	Very Strong	Strong but only for iOS	Strong

Table 3.24 Frontend Framework Comparison

Sources: [75]-[77]

3.3.2.1 *React Native*

React Native is a mobile application framework developed by Meta that allows developers to build cross-platform applications for both iOS and Android using a single codebase. React Native is built around the React development model, meaning applications are designed using reusable components and structured user interfaces. Since React Native is based on JavaScript and TypeScript, it fits well with modern web development practices and has a large ecosystem of tools and libraries.

For *SOLARIS*, React Native is a strong candidate because it supports both major phone platforms while reducing development time compared to writing separate native applications. This is especially important in a senior design environment where development time is limited. Another major advantage is familiarity. While the team may not have direct experience with React Native specifically, the team does have experience with React-based web development. Since React Native follows the same general component structure and workflow, this reduces the learning curve and makes it easier to build a functional application in a short timeframe.

React Native also has reliable support for Bluetooth Low Energy communication through existing libraries. Since BLE is a practical option for communicating between the robot and the phone application, React Native provides a straightforward way to implement device scanning, connecting, and sending and receiving data. Overall, React Native is a practical framework for *SOLARIS* due to cross-platform support, strong library availability, and existing team familiarity.

3.3.2.2 *SwiftUI*

SwiftUI is Apple's modern user interface framework for building iOS applications using the Swift programming language. SwiftUI is designed for creating clean, responsive

interfaces using a declarative development style, which makes it easier to build and maintain UI compared to older iOS approaches. Since SwiftUI is a native Apple framework, it integrates directly with iOS features and provides strong performance and stability.

For *SOLARIS*, SwiftUI is a valid framework to consider because it would allow the team to build a high-quality native iPhone application. SwiftUI also has strong support for Bluetooth Low Energy communication through Apple's Core Bluetooth framework, meaning the app could reliably connect to the robot and display live data. In addition, since it is a native framework, SwiftUI would generally provide the best user experience for iOS users and would have fewer compatibility concerns compared to cross-platform frameworks.

However, the main drawback of SwiftUI is that it only supports Apple devices. Using SwiftUI would require building a separate Android application using a different language and framework if we still wanted dual compatibility, which would significantly increase development time and complexity. For a senior design project where time and resources are limited, this makes SwiftUI less practical as the primary framework, even though it is a strong option for iOS-only development.

3.3.2.3 Flutter

Flutter is a cross-platform mobile application framework developed by Google that allows developers to build applications for both iOS and Android using a single codebase. Flutter applications are primarily written in Dart, and the framework provides a large set of built-in UI components for designing modern interfaces. Flutter is known for allowing fast development through features such as hot reload, which makes it easier to quickly test and refine the application during development.

For *SOLARIS*, Flutter is a strong option because it supports both major phone platforms while reducing the need to build separate native applications. This is a major advantage for a senior design project since it lowers development time and simplifies long-term maintenance. Flutter also has support for Bluetooth Low Energy communication through existing libraries, meaning it can be used to connect to the robot and display real-time data.

One drawback of Flutter is that it introduces an additional learning curve compared to other options. Since the team is more familiar with JavaScript and React-based development, React Native may be easier to adopt than learning Dart and the Flutter ecosystem from scratch. Overall, Flutter is a strong cross-platform framework with good performance and BLE support, but its practicality depends on how quickly the team can learn and integrate it into the project timeline.

3.3.2.4 Frontend Phone App Framework Selection

React Native: React Native was selected as the mobile application framework for *SOLARIS*. React Native provides cross-platform support for both iOS and Android using a single codebase, which is a major advantage for a senior design project where development time and resources are limited. This allows the team to build and maintain one application rather than developing separate native apps for each platform.

Another major reason for selecting React Native is familiarity. React Native uses JavaScript and TypeScript, which the team is already comfortable with, and it follows the same general component-based development model as React for web development. Since the team already has experience with React-based projects, the learning curve for React Native is significantly lower compared to SwiftUI and Flutter, which would require learning either the Swift ecosystem or Dart.

React Native also provides strong library support for Bluetooth Low Energy communication, which is important since BLE is the selected wireless protocol for *SOLARIS*. This makes it straightforward to implement key features such as scanning for the robot, connecting to it, sending and receiving control and telemetry data. Overall, React Native provides the best balance of cross-platform support, ease of development, and practical ecosystem support for the *SOLARIS* mobile application.

3.3.3 Backend Phone App Frameworks

	<i>FastAPI</i>	<i>Django</i>	<i>Spring Boot</i>	<i>Express.js</i>
Best DB framework pairing	Postgres	Postgres	MySQL	MySQL
Language	Python	Python	Java	Javascript / Typescript
Performance	High	Medium	Low to Medium	Medium
Team Familiarity	High	Medium	Low to Medium	Low to Medium
Library Ecosystem	Strong	Very Strong	Low to Medium	Medium

Table 3.25 Backend Framework Comparison

Sources: [78]-[81]

3.3.3.1 *FastAPI*

FastAPI is a modern Python backend framework designed for building APIs quickly and efficiently. It is commonly used for creating REST APIs and is known for having strong performance compared to many other Python web frameworks. One of the major benefits of FastAPI is that it is easy to develop with while still being scalable. It also includes automatic request validation and automatically generates interactive API documentation, which is helpful for development and testing.

For *SOLARIS*, FastAPI is a strong option because it allows the backend to be built quickly while still supporting the type of functionality the project would require. This includes storing system data, handling user requests from the phone application, and providing an organized way to access stored robot metrics from other interfaces such as a website.

FastAPI also works well with common database tools such as SQLAlchemy and Alembic, making it easier to design and manage a database for storing collected robot information.

Another major advantage of FastAPI is familiarity. Since the team already has experience using FastAPI, it reduces the learning curve and makes development more efficient. Overall, FastAPI provides a strong balance of performance, ease of development, and compatibility with modern backend tools, making it a practical choice for *SOLARIS*.

3.3.3.2 Django

Django is a Python backend framework that is designed to support structured and scalable web development. Django provides many built-in tools that are commonly needed for backend systems, such as routing, authentication support, and database integration. One of the biggest advantages of Django is that it includes its own ORM and migration system, which makes it easier to design, manage, and update a database throughout development. Django is also widely used and has a very large ecosystem, meaning there is strong documentation and long-term community support.

For *SOLARIS*, Django is a valid framework to consider because it could support a backend system that stores robot metrics, organizes collected data, and serves that data to other interfaces such as the phone application or a website. Django is also useful if the project expands into a more feature-heavy platform, since it provides a strong built-in structure for managing data models and backend logic.

However, Django may be less ideal if the backend is intended to remain lightweight and focused primarily on API endpoints. Since Django is designed to support full web applications, it can introduce additional overhead compared to frameworks that are specifically optimized for API development. Overall, Django is a strong option for backend development, but it is best suited when a project benefits from its larger set of built-in features.

3.3.3.3 SpringBoot

Spring Boot is a backend framework built on the Java programming language and is widely used for developing large and structured web applications. It is especially common in industry for building REST APIs and backend services due to its reliability, long-term support, and strong ecosystem. Spring Boot provides a clear project structure and includes many built-in tools that make backend development more standardized, especially for larger codebases.

For *SOLARIS*, Spring Boot is a framework worth considering because it has strong support for database integration and backend architecture. It commonly works with tools such as Spring Data JPA and Hibernate for database management, and it supports database migration tools like Flyway or Liquibase. This makes it a strong option for building a backend system that stores robot metrics, serves data to a phone application, and scales well if the project expands.

However, Spring Boot can also be more complex than other backend frameworks, especially for smaller projects. Since it is designed with enterprise-level applications in

mind, it can introduce more setup and configuration compared to lighter frameworks such as FastAPI. In addition, the team's familiarity with Java backend development may be lower compared to Python-based frameworks. Overall, Spring Boot is a strong and professional backend option, but it may not be the most efficient choice for a time-limited senior design project.

3.3.3.4 *Express.js*

Express.js is a backend framework for Node.js that is commonly used to build REST APIs and web servers using JavaScript or TypeScript. Express is known for being lightweight and flexible, meaning it does not force a strict project structure and allows developers to design the backend in a way that fits the needs of the application. Because of how widely used it is, Express also has a very large ecosystem of libraries and community support.

For *SOLARIS*, Express.js is a valid framework to consider because it can be used to build a backend that stores robot data, serves API endpoints to the phone application, and supports integration with a database. Express is often paired with database tools such as Prisma, Sequelize, or TypeORM, which provide structured ways to store and query system information. Since JavaScript and TypeScript are already being considered for the mobile application, using Express could also allow the project to share the same language across both frontend and backend development.

However, one drawback of Express is that it is relatively minimal by default. While this can be an advantage for small projects, it also means that more features such as validation, authentication, and project structure must be built manually or added through third-party libraries. Overall, Express.js is a strong backend option due to its simplicity, flexibility, and ecosystem, but it may require more setup and organization compared to frameworks that provide more built-in structure.

3.3.3.5 *Backend Phone App Framework Selection*

FastAPI: FastAPI was selected as the backend framework for *SOLARIS*. The biggest reason for this choice is familiarity. The team is already comfortable with Python, and some team members have direct experience working with FastAPI specifically. This reduces the learning curve and allows development time to be spent on implementing system features rather than learning a new backend framework from scratch.

FastAPI is also a strong fit for *SOLARIS* because the backend is primarily intended to serve as an API layer for storing and accessing robot telemetry, system status, and logged metrics. FastAPI is designed specifically for building lightweight, efficient APIs, which makes it more practical for this project than heavier full-stack frameworks. In addition, FastAPI pairs well with PostgreSQL through widely used Python libraries, making database integration and long-term data management straightforward. Overall, FastAPI provides the best balance of simplicity, performance, and development efficiency for the *SOLARIS* backend.

3.3.4 *Embedded Development Frameworks*

With the selection of the ESP32-S3 microcontroller (as discussed in Section 3.1.2), the embedded development framework options for *SOLARIS* can be narrowed down

significantly. For this project, both PlatformIO and ESP-IDF will be used. The following sections describe each framework and how they support development for the *SOLARIS* embedded system.

3.3.4.1 Platform I/O

PlatformIO will be used as the primary development environment for the *SOLARIS* embedded firmware. It provides the overall project structure, build system, and file organization needed to keep the codebase clean and manageable throughout development. PlatformIO also simplifies the process of compiling and flashing firmware onto the ESP32-S3, which is important for fast iteration during testing.

Another major advantage is that PlatformIO integrates directly into Visual Studio Code, allowing the team to develop, build, and upload firmware from a single environment. This improves workflow efficiency and makes it easier to manage libraries, serial monitoring, and debugging tools during development.

3.3.4.2 ESP IDF

ESP-IDF (Espressif IoT Development Framework) will be used alongside PlatformIO because it provides the official software development framework for the ESP32-S3. ESP-IDF includes the core drivers, hardware libraries, and system-level tools needed to fully access the capabilities of the microcontroller. It also provides a large collection of example projects and reference implementations, which is extremely useful for learning and integrating features such as Bluetooth Low Energy, peripherals, and power management.

ESP-IDF is especially important for *SOLARIS* because it allows the project to rely on well-supported vendor libraries rather than requiring low-level code to be written from scratch. This reduces development risk and makes it easier to implement embedded features reliably.

3.4 Wireless Communication Protocols

Table 3.26 *Wireless Communication Protocol Comparison*

	<i>Wi-Fi</i>	<i>Bluetooth Standard</i>	<i>Bluetooth Low Energy</i>	<i>LoRaWAN</i>	<i>Zigbee</i>
Range (m)	20-100	Up to 30	Up to 50	10,000	1000
Data Rate	Up to 46 Gbps	1-3 Mbps	.5-1 Mbps	27 Kbps	500 Kbps
Power Consumption	Highest	Medium-Low	Low	Lowest	Low
Frequency Bands	2.4GHz/5GHz/6GHz	2.4GHz	2.4GHz	Medium	915MHz / 2.4GHz
Phone Compatibility	Yes	Yes	Yes	No: separate hub required	ZigBee direct: Uses Bluetooth to connect to Zigbee IOT devices

Table 3.26 Wireless Communication Protocol Comparison

Sources: [82],[83],[84]-[88]

3.4.1 Wi-Fi

Wi-Fi is a common wireless protocol that provides strong compatibility with smartphones and relatively high usable range. It operates in the 2.4 GHz and 5 GHz bands, with newer standards also supporting 6 GHz. Although Wi-Fi is capable of very high data rates, *SOLARIS* does not require large bandwidth for normal operation, since communication is primarily limited to low-rate telemetry and control commands.

Wi-Fi is still a strong candidate because it can be straightforward to integrate into a phone application and typically provides greater range than Bluetooth-based communication, especially in open outdoor environments. Wi-Fi also does not strictly require a router. The robot can be configured to create its own local Wi-Fi network, allowing a phone to connect directly without relying on external infrastructure.

The primary disadvantage of Wi-Fi is power consumption. Wi-Fi communication generally requires more energy than Bluetooth-based alternatives, which can reduce battery life and negatively impact a solar-powered system's net energy gain. This drawback can be partially mitigated by enabling Wi-Fi only when necessary and using low-power sleep modes when communication is idle, but Wi-Fi is still typically less power-efficient than BLE for continuous operation.

3.4.2 Bluetooth Classic

Bluetooth Standard (Bluetooth Classic) is a widely supported wireless protocol that allows direct communication between the robot and a smartphone without requiring external

infrastructure. It typically provides higher data rates than Bluetooth Low Energy, making it well suited for continuous streaming applications such as audio or sustained real-time data transfer.

For *SOLARIS*, Bluetooth Classic was considered because it offers native phone compatibility and does not require a router or access point. However, the primary communication needs of *SOLARIS* are low-bandwidth telemetry and control commands rather than continuous high-rate streaming. As a result, Bluetooth Classic's higher throughput does not provide a meaningful advantage for this system.

Additionally, Bluetooth Classic generally consumes more power than BLE due to its connection model and design goals. Since *SOLARIS* is a solar-powered robot where energy efficiency is a key constraint, Bluetooth Classic is less attractive than BLE for long-term operation.

3.4.3 Bluetooth Low Energy (BLE)

Bluetooth Low Energy (BLE) is a wireless protocol designed specifically for low-power, battery-operated devices. BLE supports direct smartphone connectivity without requiring external infrastructure and is well suited for applications that transmit small amounts of data intermittently. This matches the needs of *SOLARIS*, where communication primarily consists of control commands, mode selection, and periodic telemetry such as battery level, light sensor readings, and system status.

BLE is a major contender for *SOLARIS* because it offers a strong balance of energy efficiency, reliability, and ease of integration into a mobile application. Compared to Wi-Fi and Bluetooth Classic, BLE typically requires significantly less power, making it more appropriate for a solar-powered robot where minimizing energy usage is critical to meeting overall system goals.

The primary drawback of BLE is range. BLE communication generally has a shorter usable range than the other protocols, especially in outdoor environments where signal strength may vary due to obstacles, orientation, or interference. This limitation may restrict the distance at which a user can reliably connect to the robot. However, because *SOLARIS* is intended to be controlled and monitored locally by a nearby user, BLE range is considered an "acceptable" tradeoff given its power advantages. However, it may be better to use a protocol like Wi-Fi for that added benefit.

3.4.4 LoRaWAN

LoRaWAN is a low-power wide-area networking (LPWAN) protocol designed for long-range communication with very low data rates. It is commonly used in IoT applications such as remote sensing, environmental monitoring, and asset tracking. LoRaWAN is attractive for solar-powered systems because it offers extremely low power consumption and can achieve communication ranges of several kilometers under favorable conditions. It also aligns well with *SOLARIS* in terms of data requirements, since the robot primarily needs to transmit small telemetry packets rather than high-bandwidth data.

LoRaWAN was considered as a potential option because it provides the best combination of range and power efficiency among the protocols evaluated. In theory, it would allow *SOLARIS* to report status information over long distances while minimizing energy usage, which is desirable for an outdoor autonomous robot.

However, LoRaWAN has a major limitation for this project: it does not provide native smartphone compatibility. A phone cannot directly connect to LoRaWAN in the same way it can connect to Wi-Fi or Bluetooth. Supporting LoRaWAN would require an intermediary device such as a gateway or hub to relay data between the robot and the phone application. This additional infrastructure increases cost, setup complexity, and development overhead, making LoRaWAN impractical for a system intended to be controlled directly through a phone.

3.4.5 Zigbee

Zigbee is a low-power wireless protocol commonly used in IoT and smart home applications. It is designed for reliable communication between many devices and is especially effective in mesh network topologies, where multiple nodes relay messages to extend coverage and improve robustness. Zigbee operates in the 2.4 GHz band globally and can also operate in sub-GHz bands in some regions. Its low power consumption and support for multi-device networking make it a strong option for distributed sensor systems.

For *SOLARIS*, Zigbee was considered because it offers low power consumption and can provide longer range than BLE in some configurations, particularly when mesh networking is used. Zigbee is also widely used in commercial IoT ecosystems, which makes it a proven technology for low-power telemetry.

However, Zigbee has a major limitation for this project: it does not provide direct, native smartphone compatibility in the same way Wi-Fi and Bluetooth do. Zigbee Direct is a newer approach that allows a phone to communicate with Zigbee devices, but it does so by using Bluetooth as the communication method between the phone and the Zigbee device. This defeats much of the purpose of choosing Zigbee for a one-to-one robot-to-phone link, since Bluetooth is still required for the phone connection.

3.4.6 Wireless Communication Protocol Selection

Bluetooth Low Energy (BLE): After comparing the available wireless communication protocols, *SOLARIS* will use Bluetooth Low Energy (BLE). None of the protocols are truly “perfect” for this project, but BLE provided the best overall balance of advantages while having the fewest drawbacks that actually matter for our system.

One of the biggest reasons BLE stands out is native smartphone compatibility. *SOLARIS* is intended to be controlled directly through a phone, and BLE supports this without requiring any external infrastructure. Protocols like LoRaWAN and Zigbee offer strong range and power advantages, but they do not provide direct phone connectivity. Supporting them would require an additional hub or gateway to relay communication between the robot and the phone application. That added hardware and development overhead would significantly increase cost, complexity, and setup requirements, which is not justified for a

one-to-one robot-to-phone link. Because of this, the only realistic contenders for *SOLARIS* were Wi-Fi and Bluetooth.

BLE is also a strong choice because it matches the communication needs of the system. *SOLARIS* does not require high bandwidth, since the data being transmitted consists of low-rate telemetry and short control commands. BLE supports this easily while staying efficient. Bluetooth Classic was considered, but its higher data rate does not provide a meaningful advantage for this project. Since Bluetooth Classic generally consumes more power than BLE, the increased throughput would be wasted for our use case.

Wi-Fi's main advantage is range, and it can also support direct backend communication when connected to a router. In theory, this would allow the robot to send data directly to a database without relying on the phone as an intermediary. However, this assumes reliable outdoor router coverage, which is not realistic for most expected use cases, especially when the robot is operating autonomously. Even with Wi-Fi, disconnects would still need to be handled, meaning *SOLARIS* would still require local storage and synchronization logic. The other Wi-Fi option is for the robot to act as its own access point so a phone can connect directly, but in that scenario neither the phone nor the robot would have internet access, making direct backend communication impossible.

Ultimately, BLE provides the most practical solution for *SOLARIS*. While it sacrifices some range compared to Wi-Fi and some power efficiency compared to LoRaWAN and Zigbee, it offers the best overall combination of low power, simplicity, and direct smartphone usability. Because *SOLARIS* is expected to be controlled locally by a nearby user, BLE range is considered an acceptable limitation, making it the most appropriate wireless communication protocol for this project.

3.5 Power Supply

Up until this point we have shared our goals and requirements for this project. Mentions of hardware usage have been stated but the reasoning for their choice have not been stated in detail. In this section we look over the valuable research used in making those decisions. It is not only valuable to look into the technologies we want to use but to also be open-minded as to why we did not choose others given sufficient investigation. This presents itself as a strong opportunity to deepen our understanding of the hardware technologies available to us for best fit within our project *SOLARIS*.

3.5.1 Brief

The *SOLARIS* battery subsystem is designed to balance weight, safety, and longevity. By prioritizing high energy density, the system maximizes runtime while maintaining a low mass profile. Safety is integrated through carefully selected chemistries and protection circuits, while cost and durability are managed by analyzing cycle life and Depth of Discharge (DoD). This multi-criteria approach ensures the robot remains operable and cost-effective over its intended lifespan.

3.5.2 Battery Considerations

SOLARIS utilizes a 12 V power supply architecture. To finalize the energy storage system, two primary factors must be determined: the optimal battery chemistry and the required capacity. The chemistry will be evaluated based on metrics such as specific energy, depth of discharge (DoD), and cost-effectiveness. Simultaneously, the capacity must be optimized to achieve a balance between system robustness and the avoidance of unnecessary weight penalties.

3.5.3 Power Distribution

	<i>Voltage</i>	<i>Current Draw</i>	<i>Quantity</i>	<i>Power Consumption</i>
ESP32-S3	3.0-3.6V	0.108A	1	0.356W
GV-5-LI-14.2V	14.2V	0.190mA	1	2.698mW
TEPT5600	1.5-6V	0.035A	4	0.462W
INA228	2.7-5.5V	0.64mA	1	2.112mW
BNO085	2.4-3.6V	0.014A	1	0.0462W
MB1020-000	2.8-5.5V	0.004A	1	0.0132W
goBilda 5203-2402-0188	12.8V	0.3A	2	7.68W
goBilda 5204-8002-0139	12.8V	0.3A	4	15.36W
Total Movement Power Consumption				23.04W
Total Sensors and Controllers Power Consumption				0.88W
Total System Power Consumption				23.92W

Table 3.27 Power Consumption Table

The power consumption for all non-locomotion components is approximately 0.9 W, totaling 10.8 Wh over a 12-hour operational window. Based on a conservative estimate of one hour of daily movement, locomotion adds 23.1 Wh to the energy budget. Combined with the 15 Wh auxiliary power allocation specified in the project objectives, the total daily energy requirement is 48.9 Wh. To ensure multi-day autonomy during periods of limited solar irradiance, a safety factor of four was applied to the daily consumption. This results in a minimum required capacity of 195.6 Wh; therefore, the system will utilize a battery with a capacity of at least 200 Wh.

3.5.4 Battery Technology Comparison

Battery Technology	Specific Energy (Wh/kg)	Depth Of Discharge	Pros	Cons	Cost (\$/kWh)
Lithium Ion (NMC)	~150-250	80-90%	Very high energy density, good power capability	Moderate safety risk if abused, requires BMS, shorter cycle life than LiFePO4, thermal runaway possible	150-300\$
Lithium Iron Phosphate (LiFePO4)	~90-120	80-100%	Excellent safety and thermal stability, very long cycle life (2000–5000+), good for deep cycling and solar storage	Lower energy density (heavier than NMC), higher upfront cost than lead-acid	200-400\$
Lead-acid	~30-50	50%	Low upfront cost, simple charging, highly recyclable	Very heavy, short cycle life, poor deep-discharge capability, lower efficiency	60-150\$

Table 3.28 Battery Technology Comparison

Sources: [89],[90]

The provided table evaluates the primary characteristics of three of the most prevalent secondary, or rechargeable, battery chemistry. To determine the most effective energy storage solution for the *SOLARIS* project, multiple battery parameters were explored. Specific Energy is a metric that defines the energy density of the chemistry, representing the amount of energy stored per unit of mass (Wh/kg). High specific energy is critical for mobile robotics, as it directly influences the weight-to-power ratio and overall payload capacity. Depth of Discharge (DoD) indicates the percentage of the total battery capacity that can be safely discharged before recharging is required. A higher allowable DoD increases the "usable" energy available in a single cycle, effectively maximizing the utility of the battery's rated capacity without compromising its cycle life. While the base cost of raw battery cells is a factor, the total system cost is variable and includes necessary overhead such as integrated Battery Management Systems (BMS), protection circuitry, and specialized charging hardware. These elements are especially essential for maintaining the safety and longevity of high-density lithium battery chemistries.

3.5.4.1 Lithium Ion (NMC)

Lithium Nickel Manganese Cobalt (NMC) batteries are renowned for their high specific energy, typically ranging from 150 to 250 Wh/kg, making them the standard choice for

weight-sensitive applications such as electric vehicles and high-end consumer electronics. While they offer a compact footprint and high power output, NMC chemistries are more susceptible to thermal runaway if punctured or overcharged, necessitating sophisticated battery management systems (BMS) for safety [89]. Their cycle life is generally moderate, often rated for 1,000 to 2,000 full discharge cycles before significant capacity degradation occurs.

3.5.4.2 Lithium Iron Phosphate

Lithium Iron Phosphate (LFP) batteries prioritize safety and longevity over absolute energy density. With a specific energy typically between 90 and 120 Wh/kg, they are heavier than NMC but offer exceptional thermal and chemical stability, significantly reducing the risk of combustion even under extreme stress [90]. LFP is characterized by its remarkable cycle life, frequently exceeding 2,000 to 5,000 cycles at high depths of discharge. Additionally, their cells exhibit a very flat discharge curve, maintaining a consistent voltage across the majority of the discharge cycle, which simplifies power regulation for sensitive electronics.

3.5.4.3 Lead-Acid

Lead-acid batteries are a mature technology primarily utilized for their low initial cost and high surge current capabilities. However, they possess the lowest specific energy of the three chemistries (30–50 Wh/kg), resulting in a much heavier and bulkier footprint for the same energy capacity [90]. Their cycle life is relatively short, often limited to 300–500 cycles, and they are sensitive to deep discharges, which can lead to permanent capacity loss. Furthermore, lead-acid batteries require longer charging times and have higher self-discharge rates compared to lithium-based alternatives.

3.5.4.4 Battery Technology Selection

Lithium Iron Phosphate: For outdoor solar-powered mobile robot applications, LFP is the superior chemistry because its operational characteristics align perfectly with the challenges of fluctuating environmental conditions and repeated daily cycling. The exceptional thermal stability of LFP is a critical safety requirement for robots exposed to direct sunlight and high ambient temperatures, as it remains stable in heat that could pose a risk to NMC cells. Its high cycle life ensures that the daily charge-discharge routines dictated by solar availability do not prematurely degrade the battery, allowing for years of field deployment. Furthermore, the high charge-acceptance rate and low internal resistance of Lithium Iron Phosphate enable the system to efficiently capture and store the varying energy levels provided by solar panels, while its flat discharge curve ensures that motor performance and sensor accuracy remain consistent as the battery depletes.

3.5.5 Battery Part Comparison

<i>Solar Panel Charge Controller Part</i>	<i>Rated Capacity</i>	<i>Operating Temperature</i>	<i>Size</i>	<i>Weight</i>	<i>Cost</i>
Dr.Prepare DBT1220LFP	20Ah (256Wh)	-10-45°C	7.13 × 3.03 × 6.61 in	5.95lbs	\$79.99
Bioenno Power BLF-1220AS	20Ah (240Wh)	-10-60°C	7.17 × 3.03 × 6.73 in	5.73lbs	\$204.99
Renogy RBT1220LFP-TM- G1	20Ah (256Wh)	-20-60°C	7.13 × 3.03 × 6.61 in	5.95lbs	\$89.99
Ultralife URB12200	21.6Ah (276Wh)	-20-60°C	7.12 × 3.03 × 6.57 in	5.97lbs	\$285.13

Table 3.29 Battery Part Comparison

Sources: [91]-[94]

The candidate batteries identified for this comparison each provide a minimum capacity of 240 Wh. While this exceeds the initial 200 Wh requirement established in the power budget, the selection was dictated by the discrete capacities available in the LiFePO₄ market. This 20% margin serves as a beneficial "buffer" for extended autonomy while remaining the closest fit for the system's design constraints. To ensure the reliability of the *SOLARIS* unit, several battery parameters were examined. The operating temperature range is a critical constraint due to the intended deployment in Florida. With summer ambient temperatures commonly reaching 38.4°C (101°F), the battery must be capable of safe operation at these levels, especially when accounting for the internal heat generated by the power electronics and motor controllers. Additionally, since *SOLARIS* is a mobile platform, minimizing the physical footprint and total mass of the energy storage system is essential. The four batteries under consideration exhibit similar size and weight profiles, but choosing a high energy density battery is optimal for the robot's mobility and operational efficiency.

3.5.5.1 Dr.Prepare DBT1220LFP

The Dr. Prepare 12V 20Ah battery prioritizes accessible operational parameters and clear charging instructions for general-purpose solar and DC-DC systems [91]. It specifies an optimal charging current of 10 A with a 20 A maximum limit and maintains a wide operating temperature range from -20°C to 60°C. The documentation provides a detailed overview of charging methods and safety guidelines, ensuring that users can easily integrate the 20 A continuous discharge unit into variable energy environments without exceeding the limits of the internal battery management system.

3.5.5.2 Bioenno Power BLF-1220AS

The Bioenno Power BLF-1220AS is optimized for high-power applications, featuring a robust maximum continuous discharge current of 40 A and a peak pulse capability of 60A

for up to five seconds [92]. This high current ceiling makes it particularly suitable for small electric motor use cases where transient power surges occur. The documentation emphasizes the necessity of a 14.6 V charging profile to trigger internal balancing and maintain the health of the 20 Ah capacity.

3.5.5.3 Renogy RBT1220LFP-TM-G1

The Renogy 12.8 V 20 Ah battery is designed as a deep-cycle storage solution [93]. It operates at a nominal 12.8 V and is engineered for repeated high-depth discharge cycles, which is fundamental for solar-powered platforms. The technical literature highlights the battery's role in grid-independent systems and provides a comprehensive safety framework, positioning the unit as a reliable energy core within a larger ecosystem of solar panels and charge controllers.

3.5.5.4 Ultralife URB12200

The Ultralife URB12200 is distinguished by a slightly higher nominal capacity of 21.6 Ah and a total energy rating of 276 Wh, offering a larger energy density of 102 Wh/kg [94]. It is engineered for high reliability and is IEC 62133 compliant, featuring a rugged exterior with integrated carry handles to facilitate mobility. Technically, it supports a maximum continuous discharge current of 20.0 A and maintains uniform voltage throughout the discharge cycle, which is a critical characteristic for applications requiring consistent power delivery.

3.5.5.5 Battery Part Selection

Renogy RBT1220LFP-TM-G1: The Renogy RBT1220LFP-TM-G1 distinguishes itself as the optimal choice for the *SOLARIS* project due to its specialized focus on deep-cycle longevity and seamless integration into solar-powered ecosystems. Unlike general-purpose lithium batteries, the Renogy unit is specifically engineered to withstand the frequent, varying charge-discharge cycles typical of a robot that relies on fluctuating solar intensity. This resilience is backed by a sophisticated internal Battery Management System (BMS) that offers comprehensive protection against overcharging from high-output panels and over-discharging during peak motor draws. While other batteries might offer higher burst currents, the Renogy provides a more stable and predictable energy density of 12.8 V and 20 Ah, which perfectly aligns with the project's requirement for a robust, high-cycle-life core. Its industry-leading reputation for durability in off-grid applications makes it the most reliable battery for ensuring the robot remains operational through years of field deployment.

3.5.6 Voltage Regulator Technology Comparison

3.5.6.1 Voltage Regulator Considerations

To properly operate the peripherals on *SOLARIS*, there must be two voltage regulators to buck the 12 V battery voltage to a usable 5 V and 3.3 V rails for usage in the smaller electronics in the robot. The dedicated 5 V rail will be used by the USB-A to charge the robot manually and the 3.3 V rail will be used to support the ESP32 and the multiple onboard sensors such as the ultrasonic sensors that will be used for ranging during navigation. The 5 V regulator must be able to support moderate current draw while having high thermal and energy efficiency to preserve power efficiency.. The 3.3 V rail mainly requires good voltage stability and low ripple to facilitate reliable ESP32 and sensor

operation. One of the main goals with *SOLARIS* is power efficiency. Because of this, regulator efficiency directly impacts day-to-day operational runtime and overall energy usage. Long term component reliability can be affected if there is excess dissipation in the regulator leading to excess heat generation. Finally, noise generated by the switching regulators must be considered, especially the ones near the sensors. For these concerns, both switching and linear regulator technologies were evaluated to determine what the optimal balance was for *SOLARIS* to optimize, efficiency, thermal performance, noise, integration complexity and cost.

<i>Voltage Regulator Technology</i>	<i>Power Efficiency</i>	<i>Daily Energy Capture</i>	<i>Safety</i>	<i>Complexity</i>	<i>Cost</i>
Switching	High efficiency, lower power dissipation than LDO for large voltage drops.	Maximizes battery energy, Reduced thermal loss.	Lower heat generation reduces thermal stress, switching action produces ripple.	Moderate. Requires inductor, MOSFET, a rectifier, output filter and careful PCB layout.	Moderate. More components than LDO means higher cost.
Linear	Efficiency = V_{out}/V_{in}. Larger voltage differences result in heat dissipation	Reduces usable battery energy due to excess voltage dissipated as heat.	Low ripple with lower noise and EMI, needs thermal management at higher currents.	Low. Requires minimal support components and simpler PCB layout.	Low, Fewer external components and simple architecture.

Table 3.30 Voltage Regulator Technology Comparison

Sources:[95]-[97]

3.5.6.2 Switching Regulator

Switching regulators regulate voltage by rapidly switching a MOSFET on and off, using an inductor to transfer energy through the circuit [96]. Because the MOSFET fully switches on and off each cycle, power dissipation across the transistor is minimized, which allows for increased efficiency depending on the load and operating conditions. This makes switching regulators particularly well suited for applications with a large voltage difference between input and output, such as stepping down from 12 V to 5 V. The higher efficiency directly reduces heat generation and improves overall energy efficiency, which is critical for maximizing battery runtime in an application like *SOLARIS*. The main drawback of switching regulators is that the MOSFET switching produces high ripple and electromagnetic interference (EMI). However, with careful PCB design and layout along with properly placed filtering capacitors, these effects can be minimized to ensure stable, low-noise operation for the onboard sensors.

3.5.6.3 Linear Regulator

Linear regulators regulate voltage by operating a pass transistor in its linear region to drop the excess voltage as heat [95]. Unlike switching regulators, linear regulators do not store and transfer their energy via an inductor. Instead, the voltage difference between the input and output is dissipated across the regulator. This makes the theoretical efficiency = V_{out}/V_{in} . This means that stepping down from 12 V to 5 V would yield approximately a 41.7% efficiency rating meaning roughly half the input would be dissipated as heat. This makes linear regulators very inefficient for larger voltage drops in battery-powered systems such as *SOLARIS*. The advantage of linear regulators is that they have much lower ripple and EMI than switching regulators which makes them very good for powering microcontrollers and analogue sensors [97]. They also have greater simplicity and low component count which makes them a fantastic choice for applications that can handle their efficiency drawbacks.

3.5.6.4 Voltage Regulator Technology Selection

Switching Regulator: With *SOLARIS* having a 12 V battery system and the desire to have a 5 V USB-A charging rail and the 3.3 logic rail for the MCU and sensors, the selected regulator must be able to balance efficiency, thermal performance, electrical noise, cost, and ease of integration. The 5 V rail needs to be able to support a moderate current draw and has a very sizable voltage drop from the input voltage to the output voltage from the battery, thus, a switching buck regulator was selected to ensure the desired efficiency and minimize the thermal losses on the PCB. This will ensure that the daily energy captured by the solar panel and charging system will be able to be used effectively and not simply lost as heat. The 3.3 V rail is for powering the MCU and multiple sensors. While there is higher ripple on the switching buck regulators, there will be adequate components in place on the PCB to ensure proper stable operation for these more sensitive components.

3.5.7 Voltage Regulator Part Comparison

<i>Voltage Regulator or Part</i>	<i>Conversion Efficiency ($V_{in} = 12\text{ V}$, $V_{out} = 5\text{ V}$)</i>	<i>Output Channels</i>	<i>Key Features</i>	<i>Max Current Per Channel</i>	<i>Quiescent Current</i>	<i>Cost</i>
LM2679	84% ($I_{Load} = 5\text{ A}$)	1	High current (5A); Current limit adjustable	5A	~7mA	\$7.42
LM2576	77% ($I_{Load} = 3\text{ A}$)	1	Robust; Industry standard	3A	~5mA	\$1.42
LM2596	73% ($I_{Load} = 3\text{ A}$)	1	Common; Easy to implement	3A	~5mA	\$2.34
LT8650S	93% ($I_{Load} = 3\text{ A}$)	2	Dual outputs; Synchronous; Silent Switcher 2 (Low EMI); Compact	4A	6.2μA	\$18.31
LTC3633	93% ($I_{Load} = 3\text{ A}$)	2	Dual outputs; Compact	3A	2mA	\$12.69

Table 3.31 Voltage Regulator Part Comparison

Sources: [98]-[102]

The *SOLARIS* power distribution network must step down the nominal 12.8 V LiFePo₄ battery voltage to two distinct rails: a 3.3 V rail to power the microcontroller and sensor suite, and a 5 V rail dedicated to the USB-A peripheral interface. To ensure the system meets its high-efficiency design goals, the above-listed table characteristics were considered.

As the primary objective of *SOLARIS* is to maximize solar energy harvesting and minimize parasitic losses, high conversion efficiency is the most critical metric. Synchronous buck regulators boast an increased efficiency of about 15% compared to asynchronous buck regulators.

The architecture favors dual-channel regulators. Utilizing a single Integrated Circuit (IC) to provide both 3.3 V and 5 V outputs reduces the total PCB footprint, simplifies the Bill of Materials (BOM), and streamlines the power routing logic compared to utilizing two discrete single-channel regulators. Based on power budget calculations, the 3.3 V logic rail requires a continuous average current of approximately 0.3 A. All candidate regulators

provide at least an order of magnitude of current headroom (typically 3 A or greater), ensuring high reliability, lower operating temperatures, and the ability to handle transient current spikes without voltage sagging. For a solar-powered autonomous system, minimizing quiescent current is essential. Quiescent current is the power consumed by the regulator in a non-switching or no-load state. Selecting a regulator with a low I_Q ensures that the system remains efficient during idle periods.

3.5.7.1 LM2679

The LM2679 is part of the "Simple Switcher" family and is characterized by its high-current capability, providing up to 5 A of output current [98]. Its primary advantage is a resistor-programmable peak current limit, allowing you to protect sensitive downstream loads or prevent battery surges by capping current between 3 A and 7 A. While it is highly efficient (up to 92%) and easy to design with, it is a single-channel device, meaning two separate units would be required to provide both 5 V and 3.3 V rails.

3.5.7.2 LM2576

The LM2576 is an older but extremely robust "Simple Switcher" capable of driving 3 A loads with excellent line and load regulation [99]. It operates at a relatively low switching frequency of 52 kHz, which results in very large filter components but offers high stability and ease of layout for beginner PCB designs. While it is commonly used in the industry, its larger footprint and lower efficiency make it less ideal for high-efficiency solar-powered robotics compared to the modern synchronous regulators.

3.5.7.3 LM2596

The LM2596 is a widely used, industry-standard step-down regulator that operates at a fixed switching frequency of 150 kHz [100]. Its key aspect is simplicity; it is internally compensated and requires very few external components, making it a reliable choice for general-purpose power regulation. However, because of its lower switching frequency, it requires larger inductors, and its efficiency (typically around 73%–80%) is notably lower than the synchronous alternatives.

3.5.7.4 LT8650S

The LT8650S is a high-performance, dual-channel synchronous regulator that is ideal for noise-sensitive applications due to its integrated "Silent Switcher 2" architecture, which drastically reduces electromagnetic interference (EMI) on the PCB [101]. Its most standout feature is an ultralow quiescent current of 6.2 μA while regulating, ensuring that the logic rail does not become a significant parasitic drain on the 20 Ah battery. By providing two 4 A channels in a single small package, it efficiently generates both the 5 V and 3.3 V rails while maintaining up to 94.6% efficiency even at high switching frequencies.

3.5.7.5 LTC3633

The LTC3633 is a dual-channel monolithic synchronous buck regulator designed for high-frequency operation up to 4 MHz, which allows for the use of very small external inductors and capacitors [102]. It supports a 3 A output current per channel and features a unique selectable phase shift between the two channels to reduce input and output ripple. This makes it a robust choice for compact designs where power density is a priority, though its quiescent current is higher than the LT8650S at 2 mA.

3.5.7.6 Voltage Regulator Part Selection

LM2679: LM2679 is selected for its ease of implementation and proven reliability, supported by extensive documentation. While it does not offer the highest efficiency among modern switching regulators, its larger pin pitch and through-hole package significantly simplify PCB assembly and prototyping. Its fixed 260 kHz switching frequency enables the use of smaller passive components, contributing to a more compact design. Furthermore, its 5 A maximum output current provides ample margin above the expected load requirements, ensuring robust and reliable operation under varying conditions.

3.5.8 Circuit Safety

Implementing a robust circuit safety strategy for *SOLARIS* is essential to protect the LiFePo₄ battery and the ESP32-S3 control logic from the significant current demands of the drivetrain. A 40 A inline fuse positioned as close to the battery's positive terminal as possible serves as the final line of defense against catastrophic short circuits. Given that the four goBILDA 5204 motors have a combined maximum stall current of 36.8 A, a 40 A rating provides a functional buffer that prevents "nuisance tripping" during high-torque maneuvers on Florida's uneven terrain.

Complementing the passive fuse, a manual disconnect switch provides a reliable mechanical "hard" isolation point for both emergency situations and long-term storage. This switch should be rated for a continuous current exceeding 40 A to minimize resistive voltage drops and prevent localized heating at the contact points. From a structural standpoint, mounting this switch on an accessible exterior portion of the robot chassis allows for an immediate "kill" command that physically breaks the battery's connection, bypassing all software logic. This dual-layer combined fuse and disconnect switch approach ensures that *SOLARIS* can safely manage its high-density energy reserves while protecting the sensitive navigation sensors and processing units from back-EMF or overcurrent events.

4 Standards and Design Constraints

4.1 Standards

4.1.1 Brief

The development of an autonomous, solar-powered robotic system requires strict adherence to established engineering standards to ensure hardware interoperability, data integrity, and operational safety. This section outlines the communication and wireless protocols selected for the *SOLARIS* platform. By utilizing standardized bus architectures and RF specifications, the system can reliably interface various high-precision components- such as the ESP32-S3 controller, the BNO085 IMU, and the INA228 power monitor- within a unified framework. The following subsections detail the physical and logical constraints of the I2C, UART, and SPI wired protocols, as well as the Bluetooth

Low Energy (BLE) wireless standard, providing the technical justification for their implementation in the final design.

4.1.2 Communication Protocol Standards

4.1.2.1 I2C

The I2C communication protocol is a software-based standard used for communication across digital systems. This inter-IC bus is defined as a simple bidirectional 2-wire bus for efficient inter-IC control, originally developed by Philips Semiconductors [103]. At a high level, the protocol allows multiple integrated circuits (ICs) to communicate using only two shared bus lines known as the serial data line (SDA) and the serial clock line (SCL) [103]. The bus is a true multi-controller system with built-in collision detection and arbitration procedures to prevent data corruption when more than one controller attempts to initiate a transfer simultaneously. Each device on the bus is recognized by a unique software address and operates within a controller/target relationship, meaning a controller can function as either a controller-transmitter or a controller-receiver [103]. By integrating the addressing and data transfer protocol directly on-chip, I2C enables software-defined systems and allows designers to prototype rapidly by simply connecting or disconnecting compatible ICs along the shared bus lines [103].

As discussed above, I2C is primarily a bidirectional bus but can also operate as a unidirectional bus depending on the communication mode. Bidirectional modes include Standard-mode (Sm), Fast-mode (Fm), Fast-mode Plus (Fm+), and High-speed mode (Hs-mode) [103]. I2C also supports an Ultra Fast-mode (UFm) which takes advantage of a unidirectional bus for applications such as LED controllers [103]. The different communication modes are compared below by their speed, directionality, driver type, and compatibility:

<i>Feature</i>	<i>Sm / Fm / Fm+</i>	<i>Hs-mode</i>	<i>Ultra Fast-mode</i>
Max Speed	100kHz / 400kHz / 1MHz	3.4 Mbit/s	5 Mbit/s
Direction	Bidirectional	Bidirectional	Unidirectional
Driver Type	Open-drain	Open-drain/Current Source	Push-pull resistors
Compatibility	Fully downward compatible	Downward compatible	Not compatible

Table 4.1 Communication Modes for I2C

Source: [103]

The five operating modes discussed above represent the bit-rate modes defined in the I2C standard. Beyond these, I2C protocol must also comply with other mandatory features and schemes. One such requirement is support for the START condition, STOP condition, and Acknowledge [103]. Combined with 8-bit oriented transactions using most significant bit (MSB) first ordering, these features make up a large portion of the standard [103]. All transactions begin with a START condition and are terminated by a STOP condition. A HIGH to LOW transition on the SDA line while SCL is HIGH defines a START condition, while a LOW to HIGH transition on the SDA line while SCL is HIGH defines a STOP condition [103]. These conditions are always generated by the controller. Since data is transferred in byte format, every byte placed on the SDA line must be eight bits long, and the controller must issue a START condition if it wishes to continue transmitting between individual byte messages [103]. As previously described, data is always transferred with the MSB first.

The START and STOP conditions only cover the transmitter side of the transaction. To complete the communication cycle, I2C also uses an Acknowledge (ACK) and Not Acknowledge (NACK) after every byte of data [103]. Through the ACK and NACK signals, the receiver is able to communicate back to the transmitter whether data was successfully received or not. In terms of physical bus behavior, during a transaction the transmitter releases the SDA line during the ACK clock pulse, allowing the receiver to pull SDA LOW and hold it stable through the HIGH period of the clock pulse [103]. When SDA remains HIGH during this ninth clock pulse, it is defined as a NACK signal [103]. Upon receiving a NACK, the controller has the opportunity to either generate a STOP condition to end the transaction or issue a repeated START condition to begin a new transfer [103]. Being largely software-defined, I2C remains one of the most versatile communication protocols available, supporting both bidirectional communication between two systems and unidirectional communication for simpler peripherals.

4.1.2.2 UART

Universal Asynchronous Receiver-Transmitter (UART) is a hardware-based communication protocol generally used for device-to-device data transfer. UART operates as a serial protocol where data is transferred bit by bit over a single wire, typically requiring only two wires total [104]. Unlike other communication methods, UART does not use a clock signal for synchronization. Instead, its asynchronous nature relies on a pre-configured baud rate to manage timing [104]. Both devices in the setup must be configured to the same baud rate, with an allowable difference of up to 10% before timing errors occur [104]. Unlike I2C, UART is hardware-specific and must be implemented at the hardware level to function properly. Despite this, UART remains one of the most commonly used protocols in embedded systems, microcontrollers, and computers due to its low cost and minimal wiring requirements. The standard defines a physical interface consisting of two wires, a transmitter (Tx) and a receiver (Rx), and supports only one controller and one peripheral in a full-duplex communication arrangement [104]. Data is delivered in a standard structure known as a packet, which consists of different components that work together to properly transmit and receive information.

Inside a device like a microcontroller, UART hardware acts as an intermediary between parallel and serial communication lines [104]. It receives data in parallel and wraps it into a packet containing key components such as a start bit, data frame, parity bit, and stop bit [104]. These components make up what is known as the standard UART frame structure. Packets are sent serially over the Tx wire, where the receiving UART performs the reverse process by stripping the frame bits and converting the serial data back into parallel [104]. For both devices to stay synchronized, several configuration requirements must be met: baud rates must match, an internal clock must be used to sample incoming data, and the hardware must be properly configured to define the communication parameters [104]. In a standard setup, the Tx pin of the first UART is connected directly to the Rx pin of the second UART and vice versa [104]. When not actively transmitting, UART lines are held in a high voltage state known as the idle state. This idle state allows one UART to take the role of the controller while the other acts as a peripheral to the transmission [104]. As mentioned, all of this is accomplished using just two wires to achieve reliable full-duplex communication.

Synchronization and timing play a key role in correctly communicating data between two pieces of hardware. UART achieves synchronization through start bits, baud rate matching, and timing tolerance [104]. As discussed, UART wraps communication data into a data frame that is assembled and disassembled by their respective devices. The key feature of this frame is the start bit, which signals when communication is allowed to begin [104]. During idle, when the transmission line is driven high, the controller device may drive Tx low for one clock cycle, alerting the receiver to begin reading bits. For this to work, both devices must have matching baud rates so that the sampling frequency matches the data transfer frequency [104]. The allowable difference between baud rates is limited to 10% before bits become too misaligned to be read correctly. To configure this timing, four major parameters must be determined: the Peripheral Clock (PCLK), oversample rate (OSR), baud rate divider (DIV), and fractional baud rate parameters (DIVM/DIVN) [104]. PCLK is defined as the input system clock speed at which an incoming signal can be sampled [104]. The general formula for calculating baud rate is shown below:

$$\text{Baud rate} = \frac{PCLK}{\left(\left(DIVM + \frac{DIVN}{2048}\right) \times 2OSR + 2DIV\right)}$$

4.1.2.3 SPI

The Serial Peripheral Interface (SPI) is a synchronous, high-speed, full-duplex communication protocol used to transfer data between a controller device and one or more peripherals [105]. Originally developed by Motorola in the 1980s, it has since become a de facto standard in embedded systems due to its simplicity and efficiency [105]. One of the defining characteristics of SPI is its shared clock signal, which ensures precise timing during data transfer. While it lacks the rigid standard format found in protocols like I2C, this actually gives developers the flexibility to tailor the protocol to their specific application [105]. A key strength of SPI is its bus architecture, which supports several standard configurations including independent chip select, daisy chain, and 3-wire SPI [105]. Each of these configurations highlights the customizability of SPI as a communication protocol for dynamic embedded applications. Rather than being defined by a strict data format, the standard is built around primary signals and four specific modes

of operation. SPI also benefits from being an exceptionally high-speed protocol, allowing for faster data transfer compared to other common serial protocols. Typical clock frequencies reflect this, as devices like the TXE8116/24 can operate with an SPI clock period of 100ns, corresponding to a frequency of 10MHz [105].

While SPI does not define a single rigid communication format, it does establish specific hardware signals, bus architecture, and data transfer modes that serve as operational standards. SPI is primarily defined by four hardware signals. Peripheral Input Controller Output (PICO) transmits data from the controller to the peripheral [105]. Peripheral Output Controller Input (POCI) transmits data from the peripheral back to the controller [105]. Serial Clock (SCLK) is a shared signal generated by the controller to synchronize data transfer [105]. Chip Select (CS) uses an active low signal from the controller to select and enable specific peripherals [105]. These four signals make up the foundation of the SPI bus architecture in a standard four-wire setup. Within this architecture there are a few notable configurations. In a daisy chain configuration, peripherals are connected in series, allowing the entire chain to be controlled by a single CS line while data shifts through each device sequentially [105]. A simplified three-wire SPI configuration combines PICO and POCI into a single bidirectional line called Data Input/Output (DIO), reducing the total wire count [105]. Because SPI allows for tailored formats, individual devices often establish their own data frame standards. One example is the TXE8116/24, which transmits the most significant bit (MSB) first over a fixed 24-bit transaction format [105].

4.1.3 Wireless Communication Standards

4.1.3.1 *Bluetooth Low Energy (BLE)*

[106] Bluetooth Low Energy (BLE) was introduced in Bluetooth Core Specification version 4.0 as a low-power alternative to Bluetooth Classic (BR/EDR). Bluetooth Classic provides higher bandwidth but was designed for continuous data streaming applications like audio, making it overkill for devices that only need to transmit small amounts of data intermittently. BLE was built for battery-operated devices that communicate in bursts, prioritizing energy efficiency over raw throughput.

[106] BLE divides the 2.400 to 2.4835 GHz ISM band into 40 channels, each 2 MHz wide. To improve reliability in environments with RF interference, BLE uses a spread spectrum technique called Adaptive Frequency Hopping (AFH). Three of the 40 channels (37, 38, and 39) are reserved as advertising channels and are used for device discovery and connection establishment. The remaining 37 channels are used for data transmission. BLE maintains a channel map that classifies each data channel as used or unused, and the frequency hopping algorithm avoids channels that are performing poorly due to interference, ensuring more reliable communication.

[106] GAP or the Generic Access Profile, defines how BLE devices discover each other and establish connections. GAP assigns two roles to communicating devices, the Peripheral and the Central. The Peripheral is typically an embedded device that advertises its presence by broadcasting packets on the three dedicated advertising channels. The Central is typically a smartphone or other host device that scans for these advertising packets and initiates a connection request when it finds a device of interest.

[106] Once a connection is established, communication moves off the advertising channels and onto the 37 data channels controlled by the Link Layer. At this point the two devices have formed a connection and can begin exchanging data. A single Central device is also capable of maintaining connections to multiple Peripherals simultaneously, which is managed at the Link Layer through the channel hopping schedule assigned to each individual connection.

[106] GATT, or the Generic Attribute Profile, defines how data is organized and exchanged once a connection has been established between two devices. GATT uses a client and server model where the Peripheral acts as the server, hosting data, and the Central acts as the client, reading or writing that data. Data on the server is organized into a hierarchy of Services and Characteristics. A Service represents a functional grouping of related data and each Service contains one or more Characteristics, which are the individual data values that can be read, written, or subscribed to by the client.

[106] The Bluetooth SIG predefines UUIDs for common Services and Characteristics such as battery level or device information, allowing devices from different manufacturers to interoperate without any custom configuration. For applications with unique data requirements, custom Services and Characteristics can be defined using application specific UUIDs, giving developers the flexibility to represent any data their application needs to exchange.

BLE has become one of the most widely adopted wireless communication standards in embedded and IoT applications due to its combination of low power consumption, native smartphone connectivity, and a flexible data model. Its layered architecture, from the physical layer up through GAP and GATT, allows developers to build wireless systems without relying on external networking infrastructure, making it well suited for a wide range of device-to-device communication uses.

4.1.3.2 IEEE 802.11

The IEEE standard 802.11 represents a comprehensive revision of the standard for wireless local area networks (WLANs). It consolidates technical corrections, clarifications, and enhancements to the Medium Access Control (MAC) and Physical Layer (PHY) functions of the previous versions [107]. The standard defines an architecture for interoperable wireless connectivity, addressing diverse operational environments including 2.4 GHz, 5 GHz, 60 GHz, and television white spaces [107]. Some key takeaways include expanded support for high-throughput (HT) and very high-throughput (VHT) stations, robust security, and sophisticated traffic management via Quality of Service (QoS) [107]. The standards are developed by the LAN/MAN Standards Committee of IEEE Computer Society [107].

The architectural framework of the 802.11 standard is composed of several fundamental building blocks that define how stations interact within a wireless network. The building block of a 802.11 LAN is known as the Basic Service Set (BSS) [107]. Within this there are many types but the first is Independent BSS which uses an ad hoc network where stations (STAs) communicate directly without an Access Point (AP) [107]. The next major

kind is an Infrastructure BSS where an AP provides local integration and access to a Distribution System (DS) [107]. Another kind is the Personal BSS (PBSS) where STAs can communicate directly and this is often seen utilized with multidirectional communication networks [107]. Lastly we have the Mesh BSS (MBSS) where a BSS forms a self-configuring network of mesh STAs [107]. While BSSs are a building block of this dynamic wireless network a proper Wi-Fi link must also use an interconnect and must be capable of larger coverage possibly formed by multiple BSSs. A Distribution System (DS) is used to interconnect BSSs to enable mobility but also integration with non-IEEE-802.11 LANs [107].

Enough about the architecture that makes up the 802.11 standard and now for a bit more about the logical services. The 802.11 standard defines connectivity and mobility required to be offered as a service. They must include association and reassociation which establishes a link to a STA and an AP but also allows for a STA to transfer to another AP when possible [107]. With the connectivity comes security concerns which are also a large portion of the standard. The standard provides a Robust Security Network Association (RSNA) framework which authenticates an STA, employs protocols to protect data from unauthorized users, and ensures that received data is from the claimed source and has not been sent maliciously [107]. The standard also defines measures for controlling power levels to meet regulatory requirements and minimize interference as well as enabling operation to remain within bands where radar interference may occur by dynamically switching channels [107].

4.1.4 Health & Safety Standards

4.1.4.1 IPC-2221

IPC-2221 standard is a general standard focused on Printed Circuit Board (PCB) design. There are several critical areas that the documentation covers in terms of design, component mounting, and interconnecting structures. These serve as foundational guidelines to produce detailed designs for active and passive components. The first critical area covered by this standard is design layout and requirements which outline the general parameters for physical features, including information hierarchy, end-product performance classes, and producibility levels [108]. Also covered by the standard is material selection and physical properties like that of guidance when selecting base materials and conductive materials like copper foils and various platings [108]. This also falls under the discussion about electrical properties of the power distribution system with electrical clearances being a key concern [108].

In general the IPC's standardization efforts are guided by principles adopted in 1995 [108]. These standards are designed to focus on end-product performance and minimize time-to-market. They do this by using simple data backed language while innovating rather than inhibiting newer manufacturing methods. It also stands as a method to standardize efforts to educate the potential relationships between purchasers and manufacturers. IPC-2221 is actually a "parent" document with a list of generic stands in conjunction for each specific technology [108]. These descriptions are then followed by an engineering description.

<i>Document Number</i>	<i>Focus Area</i>
<i>IPC-2221</i>	<i>Generic Design</i>
<i>IPC-2222</i>	<i>Rigid Organic Printed Board Structure</i>
<i>IPC-2223</i>	<i>Flexible Printed Board Structure</i>
<i>IPC-2224</i>	<i>Organic, PC Card Format</i>
<i>IPC-2225</i>	<i>Organic, MCM-L (Multichip Module)</i>
<i>IPC-2226</i>	<i>High Density Interconnect (HDI) Structure</i>
<i>IPC-2227</i>	<i>Discrete Wiring</i>

Table 4.2 Documentation Hierarchy for IPC-2221

Source: [108]

Mentioned above this standard focuses on product classification and producibility. Breaking them down into three main groups a user becomes responsible for identifying the class to which their product belongs along with the level of inspection and perfection required [108]. These classes are class 1 for general electronic products, class 2 for dedicated service electronic products, and lastly class 3 for high reliability electronic products. For example these products that you would find in class 1 are relative to function as that is the main requirement whereas cosmetic damage is acceptable. In reference to a class 2 item these products require high levels of performance and extended life where uninterrupted service is desired but not critical. On top of these classifications also falls different producibility levels. These can similarly be broken into three separate levels being level a, b, and c. In level a general design complexity is preferred since it is tied to lowest cost [108]. In level c we see the moderate design complexity desired since this is more of a standard acceptability level. Lastly we have level c which relies on high design complexity where producibility may be reduced and the penalty of having the highest cost [108].

One key consideration covered by this standard is an emphasis on testability review where designs must be considered for balance between cost and performance. Basic testing types may include functional tests where electrical design functionality is tested through connectors and test points [108]. Another type of basic testing type is known as the In-Circuit Test (ICT) which uses a “bed-of-nails” approach to find manufacturing defects which can be shorts, opens, and just simply using the wrong parts [108]. ICT has physically defined requirements due to its fixturing based mechanism. As defined by the IPC-2221 standard these fixtures must have test lands no smaller than 0.9 mm, grid placements of 2.5 mm centers, side access with secondary solder sides, and then clearance equal to 80% the components height at a minimum of 0.6 mm [108]. These test points must also be covered

with a solder-coating which is conductive, non-oxidizing coating, free of solder resist and markings [108].

This leads into another common standard described by IPC-2221 which defines materials selection and physical properties. For example material choice is driven mainly by temperature during operation and soldering, electrical properties, and structural strength [108]. Conductive materials and plating are generally between the choice of copper plating, gold plating, and nickel plating each with their own constraints. Copper plating must be the final surface and hole wall plating which requires 20 micro meters to 25 micro meters depending on performance class [108]. Gold plating can be used for edge-board connectors but excess gold above 0.8 micro meter in areas to be soldered can cause embrittled solder joints [108]. Lastly nickel plating acts as a barrier to prevent copper-tin diffusion and provides what's known as the “anvil effect” under gold contacts [108]. This also bleeds into thermal management and constraining materials. This is standardized where symmetrical construction is highly recommended to prevent warping. A symmetrical cored board uses two restraining cores as power and ground planes generally [108]. The drawback led by asymmetrical cores may cause distortion during soldering due to differing coefficients of thermal expansion [108].

4.1.5 Documentation Standards

4.1.5.1 *ISO-5807*

The ISO-5807 standard is in a grammatical guideline for technical diagrams which conveys standards for documentation symbols and conventions for flowcharts [109]. Though it was published in 1985 it still stands as a definitive guide for ensuring that a flowchart drawn across the world can be understood by many different engineers. The core components of the standard rely on the way it defines symbol usage and they can be broken into several major categories. This could include specialized data symbols for components and structures. It also characterizes connections and flow symbols to standardize how flowlines should be read with a top-to-bottom approach being inherent when known lines are present [109]. Other key conventions and rules introduced by the standard involve clarity over complexity where symbols should be uniform in size as much as possible, text should be inside symbols, and crossing lines should be avoided at all cost [109].

<i>Symbol</i>	<i>Meaning</i>
<i>Rectangle</i>	<i>Process; represents a single step or an operation that changes value, form, or location of data</i>
<i>Diamond</i>	<i>Decision; a point where a choice is made, usually resulting in a “Yes/No” path</i>
<i>Paralelogram</i>	<i>Data; represents input or output of information</i>
<i>Oval</i>	<i>Terminator; marks the start or end point of a process</i>

Table 4.3 General Symbols Guide for ISO-5807

Source: [109]

Above mentioned are some general guidelines for how flowcharts symbols can be standardized to be readable by many different people. Most modern flowcharting software like Visio have an ISO standard library pre-loaded making it one of the most universal standards for creation.

4.2 Design Constraints

4.2.1 Brief

The realization of the SOLARIS platform is governed by a complex set of design and practical constraints that balance theoretical performance with real-world engineering limitations. This chapter details the technical boundaries of the project, specifically focusing on the power distribution logic, signal integrity, and the economic factors influencing component selection. By identifying these constraints early in the design cycle, the team can ensure that the final autonomous system remains energy-positive, structurally sound, and financially viable within a student-oriented budget. The following sections outline the specific challenges associated with stepping down high-voltage battery power, managing electromagnetic interference from drive motors, and optimizing the cost-to-performance ratio of the mechanical chassis and electrical subsystems.

4.2.2 Power Constraints

The *SOLARIS* architecture must interface various peripherals and controllers that operate at various voltage rails, differing from the 12.8 V nominal output of the LiFePO₄ battery. To maintain a compact form factor without sacrificing performance, the system utilizes a dual-channel synchronous voltage regulator. This integrated approach provides the required 3.3 V and 5 V outputs from a single IC, achieving high conversion efficiency and minimizing thermal dissipation. Maximizing efficiency is a foundational constraint; to ensure the highest possible energy yield, the design incorporates oversized conductors to

mitigate I^2R (ohmic) losses and utilizes low-power sleep modes for non-critical components during idle periods.

Central to the energy harvesting strategy is the Genasun charge controller, which employs Maximum Power Point Tracking (MPPT). Unlike standard PWM controllers, MPPT dynamically adjusts the electrical operating point of the photovoltaic cells to extract the maximum available power under varying irradiance and temperature conditions. This technology is critical for ensuring that the battery gains consistently outweigh the system's baseline consumption.

Rigorous power budgeting is essential for mission success. A sensor or peripheral is only considered a net benefit if its operational value outweighs its energetic cost. To manage this, the architecture evaluates the use of MOSFET-based power switching to programmatically isolate high-draw components. While this allows for granular power control, it introduces specific engineering trade-offs: the addition of a MOSFET increases the series resistance (R_{DS}) of the power path and requires the allocation of limited General Purpose Input/Output (GPIO) pins from the microcontroller unit (MCU).

The actuation system introduces significant constraints regarding signal integrity. Motors generate substantial current spikes and inductive back-EMF during operation, which can inject electrical noise into shared power planes. To prevent interference with sensitive I^2C sensor data or RF communications, the motor power bus must be logically and physically isolated from the logic-level power pins of the sensors and MCU.

A primary constraint of the solar subsystem is the trade-off between solar alignment and actuation cost. While the photovoltaic effect is maximized when the panel is perfectly perpendicular to the sun, constant adjustments by the pan-and-tilt motors consume significant energy. The system is therefore designed to optimize net energy gain—the difference between energy harvested and energy consumed. Rather than continuous tracking, the motors are actuated only when the predicted increase in solar harvest significantly outweighs the power cost of the maneuver, ensuring the system remains energy-positive throughout the day.

4.3 Practical Constraints

4.3.1 Economic

One of the major constraints involved in designing and building *SOLARIS* has been analyzing and minimizing the overall cost of the system build that falls within normal student budget. While we had a full-sized team to split costs, the *SOLARIS* design involves having several components that, especially if accidentally overengineered, can incur a high cost. Because *SOLARIS* is designed for off-road movement, ensuring that our components meet our technical requirements, will not fail under stress and have sufficient margin became a significant factor when selecting and procuring items.

The biggest single cost item that we've worked to minimize is the overall *SOLARIS* chassis system. The chassis needs to meet several design requirements to ensure that the chassis

will not give under significant torque loads thus causing the motors to potentially spike in current usage unnecessarily thus causing excessive heating and power inefficiencies. The chassis must also be made of a material light enough to not incur an excessive contribution to the overall weight of the robot which would require larger motors, motor drivers, lower energy efficiency and even higher cost. The difficulty is in finding a pre-built frame that meets our design specifications for a low price. Research indicated that this was impossible. The economic compromise was designing the frame from scratch and using aluminum sheets to ensure *SOLARIS* had the necessary rigidity and robustness without incurring heavy costs. The cost, instead, was shifted over to the time required to research and design the chassis from scratch.

The second biggest cost factor for *SOLARIS* was the DC motor footprint on the project. To ensure that the robot could reliably navigate in the desired terrain with the expected weight, the drive DC motors were a significant design factor to attempt to minimize the cost of. Finding DC motors that were well-documented that met design requirements and were not >\$80 proved challenging. The robot required 4 motors, one for each wheel and needed to be able to support the estimated 20 lb weight of the overall robot as well as navigate it over the desired terrain. The desired terrain of *SOLARIS* is grass which, when evaluating practical environmental hazards, implies that there will be roots, patches of thick grass and sticks/twigs that the robot must successfully negotiate without stalling. The other two motors were the solar panel tilt and pan motors. These motors did not have such steep requirements but were selected to ensure that, in particular, the pan motor could support the slip ring arm, solar panel, and the tilt motor. To support these larger motors, higher cost drivers were required to support the current load.

The next area of concern was the overall power and power management system, particularly the voltage regulation and energy management. Since *SOLARIS* primarily operates from a battery charged from a solar panel, efficiency is an essential consideration. This consideration meant that we chose switching regulators for both the 3.3 V and 5 V regulators to minimize complexity and cost since we would only be purchasing a single dual-buck chip. The other consideration was minimizing the solar panel cost. To meet this goal we chose a material and model that has high efficiency, did not have an excessive size, and good thermal response while also not weighing an unnecessary amount which proved to be a challenge due to the demand of solar energy and the specialty material. Additionally we needed to purchase a battery that was small enough to fit in a reasonable robot size while also having the desired capacity and ease of integration.

5 Comparison of ChatGPT with other Similar Platforms

5.1 Brief

As generative artificial intelligence becomes an increasingly integrated tool in the engineering workflow, evaluating the efficacy and accuracy of different Large Language Models (LLMs) is essential for technical validation. This chapter presents a series of case

studies comparing OpenAI's ChatGPT against other industry-leading platforms, including Google's Gemini, xAI's Grok, and Anthropic's Claude. The objective of this comparative analysis is to determine which platform provides the most actionable, well-cited, and technically sound data for a complex robotics project like SOLARIS. By subjecting these models to prompts regarding battery chemistry, motor selection, and sensor specifications, the following sections analyze the variations in response depth, formatting, and the utility of the provided design frameworks.

5.2 Case Study 1

Please refer to Appendix B, [1a], [1b], and [1c] for the Case Study 1 prompt and the corresponding outputs generated by OpenAI's ChatGPT and Google's Gemini, both operating in deep research mode.

ChatGPT's response was concise and well-structured. The introductory paragraph clearly identified the five most common battery chemistries used in renewable solar applications, along with their key strengths and typical use cases at a high level. This was followed by one paragraph per chemistry, where each was analyzed systematically in terms of performance parameters and primary drawbacks. Notably, each paragraph included at least four in-text citations which made verification of the information straightforward and efficient.

The response by ChatGPT concluded with a horizontal bar chart comparing typical battery lifespans over time. This visualization reinforced ChatGPT's implicit conclusion that lithium iron phosphate (LiFePO_4) is the most favorable option, as it clearly demonstrated its superior lifespan relative to the other chemistries discussed.

In contrast, Gemini's response was significantly more detailed- approximately five times longer than ChatGPT's. It adopted a broader, big-picture perspective by examining the evolution of battery selection for renewable energy applications, including trends in the market from 2010 to 2024. While the introduction similarly listed the five chemistries covered, it remained more general and did not immediately compare them. Although this level of abstraction was not explicitly requested in the prompt, it did not detract from the overall quality of the response.

Each section in Gemini's response included tables that provided additional context, either by comparing chemistries directly or by evaluating their suitability for different applications. Furthermore, Gemini incorporated sections not present in ChatGPT's response, including a comprehensive comparison of economic factors, safety considerations, environmental impact, and installation requirements. The response concluded with a "Strategic Deployment Summary," which identified LiFePO_4 as the industry benchmark.

Both responses included a total of approximately 20 cited sources, offering strong opportunities for further research into battery chemistries. Although Gemini's response was less concise and included tables contrary to the prompt's instructions, it was ultimately

preferred due to its more comprehensive comparative analysis and its clear articulation of the advantages and disadvantages across all five chemistries.

5.3 Case Study 2

Please refer to Appendix B, [2a], [2b], and [2c] for the Case Study 2 prompt and the corresponding outputs generated by ChatGPT and Grok.

ChatGPT's response was extremely expansive and detailed in its analysis. It began by summarizing why DC motors are desirable in mobile robotics platforms. It went on to break down various motor types and their use cases citing controllability, efficiency and compatibility as a main reason. ChatGPT then went on to explain how DC motors are integrated into, and how they fit into the larger overall system.

ChatGPT then concluded its research by discussing various motor control technologies and methods commonly used in mobile robotics. It described basic controls such as PWM and H-Bridges for brushed DC motors and inverters for brushless motors. It also discussed more advanced control techniques such as using closed loop feedback control to allow the robot to maintain speed over any kind of terrain or laid that hinders its movement.

In contrast, Grok had a much shorter response but with far more pictures included in its explanation. It began by giving a succinct and directly to the point overview of the types of DC motors used in mobile robotics platforms but included a more technical and low level approach that included the use of schematics and drawings to aid understanding. The main difference here was Grok got more into the details on construction and technical operation rather than wider robot design.

Grok then talked about motor control methods and technologies. Grok spoke about H-Bridge and PWM control for Brushed motors and inverter control for brushless DC motors like ChatGPT did as well has discussed the use of PID, FOC and sensorless controls. THE main difference was Grok treated it as more of an assorted list of methods. One thing that Grok did include was a discussion of the use of AI based tuning for control systems.

Overall, ChatGPT's explanation was more generalized system level and application focused explaining how motor choices will affect a given project's performance. Grok's response was more based on component level descriptions, principles of operation and how they apply to the scope of mobile robots.

5.4 Case Study 3

Please refer to Appendix B, [3a] and [3b] for the Case Study 3 prompt and the corresponding output generated by Claude.

Claude [3b] provided a great starting point on what should be considered while looking for ultrasonic sensors for a design project. The response provided a comprehensive breakdown of the most critical selection criteria for an ultrasonic sensor including detection range, beam width, resolution, output type, (UART, I2C, PWM, etc..) update rate, environmental

sensitivity, power consumption, cross-talk, physical form factor, and cost and availability. While all of these factors were acknowledged and considered during the selection process, the response helped the team establish a clear prioritization, identifying range and beam width as the two most impactful criteria for *SOLARIS*' use case, while confirming that the remaining factors, though still relevant, were secondary concerns that the chosen sensor addressed.

Detection range was the first major filter applied. *SOLARIS* is a robot intended to navigate in the real world, meaning the sensor needed to provide enough advance warning for it to reroute before colliding with said obstacle. The MB1020-000 offers a maximum range of 6.5 meters which gives ample time for *SOLARIS* to see obstacles and change course. Claude framing range as a top tier concern reinforced that this should be a hard requirement rather than a preference.

Beam width was the second primary consideration and required the most analysis. The MB1020-000 has a beam width of 42 degrees, which the team evaluated mathematically to confirm it was appropriate for *SOLARIS*' geometry and obstacle detection goals. A beam that is too narrow risks missing obstacles that fall outside the sensor's cone, while a beam that is too wide can introduce picking up on objects not in front of *SOLARIS*. The 42 degree requirement was a good middle ground to choose from. Claude's response's called out the beam width as another top tier priority, which again reinforced this as another hard requirement.

Beyond the two primary criteria, Claude's response also prompted the team to verify other things such as the update rate, which turned out to be a non-issue for the MB1020-000. *SOLARIS* is designed to drive at slow speeds, and the minimum read time of 50ms means we should reasonably be able to read from each sensor multiple times per second. Although this was a small consideration, all of Claude's responses helped make sure we didn't miss any specification during the verification process.

Overall, Claude's [3b] response served as a useful decision making framework that aided in the sensor selection process.

5.5 Case Study 4

For this case study I chose to work with Google's Gemini 3 and OpenAI's ChatGPT. To test their functionality and capability they were each given the exact same prompt [4a] then after their response they were fed prompt [4d]. What can be discussed about their responses is a difference in the styles by which they answer the question. While it can be agreed that both Gemini and ChatGPT offer different structural approaches and focus areas to prompt [4a] they do both agree on fundamental strategies. These propositions are seen to be similar in response [4b] and [4c] where they each recommend similar features. Both make mention of the ULP co-processor as well as using strategies found across all embedded systems like sleep modes for reducing power draw during times of low usage. The main categories on which they differ can be seen in a few main areas. These main areas of interest are their presentation and tone, hardware and software focus, and communication strategies for implementation.

To start we look at Gemini 3 via its response [4b] to prompt [4a]. Gemini uses a narrative and analogical approach which compares the ESP32-S3 to a “high-performance sports car” which requires “careful management to avoid draining the tank”. It organizes information into five broad conceptual categories and includes a structured table for viewing them. This table compares sleep modes (Active, Light Sleep, Deep Sleep, and Hibernation). It also mentions the respective current draws and the use cases for each of these modes. In a similar fashion ChatGPTs response [4c] seems to provide a more granular, list-based, instructional guide with the way it responds. Rather than focusing on five major categories ChatGPT breaks the advice down into many smaller sections which include a “Typical Low-Power Workflow” and a list of “Common Mistakes”. The tone of it is directly focused on a more practical application and its implementation steps.

Moving on Gemini 3 communicates quite differently compared to ChatGPT in that it chooses more broad suggestions for how to communicate with other devices in its response [4b]. Gemini suggests a specific protocol alternative known as ESP-NOW as a wireless way to bypass the power heavy Wi-Fi handshake process. It also provides technical details on Wi-Fi Modem Sleep by explaining how it shuts down the radio between DTM beacons. Rather than continuing this process ChatGPT chose a much more explicit approach to solving an energy issue involved with communication. ChatGPTs response [4c] focuses on an explicit management of the radio stack in a code oriented manner. Some examples of this are providing specific API calls like `esp_wifi_stop()` and `esp_bt_controller_disable()` to ensure that the radios are off when not in use.

This behavior by ChatGPT is continued as it provides more code-level snippets for general GPIO optimization. ChatGPT provides function calls like `gpio_set_direction()` and `gpio_pullup_en()` to prevent floating pins from drawing power. In opposition to this Gemini highlights more of the software abstraction by zoning in on batching data as a strategy. Overall, this first prompt showed a key difference in structure where Gemini gives more narrative with a summary table while ChatGPT provides a detailed list with workflow and mistakes. Continuing, Gemini mentions hibernation mode for sleeping while ChatGPT focuses more on the Deep, Light, and Modem Sleep modes that the ESP32-S3 is capable of. ChatGPT goes as far as to mention specific tools to be used like shut resistors and a focus on voltage regulators with external power gating. Gemini misses the mark in this area by focusing on a more internal based solution which does not come across as complete.

As a follow up to prompt [4a] both Gemini and ChatGPT were given a second prompt [4d] to see how they would each continue on a talking point. Here the prompt regards C/C++ programming for low-power design on the ESP32-S3 in particular. Gemini’s response [4e] and ChatGPTs response [4f] both emphasize a shift from a continuous execution style to a more event-driven, sleep-oriented programming style. However they do both differ significantly in their level of detail, specific technical recommendations, and overall structure. These categories that they differ in can first be defined as the way they frame their argument with a different approach in depth and breadth of responses. Each model also provides a unique focus with separate technical recommendations and programming practices.

Gemini poses its response [4e] as a challenging transition from “Linear Programming” to “Event-Driven Programming” by focusing on the idea of a chip spending minimal time in an active state. ChatGPT follows the same train of thought with its response [4f] by focusing on a program that follows a wake sleep cycle for maximum efficiency. ChatGPT deepens this understanding by offering a list of 15 concise points. It covers a wider range of practical firmware tasks, such as compile-time optimizations and reducing UART logging to save power. Gemini responds in a more detailed manner [4e] by including mathematical formulas for power saving offering a solution like frequency scaling and a detailed explanation as to why ESP-NOW is more efficient than standard Wi-Fi due to bypassing handshakes.

Both AI models recommended different practices and also API calls that would both be useful for running power responsible code. In ChatGPT’s response [4f] it focuses on a logging management system where it suggests disabling logs in production because UART communication can be power-hungry. It follows up this idea with a fast boot strategy where it advises skipping unnecessary initializations and caching configurations in RTC memory to speed up the boot process after deep sleep reset. Lastly it recommends disabling modules during no use or sleep scenarios. Gemini’s response [4e] tries to drive the importance of ESP-NOW even though it is irrelevant for communication with client devices and only works for ESP to ESP communication. It mentions allowing FreeRTOS to trigger low-power states automatically during idle tasks as well as isolating pins to prevent current leakage during sleep.

Overall the main categorical differences between both models breaks down to organizational style, technical depth and breadth, and specific programming focus. Gemini prefers a more event-driven programming style as opposed to ChatGPT’s sleep as the default state approach. ChatGPT prefers to provide a comprehensive checklist of actions while Gemini likes to explain the why with physics and math involved. Each model provides its own unique insights to the project but overall ChatGPT offers more usable and user friendly advice as it comes from the perspective of a firmware developer choosing to approach each problem with a software solution as opposed to Gemini’s frequent physical and hardware applications.

6 Hardware Design

6.1 Brief

The schematic design for *SOLARIS* is centered on a modular architecture that integrates high-precision sensing, efficient power regulation, and optimized motor control through a central microcontroller hub. At the core of the system, an ESP32 manages all peripheral communication and hosts a web server for mobile interaction, utilizing a partitioned layout that includes a dedicated USB-to-UART programming bridge and multiplexed analog inputs to conserve I/O resources. The sensing suite employs an INA228-Q1 for ultra-precise battery telemetry and a network of phototransistors for solar tracking, while the power stage utilizes LM2679-based regulators to provide stable 3.3V and 5V rails. To drive

the robot's mobility and solar-tracking motors, the design incorporates VNH5019A-E H-bridge drivers featuring a hardware-inverted logic scheme that ensures bidirectional movement while further minimizing the microcontroller pin count. Together, these schematics form a cohesive electrical ecosystem engineered for maximum energy efficiency and reliable autonomous operation. Furthermore, the report examines the mechanical chassis, detailing its structural design and physical integration within the overall system.


6.2 Schematic Diagrams

6.2.1 MCU

The MCU schematic serves as the central hub for the *SOLARIS* system, managing all peripheral interconnects and hosting the web server required for the integrated mobile application. To ensure clarity, the schematic is partitioned into three distinct sections within this report: the primary microcontroller and its supporting circuitry, the dedicated USB-to-UART programming bridge, and the comprehensive array of connectors that interface with other *SOLARIS* subsystems.

The microcontroller schematic features an integrated circuit with grounding pads on its underside connected to the system ground plane to ensure electrical stability and heat dissipation. Decoupling capacitors are placed in close proximity to the voltage input pins to filter out high-frequency noise and provide a stable supply for the module. A dedicated manual switch is connected to the enable pin to allow for physical control over the power state of the unit. Additionally, the communication lines used for the data bus are both pulled up by ten thousand ohm resistors to maintain the necessary signal levels for reliable inter-chip communication. Pin IO46 serves as an active-low interrupt source, connected through J30 to a manual push button for physical activation.

Two test circuits featuring light-emitting diodes are included to assist with hardware and software diagnostics across the system. The first circuit consists of a manual button that serves as a verification point for the presence of the primary 12.8 V battery supply on the main board. The second circuit is an indicator controlled by a transistor that can be driven by the microcontroller to confirm successful firmware uploads through a simple blinking sequence. These visual indicators provide essential feedback during the initial assembly and testing phases of *SOLARIS*.



The schematic for the USB to UART bridge is centered on a single chip that facilitates data transfer between the host computer and the controller. The hardware interfaces with the host via a USB AB port, utilizing standard differential data lines for signal transmission. This connection also includes a dedicated sense line designed to detect the presence of the bus and verify that the device is actively linked to the host system. Power for the logic and internal oscillator is derived from the bus voltage using an integrated regulator, which provides a stable supply for the system components. To ensure reliable performance, grounding pads on the underside of the package are tied to the shared ground plane, and decoupling capacitors are strategically placed at the power input pins to eliminate electrical noise.

Data communication with the microcontroller is handled through the transmit and receive lines, while hardware flow control is supported by a single control signal. For this

implementation, the four general purpose input and output pins identified as GPIO.0, GPIO.1, GPIO.2, and GPIO.3 are designated as no connects and left floating. Additionally, several modem control pins including the ring indicator, data carrier detect, data set ready, and clear to send lines are also left unconnected. This configuration simplifies the circuit layout while focusing on the essential signals required for high-speed serial communication between the controller and the host computer.




Figure 6.2 MCU Programming Bridge Schematic

The microcontroller connectors schematic incorporates a wide array of connectors to facilitate communication between the various peripheral boards in the system. A dedicated IMU component serves as the interface for a pre-assembled inertial measurement unit PCB, ensuring precise orientation data. Because the microcontroller has limited I/O pins, the design utilizes two multiplexers to manage analog signals. This configuration enables the system to independently sample eight light sensors using four I/O pins and four ultrasonic sensors using only three I/O pins, effectively conserving the microcontroller's hardware resources.

Physical routing for power and data is handled through a combination of headers and high-current connectors. Standard header pins provide the necessary power and signal links for the ultrasonic sensors, light sensors, voltage regulators, and energy monitor. For the main power input, a terminal block connector securely brings the 12.8V supply into the circuit. This central power rail feeds directly into the voltage regulators, which stabilize the energy for the rest of the electronic components.




Figure 6.3 MCU Connectors Schematic

6.2.2 Sensors

The energy monitoring subsystem is implemented using the Texas Instruments INA228-Q1, a 20-bit ultra-precise power and energy monitor communicating over the I2C protocol. The device is configured in a high-side sensing topology, with a 15 milliohm shunt resistor placed on the main power line between the LiFePO4 battery and the system load. This resistor value was selected to utilize 73 percent of the device's narrower ADC input range of plus or minus 40.96 millivolts at the system's expected maximum current draw of 2 Amperes, which accounts for four drive motors drawing 0.3 Amperes each along with 0.3 Amperes of overhead from the main controller and sensing circuitry. Operating at 73 percent of full scale ensures strong measurement resolution while maintaining adequate headroom for transient motor startup currents that may briefly exceed nominal operating conditions.

The INA228-Q1 interfaces with the system microcontroller via a two-wire I2C bus, with both the SDA and SCL lines pulled up to the 5-volt supply rail through 10 kilohm resistors, which satisfies the open-drain requirements of the I2C protocol and ensures reliable logic level signaling. A 100 nano-farads decoupling capacitor is placed directly at the VS supply pin of the device, following the manufacturer's layout recommendations, to suppress high-frequency noise introduced by the switching motor loads in the system. The device continuously accumulates charge and energy data in its internal 40-bit hardware registers, enabling real-time coulomb counting for state of charge estimation of the battery without requiring significant computational overhead from the main microcontroller. This configuration allows the system to accurately track both energy flowing out of the battery during motor operation and energy returning during solar recharge cycles.




Figure 6.4 Energy Monitor Schematic

The light sensing subsystem utilizes 8 Vishay TEPT5600 silicon NPN phototransistors, which are sensitive to visible light across a spectral range of 440 to 800 nanometers with a peak sensitivity at 570 nanometers, closely matching the responsivity of the human eye. The device is configured in a transimpedance topology where the collector is connected to the 3.3-volt supply rail, and the emitter is connected to ground through a 10-kilohm load resistor. When incident light strikes the phototransistor, it generates a collector current proportional to the illuminance level, which at 100 lux under standard conditions produces a typical photo current of 630 microamperes. This current flowing through the 10-kilohm emitter resistor develops a voltage that rises proportionally with increasing light intensity, providing an analog output signal suitable for digitization.

The voltage developed across the 10-kilohm resistor is routed directly to an analog to digital converter input on the system microcontroller, allowing the firmware to sample and interpret the ambient light level in real time. The TEPT5600 has a narrow acceptance angle of plus or minus 20 degrees at half sensitivity, which helps the sensor reject off-axis ambient light and focus its response on the light source directly in front of it, making it well suited for detecting the intensity of solar irradiance in the context of *SOLARIS*. The 3.3 volt supply rail was chosen to remain within the 6 volt collector emitter breakdown voltage limit specified in the datasheet, and the 10 kilohm resistor was selected to keep the maximum collector current well below the 20 milliamp absolute maximum rating under any foreseeable illumination condition the sensor may encounter during outdoor robot operation.




Figure 6.5 Light Sensor Schematic

6.2.3 Power

The power system for *SOLARIS* is a critical architecture designed to maximize battery charging efficiency while strategically managing energy for all other system components. To provide the necessary supply levels for various subsystems, the design features 3.3 volt and 5-volt power rails regulated by the LM2679 adjustable integrated circuit. Both regulator stages were developed using Texas Instruments WEBENCH software and are integrated onto a dedicated daughter board that interfaces with the main PCB via a four-pin header.

The LM2679 utilizes a specialized architecture to provide efficient step-down regulation, starting with the Switch Output. This serves as the primary switching node and operates alongside an external inductor and Schottky diode to manage energy storage and current return paths. The Input provides the supply voltage and requires bypass capacitors to filter high frequency switching noise, while the Boost connection employs a small ceramic capacitor as a charge pump to ensure proper gate drive.

For electrical stability, the Ground provides a low impedance common reference and thermal dissipation path, while the Current Adjust allows for programmable safety thresholds. Both regulators use a 12.54k ohm resistor to establish a 3-ampere peak current limit. Regulation is achieved via the Feedback connection, which monitors the output voltage through a resistor divider to inform the internal control loop. The resistor voltage divider is set to a ratio of 1.74 to maintain the target 3.3V output. Finally, the Soft Start uses a 15nF capacitor to ground to define a 3.1 millisecond startup ramp, which is essential for preventing excessive inrush currents from sagging the input rail or tripping protection fuses. An integrated LED indicator circuit is tied directly to the regulator's output stage to serve as a diagnostic verification tool. This feature provides immediate visual confirmation that the regulated power rail is active and delivering the intended voltage, hastening the troubleshooting process.




Figure 6.6 3.3V Regulator Schematic

The fundamental design characteristics and pin behaviors detailed for the 3.3V regulator remain consistent when implementing the 5V regulator version. Since the underlying switching architecture is the same, the main system changes involve the feedback loop, where the resistor voltage divider is set to a ratio of 3.16 to maintain the target 5V output. Additionally, some peripheral components, such as the output filter capacitors, are resized or selected for higher voltage tolerances to ensure the rail remains stable and noise-free under different load conditions.




Figure 6.7 5V Regulator Schematic

The motor driver schematic centers on the VNH5019A-E integrated circuit, which serves as the primary interface between the high-power battery and the logic controller. This architecture integrates the main 12.8V battery with the motor while utilizing a dedicated 5V logic rail to power the driver's internal circuitry. To ensure a stable electrical reference and signal integrity across the entire system, it is vital to tie all grounds from the battery and the controller together. Motor speed is then precisely regulated through a PWM signal transmitted directly from the controller to the driver hardware.

Directional control is achieved via the driver's two input pins. While the driver is typically capable of a dedicated braking mode when inputs are matched, this specific design omits that function to optimize microcontroller resource allocation. By tying the primary direction signal to an inverter that feeds the secondary input, the two pins are hardware-locked into opposite states. This configuration ensures the consistent polarity required for bidirectional movement while significantly reducing the required I/O pins on the microcontroller. The following table outlines the specific operating modes derived from these input configurations.

IN _A	IN _B	DIAG _A /EN _A	DIAG _B /EN _B	OUT _A	OUT _B	CS (V _{CS} D = 0 V)	Operating mode
1	1	1	1	H	H	High imp.	Brake to V _{CC}
1	0	1	1	H	L	I _{SENSE} = I _{OUT} /K	Clockwise (CW)
0	1	1	1	L	H	I _{SENSE} = I _{OUT} /K	Counterclockwise (CCW)
0	0	1	1	L	L	High imp.	Brake to GND

Figure 6.8 Motor Driver Operating Truth Table

The diagram illustrates a motor driver circuit using the VN75019ATR-E IC. The IC is a 16-pin component with the following pin configuration:

- 1: OUTA
- 2: NC
- 3: VCC
- 4: INA
- 5: ENA/DIAGA
- 6: CS_DIS
- 7: PWM
- 8: CS
- 9: ENB/DIAGB
- 10: INB
- 11: CP
- 12: VBAT
- 13: VCC
- 14: NC
- 15: OUTB
- 16: OUTB

The circuit includes the following components and connections:

- Power Supply:** A 12.8V supply is connected to the VCC pin (pin 3) and the VBAT pin (pin 12). A 5V supply is connected to the VCC pin (pin 13).
- Control Header (U3):** A 4-pin header with pins 1, 2, 3, and 4. Pin 1 is connected to 12.8V, pin 2 to GND, pin 3 to DIR, and pin 4 to PWM_MOTOR.
- Motor Driver IC (U1):** VN75019ATR-E. The OUTA pin (pin 1) is connected to the motor's MOTOR+ terminal. The OUTB pin (pin 16) is connected to the motor's MOTOR- terminal.
- Passive Components:**
 - Resistors: R1 (1K), R2 (1K), R3 (3.3k), R4 (1K), R5 (3.3k), R7 (3.3k).
 - Capacitors: C1 (0.1uF), C2 (0.1uF), C3 (470uF), C4 (0.1uF), C6 (0.1uF), C7 (10nF).
 - MOSFET: Q1 (N-Channel, 30A).
- Other Connections:**
 - The DIR pin (pin 4) is connected to the gate of the MOSFET (Q1).
 - The PWM_MOTOR pin (pin 4) is connected to the gate of the MOSFET (Q1).
 - The GND pin (pin 2) is connected to the common ground.
 - The MOTOR+ and MOTOR- terminals are connected to the motor.

6.3 Mechanical Design

The mobility system consists of four drive motors mounted at each corner of the robot base, providing a stable four-point drive that allows the robot to navigate varied outdoor terrain

while maintaining a low center of gravity beneath the solar panel payload. Atop the chassis, a two-motor pan and tilt mechanism is responsible for orienting the 20 watt solar panel, with one motor driving azimuthal rotation and a second motor controlling the elevation angle of the panel relative to the sun. This dual-axis articulation allows the system to actively track solar position throughout the day, maximizing the energy harvested and delivered to the onboard LiFePO4 battery. The overall chassis architecture prioritizes accessibility and repairability, as the removable PETG panel covers allow any internal subsystem including the battery, power electronics, and sensing hardware to be accessed and serviced without disassembling the structural frame.



Figure 6.10 SOLARIS Chassis Without Covers




Figure 6.11 SOLARIS Chassis With Covers

7 Software Design

7.1 Brief

The *SOLARIS* software architecture is designed to integrate autonomous environmental sensing, real-time manual control, and robust mobile user interface. At its core, the software manages the transition between distinct operational states which can range from the efficiency of hibernation to the complex auto solar tracking used to ensure maximized energy acquisition. The chapter is broken into four primary areas which cover the many facets of the software's reach. It covers state and case diagrams which deep dive into the behavioral logic of the robot which include the “Wandering” and “Alignment” routines with a logical user interface (UI). We also cover the autonomous navigation to break down the process of analysis given sensor driven decision making which allows the device to avoid obstacles and seek optimal light sources. It is also important to define the data transfer that takes place over Bluetooth Low Energy (BLE) which bridges the robot's telemetry data with a user device. Lastly we walkthrough the mobile application which focuses on user authentication, real-time telemetry visualization, and manual control systems used to drive the robot. By mapping these software components, we establish a blueprint for a system that is not only reactive to its environment but also remains transparent and controllable for the end-user.

7.2 Case and State Diagrams

7.2.1 User State Diagram




Figure 7.1 Software State Diagram with User Interaction

The basic logical state diagram above describes the base level software architecture for *SOLARIS* functionality. Behavioral logic of the device can be broken up into four main categories being Power Off, Searching, Manual, and Autonomous states. In the first state of Power Off the ESP32 is booted at the toggle of a physical switch and Bluetooth begins to advertise to a potential user. After this we enter the Searching phase where the remote

device has been paired and motor drivers are enabled. Staying in manual state allows for sensors to be read to update the app and receive movement data to move the robot. If during the manual control period autonomous mode is selected we then move to autonomous functionality. The Autonomous state is characterized by its lack of manual overrides and its sensor based pathing. During the Autonomous state a toggled manual can be asserted where the autonomous routine is topped and user input is listened for. From any state the toggle of a physical switch will move the device into a power off state where all operation ceases and power is cut to the ESP32.




Figure 7.2 Software State Diagram with Phone User Interface

Figure 7.2 shows the control flow of the *SOLARIS* companion app. On launch, the app checks whether a cached Clerk token exists. If one is found, the user is brought directly to the sign-in page to validate their credentials. If not, the app presents the home screen where the user can choose to sign in or sign up. The sign-up flow collects an email and password, sends a validation code to the user, and registers the account in Clerk upon successful verification. Sign-in through Google or Facebook is also supported.

Once credentials are validated, the app attempts to connect to *SOLARIS* over BLE. If the connection succeeds, the app begins receiving telemetry data from the robot and forwarding it to the backend, then navigates to the Telemetry screen. If the BLE connection fails, the user can still view previously cached telemetry data, but mode switching is disabled.

From the main interface, the user can navigate between the Telemetry and Mode tabs. The Telemetry tab displays all robot data. The Mode tab fetches the current operating mode

from *SOLARIS* over BLE and reflects it as the selected option. If the user changes the mode, the new selection is sent to the robot over BLE. If Manual mode is selected, two joysticks appear on screen and throttle and steering data are sent to *SOLARIS* in real time. Pressing the sign out button from any state returns the user to the home screen.

7.2.2 Autonomous Navigation State Diagram




Figure 7.3 Autonomous Navigation Sub-User State Diagram

The table above shows exactly how the autonomous functionality of *SOLARIS* is broken into many logical states. We begin with Clock Syncing where we see our first fork in the road after receiving a Bluetooth time packet from the app. Based on the current time the robot chooses either to wander for better light or hibernate until a proper time has arrived. Then as we begin to wander, we are within a safety loop to keep the robot from running into objects but also navigate dynamically. The robot searches for light but is interrupted by an ultrasonic sensor reading when approaching an obstacle. When the phototransistor readings become satisfied or the search timer threshold has been reached the robot will stop and begin to align its solar panel automatically. After angling its solar panel correctly or running out of time it will go back into hibernation where it will wait for the next wake signal to bring it back out of hibernation.

7.2.3 Software Case Diagram

Case Category	Specific Condition (Input)	Resulting State	Hardware Behavior
I. Connectivity	No Bluetooth paired	Search/ Advertising	ESP32 LED blinks; Motors locked; System initialized on startup
	Bluetooth paired	Standby	Await Mode command (Manual/Auto)
II. Control Mode	Mode = “Manual” (via App)	Manual	Listen to controls; Service BLE with telemetry
	Mode = “Stationary” (via App)	Hibernation	Go to sleep in this spot
	Mode = “Auto” (via App)	Autonomous	Move to Autonomous case
III. Autonomous	Time = Night (from RTC)	Hibernation	Enter Deep Sleep immediately
	Time = Day & Obstacle < 20cm	Obstacle Avoidance	Stop; Pivot using IMU and ultrasonic sensors; clear path
	Time = Day & Obstacle > 20cm	Solar Search	Move to Goal Evaluation case
IV. Goal Evaluation	Light < Threshold	Wandering	Moving forward; Keep IMU heading
	Light > Threshold	Solar Alignment	Stop wheels; Activate Pan/Tilt motors
	Alignment = Success	Hibernation	Shutdown to charge; Set wake timer
V. Safety/System	Battery < 10%	Critical Stop	Send alert to app; kill all motors
	Robot is unable to move	Critical Stop	Send alert to app; kill all motors

Table 7.1 SOLARIS Software Case Overview

The case logic for *SOLARIS* can be broken into five main categories as shown above. Our system starts with its very first decision which determines if the system will begin to run. This is decided during the Connectivity case where the system must pass through an initialization state and then fall into an autonomous mode or manual mode for continued control of the robot. If the robot is put into autonomous mode it does not just drive but it enters a sub-case structure defined by the time and environment settings. Here the robot will check for the time gate to make sure it is not moving during the night

time or to make sure that it does not run into obstacles. These states are followed by a goal to reach the appropriate light intensity for the robot to stop moving. When the robot finally comes to a stop the robot begins to angle the phototransistor array to find an equalizing of electrical factors across each phototransistor. When the threshold is met the solar alignment phase finishes and moves into a hibernation stage. While this is all happening a data feedback loop runs in parallel to output data via a BLE request to a connected phone client. This case diagram illustrates the main loop on a higher level to understand what phases the robot goes through in a means of functionality.

7.3 Data Transfer

7.3.1 BLE Data Transfer




Figure 7.4 Bluetooth Low Energy (BLE) Data Transfer (Client to ESP)

The BLE data transfer serves as a digital nervous system for our project by facilitating high-speed and low power bridge between the user client and the machine logic. During the BLE exchange a user client may push three primary data packets to the robot to dictate its behavior and synchronize its internal clock. The robot can be controlled via a mode selection, moved via a joystick functionality, and lastly is updated with a Unix time packet to synchronize a real time clock. On the other side the ESP32 pushes real time data from internal storage to the mobile interface. The levels to be monitored involve temperature data for ESP32 CPU temperature to prevent thermal throttling. Another key insight delivered by the ESP32 is a battery level as a percentage to analyze critical low points. Lastly we look at navigational and performance metrics like distance traveled by the robot given its last shutdown and a net power draw which gives key insights into losing or gaining energy.

7.4 User Interface



Figure 7.5 SOLARIS Phone App Home Screen

The home page presents two options: Sign In and Sign Up. Tapping either button navigates to the respective authentication screen. If the app detects an existing session, it skips the home page entirely and routes the user directly to the main dashboard.




Figure 7.6 SOLARIS Phone Sign Up Screen

The sign-up screen offers two paths, OAuth via Google or Facebook, or manual registration with an email and password. For manual sign-up, the user enters their email, creates a password, and confirms it. Authentication is handled by Clerk. Dynamic inline alerts provide real-time feedback for input errors such as invalid email formats or passwords that don't meet minimum length requirements.




Figure 7.7 SOLARIS Phone Verification Functionality

After submitting valid credentials, the user is prompted to enter a 6-digit verification code sent to their email. The verify button remains disabled until all 6 characters are entered. If the code wasn't received, a resend option is available, though it's rate-limited once every 2 minutes. Once a valid code is submitted, Clerk verifies it server-side and the user is granted access as a registered account.




Figure 7.8 SOLARIS Phone Sign In Page

The sign-in screen mirrors the sign-up flow, offering OAuth via Google or Facebook or manual entry with a registered email and password. Dynamic alerts surface credential errors in real time, such as an incorrect password or unrecognized email. Once Clerk verifies the credentials, the user is routed into the authenticated portion of the app.



Figure 7.9 SOLARIS Telemetry Page

The Telemetry screen is the primary monitoring interface of the *SOLARIS* companion app, accessible from the bottom navigation bar. A 3D render of the robot sits at the top of the screen, with the four telemetry cards arranged in a 2x2 grid below it, showing Battery Level, Distance Traveled, Net Power, and CPU Temp. Each card is independently swipeable, letting the user cycle through different time window views for that metric. A row of dot indicators at the bottom of each card tracks the current position. The available time windows differ depending on the metric. Battery Level has two views, current and past day. Distance Traveled has five, covering past hour, past day, past week, past month, and all time. Net Power has four, covering current, past hour, past day, and past week. CPU Temp has three, covering current, past hour, and past day.

Current-value views show a large numeric readout with a label above it. The Battery Level card also includes a horizontal progress bar that is based on the current charge percentage. Both the bar and the percentage text change color dynamically, green above 90%, yellow between 21% and 89%, and red at 20% or below to signal a low battery. The other current views display their values in their respective units, watts for Net Power and degrees Celsius for CPU Temp.

Historical views replace the numeric readout with a bar chart, where each bar represents a time-bucketed average with timestamps or day labels along the x-axis. The Net Power chart spans both positive and negative values on the y-axis, since net power can go negative when consumption exceeds solar generation. All charts share a consistent dark blue style matching the overall theme of the app.




Figure 7.10 SOLARIS Mode Switch Page

The Mode screen allows the user to switch between the three operating modes of *SOLARIS*. Stationary, Automatic, and Manual. The active mode is highlighted with an orange border and orange text, while inactive modes use a plain dark button style.

Selecting Stationary or Automatic presents no additional controls on screen, as those modes are managed by the robot itself. Selecting Manual reveals two joysticks in the lower half of the screen. The left joystick is constrained to the y-axis and controls throttle, and the right joystick is constrained to the x-axis and controls steering. Each joystick displays its current value as a decimal readout below it, defaulting to 0.00 at rest. When the user interacts with the joysticks, the throttle and steering values are sent to the robot over BLE, allowing direct manual control from the app.

8 System Fabrication/Prototype Construction

8.1 Brief

The PCB design for the *SOLARIS* platform is built on a modular architecture that prioritizes electrical efficiency, signal integrity, and structural rigidity across its various subsystems. At the heart of the system is the primary MCU board, which manages central processing via an ESP32 and an inertial measurement unit while utilizing multiplexing to coordinate a high density of peripheral sensors. This central hub interfaces with a suite of specialized daughterboards, including phototransistor-based light sensors for autonomous

navigation, an INA228-equipped energy monitoring module for comprehensive power telemetry, and unified voltage regulators that support both 3.3V and 5V rails. High-current tasks are handled by custom VNH5019A-E motor driver boards, which incorporate conservative trace widths and enhanced thermal management to ensure reliable motor control. By utilizing a strategic combination of header connections, robust terminal blocks, and mechanical standoffs, the hardware suite ensures seamless communication and power delivery while remaining scalable for future system development.

8.2 Sensor Boards

The *SOLARIS* platform utilizes two specialized sensing boards designed to optimize system performance and autonomy. The first is a dedicated energy monitoring module tasked with tracking the power consumption and health of the onboard LiFePO₄ battery, while the second monitors ambient sunlight intensity to provide data for the system's autonomous movement algorithms. The energy monitor functions by routing battery current through a 20-milliohm shunt resistor to eight independent output ports. To ensure electrical efficiency and safety, high-current copper traces were designed with a minimum width of 2 mm, a specification verified by trace-width calculations to support a continuous current of approximately 4 amperes.

The system measures current consumption by detecting the small voltage drop developed across a four-terminal shunt resistor, which is then processed by an INA228 monitoring IC. Once configured via I2C, the IC provides the microcontroller with a comprehensive suite of electrical telemetry, including shunt and battery voltages, die temperature, current, power, and accumulated energy and charge. This module operates as a daughterboard to the primary MCU, interfacing through a header connection that provides 3.3V logic power, ground, and the I2C communication bus.




Figure 8.1 2D Energy Monitor Board Layout




Figure 8.2 3D Energy Monitor Board Layout

The light sensor PCB features a streamlined design optimized for high-density integration and reliable data acquisition. Each module utilizes a phototransistor and a SMD resistor in a compact configuration that facilitates a small physical footprint. This minimalist

approach is intentional, allowing the sensors to be easily fitted into the various constrained locations required by the platform's mechanical architecture without interfering with other subsystems.

The deployment strategy involves eight independent boards strategically mounted to maximize the system's sensing capabilities. Four PCBs are positioned at the top corners of the rectangular robot chassis, while the remaining four are secured to the sunlight tracker module using M3 bolts. Each board interfaces with the central microcontroller via a three-pin header providing 3.3V power, ground, and analog voltage output. This PCB ensures seamless communication and processing of light intensity values back to the MCU.




Figure 8.3 2D Light Sensor Board Layout




Figure 8.4 3D Light Sensor Board Layout

8.3 Power Boards

The 3.3V and 5V voltage regulator modules utilize a common integrated circuit architecture, allowing for a unified PCB design across both voltages. This approach optimizes the project's bill of materials, as both regulators are assembled from the same board layout by simply adjusting the values of specific passive components. To ensure stable operation and mitigate high-frequency noise, decoupling capacitors for the 12.8V input are positioned near the IC, strictly adhering to the manufacturer's datasheet recommendations.

Visual telemetry is provided by an indicator LED located at the lower-left corner of the PCB, which confirms that the regulator is maintaining the target output voltage. The electrical interface is managed via a four-pin header—consisting of the voltage output, dual ground pins, and the primary input voltage. The regulators boards are fastened to the main MCU board using dedicated mounting holes. This configuration uses a combination of header pins for electrical connectivity and standoffs to ensure structural rigidity and support within the *SOLARIS* platform.




Figure 8.5 2D Regulator Board Layout




Figure 8.6 3D Voltage Regulator Board Layout

The motor driver PCB design for *SOLARIS* utilizes the VNH5019A-E, a fully integrated H-bridge capable of delivering 30 A of continuous output current. We used a minimum motor line trace width of 1.524 mm for a maximum anticipated current of 3.3 A. We anticipate the motors operating around 0.3 A under normal load. This conservative trace sizing helps minimize resistive losses and thermal buildup across the six-motor drive system, ensuring reliable performance during prolonged operation. To further manage heat, the design leverages the package's exposed die pads, which are specifically engineered for improved thermal performance in high-current applications. The interface between the motor driver and the microcontroller is designed for both precision control. Speed regulation is achieved through PWM operation at frequencies up to 20 kHz, while direction of motors is managed via input signals.




Figure 8.7 2D Motor Driver Board Layout




Figure 8.8 3D Motor Driver Board Layout

8.4 MCU Board

The main MCU board serves as the central nervous system of the *SOLARIS* platform, coordinating high-power energy distribution alongside sensitive logic-level communication. At the power entry point, a robust terminal block accepts the twelve-point eight-volt supply, which is immediately routed to the voltage regulation stages. These stages must be physically isolated from the high-speed logic to prevent switching noise from interfering with the microcontroller or the serial programming bridge.

The processing core is anchored by the ESP32, which is positioned to minimize the length of data traces to the adjacent USB-to-UART programming circuit. The design incorporates

a pre-built inertial measurement unit; this component is mounted in the geometric center of the board to ensure that its orientation and movement data accurately reflect the robot's center of mass. To manage the high density of sensors, the board utilizes multiplexers to consolidate signals from the light and ultrasonic arrays, allowing a wide range of data to be processed using a limited number of input pins.

The board's perimeter is dedicated to a comprehensive array of connectors that interface with various daughterboards throughout the chassis. These connections range from standard headers for low-power sensors to a more substantial terminal block interface for the energy monitoring module. This strategic organization minimizes cable clutter while allowing the structural standoffs to provide both electrical connectivity and mechanical rigidity for the entire assembly. Ultimately, this modular architecture facilitates efficient troubleshooting and ensures the future scalability of the individual sensing and drive systems.




Figure 8.9 2D MCU Board Layout




Figure 8.10 3D MCU Board Layout

9 System Testing and Evaluation

9.1 Brief

Successful implementation of an autonomous solar-tracking vehicle is not only reliant on the selection of high-performance components but also in the integration of a reliable system. The testing and evaluation protocols used to verify each subsystem are created to validate the design requirements established by the project in the initial phases. Testing was approached from an initial hardware level where the electrical integrity of each component was tested. Following validation, we conducted integration testing to ensure that each component would not only work with the controller but also that the controller could manage the concurrent demands of the system. The evaluation metrics serve as a baseline for the transition into Senior Design 2 and how we hope to continue verifying functionality and feasibility of *SOLARIS* moving forward.

9.2 Hardware and Software Integration

9.2.1 ESP32-S3-WROOM-1/1U MCU Dev Kits

The ESP32-S3-WROOM-1 (PCB antenna) and 1U (U.FL connector for external antenna) modules were evaluated for their ability to manage concurrent tasks: motor PID loops, sensor fusion, and power management. Testing of these devices focuses on the stability of the development kit under the heavy I/O load required by *SOLARIS*. The N16R8 variant comes with 16 MB of Flash and 8 MB of PSRAM memory. This provides plenty of memory overhead which is necessary for complex multi-tasking. Testing must validate the extended memory partitions and ensure that the high-speed PSRAM does not conflict with the GPIO requirements of the motors and sensors.

Hardware Setup:

- **Memory Configuration:** 16 MB Flash / 8 MB OPI PSRAM
- **Interface:** The S3 utilized its on-board USB-UART for easy programming and logging.

The N16R8 variant uses specific internal pins for the Octal PSRAM (GPIO 33-37). These were verified through continuity testing that our external sensors were not assigned to these reserved internal pins. Data can be maintained by PSRAM as a buffer of readings for all sensors which allows for filtering without slowing down the main control loop. Due to its 240 MHz dual-core processor context switching between each task is fast and does not prove to be a problem with no memory contention errors.

Integration of this system is vital as it acts as a nexus for all subsystems. On board storage must be utilized for data logs from the NPA20S-I2J Solar Panel and the main power line from our power supply. It acts as a computational hub for the solar tracking algorithms and for movement control. The core also runs high-priority safety tasks which monitor the MB1020 ultrasonic sensors during driving to override all other tasks if a collision is imminent. To build upon this groundwork for Senior Design 2 we hope to utilize the onboard memory to cache the solar harvest data and write to the Flash for low power operation. We would also like to implement a mode where the PSRAM is powered down during long periods of darkness to extend the battery life of the system if necessary.

9.2.2 TEPT5600 Phototransistors

For the light sensing capabilities of the vehicle we utilize eight TEPT5600 phototransistors divided into two functional groups. One four-unit “quadrant” array is used for solar alignment while the other four-unit “navigation” array is used for seeking peak solar intensity while driving. Each phototransistor is configured in a common-emitter circuit with a 10 k Ω pull-down resistor, creating a voltage divider that interfaces with the CD4051 MUX.

Hardware Setup:

- **Load Resistor:** 10 k Ω (selected to balance sensitivity in ambient light with saturation limits in direct sunlight)
- **MUX Interface:** Inputs and Selector

- **Spectral Match:** The TEPT5600 peak sensitivity (570nm) is verified against the solar spectrum to ensure efficient tracking.

Testing of this equipment consists of hardware voltage output measurements to verify there are no false positives in shadows. Similarly when the sensors are exposed to sunlight we verify the output voltage remains saturated around the input voltage level but also below 3.1V to avoid “clipping” of the ADC on the ESP32-S3. Then we verify the minimum time needed to sample all sensors considering we must address the MUX to switch channels.

Integration of this system relies heavily on the closed-loop system where the TEPT5600 data drive the goBILDA 5203-series motors. The panel array is arranged in a “Shadow Box” configuration where four sensors report equal voltage when the panel is perfectly perpendicular to the sun. Similarly the navigational array is mounted on all four major directional headings to verify which direction has the clearest path to the sun. A searching algorithm must be implemented to simulate sun movement simply by moving a high-intensity light. The ESP32-S3 must successfully process the MUX data and command the motors to re-center the panel within a limited time. Similarly the motors must avoid jittering due to obstructions with a software implemented threshold.

During Senior Design 2 we may experiment with optical diffusers which would allow for the phototransistors to have a wider field of view. Similarly, the code could be updated to “average” the reading from all of the eight sensors to determine if it is a cloudy day, which would tell the system to stay in a power saving mode. Lastly we hope to integrate a pulldown network onto the PCB which would help to reduce lead inductance and noise.

9.2.3 MB1020-000 Ultrasonic Sensors

SOLARIS utilizes a multi-sensor array of MB1020-000 ultrasonic sensors for short range collision detection. To eliminate ultrasonic interference, the sensors are configured in a Daisy Chain command sequence controlled by the ESP32-S3. Testing is used to verify the timing of the chaining pulse and the successful analog read from the ultrasonic module.

Hardware Setup:

- **Chaining Logic:** The TX pin of Sensor 1 is connected to the RX pin of Sensor 2
- **Master Trigger:** GPIO is connected via the ESP32-S3 to the RX pin of Sensor 1
- **Data Interface:** The UART TX pins of all sensors are tied to a single UART RX pin on the ESP32-S3
- **Baud Rate:** Baud rate is set to 9600 bps.

To test the hardware we look at the trigger timing which can be verified using an oscilloscope that a high signal is sent to the first sensor’s RX pin to start the chain. The chain propagation then begins and it is verified that Sensor 1 completes its reading, then it pulls the TX pin high to trigger Sensor 2. Lastly we verify that the data received over UART can be correctly converted from an ASCII string sent by each sensor and converted into an integer value in inches.

As a safety component, the MB1020 is evaluated for its “Field of View” and its ability to detect various materials. Minimum distance of the sensor should reliably detect objects up to 6 inches away. Below this the sensor reports “R006” which is handled as a critical proximity zone. With a minimum of 6 inches the maximum distance should be a

standard wall up to 254 inches away. The resolution of these measurements is always 1-inch meaning that across the full operating range we should be able to monitor at minimum an inch change. To test the wide beam we can use something of slim diameter to cover the beam and analyze the output from the TX.

The system must be integrated into the vehicle's main control loop alongside the BNO085 and the goBILDA motors. This means that factors of latency for approximate time of scanning must be sufficient for the robot to be issued a stop command before impact. This may be an application of a high priority interrupt where motors must be stopped immediately and override any navigation controls. Similarly motor noise must not trigger false positives on the ultrasonic sensors. For Senior Design 2 a filter can be added to the ESP32-S3 code to ignore reflections that or simply spikes in the sonar data. Similarly testing the sensors in outdoor conditions with wind and temperature changes may show variation that could impact the 1-inch resolution. Lastly, implementing a dynamic break rather than a hard stop would allow the vehicle to slow down as it gets closer to an object and then use an encoder for a smooth deceleration curve.

9.2.4 BNO085 IMU

The BNO085 9-axis IMU provides the orientation data necessary for *SOLARIS* to properly track the heading of travel and direction of the sun in conjunction with timing data. Due to known stability issues with I2C clock stretching on the ESP32-S3 platform, the SPI (Serial Peripheral Interface) protocol was selected to ensure high-speed and reliable data transmission. Testing of this hardware focuses on validating the 6-wire SPI interface which is required by the Hillcrest Labs SH-2 protocol (MISO, MOSI, SCK, CS, WAK, and INT). The BNO085 was interfaced with the ESP32-S3 WROOM-1 using the VSPI/FSPI bus, The ESP32-S3's internal pull-up resistors were disabled in favor of external 10 k Ω resistors on the Chip Select (CS) and Reset lines to ensure stable logic states during MCU boot-up.

Hardware Setup:

- **SPI Clock Speed:** 3 MHz
- **Mode:** SPI Mode 3 (CPOL=1, CPHA=1)
- **Pins:** MOSI (GPIO 11), MISO (GPIO 13), SCK (GPIO 12), CS (GPIO 10)

To test this hardware we first verify our protocol selection which requires us to strap our PS0 and PS1 pins for correct SPI communications. Using an oscilloscope we can validate the handshake by monitoring the HINT (Host Interrupt) pin. When data becomes ready for the ESP32-S3 the BNO085 pulls the pin low and the SPI transaction initiates. We also can confirm the sensor successfully exits the reset state and sends the "SHTP Advertisement" packet over SPI. Verifying that these three test cases pass we can then move on to a larger performance evaluation where we can look at dynamic conditions of the systems interaction. For our application we look into the magnetic heading and relative orientation tracking capabilities.

We test these by executing a turn test which not only tests the magnetic heading but also tests the gyroscopes rotation vector fusion. This system must be integrated into the vehicle's PID Control Loop. During this loop the SPI bus orientation updates at around 100 Hz which allows the ESP32-S3 to calculate motor corrections every

10 ms. Similarly because the BNO085 performs internal sensor fusion we can verify that the magnetic heading remains stable even when the vehicle is titled at an angle. These tests show the critical advantage that a self contained unit like the BNO085 has against other 9-axis IMUs. During Senior Design 2 we will execute these tests in more detail as we reveal access to testing equipment like an oscilloscope, bench power supply, and other solidified details of the project.

9.2.5 INA228EVM Energy Monitor Unit Dev Kit

The INA228 was integrated into this project to allow for real-time power analytics and State of Charge (SoC) estimation for the *SOLARIS*'s battery pack. In particular the INA228 uses an integrated Charge Accumulator (Coulomb Counter) and Energy Accumulator. This is a surefire way to monitor the battery of *SOLARIS*. To communicate with the device we use a I2C connection which provides a stable and fast software based communication method. Testing of the device relies on the accuracy of the shunt resistor interface and the reliability of the I2C register mapping. The INA228EVM was configured in a High-Side Sensing configuration. The power supply terminal was connected to the IN+, and the load was connected to the IN-.

Hardware Setup:

- **Shunt Resistor:** 15 mΩ (On-board EVM default)
- **Interface:** I2C at 400 kHz
- **Address:** 0x40 (A0 and A1 tied to GND/Hanging)

The testing procedure went as follows starting with the Calibration Register setting. To convert the raw 20-bit ADC values into Amps and Watts, the SHUNT_CAL register must be programmed. Using a maximum expected current of 10A, the Current_LSB was calculated as follows:

$$Current_{LSB} = \frac{Max_Expected_Current}{2^{19}}$$

With the load disconnected, we verified that the current register reported 0.00mA. We then performed a “Stress Read” test where the ESP32-S3 polled the Bus Voltage register 1,000 times in a loop to ensure no I2C timeout or bus errors occurred. We then used three types of resistor as our main loads being a 1 kΩ, 100 Ω, and 220 Ω resistors. Starting the system with each connected we verified the V_BUS, CURRENT, and WATTS register readouts to compare them for accuracy. It is important to note that all values delivered from the IC must be multiplied by the *Current_LSB* to reveal the true value. In each instance all load values produced expected test values and were compared using a handheld multimeter. During Senior Design 2 we plan to implement the following improvements to the battery monitoring system. We can combine the Coulomb counting with a Voltage-to-SoC look-up table to keep the battery percentage conversion simplified. We also can use the internal thermal monitoring sensor to monitor the PCB temperature near the shunt resistor to compensate for resistance changes during high current draw. Lastly we would like to implement data logging features for analysis after a run.

9.2.6 Multiplexer

To maximize the I/O capabilities of the ESP32-S3, an eight channel CD4051BM96 and a four channel CD4052BE analog multiplexer are to be employed. The CD4051 will be used for the optical array of phototransistors while the CD4052 will be used for motor-related signal routing. The CD4051BM96 serves as the primary interface for the 8-unit phototransistor array. It allows the ESP32-S3 to scan the entire array via a single ADC pin.

Hardware Setup:

- **Select Lines (S0, S1, S2):** Connected to GPIOs
- **Common Out (COM):** Connected to ESP32 ADC1 via GPIO
- **Testing Tool:** Oscilloscope used to verify timing and signal leakage

First we confirm that the hardware when powered by a 3.3V source is capable of sensing a 3.3V logical high from the ESP32-S3 correctly. When triggered the CD4051's address pins will switch between inputs. A test loop cycles through all eight channels. During this test loop a 1 kHz sine wave can be applied to each input which can ensure the signal was passed to the output with zero distortion or clipping. This concludes the testing for the CD4051 with both the isolation test and linearity test showing a pass.

The CD4052BE is utilized for the motor subsystem, specifically for control signals and PWM signal switching. Because it is a Dual 4-channel MUX, it switches two independent signals (X and Y) simultaneously using the same select lines (A, B). To test this we first must verify that when Select Bit A is toggled, both the X-channel and Y-channel switch to their respective inputs at the exact same time. This is critical for maintaining phase synchronization in signals. Then we also want to test the MUX bandwidth to verify that a 20 kHz PWM signal can pass through the CD4052 without significant rounding of square wave edges. This will also use a very similar hardware setup as shown above with one less select line in this instance.

These final tests are used to ensure that the integration of both multiplexers do not interfere with one another. On the final PCB the CD4051 and CD4052 will be placed as close to the sensors as possible to minimize the length of analog traces, which reduces the pickup of electromagnetic noise from the motors. We must also investigate using the ESP32-S3s RMT (Remote Control) peripheral or LEDC to automate the MUX address switching, further reducing CPU overhead. It may also be important to look into a buffer after MUX outputs to eliminate the effects of the internal resistance and provide a low impedance signal to the ESP32 ADC.

9.2.7 Motors

Project *SOLARIS* uses two distinct classes of goBILDA 5200-series planetary gear motors to meet the varying torque and speed requirements of the chassis and the solar tracking systems. Testing must be conducted on two key parts of these motors being the PWM response and the integration of built-in Hall-effect encoders with the ESP32-S3.

Hardware Setup:

- **Propulsion:** Four 5204-8002-0100 (99.5:1 ratio, 60 RPM) motors
- **Solar Tracking:** Two 5203-2402-0188 (188:1 ratio, 30 RPM) motors
- **Encoders:** Interfaced with the ESP32-S3 using the PCNT (Pulse Counter) peripheral to track position without CPU interrupts.

To test these we must test each motor for directional consistency. This means a “forward” command in software is verified to produce a clockwise rotation on all four drive motors. Using a multimeter we monitor the current draw of the 5024-series motor under a block condition to analyze current spikes. We then analyze the minimum start voltage required to overcome the internal friction of the planetary gearboxes by increasing the PWM duty cycle from 0%.

The primary concerns for the project were always how the motors fulfill specific mechanical requirements of the project. In this concern we analyze the drive motors ability to maintain a straight heading. All four motors must be fed an identical PWM signal while the deviation in RPM between the slowest and fastest motors are observed. Similarly the solar tracking system requires the motors to have high torque but have high precision to angle the panel at the sun. We verify this by applying a minimum “back-drive” resistance to hold the solar panel in place without power. We also use the phototransistor array and the CD4051 MUX, we test the motors ability to turn in steps. The system must achieve high angular resolution and exceed a minimum requirement for efficient solar harvesting. The integration of six high-power motors presents significant challenges for the sensitive sensors like the INA228 and the BNO085. Using the INA228 we monitor the battery rail during a “High-Torque Start” of all four drive motors. The motor leads must also be twisted and kept away from the SPI lines of communication which ultimately may cause “bit-flips” or communication resets during peak operation. To continue in this direction we look to follow up with mechanical and electrical refinements. These refinements consist of using acceleration and deceleration curves which reduce the mechanical stress on the goBILDA motors during sudden starts and stops. We also want to implement a closed loop tracking algorithm which allows the software line between the phototransistors and the MUX data to autonomously track the brightest light source in real-time. Finally using the INA228 we may be able to detect if the solar panel has hit a physical limit by monitoring for a sudden spike in current which could act like a virtual limit switch.

9.2.8 NPA20S-I2J Solar Panel

The NPA20S-I2J serves as the primary energy source for the *SOLARIS* charging network. Testing is designed to verify the panel’s electrical characteristics under varying light conditions and ensure its structural integrity when mounted to the dual-axis tracking system. Before integration with the charging circuit, the panel’s raw output must be benchmarked against the datasheet specifications:

$$(P_{max} = 20W, V_{mp} = 20.0V, I_{mp} = 1.0A).$$

We test the panel in an open-circuit configuration in direct noon sunlight using a multimeter to ensure the panel can reach its rated ~24 V. We then measure the maximum current output by shorting the leads through an ammeter to verify a value of approximately 1.0A. Then lastly we verify that the MC4 connectors and wiring gauge (12 AWG) were sufficient to handle the peak current without resistive heating.

Our panel is integrated with the Genasun GV-5-LI-14.2V MPPT and the battery storage system. We verify for the sake of integration that the charge controller correctly identifies the 20V input from the NPA20S-I2J and speeds it down to the appropriate charging voltage for the battery. The INA228 can also be placed between the solar panel and the charge controller. This allows for the ESP32-S3 to calculate the “Harvested Energy” in Watt-hours, and provides a metric for system autonomy. During Senior Design 2 we hope to refine the solar tracking algorithm so the motors only move when a significant increase in voltage is predicted, which reduces the energy cost of moving the panel. It is also worth testing the mounting brackets under wind loads to ensure that the motors do not strip mounting hardware. Lastly it would be important to create a software-based “Sun Map” which uses the time-of-day data received over BLE to predict the optimal heading of the panel angles even if the phototransistors are momentarily obstructed.

9.2.9 Charge Controller

The Genasun GV-5-LI-14.2V is a high-efficiency MPPT (Maximum Power Point Tracking) controller that is specifically designed for lithium based batteries. It is integrated to bridge the NPA20S-I2J solar panel and the lithium battery pack. This component is responsible for high-speed tracking of the panel’s maximum power point and regulating the 14.2V charging profile. Since the GV-5 is a self-contained hardware controller with no digital communication port, testing must be performed indirectly by the ESP32-S3 via the INA228 energy monitor to verify the working functionality.

Hardware Setup:

- **Input:** Connected to the MC4 leads of the NPA20S-I2J panel
- **Output:** Connected to the LiFePo4 battery via a fuse
- **Monitoring:** One INA228 monitor is placed on the solar input side and a second on the battery output side to measure conversion efficiency.

For the test we can measure the power consumed by the GV-5 in a dark condition to confirm it is not a significant drain on the system. We must also verify that there is safety on an improper connection which can be verified by briefly reversing the panel leads to ensure no component damage occurs. Lastly we measure the charging cycle until the battery reaches its peak. This can be verified by watching the GV-5 transition to “Float/CV” mode at around 14.2 V.

One core metric for the success and valid functionality of this MPPT is its efficiency. Efficiency can be calculated using the two-point measurement system of the INA228, efficiency is calculated as:

$$\eta = \left(\frac{V_{bat} \times I_{bat}}{V_{sol} \times I_{sol}} \right) \times 100$$

Under a max charging current the GV-5 must remain cool to the touch remaining within safe operating limits. One concern of this high power converter is the introduction of high-frequency noise that could be introduced to the system. We must use an oscilloscope on the ESP32-S3’s 3.3V rail to ensure the GV-5’s switching frequency does not interfere with components like the BNO085 IMU or the CD4051 MUX signals. For Senior Design 2 we should implement a charge cycle log where the ESP32-S3 can be programmed to create a

total record of mAh harvested by the GV-5 per day. Similarly we could try to provide passive venting on the final product to ensure that all components including the MPPT remain cool during operation in Florida summers.

9.3 Phone Application Stack

9.3.1 NimBLE Setup on ESP32-S3-WROOM-1/1U

The starting point for the NimBLE stack on the ESP32-S3 came from the ESP-IDF NimBLE peripheral example in the Espressif docs. That example gave us a working GAP and GATT skeleton with advertising, connection handling, and a basic service registration flow. We stripped out the parts we did not need and restructured it into its own components so the *SOLARIS* firmware could add its own services on top of it without touching the init code.

The main changes we made on top of the example were writing our own GATT characteristics for the stuff *SOLARIS* needs, telemetry upload, mode switching, manual control, time syncing, and weather data. We also renamed the device to *SOLARIS*, so it shows up properly when scanning.

The telemetry upload itself works like a little TCP style service on top of BLE. The phone asks for a chunk of log records, the ESP32 sends it, the phone ACKs, and only then does the ESP32 move on to the next chunk. If an ACK does not come back in time, the ESP32 resends the current page. Once the phone has everything, it sends a clear command, and the ESP32 wipes the log.

To confirm the stack was actually up and behaving, we tested in two stages. First, we used nRF Connect on a phone to scan for the device, connect to it, and walk the GATT table. This was the fastest way to check that advertising was working, the device name showed up as *SOLARIS*, and every service and characteristic we had registered was discoverable with the right UUIDs and properties. nRF Connect also let us manually subscribe to notify characteristics and write raw bytes to the mode, manual control, and telemetry reads, which is how we verified the write handlers were parsing payloads correctly before the companion app was ready.

Once the companion app had its BLE client wired up, we ran the same checks from the phone app side. We confirmed it could discover the device, connect, and read and write to each of the characteristics. We watched the serial monitor the whole time to make sure the ESP was receiving everything properly and that the log output lined up with what the app was sending.

9.3.2 Phone App

9.3.2.1 Frontend Integration

The phone app is built in React Native with Expo and runs the full BLE stack on iOS using react-native-ble-plx. All BLE logic lives in a single context provider so any screen can pull in the connection and the read, write, and sync helpers without managing its own state.

Auth is handled through Clerk and every backend request goes out with a bearer token attached.

We tested on a physical iPhone the whole time. Expo Go does not support custom native modules like the BLE library, so we used a development build installed on the phone and connected to the ESP32 over real Bluetooth for every test.

The BLE connection flow runs as soon as the user signs in. It asks for permissions, scans for a device named *SOLARIS*, connects, then syncs time, pulls weather from the backend and writes it to the ESP, and runs a telemetry upload. We tested by watching the phone console and the ESP serial monitor side by side, confirming each step completed on both ends.

The telemetry home page pulls from the GET telemetry endpoint and feeds the JSON into a grid of swipeable cards for battery, distance, net power, and CPU temp. We tested by loading the page with real data in the database and confirming every card showed the right numbers and charts. Pull to refresh triggers a fresh BLE sync if connected then re-fetches, which we verified by watching new data come in from both sides.

The mode change screen switches between stationary, automatic, and manual, and shows throttle and steering joysticks when manual is selected. On mount it reads the current mode off the ESP. Mode changes get written back over BLE, and in manual mode the joystick positions stream down as signed bytes while the user moves them. We tested by flipping through all three modes and watching the ESP logs confirm each one, and by dragging the joysticks and confirming the throttle and steering values updated on the ESP in real time.

The weather flow is tied into the BLE sync. The app grabs the phone's GPS location, hits the backend weather endpoint, and writes the forecast down to the ESP as a 152 byte payload. We tested by triggering a sync and confirming the values the backend returned matched what the ESP logged when it parsed the write.

The timesync characteristic is the first triggered during BLE synchronization. It retrieves the current datetime from the phone, converts it to a 32-bit Unix timestamp, and transmits the payload to the ESP. Once received, the ESP uses the time library to set its onboard clock, making the current time available to any other module that needs it.

All testing was manual, done with the phone in hand and the ESP on the bench with the serial monitor open. Between the phone console and the ESP logs we could follow every message end to end and catch any place the two sides disagreed.

9.3.2.2 Database Setup and Connection

The database for *SOLARIS* is a PostgreSQL instance hosted on Supabase. We keep the connection string in an environment variable so the backend can pick it up without having any credentials hardcoded. Alembic handles the migrations, so whenever we changed a table we generated a new migration and ran it against the database instead of editing anything by hand.

The schema is simple. There is a users table keyed on the Clerk user ID, and a telemetry table that references it. Each telemetry row has a timestamp, battery level, CPU temp, distance traveled, and net power, and the primary key is the user ID plus the timestamp so every record is tied back to who sent it.

On top of that we have two services, one for posting telemetry and one for getting it back out. The POST service takes the list of records the phone uploaded, makes sure the user exists in the users table, then inserts all the rows at once. If a record with the same user and timestamp is already in there it just skips it so we do not get duplicates. The GET service pulls every telemetry record for the user, sorts them newest first, and builds out the JSON the app needs for the home page. That means current values for battery, CPU temp, and net power, running totals for distance over the past hour, day, week, month, and all time, and time bucketed averages for the charts so the app does not have to do that work itself.

We tested the POST side by sending telemetry from the backend and checking the Supabase table editor to confirm the rows actually showed up with the right values. The GET side got tested through the companion app directly, since the telemetry home page pulls from that endpoint and will not render correctly if anything is off. If the numbers and charts on the home page looked right, we knew the endpoint was returning the data in the shape the app expected.

9.3.2.3 *Backend Weather Services*

The weather endpoint takes a latitude and longitude from the phone, calls the Open-Meteo API under the hood, and returns a trimmed down version of their response that fits what SOLARIS actually needs. That means 24 hours of hourly forecast with cloud cover and precipitation probability, plus the sunrise and sunset times for the day. Everything is returned as Unix timestamps so the ESP32 does not have to parse date strings.

The flow is pretty simple. The phone app grabs the user's location, hits this endpoint, gets the JSON back, and then forwards that weather data down to the ESP32 over BLE.

To test it we temporarily turned off the auth requirement and called the endpoint directly with a known latitude and longitude, then checked that the JSON coming back had the right shape, 24 forecast entries, valid sunrise and sunset timestamps, and numbers that matched what Open-Meteo was showing for that location. Once the payload looked right we turned auth back on and tested it through the full stack, phone app pulling the weather, sending it down to the ESP32 over BLE, and confirming the ESP32 parsed it correctly on the other end.

10 Administrative Content

10.1 Budget

The total budget allocated for the *SOLARIS* project is approximately \$1,000. This budget is intended to cover all major system components, including mechanical hardware, electronic components, sensing and actuation devices, power management hardware, and

prototyping materials. Funds are also reserved for iterative development and testing, allowing for component replacement or redesign as needed throughout the project lifecycle. Establishing a defined budget constraint ensures that design decisions remain practical and encourage cost-effective engineering tradeoffs while still meeting the project's functional and performance objectives.

10.2 Bill of Materials

Part Name	Part Number	Market Availability		Total # Ordered		Cost	
		Market Price	Min Qty.	Min Need	Total Bought	Min Cost	Total Cost
ESP-32-S3 Microcontroller Chip	ESP32-S3-WROOM-1-N16R8	\$6.56	1	1		\$6.56	
ESP-32-S3 (Dev Kit)	ESP32-S3-WROOM-1-N16R8(Dev Board)	\$23.74	3	0	3	\$0	
DC Drive Motor	5204-8002-0100	goBILDA: \$56.99	1	4		\$227.96	
Pan/Tilt Motor	5203-2402-0188	goBILDA: \$54.99	1	2		\$109.98	
Motor Driver	VNH5019ATR-E	\$9.22	1	3		\$27.66	
LiFePo4 Battery	RBT1220LFP-TM-US	Renogy: \$89.99	1	1		\$89.99	
Newpowa Solar Panel	NPA20S-12J	\$37.99	1	1		\$37.99	
Genasun Multi-Power-Point-Tracker	GV-5-LI-14.2V	\$99.00	1	1		\$99.00	
TEPT Phototransistor	TEPT5600	\$0.70	1	8	16	\$5.60	\$9.25
Energy Monitor Unit (Dev Kit)	INA228EVM 5832 var.	\$14.95	1	0	2	\$0.00	\$40.48
TI Energy Monitor Unit (Automotive Grade)	INA228AQDGSRQ1	\$5.27	1	1		\$5.27	

Inertial Measurement Unit (IMU)	BNO085	\$24.95	1	1	2	\$24.95	\$49.90
Voltage Regulator 3.3V & 5V Separate	LM2679T-ADJ/NOPB	\$7.42	1	2		\$14.84	
Ultrasonic Sensor	MB1020-000	\$19.95	1	4	6	\$79.80	\$119.70
TOTAL:		ex. (total cost is total spent so far while min. is based on parts)				\$729.60	\$219

Table 10.1 Bill of Materials

10.3 Distribution of Work

<i>Electrical Engineering</i>	<i>Responsibilities</i>
Rafael Puig	Input and output port configuration from the user
	Battery safety and power system design
	Lead PCB design
<i>Electrical Engineering</i>	<i>Responsibilities</i>
Nathan Hammond	Servo and motor technology
	Design of the power delivery system
	Co-Lead PCB design for drivetrain control
<i>Computer Engineering</i>	<i>Responsibilities</i>
Ryan Markowitz	Phone app design for interface with the MCU
	Software lead for IoT connection
	Co MCU firmware design
<i>Computer Engineering</i>	<i>Responsibilities</i>
Garrett Fortier	Lead MCU firmware design and debug
	Automation for the self-driving algorithm and modularity
	Hardware/Software integration

Table 10.2 Distribution of Work

10.4 Milestones

<i>Milestone Number</i>	<i>Milestone Description</i>	<i>Start Date</i>	<i>Anticipated End Date</i>	<i>Dependencies & Requirements</i>
#1	Recruitment	01/12/2026	01/16/2026	Scout for potential team members interested in working together
#2	Brainstorming	01/16/2026	01/22/2026	Talk with one another briefly over a week to discuss ideas
#3	D&C Document	01/22/2026	01/30/2026	Divide up the task and work independently to add relevant information
#4	Mentor Meeting	02/04/2026	02/04/2026	Meet with Dr. Suboh for D&C revision
#5	60-Page Milestone	02/06/2026	03/27/2026	Work will be divided up, and then relevant portions of the paper will be written
#6	Final Document	03/01/2026	04/28/2026	All group written portions will be formatted onto one document

Table 10.3 Milestone Progress

11 Conclusion

The Solar-Optimized Land Autonomous Recharge and Intelligence System, known as SOLARIS, was developed to expand the domain of renewable energy solutions by investigating the feasibility of mobile solar panel platforms. Although the twenty watt rated solar panel is modest in scale, the design of the system represents a significant step toward proving the viability of such robotics and uncovering the technical challenges associated with mobile energy harvesting. Through this project, our team of electrical and

computer engineers at the University of Central Florida sought to gain a deeper understanding of power electronics and autonomous systems while creating a functional blueprint for future solar driven mobile technologies.

This report has provided a comprehensive overview of the engineering lifecycle for the project, beginning with the initial goals and objectives set by the team. We conducted extensive research into hardware and software selections and evaluated multiple engineering standards and constraints relevant to the system. Additionally, the project included a case study on the use of large language models in engineering and administrative workflows to enhance productivity. The process transitioned into a realized design through the development of complex schematics and printed circuit board layouts, alongside the creation of autonomous software algorithms and hardware demonstrations to mitigate integration risks.

The physical embodiment of the robot relies on a modular and robust chassis constructed from aluminum extrusions for structural support and custom three dimensional printed panels for the containment of internal electronics. Movement is facilitated by a four wheel drive system using skid steering for navigation across various terrains. For the solar array itself, a two axis arm was designed using two motors in an arrangement of bevel gears and specialized fasteners to allow for precise pan and tilt motion. The electrical backbone consists of a central ESP32 S3 microcontroller board that hosts a bluetooth phone app server and interfaces with specialized daughterboards for energy monitoring, power regulation, and motor driving.

Since the project has not been fully constructed at this time, the experimental data required to fully validate the concept remains to be collected. The results gathered during future testing will assist organizations interested in expanding mobile solar energy projects by defining the practical limits and potential audience for such systems. Our team plans to spend the coming months procuring the remaining circuit boards, electronic hardware, and chassis materials to begin the final assembly, programming, and debugging phases. This transition will provide a definitive starting point for measuring specifications and making any necessary modifications to ensure that the completed system meets its set objectives and performance goals.

Appendix A – Reference

- [1] 曾舒婧, “News,” *Jackery*, Jan. 11, 2024.
<https://www.jackery.com/blogs/news/jackery-debuts-revolutionary-solar-generator-mars-bot-at-ces-2024?srltid=AfmBOoqCnd6ULgSK9d6mTIQqmZTIWY50-8vI0sHEHEDs32RMOLLjCmBs>
- [2] pico bricks, “SolarX: DIY STEM Solar Powered Robot,” *Picobricks.com*, 2026.
https://picobricks.com/products/solar-tracker-robot?srltid=AfmBOoqkf9bFY7E0rjGaOA1-cR2sDuRMQcqBFjFm8LCpA2Kn5h_ZYap7
- [3] J. Schneider and I. Smalley, “What is a microcontroller? | IBM,” *www.ibm.com*, Jun. 04, 2024. <https://www.ibm.com/think/topics/microcontroller>
- [4] GeeksforGeeks, “Microcontroller and its Types,” *GeeksforGeeks*, Feb. 20, 2022.
<https://www.geeksforgeeks.org/digital-logic/microcontroller-and-its-types/>
- [5] Espressif, “ESP32-P4 Chip Revision v1.3 Datasheet,” *Espressif.com*, 2026.
https://documentation.espressif.com/esp32-p4-chip-revision-v1.3_datasheet_en.html
(accessed Apr. 27, 2026).
- [6] Espressif, “ESP32-S3 Datasheet,” *Espressif.com*, 2025.
https://documentation.espressif.com/esp32-s3_datasheet_en.html
- [7] Espressif, “Espressif Documentation,” *Espressif.com*, 2025.
https://documentation.espressif.com/esp32_datasheet_en.pdf
- [8] T. Instruments, “MSP430FR6989 MSP430FR698x(1), MSP430FR598x(1) Mixed-Signal Microcontrollers 1 Device Overview,” 2014. Accessed: Apr. 27, 2026. [Online]. Available:
<https://www.ti.com/lit/ds/symlink/msp430fr6989.pdf?ts=1771368753952>
- [9] STMicroelectronics, “STM32U0 series advanced Arm-based 32-bit MCUs,” *www.st.com*, Jul. 2025.
https://www.st.com/resource/en/reference_manual/rm0503-stm32u0-series-advanced-armbased-32bit-mcus-stmicroelectronics.pdf

- [10] STMicroelectronics, “RM0394 Reference manual STM32L41xxx/42xxx/43xxx/44xxx/45xxx/46xxx advanced Arm ® -based 32-bit MCUs,” Dec. 2024. Available: https://www.st.com/resource/en/reference_manual/rm0394-stm32l41xxx42xxx43xxx44xxx45xxx46xxx-advanced-armbased-32bit-mcus-stmicroelectronics.pdf
- [11] J. Schneider, “Microcontroller Vs Microprocessor,” *Ibm.com*, Jun. 13, 2024. <https://www.ibm.com/think/topics/microcontroller-vs-microprocessor>
- [12] J. Schneider, “FPGA Vs Microcontroller,” *Ibm.com*, Jun. 03, 2024. <https://www.ibm.com/think/topics/fpga-vs-microcontroller>
- [13] eMotors Direct, “The Difference Between AC and DC Motors | eMotors Direct,” *www.emotorsdirect.ca*, Sep. 08, 2023. <https://www.emotorsdirect.ca/knowledge-center/article/the-difference-between-ac-and-dc-motors>
- [14] P. Millett, “Brushless Vs Brushed DC Motors: When and Why to Choose One over the Other,” *Monolithicpower.com*, 2024. <https://www.monolithicpower.com/en/learning/resources/brushless-vs-brushed-dc-motors>
- [15] J. Hertz, “H-bridge DC Motor Control Using Complementary PWM, Shoot-through, and Dead-time - Technical Articles,” *All About Circuits*, Aug. 29, 2022. Available: <https://www.allaboutcircuits.com/technical-articles/h-bridge-dc-motor-control-complementary-pulse-width-modulation-pwm-shoot-through-dead-time-pwm/>
- [16] D. Birks, “Brushed DC Motors and How to Drive Them,” *Diodes Incorporated*, Sep. 02, 2020. <https://www.diodes.com/design/support/technical-articles/driving-brushed-dc-motors>
- [17] R. Shaft, “5204 Series Yellow Jacket Planetary Gear Motor (99.5:1 Ratio, 80mm Length 8mm REX® Shaft, 60 RPM, 3.3 - 5V Encoder),” *goBILDA®*, 2026. <https://www.gobilda.com/5204-series-yellow-jacket-planetary-gear-motor-99-5-1-ratio-80mm-length-8mm-rex-shaft-60-rpm-3-3-5v-encoder/> (accessed Apr. 27, 2026).
- [18] StepperOnline, “Brushed 24V 60W 6.27Nm(887.91oz.in) 60RPM w/ 50:1 Spur Gearbox DC Gear Motor - F5D60-24GN-30S/5GN50K | StepperOnline,” *Omc-stepperonline.com*, 2026. <https://www.omc-stepperonline.com/brushed-24v-60w-6->

27nm-887-91oz-in-60rpm-w-50-1-spur-gearbox-dc-gear-motor-f5d60-24gn-30s-5gn50k (accessed Apr. 27, 2026).

[19] RobotShop, “36D DC Planetary Gearmotor w/ Encoder - 24V 120RPM,” *www.robotshop.com*, 2023. <https://www.robotshop.com/products/e-s-motor-36d-dc-planetary-gearmotor-w-encoder-24v-120rpm>

[20] R. Shaft, “5203 Series Yellow Jacket Planetary Gear Motor (188:1 Ratio, 24mm Length 8mm REX® Shaft, 30 RPM, 3.3 - 5V Encoder),” *goBILDA®*, 2026. <https://www.gobilda.com/5203-series-yellow-jacket-planetary-gear-motor-188-1-ratio-24mm-length-8mm-rex-shaft-30-rpm-3-3-5v-encoder/> (accessed Apr. 27, 2026).

[21] RobotShop, “High Torque Gearmotor Self-Locking Worm Gear Motor w/Encoder -12V, 105RPM,” *www.robotshop.com*, 2023. <https://www.robotshop.com/products/e-s-motor-high-torque-gearmotor-self-locking-worm-gear-motor-w-encoder-12v-105rpm>

[22] RobotShop, “FeeTech 180 Degrees Digital Servo 7.4V 35kg/cm FT5330M,” *RobotShop USA*, 2023. <https://www.robotshop.com/products/feetech-180-degrees-digital-servo-74v-35kg-cm-ft5330m>

[23] Mouser Electronics, “STMicroelectronics VNH5019ATR-E,” *www.mouser.com*. <https://www.mouser.com/ProductDetail/STMicroelectronics/VNH5019ATR-E?qs=T%2FOtf55vL7cDceYVZzSNKw%3D%3D>

[24] U.S. Department of Energy, “How Does Solar Work?,” *Energy.gov*, 2025. <https://www.energy.gov/eere/solar/how-does-solar-work>

[25] U.S. Energy Information Administration, “Photovoltaics and Electricity,” *Eia.gov*, May 24, 2024. <https://www.eia.gov/energyexplained/solar/photovoltaics-and-electricity.php>

[26] E. Faur and C. Speranza, “Introduction,” *ENERGEIA. ONLINE JOURNAL FOR LINGUISTICS, LANGUAGE PHILOSOPHY AND HISTORY OF LINGUISTICS*, pp. 1–14, Oct. 2025, doi: <https://doi.org/10.55245/energeia.2025.01>.

[27] Victron Energy, “Which solar charge controller: PWM or MPPT?,” *www.victronenergy.com*, Jan. 20, 2020. <https://www.victronenergy.com/upload/documents/Technical-Information-Which-solar-charge-controller-PWM-or-MPPT.pdf>

- [28] SunForgeLLC, “GV-5 – 5A MPPT Solar Charge Controller,” *Sunforgellc.com*, 2026. <https://sunforgellc.com/product/gv-5/> (accessed Apr. 27, 2026).
- [29] SunForgeLLC, “GV-5-MOD – 5A MPPT Solar Charge Controller,” *Sunforgellc.com*, 2026. <https://sunforgellc.com/product/gv-5-mod/> (accessed Apr. 27, 2026).
- [30] Victron Energy, “SmartSolar MPPT 75/10, 75/15, 100/15 & 100/20,” *Victron Energy*. <https://www.victronenergy.com/solar-charge-controllers/smartsolar-mppt-75-10-75-15-100-15-100-20>
- [31] Renogy, “Li 20/30/40 Amp Rover MPPT Solar Charge Controller | Renogy Solar,” *Renogy US*, 2025. <https://www.renogy.com/products/rover-li-40-amp-mppt-solar-charge-controller?variant=43178688675955> (accessed Apr. 27, 2026).
- [32] electronics notes, “Phototransistors vs Photodiodes: Choosing the Right Sensor for Your Light-Detection Application» Electronics Notes,” *Electronics-notes.com*, 2025. https://www.electronics-notes.com/articles/electronic_components/transistor/phototransistors-vs-photodiodes.php?utm_source=chatgpt.com
- [33] EEPower, “Photoresistor | Resistor Types | Resistor Guide,” *eepower.com*. <https://eepower.com/resistor-guide/resistor-types/photo-resistor>
- [34] H. Lastya *et al.*, “Performance Comparison Between LDR and Phototransistor Sensor for Dual-Axis Sun Tracker Sensor Based on Tetrahedron Geometry,” *scispace.com*, Feb. 25, 2023. <https://scispace.com/pdf/comparative-analysis-of-support-vector-machine-and-lusnsqph2b.pdf>
- [35] Adafruit Industries, “Adafruit TSL2591 High Dynamic Range Digital Light Sensor,” *www.adafruit.com*, 2023. <https://www.adafruit.com/product/1980>
- [36] S. C. | Arduino | 0, “How to Use Photoresistors to Detect Light on an Arduino,” *Circuit Basics*, Sep. 30, 2021. <https://www.circuitbasics.com/how-to-use-photoresistors-to-detect-light-on-an-arduino/>
- [37] “GL55 Series CdS Photoresistor Manual.” Available: <https://www.kth.se/social/files/54ef17dbf27654753f437c56/GL5537.pdf>

- [38] O. Nydal Dahl, "Phototransistor - A Newbie's Guide," *Build Electronic Circuits*, May 16, 2023. <https://www.build-electronic-circuits.com/phototransistor/>
- [39] N. Das, "Photodiode - A Beginner's Guide," *Build Electronic Circuits*, May 05, 2023. <https://www.build-electronic-circuits.com/photodiode/>
- [40] RFWirelessWorld, "Types of Phototransistor : Key differences & Applications," *Rfwireless-world.com*, 2026. <https://www.rfwireless-world.com/articles/phototransistor-basics-types-applications> (accessed Apr. 27, 2026).
- [41] W. Storr, "Phototransistor Basics including the Photodarlington," *Basic Electronics Tutorials*, Aug. 10, 2025. <https://www.electronics-tutorials.ws/transistor/phototransistor-basics.html>
- [42] Vishay Semiconductors, "Vishay Semiconductors Ambient Light Sensor," Jun. 2016. Available: <https://www.vishay.com/docs/81341/tept4400.pdf>
- [43] Vishay Semiconductors, "TEPT5600 Ambient Light Sensor," Nov. 2014. Available: <https://www.vishay.com/docs/84768/tept5600.pdf>
- [44] Excelitas Technologies, "VTT9812FH and VTT9814FH | IR-Bloc Ambient Light Sensor | IR-Blocking Silicon Phototransistor," *www.excelitas.com*, Jul. 01, 2016. https://www.excelitas.com/file-download/download/public/61631?filename=Excelitas_VTT9814FH_Phototransistor_specifications.pdf
- [45] Nigel, "SoC Estimation by Coulomb Counting," *Battery Design*, May 27, 2022. <https://www.batterydesign.net/soc-estimation-by-coulomb-counting/>
- [46] Renesas Electronics, "Coulomb Counting and State of Charge Estimation featuring the RAA489206/ISL94216A Battery Front End," *renesasweb.s3.us-west-2.amazonaws.com*, Jan. 23, 2023. <https://www.renesas.com/en/document/apn/coulomb-counting-and-state-charge-estimation-featuring-raa489206isl94216a-bfe>
- [47] Lithionics, "What is SOC Drift and how to avoid it?," *Lithionics*, Feb. 12, 2024. <https://lithionics.freshdesk.com/support/solutions/articles/154000127575-what-is-soc-drift-and-how-to-avoid-it->

- [48] Analog Devices, “60V Battery Gas Gauge with Temperature, Voltage and Current Measurement,” *www.analog.com*, 2017. <https://www.analog.com/media/en/technical-documentation/data-sheets/2944fa.pdf>
- [49] Texas Instruments, “INA226-Q1 AEC-Q100, 36V, 16-Bit, Ultra-Precise, I 2 C Output Current, Voltage, and Power Monitor With Alert,” Aug. 2024. Accessed: Apr. 27, 2026. [Online]. Available: https://www.ti.com/lit/ds/symlink/ina226-q1.pdf?HQS=dis-dk-null-digikeymode-dsf-pf-null-ww&ts=1777314223267&ref_url=https%253A%252F%252Fwww.ti.com%252Fgeneral%252Fdocs%252Fsuppproductinfo.tsp%253FdistId%253D10%2526gotoUrl%253Dhttps%253A%252F%252Fwww.ti.com%252Flit%252Fgpn%252Fina226-q1
- [50] Texas Instruments, “INA228-Q1 AEC-Q100, 85-V, 20-Bit, Ultra-Precise Power/Energy/Charge Monitor With I2C Interface,” *www.ti.com*, Jun. 2021. <https://www.ti.com/lit/ds/symlink/ina228-q1.pdf>
- [51] RF Wireless World, “Accelerometer vs Gyroscope: Key Differences Explained,” *Rfwireless-world.com*, 2026. <https://www.rfwireless-world.com/terminology/accelerometer-vs-gyroscope> (accessed Apr. 27, 2026).
- [52] RF Wireless World, “Magnetometer Basics: Types and Applications,” *Rfwireless-world.com*, 2026. <https://www.rfwireless-world.com/articles/magnetometer-basics-types-and-applications> (accessed Apr. 27, 2026).
- [53] EPSON, “What are accelerometers? Explain how they works, what they measures, and how they are used | Technical Column | Technical Information | Epson Sensing Device,” *Epsondevice.com*, 2025. <https://www.epsondevice.com/sensing/en/tech/column/accelerometer/>
- [54] Global Spec, “Accelerometers Selection Guide: Types, Features, Applications | Engineering360,” *www.globalspec.com*, 2026. https://www.globalspec.com/learnmore/sensors_transducers_detectors/acceleration_vibration_sensing/accelerometers
- [55] SparkFun, “Gyroscope - learn.sparkfun.com,” *Sparkfun.com*, 2019. <https://learn.sparkfun.com/tutorials/gyroscope/all>
- [56] Matan, “MEMS gyroscope | How it works, Application & Advantages,” *Electricity - Magnetism*, Oct. 26, 2023. <https://www.electricity-magnetism.org/mems-gyroscope/>

- [57] Jacob, “Getting Started With Inertial Measurement Units | Exploring Degrees Of Freedom - Tutorial Australia,” *Core Electronics*, Jul. 19, 2023. <https://core-electronics.com.au/guides/getting-started-with-inertial-measurement-units-exploring-degrees-of-freedom/>
- [58] Hillcrest Labs, “BNO08X Datasheet Revision 1.17 BNO08X Data Sheet,” CEVA, 2023. Available: https://www.ceva-ip.com/wp-content/uploads/BNO080_085-Datasheet.pdf
- [59] InvenSense Inc., “MPU-6000 and MPU-6050 Product Specification Revision 3.4 MPU-6000/MPU-6050 Product Specification,” Aug. 2013. Available: <https://invensense.tdk.com/wp-content/uploads/2015/02/MPU-6000-Datasheet1.pdf>
- [60] InvenSense, “GENERAL DESCRIPTION,” Mar. 2024. Available: <https://invensense.tdk.com/wp-content/uploads/2024/03/DS-000189-ICM-20948-v1.6.pdf>
- [61] M. Toa and A. Whitehead, “Ultrasonic Sensing Basics,” Dec. 2021. Accessed: Apr. 27, 2026. [Online]. Available: <https://www.ti.com/lit/an/slaa907d/slaa907d.pdf?ts=1771567279190>
- [62] MaxBotix, “LV-MaxSonar-EZ Datasheet,” *MaxBotix*, 2024. <https://maxbotix.com/pages/lv-maxsonar-ez-datasheet>
- [63] MaxBotix, “I2CXL-MaxSonar-EZ Datasheet,” *MaxBotix*, 2023. <https://maxbotix.com/pages/i2cxl-maxsonar-ez-datasheet>
- [64] MaxBotix, “HRLV-MaxSonar-EZ Datasheet,” *MaxBotix*, 2026. <https://maxbotix.com/pages/hrlv-maxsonar-ez-datasheet>
- [65] Simplify3D, “Ultimate 3D Printing Material Properties Table,” *Simplify3D*, 2024. <https://www.simplify3d.com/resources/materials-guide/properties-table/>
- [66] D. McClements and P. Keane, “Aluminum: Properties, Uses, and Benefits,” *Thomasnet.com*, Aug. 22, 2025. <https://www.thomasnet.com/articles/metals-metal-products/aluminum/>

- [67] S. Mirdehghan, “Carbon Fiber Properties - an overview | ScienceDirect Topics,” *www.sciencedirect.com*, 2021. <https://www.sciencedirect.com/topics/materials-science/carbon-fiber-properties>
- [68] MatWeb, “MatWeb - The Online Materials Information Resource,” *www.matweb.com*, 2026. <https://www.matweb.com/search/DataSheet.aspx?MatGUID=96d768abc51e4157a1b8f95856c49028&ckck=1>
- [69] MatWeb, “MatWeb - The Online Materials Information Resource,” *www.matweb.com*, 2026. <https://www.matweb.com/search/DataSheet.aspx?MatGUID=3a2e111b27ef4e5d813bad6044b3f318&ckck=1>
- [70] MatWeb, “MatWeb - The Online Materials Information Resource,” *www.matweb.com*, 2026. <https://www.matweb.com/search/DataSheet.aspx?MatGUID=4f19a42be94546b686bbf43f79c51b7d>
- [71] A. Medina-Santiago, L. A. Morales-Rosales, C. A. Hernández-Gracidas, I. Algreto-Badillo, A. D. Pano-Azucena, and J. A. Orozco Torres, “Reactive Obstacle–Avoidance Systems for Wheeled Mobile Robots Based on Artificial Intelligence,” *Applied Sciences*, vol. 11, no. 14, p. 6468, Jul. 2021, doi: <https://doi.org/10.3390/app11146468>.
- [72] W. Franklin, “Kalman Filter Explained Simply,” *The Kalman Filter*, Dec. 31, 2020. <https://thekalmanfilter.com/kalman-filter-explained-simply/>
- [73] G. Sarraf, “Implementing a Kalman Filter in Python for Sensor Fusion,” *ThinkRobotics.com*, Jul. 29, 2025. <https://thinkrobotics.com/blogs/learn/learn-to-design-and-3d-print-robotic-grippers-discover-materials-mechanisms-sensors-and-optimization-techniques-for-effective-robotic-arm-end-effectors>
- [74] I. Reda and A. Andreas, “Solar Position Algorithm for Solar Radiation Applications,” *Nlr.gov*, Jan. 2008. <https://docs.nlr.gov/docs/fy08osti/34302.pdf>
- [75] React Native, “React Native · A framework for building native apps using React,” *reactnative.dev*, 2024. <https://reactnative.dev/>

- [76] Apple Inc., “SwiftUI - Apple Developer,” *Apple Developer*, 2025. <https://developer.apple.com/swiftui/>
- [77] Flutter, “Flutter - Beautiful native apps in record time,” *Flutter.dev*. <https://flutter.dev/>
- [78] FastAPI, “FastAPI,” *fastapi.tiangolo.com*, 2023. <https://fastapi.tiangolo.com/>
- [79] Django, “The Web framework for perfectionists with deadlines | Django,” *Djangoproject.com*, 2024. <https://www.djangoproject.com/>
- [80] Spring, “Spring Projects,” *Spring.io*, 2019. <https://spring.io/projects/spring-boot>
- [81] OpenJS Foundation, “Express - Node.js web application framework,” *Expressjs.com*, 2017. <https://expressjs.com/>
- [82] Dell, “Wi-Fi Standards: IEEE 802.11ac, 802.11ax, and Wireless Internet Standards,” *Dell.com*, Dec. 09, 2024. <https://www.dell.com/support/content/en-ke/article/product-support/self-support-knowledgebase/networking-wifi-and-bluetooth/wi-fi-network-standards-overview>
- [83] M. E. Yuksel, “Power Consumption Analysis of a Wi-Fi-based IoT Device,” *Electrica*, vol. 20, no. 1, pp. 62–70, Feb. 2020, doi: <https://doi.org/10.5152/electrica.2020.19081>.
- [84] Y. Huang, “Complete Bluetooth Versions Guide: From 1.0 to 6.2,” *MOKOSmart #1 Smart Device Solution in China*, Nov. 05, 2025. https://www.mokosmart.com/guide-on-different-bluetooth-versions/#elementor-toc__heading-anchor-5 (accessed Apr. 27, 2026).
- [85] GeeksforGeeks, “Bluetooth vs Bluetooth Low Energy,” *GeeksforGeeks*, May 15, 2022. <https://www.geeksforgeeks.org/computer-networks/bluetooth-vs-bluetooth-low-energy/>
- [86] The Things Network, “What are LoRa and LoRaWAN?,” *The Things Network*, 2020. <https://www.thethingsnetwork.org/docs/lorawan/what-is-lorawan/>
- [87] Amazon Web Services, “LoRaWAN - Implementing Low-Power Wide-Area Network (LPWAN) Solutions with AWS IoT,” *Amazon.com*, 2025.

<https://docs.aws.amazon.com/whitepapers/latest/implementing-lpwan-solutions-with-aws/lorawan.html>

[88] Connectivity Standards Alliance, “Zigbee FAQs | Frequently Asked Questions,” *CSA*. <https://csa-iot.org/all-solutions/zigbee/zigbee-faq/>

[89] Exspenditure, “Why I’d Never Buy NMC Again: LiFePO₄ vs NMC Battery Review,” *Exspenditure*, Jan. 28, 2026. <https://exspenditure.com/lifepo4-vs-nmc-portable-power-station-battery-2026/> (accessed Apr. 27, 2026).

[90] Niclas, “Energy density and Specific Energy of Battery,” *sinovoltaics.com*, 2020. <https://sinovoltaics.com/learning-center/storage/energy-density-and-specific-energy-of-battery/>

[91] Dr.Prepare, “12V 20Ah LiFePO₄ Lithium Deep Cycle Battery,” *Dr.Prepare*, 2026. https://drprepare.com/products/12v-20ah-lifepo4-lithium-deep-cycle-battery?_pos=10&_sid=8236c4fb3&_ss=r (accessed Apr. 27, 2026).

[92] Bioenno Power, “12V, 20Ah LFP Battery (ABS, BLF-1220AS),” *Bioenno Power*, 2026. https://www.bioennopower.com/products/12v-20ah-lfp-battery-abs-blf-1220as?_pos=16&_fid=ce683e851&_ss=c (accessed Apr. 27, 2026).

[93] UltraLife, “URB12200 Technical Datasheet,” *Mouser.com*, Oct. 28, 2025. https://www.mouser.com/datasheet/3/186/1/TDS_URB12200.pdf (accessed Apr. 27, 2026).

[94] Renogy, “Mini Size 12V 20Ah Trolling Motor Lithium Iron Phosphate Battery,” *Renogy US*, 2026. <https://www.renogy.com/products/12v-20ah-trolling-motor-lithium-iron-phosphate-battery> (accessed Apr. 27, 2026).

[95] B. Lee, “Application Report Understanding the Terms and Definitions of LDO Voltage Regulators,” 1999. Available: <https://www.ti.com/lit/an/slva079/slva079.pdf>

[96] Texas Instruments, “Understanding Buck Power Stages in Switchmode Power Supplies Application Report SLVA057,” 1999. Available: <https://www.ti.com/lit/an/slva057/slva057.pdf>

[97] S. Ziel, “Tips, Tricks and Advanced Applications of Linear Regulators,” Mar. 2024. Available: <https://www.ti.com/lit/ml/slup415/slup415.pdf>

- [98] Texas Instruments, “LM2679 SIMPLE SWITCHER® 5-A Step-Down Voltage Regulator With Adjustable Current Limit,” *www.ti.com*, Jun. 2016.
https://www.ti.com/lit/ds/symlink/lm2679.pdf?ts=1624370319812&ref_url=https%253A%252F%252Fwww.ti.com%252Fproduct%252FLM2679%253FkeyMatch%253DLM2679%2526tisearch%253Dsearch-everything%2526usecase%253DGPN
- [99] EVVOSEMI, “Product Specification | LM2576/LM2576HV,” *mm.digikey.com*. https://mm.digikey.com/Volume0/opasdata/d220001/medias/docus/7158/5272_LM2576HVS-ADJ-EV.pdf
- [100] Semiconductor Components Industries, “3.0 A, Step-Down Switching Regulator,” *www.onsemi.com*, Dec. 2022.
<https://www.onsemi.com/pdf/datasheet/lm2596-d.pdf>
- [101] Analog Devices, “Dual Channel 4A, 42V, Synchronous Step-Down Silent Switcher 2 with 6.2 μ A Quiescent Current,” *www.analog.com*, Nov. 2024.
<https://www.analog.com/media/en/technical-documentation/data-sheets/lt8650s.pdf>
- [102] Linear Technology, “Dual Channel 3A, 15V Monolithic Synchronous Step-Down Regulator,” *www.analog.com*, May 2015. <https://www.analog.com/media/en/technical-documentation/data-sheets/3633fd.pdf>
- [103] NXP Semiconductors, “UM10204 I 2 C-bus specification and user manual Rev. 6 -4 April 2014 User manual Document information Info Content,” Oct. 2021.
Available: <https://www.nxp.com/docs/en/user-guide/UM10204.pdf>
- [104] E. Peña and M. Legaspi, “UART: A Hardware Communication Protocol Understanding Universal Asynchronous Receiver/Transmitter,” *VISIT ANALOG.COM*, vol. 54, no. 4, 2020, Available: <https://www.analog.com/media/en/analog-dialogue/volume-54/number-4/uart-a-hardware-communication-protocol.pdf>
- [105] Y. Pan, “Understanding the SPI Bus,” Aug. 2025. Accessed: Apr. 27, 2026. [Online]. Available: <https://www.ti.com/lit/ab/sboa621/sboa621.pdf?ts=1774451886756>
- [106] Bluetooth, “The Bluetooth® Low Energy Primer | Bluetooth® Technology Website,” *Bluetooth® Technology Website*, 2019. https://www.bluetooth.com/bluetooth-le-primer/#mcetoc_1iiprfme5 (accessed Apr. 27, 2026).

- [107] IEEE SA, “IEEE Standard for Information Technology- Telecommunications and Information Exchange between Systems Local and Metropolitan Area Networks-Specific Requirements Part 11: Wireless LAN Medium Access Control (MAC) and Physical Layer (PHY) Specifications,” 2020. Available: <https://thewifiofthings.com/wp-content/uploads/2021/08/802.11-2020-Preview.pdf>
- [108] IPC, “IPC-2221 Generic Standard on Printed Board Design The Institute for Interconnecting and Packaging Electronic Circuits,” 1998. Available: https://www-eng.lbl.gov/~shuman/NEXT/CURRENT_DESIGN/TP/MATERIALS/IPC_2221.pdf
- [109] ISO, “ISO 5807 -1985,” *Scribd*, 2026.
<https://www.scribd.com/document/519991367/ISO-5807-1985>

Appendix B – AI References

[1a] **Case Study 1 Prompt:** “Write a concise explanation comparing the five most common battery chemistries used in solar systems, using natural paragraphs instead of structured tables.”

[1b] **Case Study 1 ChatGPT Deep Research Outcome: “Solar Battery Chemistries: An Executive Summary**

Among common solar batteries, lithium iron phosphate (LiFePO₄) leads for long life and safety, while standard lithium-ion (NMC) packs more energy per weight. Traditional lead–acid types cost less upfront but age faster and need upkeep. AGM/gel (sealed lead–acid) improve on flooded types with less maintenance. NiMH batteries are rarely used in large solar systems, finding niche use in small solar lights and electronics.

LiFePO₄ (Lithium Iron Phosphate)

LiFePO₄ batteries have become the go-to choice for home solar storage because they last a long time and are very stable. A good LiFePO₄ battery usually gives 8–12 years of life (often 10–15) with a very high cycle count – typically 2,000–6,000 full cycles at about 80% depth-of-discharge. They retain voltage well across discharge, can be drained deeply daily (most systems use 20–80% each cycle), and require essentially no maintenance (no watering or cleaning corrosion). LiFePO₄ chemistry is very thermally stable and safe (far less fire risk than other lithium types). Round-trip efficiency is excellent – often above 92%, meaning most of the solar power you store comes back out. The trade-off is higher initial cost (more than lead-acid), but because they last so much longer with no upkeep, their cost per usable kilowatt-hour is low.

NMC (Nickel-Manganese-Cobalt Lithium-Ion)

Standard lithium-ion (NMC) batteries are common in electronics and some solar products. They pack more energy into a given space than LiFePO₄ (higher energy density) but have some tradeoffs. A typical NMC home battery can cycle a few thousand times (often up to 3,000–5,000 partial cycles), giving roughly a 10–12 year lifespan with good care. NMC units have very high round-trip efficiency (around 90–95% or even up to ~98% under ideal conditions), so you lose little energy charging them. However, they are more expensive upfront (especially since they use cobalt) and their chemistry is less forgiving: they require strong battery management to avoid overheating or overcharging. In practice, NMC batteries deliver large bursts of power (good C-rate) and are lighter, making them attractive when space is tight (e.g. wall-mounted storage or electric vehicles). But for typical home solar use, their slightly shorter life and greater sensitivity mean LiFePO₄ is often preferred unless weight/size is critical.

Flooded Lead–Acid

Flooded (wet cell) lead–acid batteries are the classic, lowest-cost option. You’ll often see them in off-grid cabins or diesel backup because you can get a lot of capacity for cheap. They perform adequately but have important drawbacks: their round-trip efficiency is only about 75–85% (a lot of energy is lost as heat), and you should never regularly discharge them beyond ~50% without hurting their life. Typical cycle life is low (hundreds of cycles), so in a daily solar-use scenario they usually last only 3–5 years before they fall below ~80% capacity. Maintenance is high: flooded cells must be kept topped up with distilled water, cleaned of corrosion, and equalized regularly. They also emit hydrogen gas when charged, so they need ventilation and safe handling. On the plus side, flooded lead–acid has the lowest upfront cost of all chemistries, making it a budget-friendly starter choice. But owners should expect more frequent replacements, lower efficiency, and a lifespan of just a few years unless the system is very carefully managed.


AGM/VRLA (Sealed Lead–Acid – AGM and Gel)

AGM and gel batteries are sealed varieties of lead–acid. They use absorbed glass mats or gelled electrolyte so you don’t have to add water. They bridge the gap between flooded lead–acid and lithium. Efficiency is still on the order of 80–85% (similar to flooded), but sealed construction means much lower routine maintenance (no watering, no acid spills). Sealed lead–acid typically achieves a bit longer life – roughly 800–1,500 cycles or about 4–8 years in a solar cycling setup. They handle deep discharges somewhat better than flooded, but manufacturers still often limit you to ~50–70% DoD to maximize life. Safety is fair: AGM/gel will not leak, but they can overheat if overcharged and lack the intrinsic safety of LiFePO_4 . Upfront cost is higher than flooded lead–acid (usually a medium price point), yet they remove the fuss of maintenance. In practice, VRLA batteries are often used for small to mid-size off-grid kits or backup systems where you want sealed convenience, but they will still need replacement roughly every 5–8 years under normal solar usage.

NiMH (Nickel–Metal Hydride) and Other Nickel-Based

Nickel–metal hydride batteries are unusual in home solar. They were common in older consumer electronics and hybrid cars, but rarely used for whole-house storage. NiMH cells are safer than the old NiCd (no toxic cadmium) and have decent charge efficiency (charging up to ~70% state is nearly 100% efficient). However, they have lower energy density and higher self-discharge, so they aren’t a great fit for large stationary use. NiMH cycle life is modest – on the order of 300–500 full cycles. In practical terms that means if used daily, they might only last ~1–3 years before needing replacement. Typical round-trip efficiency is lower than lithium (charging generates heat near end of charge, wasting energy), and they still require proper charging termination to avoid overheat. Because of this, NiMH tends to show up only in very small solar applications: for example, garden lights, sensors, or portable devices often use AA- or C-size NiMH packs powered by a solar panel. For a homeowner looking at a solar battery bank, NiMH

is not a competitive choice – it’s cheap to buy, but it has short life, modest efficiency, and high self-discharge, meaning it’s rarely used for home energy storage.



”

[1c] Case Study 1 Gemini Deep Research Outcome: “*The Global Transition of Energy Storage: A Comparative Analysis of Primary Battery Chemistries in Solar Systems*”

The global trajectory of renewable energy integration has reached a pivotal juncture in 2026, where the efficacy of solar power is no longer defined solely by photovoltaic efficiency but by the sophistication of the associated energy storage systems. As the decentralized energy landscape matures, the selection of battery chemistry has evolved from a secondary technical consideration into a primary strategic pillar for residential, commercial, and utility-scale stakeholders. This transformation is underpinned by a staggering 93% decline in system-level battery storage costs between 2010 and 2024, falling from \$2,571 per kilowatt-hour (kWh) to approximately $192/\text{kWh}$.¹ This economic shift, coupled with intensified research into material science, has solidified five primary chemistries as the dominant architectures for solar storage: Lithium Iron Phosphate (LiFePO_4 or LFP), Lithium Nickel Manganese Cobalt Oxide (NMC), Lead-Acid variants, Redox Flow systems, and the burgeoning Sodium-Ion (Na-ion) technology.²

The fundamental role of energy storage in a solar-integrated grid is the temporal shifting of power—capturing generation during peak solar irradiance for use during evening peak demand or providing critical resilience during grid outages.⁵ However, the physical and chemical constraints of each battery type dictate its suitability for specific operational profiles. While high energy density remains the priority for the mobile transportation sector, the stationary storage market has pivoted toward safety, cycle life, and leveled cost, leading to a significant divergence in chemical adoption.⁷ By 2025, the global residential battery storage market was valued at \$6.85 billion, with projections suggesting a rise to \$8.73 billion by 2026, reflecting a compound annual growth rate (CAGR) of 17.80%.⁵

Lithium Iron Phosphate (LiFePO_4): The Technical and Economic Benchmark
Lithium Iron Phosphate (LFP) has emerged as the unequivocal leader in the 2026 stationary storage market, capturing an estimated 85% share of utility-scale projects and a dominant portion of new residential installations.¹ The ascent of LFP is

not a result of superior energy density—it actually offers less capacity per kilogram than other lithium variants—but rather its exceptional safety profile and lifecycle economics.²

Chemical Architecture and Thermal Stability

LFP batteries utilize an iron phosphate cathode, which provides a robust olivine crystalline structure. This structure is significantly more stable than the metal-oxide layers found in chemistries like NMC or LCO.⁸ In practical terms, the $P - O$ bond in the phosphate cathode is much stronger, which prevents the release of oxygen during chemical decomposition or thermal stress.⁸ This characteristic virtually eliminates the risk of thermal runaway, a self-sustaining heating cycle that can lead to fires or explosions.²

The thermal stability of LFP is a decisive factor in the 2026 regulatory environment. LFP cells remain stable at temperatures exceeding 200°C (392°F), whereas other lithium chemistries may reach critical instability at temperatures as low as 150°C .⁹ This inherent safety simplifies permitting processes and reduces the need for expensive, high-capacity fire suppression systems, lowering the balance-of-system (BOS) costs for commercial and industrial (C&I) projects.⁸

Performance Metrics and Lifecycle Longevity

From an operational standpoint, LFP offers a superior value proposition over its lifespan. While its upfront cost can range from \$600 to \$900 per kWh for high-performance residential systems, it delivers between 6,000 and 10,000 full charge-discharge cycles before dropping below 80% of its original capacity.⁷ In a typical daily solar cycling application, this translates to a service life of 15 to 20 years, far outlasting the 10-year industry standard of the previous decade.⁷

LFP batteries also allow for a significant depth of discharge (DoD), often rated between 95% and 100%.² This means a homeowner with a 10kWh LFP system can utilize nearly all the stored energy every day without damaging the internal chemistry. When compared to lead-acid systems that are limited to a 50% DoD to preserve health, the usable energy of an LFP system is effectively double for the same nameplate capacity.¹³

Metric	Lithium Iron Phosphate (LFP)	Industry Context (2025-2026)
Cycle Life	6,000 – 10,000 cycles	Highest among lithium variants ⁷
Depth of Discharge	95% – 100%	Maximizes usable capacity ⁷
Round-Trip Efficiency	95% – 98%	Minimal energy loss ⁷
Thermal Stability	Stable up to >	Safe for indoor installation ⁹

Energy Density	120 – 160 Wh/kg	Lower than NMC, higher than lead-acid ⁷
Lifespan	15 – 20 years	Ideal for long-term solar ROI ⁷

Lithium Nickel Manganese Cobalt Oxide (NMC): The High-Density Specialist

*Lithium Nickel Manganese Cobalt Oxide (NMC) represents the primary alternative within the lithium-ion family for solar applications. Historically favored for its high energy density, NMC has seen its share of the stationary storage market decline in favor of LFP, yet it remains a critical choice for space-constrained environments.*⁷

Energy Density and Space Constraints

*NMC batteries utilize a mixed-metal cathode that enables significantly higher energy density than LFP, typically ranging from 150 to 250 **Wh/kg**.⁷ This density allows manufacturers to pack a larger amount of energy into a smaller, lighter enclosure. For urban residential settings or commercial retrofits where physical space is at an absolute premium, NMC remains a vital technology.⁷ A standard NMC-based storage unit might be 30% to 50% more compact than an LFP unit of equivalent capacity, facilitating easier wall mounting and reducing the footprint in garages or utility rooms.*¹⁰

Trade-offs: Lifecycle and Safety

*The high energy density of NMC comes with inherent trade-offs in longevity and safety. NMC batteries generally offer a shorter cycle life, typically between 3,000 and 5,000 cycles.⁷ In a daily cycling solar application, this results in a service life of 10 to 15 years, requiring replacement sooner than an LFP alternative.⁷ Additionally, to maintain the health of the chemistry, installers often recommend limiting the depth of discharge to 80% or 90%, which reduces the "true" usable energy compared to the nameplate rating.*¹³

*Safety remains a significant point of comparison. NMC chemistries are more susceptible to thermal runaway due to the nature of the metal-oxide bonds in the cathode, which can release oxygen during decomposition and fuel a fire.⁸ In the 2026 market, this has led to increased scrutiny from insurers and local fire authorities, often necessitating more complex thermal management systems and increased physical spacing between units, which can offset some of the initial space-saving advantages.*⁸

Feature	Lithium NMC	Strategic Implication
Energy Density	150 – 250 Wh/kg	Best for space-constrained sites ⁷
Cycle Life	3,000 – 5,000 cycles	Shorter ROI window than LFP ⁷
Depth of Discharge	80% – 90%	Limited usable energy per cycle ¹³

Safety Rating	Moderate	Requires advanced BMS and cooling ⁷
Cost per kWh	\$500 – \$700	Historically cheaper, now competing with LFP ⁷

Lead-Acid Storage: Evolutionary Persistence in a Lithium-Dominated Era

Lead-acid batteries are the oldest form of rechargeable energy storage and, despite the rapid rise of lithium, they maintain a presence in the 2026 solar market.¹² Their persistence is driven by low upfront costs and a deeply established global infrastructure for manufacturing and recycling.⁷

Variations: Flooded, AGM, and Gel

The lead-acid market for solar is divided into three primary categories: Flooded (Wet Cell), Absorbed Glass Mat (AGM), and Gel.

1. **Flooded Lead-Acid (FLA):** *These are the most economical upfront but require significant maintenance. Users must regularly check and replenish distilled water levels and perform equalization charges to prevent sulfation—the buildup of lead sulfate crystals on the plates.⁷ They also emit hydrogen gas during charging, necessitating well-ventilated dedicated enclosures.⁷*
2. **AGM (Absorbed Glass Mat):** *A sealed, maintenance-free variant where the electrolyte is held in fiberglass mats. AGM batteries offer better efficiency than flooded cells and handle high discharge rates more effectively.⁷ They are common in budget-conscious backup systems that are not cycled daily.⁷*
3. **Gel Batteries:** *These use a silica additive to turn the electrolyte into a stiff gel, making them exceptionally resistant to spills and more stable in extreme temperatures.⁷ Gel batteries are particularly suited for off-grid applications in harsh climates, although they are highly sensitive to overcharging.¹²*

Efficiency and Depth of Discharge Constraints

The primary limitation of lead-acid technology in a solar context is its inefficiency and short lifespan. Lead-acid batteries generally operate at 70% to 85% round-trip efficiency, compared to over 95% for lithium.¹¹ This means a significant portion of the solar energy collected is lost as heat during the storage process.

Furthermore, to prevent permanent chemical damage, lead-acid batteries should not be discharged below 50% of their capacity.⁷ A user requiring ^{10kWh} of usable daily energy would need to install a ^{20kWh} lead-acid bank. When combined with a cycle life that often falls below 1,000 cycles for flooded cells, the frequent replacement requirements often make lead-acid the most expensive option on a levelized cost basis over 10 to 15 years.¹¹

Type	Cycle Life (50% DoD)	Maintenance Requirement	Ideal Use Case
------	----------------------	-------------------------	----------------

Flooded (FLA)	500 – 1,000	High (Water/Venting)	Budget off-grid ¹²
AGM	1,000 – 1,500	None (Sealed)	Backup/Standby ¹²
Gel	1,000 – 1,500	None (Sealed)	Extreme temperatures ¹²
Efficiency	60% – 85%	N/A	Lowest upfront cost ¹³

Redox Flow Battery Systems: Architectures for Multi-Hour Grid Resilience

Redox Flow Batteries (RFB) represent a fundamental departure from the "closed-cell" design of lithium and lead-acid. Instead of storing energy in solid electrodes, flow batteries store energy in liquid electrolyte solutions held in external tanks.² These liquids are pumped through a central membrane stack where the chemical reaction occurs.²

Decoupling Power and Energy

The unique advantage of flow batteries is the ability to decouple power capacity (determined by the size of the membrane stack) from energy capacity (determined by the volume of the electrolyte tanks).¹⁶ For a solar farm needing to bridge 10 hours of darkness, a flow battery can be scaled simply by increasing the size of its storage tanks, providing a lower marginal cost for long-duration storage.¹⁴

In 2026, the two primary flow technologies are Vanadium Redox Flow Batteries (VRFB) and Iron Flow Batteries. Vanadium is the more mature technology, offering high electrochemical stability and the ability to reuse the electrolyte indefinitely.¹⁶ Iron flow batteries are emerging as a cheaper alternative, utilizing abundant, non-toxic materials, although they are still moving toward large-scale commercial saturation.¹⁶

Lifespan and Safety Profile

Flow batteries are unparalleled in their longevity. They can be cycled over 20,000 times with zero degradation, often lasting 25 to 30 years—effectively matching the lifespan of the solar panels themselves.¹⁴ They also allow for 100% depth of discharge every single day without any damage to the system.¹⁴

Safety is another hallmark of flow technology. Because the electrolytes are typically water-based and non-flammable, they possess zero fire risk.¹⁴ This allows them to be classified as "non-hazardous" for indoor or near-building installation, saving significant costs on fire suppression and specialized insurance.¹⁴ However, their low energy density (20 to 35 **Wh/kg**) and the footprint required for liquid tanks make them unsuitable for residential garages, relegating them to utility-scale solar farms or heavy industrial sites.¹¹

Characteristic	Vanadium Flow (VRFB)	Strategic Advantage
Cycle Life	20,000 – 30,000+	"Buy once" for 25-30 years ¹⁴
Depth of Discharge	100%	Full utilization of tanks ¹⁴
Safety	Non-flammable	No thermal runaway risk ¹⁷

Degradation	Near zero	Electrolyte never wears out ¹⁹
Physical Size	Very Large	Scalable for grid-level storage ¹⁴
Efficiency	70% – 80%	Lower than lithium due to pumps ¹⁶

Sodium-Ion (Na-ion) Technology: The Frontier of Resource-Agnostic Storage

As of 2026, Sodium-Ion (Na-ion) technology has emerged as a disruptive contender to the lithium-ion monopoly. Its value proposition is centered on resource abundance and supply chain stability, as sodium is approximately 500 times more abundant than lithium and can be sourced globally from seawater or salt deposits.⁴

Material Advantage and Manufacturing Synergy

Sodium-ion batteries function on a principle similar to lithium-ion, but utilize sodium ions (Na^+) which are larger and heavier than lithium ions (Li^+).²¹ This necessitates different cathode and anode materials—typically hard carbon for the anode instead of graphite.²¹ A critical economic advantage for sodium-ion is the ability to use aluminum current collectors for both the cathode and the anode, whereas lithium cells require more expensive copper for the anode.⁴

Remarkably, sodium-ion is a "drop-in" technology that can be manufactured on existing lithium-ion production lines with only minor modifications.²² This has allowed major manufacturers, particularly in China (which accounts for over 60% of global storage production), to rapidly scale sodium-ion as a hedge against lithium price volatility.⁶

Performance in Extreme Conditions

Sodium-ion batteries exhibit superior performance in extreme temperatures compared to standard lithium chemistries. They can discharge normally in environments as cold as -40°C and maintain 80% capacity at -20°C , a range where LFP batteries often struggle or require energy-intensive heating systems.²⁰ Furthermore, they are inherently safer to transport; because they can be discharged to zero volts without damage, they can be shipped as "non-dangerous" goods without the residual fire risk associated with lithium batteries.⁴ While their current energy density (110-160 Wh/kg) is on the lower end of the lithium spectrum, they represent the most promising path toward low-cost, sustainable mass storage.²⁰

Attribute	Sodium-Ion (Na-ion)	LFP Comparison
Raw Material Cost	Very Low	Higher (Lithium/Iron) ²⁰
Low-Temp Performance	Excellent (-40°C)	Moderate ²⁰
Shipping Safety	Safe at 0 Volts	Risk at 0 Volts ⁴

Charge Speed	Up to 90% in 15 mins	Slower (typically 1-2 hours) ²⁰
Cycle Life	2,000 – 6,000 (Current)	6,000 – 10,000 ²⁰
Current Status	Early commercialization	Mature market leader ²²

Comparative Economic Analysis and Levelized Cost of Storage (LCOS)

To evaluate these chemistries for professional deployment, one must look beyond the initial capital expenditure (**CAPEX**) and calculate the Levelized Cost of Storage (**LCOS**). This metric represents the cost per unit of energy delivered over the entire lifetime of the system, accounting for degradation, efficiency, and replacement frequency.¹⁴

Upfront Price vs. Long-Term Value

In 2026, LFP systems often appear more expensive upfront (\$650-\$850/kWh installed) compared to lead-acid (\$150-\$300/kWh).¹³ However, when viewed over a 25-year horizon, the picture flips. In regions like New Zealand, the **LCOS** for LFP is estimated at 18–28 cents per **kWh**, while lead-acid exceeds 30–40 cents due to the need for multiple pack replacements and lower usable capacity.¹⁴ Flow batteries present the most dramatic long-term savings. While they have the highest upfront cost (\$900-\$1,200/kWh for smaller C&I systems), their 25-year **LCOS** drops to 11–17 cents per **kWh** because they essentially never wear out and require minimal maintenance.¹⁴ For large-scale utility projects, the **LCOS** for flow systems is increasingly the benchmark for grid stability tenders.¹⁴

Future Cost Trajectories (2030-2050)

Projections for sodium-ion suggest that it will eventually provide the lowest **LCOS** for medium-duration storage. By 2050, sodium-ion batteries are expected to deliver storage at 11–14 Euros per **MWh**, significantly cheaper than lithium-ion projections of 16–22 Euros per **MWh**.²² This trend is driven by high learning rates and the avoidance of critical mineral scarcity.²²

System Type	2026 Upfront (\$/kWh)	25-Year LCOS (\$/kWh)	Service Frequency
Lead-Acid	\$150 – \$300	\$0.25 – \$0.40	Monthly ¹³
Lithium LFP	\$600 – \$900	\$0.18 – \$0.28	Minimal ⁷
Vanadium Flow	\$900 – \$1,200	\$0.11 – \$0.17	Annual/Bi-annual ¹⁴
NMC	\$500 – \$700	\$0.20 – \$0.30	Low ⁸
Sodium-Ion	\$400 – \$600	TBD (Falling)	Minimal ²⁰

Safety, Regulation, and the Built Environment

The integration of solar batteries into residential and commercial buildings is governed by increasingly strict safety codes, notably UL 9540 (Energy Storage Systems) and NFPA 855 (Standard for the Installation of Stationary Energy Storage Systems).⁷ These regulations dictate minimum clearances from property lines, ventilation requirements, and fire suppression protocols.⁷

Thermal Management and Installation

For lithium-ion (particularly NMC) and lead-acid, thermal management is a critical design constraint. These chemistries require active cooling or heating to maintain an optimal operating range of 15°C to 35°C .¹¹ Operating outside these bounds significantly accelerates degradation and increases safety risks. LFP and sodium-ion offer more flexibility, operating effectively in wider temperature windows (-4°F to 140°F for LFP), which reduces the energy consumed by the system's own climate control—often called "parasitic load"—by 5% to 10%.⁷

Environmental and Ethical Considerations

The sustainability of battery chemistry is becoming a differentiator for corporate buyers. Lead-acid batteries are the standard for recyclability, with a mature closed-loop system.¹³ Lithium-ion recycling is a growing but complex industry, with LFP being harder to recycle profitably because it lacks high-value nickel and cobalt.¹⁰ However, LFP and sodium-ion are considered more "ethical" as they avoid the environmental and human rights issues associated with cobalt mining.² Flow batteries are perhaps the most sustainable, as the electrolyte is 100% reusable and the mechanical parts are easily salvaged.¹⁴

Strategic Deployment Summary

Choosing the correct solar battery chemistry in 2026 requires balancing physical constraints against lifecycle economic goals. For standard residential and C&I applications where safety and lifespan are paramount, Lithium Iron Phosphate (LiFePO_4) remains the industry benchmark. For space-critical retrofits, Lithium NMC continues to provide the necessary density, albeit with a higher safety and replacement overhead.

Redox Flow batteries are the definitive choice for long-duration, utility-scale projects where the goal is grid-level resilience over a 30-year timeframe. Lead-acid technologies occupy a niche for low-use backup or remote, budget-constrained off-grid sites. Finally, Sodium-Ion represents the most promising frontier for resource-secure, low-cost mass deployment, particularly in climates with extreme temperatures.

The convergence of these five chemistries allows the 2026 solar market to address a diverse range of needs, from a single home in a suburb to a multi-megawatt industrial solar park. As manufacturing continues to scale and costs continue their downward trajectory, energy storage has transformed from a luxury addition into the essential heartbeat of the global energy transition.”

[2a] Case Study 2 Prompt: Do in depth research on DC motor, and motor control technologies and their applications in mobile robotics platforms. You may use tables but minimize their use.

[2b] Case Study 2 ChatGPT Deep Research Response:

1) Why DC motors dominate mobile robotics

For mobile robots, the actuator has to do four things well: produce usable wheel torque at low speed, respond quickly to control commands, fit inside a battery-powered system, and interface cleanly with embedded electronics. That is why DC motor families—mainly brushed DC motors and brushless DC / PMSM motors—show up so often in mobile robotics. Brushed DC motors are easy to drive and model because torque is closely tied to armature current, which makes them attractive for simple wheeled robots and low-cost drivetrains. Brushless motors usually win when a robot needs better efficiency, longer service life, higher power density, smoother operation, or more advanced torque/speed control. In practice, that is why you see brushed DC on many small educational and budget differential-drive robots, while BLDC/PMSM solutions are common in higher-performance AMRs, AGVs, drones, and industrial mobile platforms. A second reason DC motors are so common is that they fit the way mobile robots are controlled. A robot controller usually thinks in layers: “I want the chassis to move at this linear and angular velocity,” then “I need each wheel at this speed,” then “I need each motor to generate this torque/current.” Differential-drive and car-like robots both map body motion into wheel motion in a very direct way, and wheel encoders can feed that back into odometry and low-level motor loops. ROS 2’s mobile robot kinematics documentation is a good reflection of how this is done in modern robotics software stacks: wheel speeds are derived from body twist, and odometry can be computed from encoder readings.

2) The main DC motor types used in mobile robots

Brushed DC motors

A brushed DC motor uses mechanical commutation through brushes and a commutator. Its biggest strengths in mobile robotics are simplicity, low controller complexity, easy reversal, and low entry cost. With a brushed motor, you can get bidirectional motion with an H-bridge and control speed with PWM. That makes brushed DC especially appealing in student robots, hobby platforms, compact service robots, and lower-cost AGV drive modules. TI’s brushed motor control application note shows a very typical embedded setup: a controller reads speed and current, runs PI/PID logic, and updates PWM duty cycle and direction through an H-bridge.

The tradeoff is that brushes wear, generate electrical noise, and become a maintenance and lifetime limitation. Brushed motors also tend to be less efficient and less scalable for very high-duty-cycle fleets than brushless alternatives. That does not make them bad for robots; it just means they are strongest where cost, design speed, and implementation simplicity matter more than maximum efficiency or maintenance-free life. In many mobile

robots, especially small differential-drive systems, brushed DC motors with gearboxes and encoders remain a completely valid choice because they are easy to integrate into closed-loop speed control.

Brushless DC motors / PMSMs

“BLDC” in practice overlaps heavily with permanent-magnet synchronous motors (PMSMs) in robotics hardware. These motors replace brushes with electronic commutation, which usually improves efficiency, power density, reliability, and noise performance. TI and ST both emphasize that advanced control methods like FOC improve efficiency, control accuracy, torque smoothness, and dynamic response in brushless drives. That makes brushless motors especially attractive in autonomous mobile robots that run for long periods from batteries, where drivetrain losses directly reduce range and usable runtime.

Brushless systems do ask more from the electronics and software. Instead of one H-bridge, a typical 3-phase BLDC drive uses a three-phase inverter, rotor position information or position estimation, commutation logic, and often a faster control MCU. That added complexity is the price you pay for better performance. On modern robotic platforms, though, that price is increasingly worth paying, especially when the platform needs precise torque control, quiet operation, regenerative braking, trajectory tracking, or high efficiency. ODrive’s robotics-oriented product line is a good real-world example of how the market has moved: it explicitly emphasizes FOC current control, torque/velocity/position modes, regenerative braking, and encoder support as core robotics features.

3) Why gearmotors are so common on mobile robots

Most mobile robots do not drive wheels directly from the raw motor shaft. They use a gearbox, often planetary or spur, because the wheel needs much more torque and much less speed than the motor naturally provides. This matters a lot in mobile robotics because the robot almost always operates near low or moderate wheel speeds where traction force, obstacle crossing, slope climbing, and start-up acceleration dominate. Real robotics case studies from maxon repeatedly show the same pattern: motors are paired with gearheads to reach practical output torque and packaging targets, whether on hull-cleaning robots, service robots, or articulated systems.

Planetary gearboxes are popular in mobile robots because they offer high reduction in compact volume and tend to handle torque better than many simple spur layouts of comparable size. Spur gear reductions can still make sense when cost, simplicity, and efficiency at lower reduction ratios matter. In mobile robotics, the gearbox choice is really a system trade: more reduction gives better low-speed torque and easier control authority, but it can also limit top speed, increase reflected inertia, add backlash, and reduce backdrivability. That becomes important if the robot relies on odometry accuracy, precise low-speed motion, or compliant interaction with terrain.

4) The control stack: from battery to wheel motion

A lot of confusion in robotics comes from mixing up motor, driver, controller, and feedback sensor. They are related, but not the same.

- The motor converts electrical energy into mechanical rotation.*
- The power driver / inverter switches current into the motor.*
- The controller computes what the driver should do.*
- The sensor reports speed, position, or current so the controller can close the loop.*

For a brushed DC wheel drive, the power stage is usually an H-bridge. NXP describes this class of device as a programmable H-bridge power IC designed to drive DC motors or bidirectional actuators, and TI's brushed-control note shows the normal software architecture: measure speed and current, run a speed loop and current loop, then update duty cycle and direction. That is the heart of modern low-level motion control.

For a brushless drive, the power stage is a 3-phase inverter. The controller must know when and how strongly to energize each phase. That can be done with simple six-step commutation, or with more sophisticated sinusoidal/FOC methods. The more advanced the commutation, the more the robot can get smooth torque, quiet operation, and accurate low-speed control.

5) Open-loop versus closed-loop motor control

Open-loop control

Open-loop control means the controller sends a command—often just PWM duty cycle—without verifying that the wheel actually achieved the intended speed or torque. This is the simplest possible approach and can work for toys, very low-cost platforms, or situations where accuracy is not important. But mobile robots rarely live in such a clean world. Terrain changes, floor friction changes, battery voltage sags, wheel wear, and left/right motor mismatch all push the robot away from the commanded motion. That is why a differential-drive robot controlled open-loop often drifts instead of moving straight.

Closed-loop control

Closed-loop control adds feedback—usually encoder speed/position and often current sensing. This is what most serious mobile robots use. A low-level controller compares commanded wheel speed to measured wheel speed and keeps correcting the PWM/inverter output. The result is more repeatable velocity, better straight-line tracking, better turning accuracy, and usable odometry. The ROS 2 kinematics documentation explicitly notes that odometry can be calculated from encoder readings, and TI's encoder note explains how incremental encoders provide pulses that indicate motion, speed, and direction.

In practice, most mobile robots benefit from nested loops:

- 1. an inner current/torque loop,*
- 2. a middle speed loop,*
- 3. and sometimes an outer position/trajectory loop.*

That structure matters because current changes motor torque quickly, while speed and position evolve more slowly. TI's recent brushed-motor note is a clear example of this thinking: speed is controlled through feedback, the current is limited or regulated, and the duty cycle is computed from those loops.

6) PWM, H-bridges, and why they matter so much in robotics

For brushed DC motors, pulse-width modulation (PWM) is the standard method for controlling average motor voltage and therefore speed/torque behavior. The H-bridge lets the system reverse polarity across the motor so the wheel can move forward or backward. This is foundational in mobile robotics because nearly every wheeled robot needs reversible traction. NXP's H-bridge documents and application material describe these drivers as the basic means to apply voltage across a motor in either direction, while TI's note shows how that switching is integrated with sensed speed/current loops.

The important engineering point is that PWM is not the same thing as closed-loop control. PWM is just the actuator command method. A robot can use PWM open-loop, which is crude, or PWM inside a closed-loop controller, which is much better. In mobile robotics, the second approach is the one that actually delivers consistent path following and odometry. This matters especially for differential-drive platforms, where tiny left/right speed differences accumulate into heading error very quickly.

7) Motor feedback technologies used on mobile robots

Incremental encoders

Incremental rotary encoders are the most common feedback device on mobile robots. TI describes them as devices that translate rotational movement into electrical signals, with pulses that can indicate speed and direction. In practice, they are popular because they are affordable, easy to read with microcontrollers, and directly useful for wheel speed control and odometry. ROS-based wheeled robots commonly use encoder counts to estimate travel distance and heading change.

The limitation is that incremental encoders tell you motion relative to a starting point, not absolute shaft angle at power-up. For many wheeled robots that is perfectly fine. You command the robot to start from a known state, then integrate wheel motion forward. But if you need instant startup position knowledge or you need to recover from resets gracefully, incremental sensing alone can be limiting.

Absolute encoders

Absolute encoders report shaft angle directly, even after power cycling. maxon's ENX encoder documentation notes that some compact encoder families can provide incremental square-wave signals and also absolute angular values via interfaces such as SSI or BiSS-C. In mobile robotics, absolute sensing is especially attractive when a platform needs precise startup alignment, accurate wheel module initialization, or robust position knowledge in higher-end servo systems.

Hall sensors

Hall sensors are common in BLDC motors for commutation and coarse position feedback. They are cheaper and simpler than high-resolution encoders, but they provide much lower angular resolution. maxon's EPOS4 application notes point out that observers become especially helpful when using low-resolution feedback like Hall sensors or relatively low-count incremental encoders, because they can produce better velocity estimation and improved low-speed regulation. That is very relevant in mobile robots that need quiet, smooth crawl speeds.

8) Brushed motor control technologies in mobile robotics

For brushed wheel motors, the most common control architectures are:

a) Voltage / PWM control

This is the simplest: command PWM and accept whatever speed results. It is cheap and easy but sensitive to load and battery variation. It is usually acceptable only for the simplest robots.

b) Closed-loop speed control with encoder feedback

This is the standard practical upgrade. The encoder measures wheel speed, and a PI or PID controller adjusts PWM so the wheel tracks the target speed. The 2022 mobile robot study on a brushed DC motor with a low-cost magnetic quadrature encoder is a good example of how useful even modest encoder-based PID control can be for overshoot, undershoot, and velocity response in a robot application.

c) Current-limited or current-controlled brushed drive

Here the system monitors motor current to protect hardware and sometimes to regulate torque more directly. This matters in mobile robots because wheel stall, curb impacts, carpet transitions, and startup acceleration can pull much higher current than steady cruising. TI's sensed brushed-motor control note explicitly uses speed and current information together, which is a strong reflection of how real robotic drives are implemented.

d) Position control

This is more common for steering actuators, pan/tilt mechanisms, and robotic appendages than for pure traction wheels, but it can also appear in wheel modules or docking mechanisms. maxon's position-control materials show how encoder interfaces and servo controllers are used when precise commanded position matters more than just raw wheel speed.

9) Brushless motor control technologies in mobile robotics

a) Six-step / trapezoidal commutation

This is the classic BLDC control method. The controller energizes the three phases in six discrete sectors per electrical cycle. Microchip's and ST's application notes describe six-step / trapezoidal / 120° commutation as a standard sensorless BLDC method. Its biggest strengths are simplicity, lower compute burden, and easier implementation compared with FOC. That makes it attractive in cost-sensitive robots, fans, pumps, and some basic traction systems.

The downside is higher torque ripple, more acoustic noise, and weaker low-speed smoothness than FOC. For a mobile robot, that can show up as jerkier crawl behavior, poorer precision at very low wheel speed, and more difficulty holding smooth motion under changing load. On a robot that only needs coarse speed control, six-step may still be enough. On a robot that needs precise indoor navigation, smooth docking, or silent operation, it is often not the best end state.

b) Sensorless BLDC control

Sensorless control removes physical rotor-position sensors and estimates commutation timing from motor signals such as back-EMF. Microchip notes that its sensorless six-step approach works on a wide range of motors, does not require detailed motor parameters, and is relatively insensitive to manufacturing variation. That is attractive for cost and wiring reduction.

The catch is that sensorless methods are usually weakest at very low speed or standstill, because back-EMF becomes hard to measure there. For mobile robots, that is a real concern: robots spend a lot of time starting, stopping, creeping, docking, and making small corrections. So sensorless BLDC can be excellent in the right operating region, but sensed systems usually behave better when a robot must deliver confident low-speed traction and repeatable startup torque.

c) Sensed BLDC with Hall or encoder feedback

Adding Hall sensors or an encoder improves startup and low-speed control. Hall sensors are often enough for commutation and moderate speed regulation, while high-resolution encoders support much better velocity and position estimation. NXP's older BLDC application note explicitly targets quadrature-encoder-based 3-phase BLDC control, which is still conceptually aligned with how robotic servo drives are built today.

d) Field-oriented control (FOC)

FOC is the most important advanced motor-control technology to understand for modern mobile robotics. TI describes FOC as making the stator and rotor magnetic fields orthogonal to achieve maximum electromagnetic torque, while improving efficiency, smoothness, and dynamic response; TI also notes that it can support operation above nominal speed through field weakening. ST similarly positions FOC as a way to improve control and accuracy for brushless motors.

Why that matters on a robot:

- smoother low-speed motion,*
- lower torque ripple,*
- better efficiency from a battery pack,*
- better torque response when load changes,*
- better suitability for torque, speed, and position loops.*

That is a big reason FOC is so common in modern robotic wheel drives, gimbals, and servo modules. ODrive's robotics products, for example, highlight FOC current control together with torque, velocity, position, trajectory control, encoder support, and

regenerative braking—exactly the feature set you expect when the motor controller is being designed for serious robot motion rather than just “spin the motor.”

10) Regenerative braking and energy management

In mobile robots, the drivetrain is part of the power system, not just the motion system. When the robot decelerates, goes downhill, or is externally driven, a motor can act as a generator. Some higher-end brushless controllers support regenerative braking, returning energy to the DC bus or safely dissipating it through a brake resistor if the battery cannot absorb it fast enough. ODrive explicitly advertises regenerative braking and brake-chopper capability for robotics applications.

This matters more as robots get heavier and faster. On a tiny classroom robot, regen may not be worth the complexity. On an AGV or AMR carrying useful payload, it can improve energy efficiency and thermal behavior. The design challenge is that regenerated energy has to go somewhere safe. If the battery is full or cannot accept charge fast enough, the DC bus can overvoltage unless the controller provides a dissipation path. That is one reason power electronics and battery management become inseparable in professional mobile robot design.

11) Applications by robot class

Here is the most useful way to think about actuator choice: the right motor technology depends on the robot’s operating regime.

<i>Robot/application class</i>	<i>Typical motor choice</i>	<i>Why</i>
<i>Small educational or hobby differential-drive robot</i>	<i>Brushed DC gearmotor + incremental encoder</i>	<i>Cheapest path to closed-loop wheel control and odometry; easy H-bridge drive.</i>
<i>Student design robot / light service platform</i>	<i>Brushed DC gearmotor or low-voltage BLDC</i>	<i>Brushed for simplicity, BLDC if efficiency/runtime/noise matter more.</i>
<i>Industrial AMR / AGV</i>	<i>BLDC/PMSM with sensored FOC</i>	<i>Better efficiency, longer life, smoother motion, stronger low-speed control.</i>
<i>Harsh-environment or long-duty-cycle mobile robot</i>	<i>BLDC + gearbox + sealed/robust feedback</i>	<i>Better reliability and lower maintenance under sustained operation.</i>
<i>Precision mobile platform or wheel module</i>	<i>BLDC/PMSM + encoder + servo drive</i>	<i>Better torque, velocity, and position regulation.</i>

That table is simplified, but the pattern is real: brushed DC tends to win on simplicity and cost; brushless tends to win on performance, efficiency, and lifecycle.

12) Motor control and robot kinematics are tightly linked

On mobile robots, motor control cannot be separated from locomotion architecture.

For a differential-drive robot, the controller independently sets left and right wheel speeds. ROS 2's equations show that body linear velocity is the average of left/right wheel velocities, while angular velocity depends on their difference divided by track width. That means wheel-speed control quality directly determines path quality. If one wheel is poorly controlled, the robot's heading will drift.

For a car-like or Ackermann robot, the traction motor and the steering actuator play different roles. The drive motor largely determines longitudinal motion, while the steering motor determines curvature. That often means the drive motor can be speed-focused while the steering motor needs tighter position control. Again, encoder placement matters: the ROS 2 documentation explicitly distinguishes odometry cases based on where encoders are mounted.

For tracked robots, drivetrain control gets harder because track-ground interaction introduces more slip and uneven resistance. In those robots, strong current limiting and good wheel/track speed feedback become even more important, because open-loop voltage control alone is usually not enough to keep motion predictable. The hull-cleaning robot case study is a good reminder that motor selection on battery-operated mobile robots is inseparable from friction, sealing, payload, and environmental loading.

13) The role of observers, estimation, and sensor fusion

A common beginner assumption is that "an encoder solves everything." It does not. Encoders give you shaft motion, but mobile robots still suffer from wheel slip, backlash, compliance, and sensor noise. That is why practical systems often combine low-level motor feedback with higher-level estimation from IMUs and localization. maxon's EPOS4 application note notes that observers can improve velocity estimation and low-speed regulation when feedback resolution is limited, and the Nav2 documentation shows how odometry is commonly fused with other sources for smoother robot state estimation.

So the real control story in mobile robotics is layered:

- *motor current control for torque,*
- *wheel-speed control for traction,*
- *wheel odometry for local motion estimate,*
- *IMU/localization fusion for a better pose estimate,*
- *planner/controller above that.*

That layered architecture is one reason motor and encoder selection matter so much more than they first appear to.

14) What usually drives actuator choice in real projects

In real mobile robot design, the choice is rarely "which motor is theoretically best?" It is usually "which motor-control stack best fits the platform constraints?" The most important constraints are:

Battery efficiency and runtime.

Brushless + FOC often wins for sustained duty because better efficiency means less heat and more runtime per charge.

Control simplicity and development time.

Brushed DC + H-bridge + encoder is often the fastest path to a working robot.

Low-speed smoothness.

If the robot must crawl, dock, or move precisely indoors, sensed BLDC with good control algorithms becomes much more attractive.

Maintenance and lifecycle.

Brushless systems usually win for long-lived fleets because they eliminate brushes as a wear item.

Cost and integration burden.

The cheapest motor is not always the cheapest drivetrain once driver complexity, sensing, firmware, and debugging are included. That is why many student teams still choose brushed gearmotors even when BLDC looks better on paper.

15) Bottom-line conclusions

If you strip away the marketing language, the practical picture is pretty clear.

Brushed DC motors are still highly relevant in mobile robotics because they are easy to drive, cheap to prototype with, and pair naturally with H-bridges, PWM, and encoder-based PID speed loops. They are especially good for educational robots, prototypes, low-cost differential drives, and systems where fast implementation matters more than peak efficiency.

Brushless DC / PMSM motors are the stronger long-term solution for many serious mobile robots because they offer better efficiency, smoother torque, better durability, and better compatibility with modern control methods like FOC. They are especially attractive in AMRs, AGVs, service robots, and battery-powered platforms that need long duty cycles, precise low-speed motion, or advanced trajectory control.

The most important control technologies to know for mobile robotics are:

- *PWM + H-bridge for brushed drives,*
- *encoder-based closed-loop speed control for repeatable wheel motion,*
- *current limiting/current control for torque and protection,*
- *six-step commutation for simpler BLDC systems,*
- *FOC for higher-performance brushless traction and servo motion,*
- *odometry from encoders tied into the robot's kinematic model,*
- *and sensor fusion/observers when raw encoder data is not enough.*

For most mobile robotics platforms, the best mental model is this:

the motor is only one part of the drivetrain; the real capability comes from the motor + gearbox + driver + sensors + control algorithm acting as a single system.

If you want, I can turn this into a report-style section with headings and references, or rewrite it specifically for a senior design mobile robot paper.


[2c] Case Study 2 Grok Deep Research Response:

DC motors serve as the workhorse actuators in mobile robotics platforms, prized for their high starting torque, variable speed control, compact size, and direct compatibility with battery power sources. Mobile robots—whether hobbyist differential-drive bots, warehouse autonomous mobile robots (AMRs), or research platforms like TurtleBot—rely on them for propulsion, precise maneuvering, and reliable operation in dynamic environments. This research draws from engineering fundamentals, control theory, and real-world robotic applications to explore DC motors, their control technologies, and integration in mobile systems.


Fundamentals of DC Motors: Brushed vs. Brushless

DC motors convert electrical energy into mechanical rotation through the interaction of magnetic fields (Lorentz force principle). Current in a conductor within a magnetic field produces torque, causing the rotor to spin. Two primary types dominate robotics: brushed and brushless (BLDC) DC motors. Both are permanent-magnet designs, but they differ fundamentally in commutation—the process of switching current to maintain continuous rotation.

Brushed DC motors use a mechanical commutator (segmented copper rings on the rotor) and carbon or precious-metal brushes to switch current direction in the armature windings as the rotor turns. The stator holds permanent magnets. This self-commutating design allows simple direct connection to a DC supply, making them inexpensive and easy to drive. However, brushes wear out over time (typically 1,000–5,000 hours), cause electrical arcing/sparking, generate electromagnetic interference (EMI), and produce audible noise. Efficiency is lower (around 60–80%) due to friction and brush losses, and they require more maintenance.



[youtube.com](https://www.youtube.com)



precisionmicrodrives.com

Brushless DC (BLDC) motors reverse the configuration: permanent magnets are on the rotor, and windings are on the stator. Electronic commutation (via an external controller) replaces brushes, using Hall-effect sensors, encoders, or sensorless back-EMF detection to time phase energization. This yields higher efficiency (often 85–95%), greater torque density, higher speeds (up to tens of thousands of RPM), quieter operation, longer life (20,000+ hours with minimal maintenance), and better thermal characteristics since heat-generating windings are stationary. Drawbacks include higher upfront cost and the need for sophisticated drivers. BLDC motors excel in battery-powered mobile robots where efficiency directly extends runtime.



lamnow.com

Brushless DC Motor Stator Winding | Lamnow

In practice, both types are often paired with gearboxes (e.g., planetary or spur) to increase torque and reduce speed for robotic traction needs. Gear motors with encoders are standard for odometry in mobile platforms.


Motor Control Technologies

Raw DC motors respond to voltage for speed (via back-EMF) and current for torque, but practical control in robotics demands precision for velocity, position, direction, and load compensation. Key technologies include:

Pulse Width Modulation (PWM) is the cornerstone for speed control. A microcontroller or driver rapidly switches full voltage on/off at a fixed frequency (typically 1–20 kHz to avoid audible noise while keeping switching losses low). The duty cycle (percentage of "on" time) sets the average voltage delivered to the motor, effectively controlling speed without wasting power as heat (unlike linear regulators). In mobile robots, PWM enables smooth acceleration and energy-efficient cruising.

H-bridge circuits handle bidirectional control (forward/reverse) and braking. Four switches (transistors or MOSFETs) form an "H" shape around the motor. Activating opposite pairs (e.g., top-left and bottom-right) reverses polarity. Modern implementations use complementary PWM with dead-time insertion to prevent shoot-through (simultaneous conduction causing shorts). Integrated drivers like the L298N (for brushed) or three-phase bridge drivers for BLDC simplify this. Braking modes (coast, slow decay, fast decay) further tune responsiveness.

hackatronic.com




circuitdigest.com

Feedback systems close the control loop. For BLDC commutation, three Hall-effect sensors (60° or 120° spacing) detect rotor position via magnetic fields. For finer control,

quadrature encoders (optical or magnetic) on the shaft provide high-resolution position/speed feedback (hundreds to thousands of counts per revolution). Encoders enable odometry—calculating robot pose from wheel rotations—in SLAM and navigation. Sensorless techniques for BLDC estimate position from back-EMF voltage during rotation, reducing wiring, cost, and failure points, though they perform poorly at very low speeds.

PID (Proportional-Integral-Derivative) control is ubiquitous for closed-loop performance. It compares measured speed/position (from encoders) against a setpoint, computing an error signal. The proportional term reacts to current error, integral eliminates steady-state offset (e.g., from friction or slope), and derivative damps overshoot. Tuned via Ziegler-Nichols or manual methods (often with ROS2 `diff_drive_controller`), PID ensures a robot maintains velocity despite uneven terrain, payload changes, or wheel slip. Cascaded loops (outer position, inner velocity) are common in advanced platforms.



keyirobot.com

Deep Dive: Understanding PID Control for Robot Motor Stabilization


For BLDC motors, advanced algorithms shine: trapezoidal (6-step) commutation is simple but causes torque ripple; sinusoidal reduces ripple; Field-Oriented Control (FOC)—also called vector control—decouples torque and flux currents (like separately excited DC motors) using Clarke/Park transforms and PI regulators. FOC delivers smooth, efficient, high-dynamic performance ideal for precise robotic motion. Modern microcontrollers (e.g., with hardware FOC peripherals) or dedicated ICs make it accessible. Sensorless FOC variants are a growing trend for compact AMRs.

Brushless Electronic Speed Controllers (ESCs) integrate these functions for hobby robotics, while industrial drivers add CAN/EtherCAT for fleet coordination.

Applications in Mobile Robotics Platforms


Mobile robots emphasize locomotion, so DC motors drive wheels, tracks, or legs with emphasis on torque for acceleration, efficiency for battery life, and controllability for

navigation. The most common configuration is differential drive: two independently controlled motors (usually geared DC or BLDC) on opposite sides, plus casters or omniwheels for stability. Varying relative speeds enables forward/backward motion, in-place rotation (zero turning radius), and curved paths—kinematics are straightforward via wheel odometry.



forum.dronebotworkshop.com

Wheeled Robot Internal Kinematics



forum.dronebotworkshop.com

Examples across scales:

- Hobby/educational platforms (e.g., Arduino or Raspberry Pi bots) use low-cost brushed gear motors with H-bridge shields and basic PWM/PID for line-following or obstacle avoidance.
- Research robots like TurtleBot employ BLDC or brushed motors with high-resolution encoders for ROS-based SLAM, mapping, and autonomous navigation.
- Industrial AMRs/AGVs in warehouses (e.g., for logistics) favor BLDC gearmotors with FOC or sensorless control for quiet, efficient 24/7 operation under varying loads. High starting torque handles inclines or sudden stops.

- *Tracked robots (e.g., search-and-rescue) use dual-motor differential steering for rough terrain.*
- *Delivery or agricultural robots benefit from BLDC's power density for longer range.*

Advantages in mobile contexts: DC motors provide instant high torque for starting heavy payloads or climbing, precise speed variation for smooth trajectories, and regenerative braking for energy recovery. Battery compatibility eliminates bulky inverters needed for AC motors. In differential setups, independent motor control simplifies software (e.g., pure pursuit or PID-based path following).

Challenges include wheel slip (mitigated by PID + IMU fusion), EMI from brushes (favoring BLDC in sensor-heavy robots), and thermal management in enclosed chassis. Gear reduction trades speed for torque but introduces backlash, addressed by high-quality encoders.

Current trends (as of 2026): The shift toward BLDC with FOC and sensorless control accelerates due to efficiency gains, falling controller costs, and integration in compact drives. Market growth reflects demand in robotics and EVs, with AI-tuned adaptive PID or neural controllers emerging for uncertain environments. Frameless or integrated gearmotors reduce weight, while functional safety features (e.g., safe torque off) appear in collaborative AMRs.

In summary, DC motors—especially BLDC variants with modern electronic controls—enable the mobility, precision, and endurance that define today's mobile robotics. From simple PWM/H-bridge hobby bots to sophisticated FOC-driven industrial fleets, these technologies balance cost, performance, and reliability, with ongoing advances in sensorless and intelligent control pushing the boundaries further. For custom platforms, motor selection hinges on torque/speed requirements, battery voltage, and control complexity, always validated through simulation and real-world testing.

[3a] Case Study 3 Prompt: When selecting an ultrasonic sensor for a design project, what are some of the most important considerations to be aware of?

[3b] Case Study 3 Claude Sonnet 4.6 Response:

Here are the key considerations when selecting an ultrasonic sensor for a design project:

Range The sensor's minimum and maximum detection distance must match your application. Most ultrasonic sensors have a blind spot (minimum range) where they can't detect objects — for close-range work, this matters a lot.

The MaxBotix sensors, for example, have variants tuned for short vs. long range.

Beam Width / Beam Pattern Ultrasonic sensors emit a cone-shaped beam. A narrow beam gives more precise directional detection; a wide beam covers more area but can pick up unintended objects (walls, floors, etc.). Know what your sensor's beam angle is and whether it suits your environment.

Resolution and Accuracy These are different things — resolution is the smallest change the sensor can detect; accuracy is how close the reading is to the actual distance. For precise obstacle avoidance or measurement tasks, both matter.

Output Type Sensors typically output via one of: PWM (pulse width), analog voltage, UART serial, or I2C. Choose based on what your microcontroller supports and how many sensors you're running simultaneously.

Update Rate How fast does the sensor report readings? For a moving robot, a slow update rate means stale distance data, which can cause missed obstacles or jittery control. Higher update rates are better for dynamic environments.

Environmental Sensitivity Ultrasonic sensors are affected by temperature (speed of sound changes with temperature), humidity, and surface material (soft/angled surfaces absorb or deflect the signal). Outdoor or variable-temp environments need more care here.

Power Consumption Relevant if you're on a battery-powered platform. Some sensors have sleep/low-power modes; others are always-on.

Cross-Talk (Multi-Sensor Interference) If you're running multiple ultrasonic sensors simultaneously, their signals can interfere with each other. You'll need staggered firing (multiplexing) or sensors designed for multi-sensor setups.

Form Factor and Mounting Physical size, mounting hole pattern, and whether the sensor faces forward-only vs. has a transducer you can aim — all affect integration into your mechanical design.

Cost and Availability For a student/prototype project, the difference between a \$4 HC-SR04 and a \$25+ MaxBotix unit is significant. Higher-end sensors earn their price in better beam control, more reliable output, and built-in filtering.

[4a] Case Study 4 Prompt: *When programming an ESP32-S3 WROOM1 how can you design for low power?*

[4b] Case Study 4 Gemini 3 Response:

Designing for low power on the ESP32-S3 is a bit like managing a high-performance sports car: it's incredibly capable, but it will drain the tank quickly if you leave the engine idling. To get the most out of a battery, you need to move beyond simple "delay" loops and utilize the chip's sophisticated power-management hardware.

1. Leverage Sleep Modes

The ESP32-S3 offers several power states. Choosing the right one depends on whether you need to maintain RAM or keep the CPU alive.

<i>Mode</i>	<i>CPU</i>	<i>RTC Memory / ULP</i>	<i>Typical Current</i>	<i>Use Case</i>
<i>Active</i>	<i>On</i>	<i>On</i>	<i>40mA – 240mA</i>	<i>Processing, WiFi/BT active.</i>

<i>Light Sleep</i>	<i>Paused</i>	<i>On</i>	<i>~240μA</i>	<i>Maintaining RAM and connections.</i>
<i>Deep Sleep</i>	<i>Off</i>	<i>On</i>	<i>~7μA – 10μA</i>	<i>Long-term sensing; periodic wakeups.</i>
<i>Hibernation</i>	<i>Off</i>	<i>Off</i>	<i>~1μA</i>	<i>Rare events; wake via RTC timer/GPIO.</i>

Implementation Tip:

Use `esp_deep_sleep_start()` for long intervals. Remember that when the chip wakes from Deep Sleep, it reboots—your code starts from the beginning of `setup()`. Use the `RTC_DATA_ATTR` attribute to store variables you want to persist through the sleep cycle.

2. The Ultra-Low-Power (ULP) Co-processor

The S3 features an FSM (Finite State Machine) and a RISC-V ULP co-processor. This allows you to monitor sensors (via I2C or ADC) while the main "power-hungry" cores are completely off.

- The Strategy: Program the ULP to check a sensor threshold. If the threshold is met, the ULP sends a wake-up signal to the main cores.*

3. Radio Management (The "Power Hogs")

WiFi and Bluetooth are the primary sources of battery drain.

- WiFi Modem Sleep: Use this if you need to stay connected to an AP but don't need to transmit constantly. It shuts down the radio between DTIM (Delivery Traffic Indication Message) beacons.*
- Protocol Choice: If your project allows, use ESP-NOW. It's a connectionless protocol that bypasses the lengthy WiFi handshake, allowing you to wake up, blast a data packet, and sleep in milliseconds.*

4. Hardware and Peripheral Optimization

Software can't fix bad hardware design. Consider these "stealth" power drains:

- Dynamic Voltage and Frequency Scaling (DVFS): Lower the CPU frequency if you aren't doing heavy math. Dropping from 240MHz to 80MHz significantly reduces current.*
- Power the Peripherals Smartly: Use a MOSFET or a dedicated Load Switch to physically cut power to external sensors or displays when they aren't in use.*

- *Internal Pull-ups: Be careful with GPIO pull-ups/pull-downs during sleep. A single 10k Ω resistor tied to ground on a high pin can leak 330 μ A—more than the entire chip uses in Light Sleep!*

5. Software Best Practices

- *Avoid delay(): In Active mode, delay() keeps the CPU spinning. Use FreeRTOS tasks and vTaskDelay() to allow the scheduler to put the CPU into a low-power state during idle periods.*
- *Batch Data: Instead of waking up every minute to send one data point, store 60 points in RTC RAM and send them all at once every hour. This minimizes the time spent in the high-current "Radio Active" state.*

Are you aiming for a device that stays "connected" (like a smart light) or one that wakes up once an hour to report data (like a soil moisture sensor)?

[4c] Case Study 4 ChatGPT Response

Designing for low power on the ESP32-S3-WROOM-1 involves a combination of hardware design and firmware strategies. The key is to minimize active time and reduce current consumption in all operating modes.

1. Use the Right Power Modes

Deep Sleep:

- *Current: ~5–10 μ A (ideal)*
- *CPU and main RAM are powered down; only RTC domain is active*
- *Wake sources: timer, GPIO, ULP coprocessor*
- *Best for infrequent operation (e.g., periodic sensor readings)*

Light Sleep:

- *Current: ~100–800 μ A*
- *CPU paused, RAM retained*
- *Faster wake-up than deep sleep*
- *Useful when responsiveness is needed*

Modem Sleep:

- *CPU stays active, WiFi/BLE duty-cycled*
- *Useful when maintaining connections but reducing power*

2. Control WiFi and Bluetooth

- *Disable radios when not needed:*
esp_wifi_stop();
esp_bt_controller_disable();
- *Only enable WiFi when transmitting data*
- *Batch transmissions instead of sending frequently*

Typical pattern:

Wake → connect → send → disconnect → sleep

3. *Reduce CPU Usage*

- *Lower CPU frequency using power management configuration*
- *Avoid busy loops; use interrupts or task delays*
- *Use FreeRTOS blocking calls like vTaskDelay()*

4. *Use RTC and ULP Coprocessor*

- *The ULP can run while the main CPU is in deep sleep*
- *It can monitor sensors and wake the CPU only when needed*

Example:

- *Periodically check a sensor*
- *Wake main CPU only if threshold is exceeded*

5. *Manage Peripherals*

- *Turn off unused peripherals (ADC, I2C, SPI, etc.)*
- *Disable or minimize LED usage*
- *Configure pins to avoid floating states*

6. *Optimize GPIO States*

- *Set unused GPIOs as inputs with defined pull states:*
`gpio_set_direction(pin, GPIO_MODE_INPUT);`
`gpio_pullup_dis(pin);`
`gpio_pulldown_en(pin);`

7. *Use Efficient Voltage Regulation*

- *Avoid high quiescent current LDOs*
- *Use low-IQ LDOs or switching regulators (buck converters)*

8. *Power-Gate External Components*

- *Use MOSFETs or load switches to cut power to sensors*
- *Only power external devices when needed*

9. *Minimize Wake Time*

- *Keep active time as short as possible*
- *Avoid unnecessary initialization*
- *Store data in RTC memory if possible*

10. *Measure Power Consumption*

- *Use measurement tools to verify real consumption:*
 - *USB power meter*

- Multimeter with shunt resistor
- Oscilloscope or power profiler

Typical Low-Power Workflow:

1. Wake from deep sleep
2. Read sensor
3. Enable WiFi
4. Send data
5. Disable WiFi
6. Return to deep sleep

Common Mistakes:

- Leaving WiFi always on
- Not disabling Bluetooth
- Using development boards with high idle current
- Ignoring regulator quiescent current
- Not using deep sleep

Typical Current Consumption:

- Deep sleep: $\sim 10\ \mu A$
- Light sleep: $\sim 0.5\ mA$
- WiFi active: 80–240 mA

If you want, I can tailor this into a design checklist or schematic-level recommendations for your specific project.

[4d] Case Study 4 Prompt: *What are some of the key ideas when programming in C/C++ to control the chip for low power design?*

[4e] Case Study 4 Gemini 3 Response:

When programming the ESP32-S3 in C/C++, your goal is to move from "Linear Programming" (doing one thing after another) to "Event-Driven Programming." You want the chip to spend the absolute minimum amount of time in an active state. Here are the key programming concepts to master for low-power design:

1. RTOS Task Management

*The ESP32-S3 runs **FreeRTOS** by default. A common mistake is using while(1) loops with empty iterations or delay().*

- **The Power-Saving Way:** *Use vTaskDelay() or **Task Notifications**. When a task calls vTaskDelay(), the scheduler puts that task into a "Blocked" state, allowing the Idle Task to run. If "Automatic Light Sleep" is enabled in the SDK*

configuration, the chip will automatically enter a low-power state during this idle time.

2. Using RTC Memory for Persistence

When you use **Deep Sleep**, the main RAM is powered off. Any global variables you've defined will be lost.

- **The Attribute:** Use the `RTC_DATA_ATTR` to store data in the RTC slow memory, which remains powered.
- **C++ Example:**
- C++

```
RTC_DATA_ATTR int bootCount = 0; // This survives deep sleep
```

-
- This is vital for **Batching**: instead of turning on the WiFi (high power) every time a sensor reads, store 50 readings in RTC memory and only trigger the radio once the memory is full.

3. Interrupt-Driven Wakeups

Instead of "polling" a pin (checking its state every few milliseconds), configure the hardware to wake the CPU only when a state change occurs.

- **External Wakeup:** Use `esp_sleep_enable_ext0_wakeup()` or `ext1` to define specific GPIOs that can trigger a wakeup from deep sleep.
- **Timer Wakeup:** Use `esp_sleep_enable_timer_wakeup()` for precise intervals.

4. Peripherals and Power Domains

The ESP32-S3 has several internal power domains (CPU, WiFi, RTC, etc.). In C/C++, you can manually toggle these to save micro-amps.

- **Isolate Pins:** When sleeping, pins can "float," causing current leakage. Use `rtc_gpio_isolate()` on pins connected to external components to ensure they don't draw power while the chip is asleep.
- **Power Down Peripherals:** If you aren't using the ADC or the Wi-Fi/BT radio, ensure they are explicitly powered down using `esp_wifi_stop()` and `adc_power_release()`.

5. Frequency Scaling (DFS)

You don't always need to run at 240MHz. If you are just reading an I2C sensor, you can downclock the CPU to 80MHz or even lower.

- **Dynamic Frequency Scaling:** You can configure the `esp_pm_config_esp32s3_t` structure to allow the system to automatically scale the frequency based on the CPU load. Lowering the frequency reduces the operating voltage, which has a squared effect on power savings: $P = V^2 \cdot f \cdot C$.

6. Efficient Communication (ESP-NOW)

If you are coding a C++ application that needs to talk to another ESP32, avoid the standard TCP/IP stack.

- **ESP-NOW:** This is a "connectionless" protocol. You can wake up, send a packet in about **3-5 milliseconds**, and go back to sleep. A standard WiFi connection can take **1-3 seconds** just to associate with an AP, consuming massive amounts of energy in the process.

Summary Checklist

Strategy	C++ Implementation
Avoid Polling	Use GPIO Interrupts / <code>esp_sleep_enable_ext0_wakeup</code>
Data Persistence	<code>RTC_DATA_ATTR</code> for variables
Radio Control	<code>esp_wifi_stop()</code> when not in use
Idle Time	Enable <code>CONFIG_PM_ENABLE</code> (Automatic Light Sleep)

Are you planning to power your device with a small LiPo battery or something like a couple of AA batteries?

[4f] Case Study 4 ChatGPT Response:

Here are the **key C/C++ programming ideas** specifically for low-power control on the ESP32-S3 (focused on firmware techniques rather than hardware):

1. Sleep as the Default State

Design your code so the chip is **almost always sleeping**.

- Main loop should quickly finish work and enter sleep:
`esp_sleep_enable_timer_wakeup(5 * 1000000); // 5 seconds`
`esp_deep_sleep_start();`

Think of your program as:

- "Wake → do minimal work → sleep again"

2. Event-Driven Instead of Polling

Polling wastes CPU cycles and power.

Avoid:

```
while(1) {
    if (sensor_ready()) { ... }
}
Use interrupts:
gpio_install_isr_service(0);
gpio_isr_handler_add(GPIO_NUM_X, isr_handler, NULL);
```

CPU sleeps until something actually happens.

3. Minimize Active Time

Keep execution short and efficient.

- *Avoid unnecessary delays*
- *Precompute values where possible*
- *Reduce logging (UART is slow and power-hungry)*

Bad:

```
printf("Debug info...\n");
```

Better:

- *Disable logs in production:*

```
esp_log_level_set("*", ESP_LOG_NONE);
```

4. Control WiFi/BLE in Code

Explicitly manage radios — never leave them on by default.

```
esp_wifi_start();
// transmit data
esp_wifi_stop();
```

- *Initialize only when needed*
- *Deinitialize after use:*

```
esp_wifi_deinit();
```

5. Use Power Management APIs

Enable dynamic frequency scaling and light sleep:

```
esp_pm_config_t pm_config = {
    .max_freq_mhz = 240,
    .min_freq_mhz = 40,
    .light_sleep_enable = true
};
```

```
esp_pm_configure(&pm_config);
```

CPU slows down automatically when idle.

6. Avoid Busy Waiting

Busy loops burn power unnecessarily.

Avoid:

```
while (!flag) { }
```

Use:

- *Task blocking:*

```
xQueueReceive(queue, &data, portMAX_DELAY);
```

- *Or delays:*

```
vTaskDelay(pdMS_TO_TICKS(100));
```

7. Use RTC Memory for Persistence

Avoid full reinitialization after deep sleep.

```
RTC_DATA_ATTR int boot_count = 0;
```

- *Keeps values across deep sleep cycles*
- *Reduces computation and wake time*

8. Shut Down Unused Peripherals in Code

Explicitly disable modules:

```
adc_power_off();
```

```
periph_module_disable(PERIPH_I2C0_MODULE);
```

- *Don't rely on defaults*
- *Turn off what you don't use*

9. Optimize Task Design (FreeRTOS)

Poor task design = wasted power.

- *Use fewer tasks*
- *Let tasks block instead of spin*
- *Use appropriate priorities*

Idle CPU = opportunity for light sleep

10. Reduce Wake Frequency

Even efficient wake cycles consume energy.

- *Increase interval between wakeups*
- *Batch sensor readings or transmissions*

11. Conditional Work (Do Less!)

Only perform work when necessary.

```
if (temperature > threshold) {  
    send_data();  
}
```

Don't wake WiFi or CPU for trivial updates.

12. Fast Boot Strategy

After deep sleep reset:

- *Skip unnecessary init*
- *Cache configs in RTC memory*
- *Avoid recalibration if unchanged*

13. Use the ULP When Possible

Offload simple tasks to the ultra-low-power coprocessor.

- *Monitor GPIO or ADC*
- *Wake main CPU only when needed*

14. Compile-Time Optimization

Use compiler settings to reduce runtime cost:

- *Optimize for size/power:*
 - *-Os instead of -O0*
- *Remove unused code*

15. Be Careful with Timing APIs

Some APIs prevent sleep.

- *Avoid tight timers that keep CPU awake*
- *Use long blocking waits instead*

Core Mindset Shift

Traditional programming:

→ *“Run continuously and react”*

Low-power embedded programming:

→ *“Sleep always, wake briefly, do the minimum, go back to sleep”*

Common Firmware Mistakes

- *Infinite loops without delays*
- *Leaving WiFi initialized*
- *Excessive logging*
- *Polling sensors constantly*
- *Not entering deep sleep at all*

*If you want, I can show you a **complete minimal low-power ESP32-S3 example (ready to compile)** that demonstrates all of these ideas in one program.*